A local rule switch destabilizes a cellular automaton market model with finite dealer capacity*

Jeffrey D. Varner
R.F. Smith School of Chemical and Biomolecular Engineering
Cornell University, Ithaca, NY 14850
jdv27@cornell.edu

Abstract

Whether activating a new local behavior can destabilize a market model that was stable immediately before the change is a causal question that separately initialized parameter sweeps cannot answer. We tested it in a deterministic cellular automaton whose agents held Buy, Hold, or Sell states and whose baseline and activated rules differed at one neighborhood sum. Each paired experiment ran both paths under the baseline rule for 1,000 steps and then switched a preselected fraction of sites in one path without resetting any dynamic state, while the matched path kept every site under the baseline rule. None of the matched no-switch controls reached a Sell majority. With price feedback and an unlimited dealer, switching 10% or 15% of sites caused no transitions, whereas switching 20% or 25% caused all 50 runs to transition. Sell-majority transitions persisted when boundary traders were held in the Hold state, showing that the local rule dynamics did not require price feedback. Across lattice sizes 32–128, the activated fraction associated with 50% transition probability fell from 0.218 to 0.166, and the transition width narrowed from 0.0275 to 0.0082. This finite-time crossover sharpened with size but did not establish an equilibrium critical point. Dealer capacity primarily controlled the price response. At 20% activation, median drawdown rose from 0.40% with capacity of 0.1 unit per trader to 3.92% with one unit per trader and 25.3% with unlimited capacity. Separate cash and inventory interventions confirmed that trader eligibility and dealer capacity constrained different stages between a desired order and its effect on price.

*Repository: https://github.com/varnerlab/Wolfram-markets-abm-paper.git.

Preprint.

1 Introduction

Large price movements arrive more often than a Gaussian return process predicts, and volatility clusters in time, so that a calm market can turn turbulent within a few sessions [1–8]. Some of the largest movements have followed a change in the behavior of a subset of participants rather than a change in fundamentals. Portfolio insurers were a minority of institutions whose rule-driven selling concentrated in the October 1987 decline [9, 10], and the forced unwind of a few leveraged funds moved the entire quantitative equity sector in August 2007 [11]. A single large automated sell program, amplified by the response of other traders, produced the 2010 Flash Crash [12–14], and retail buying coordinated on social media moved individual stocks in 2021 [15]. Financial constraints can also turn a local sale into a market-wide decline, because forced sales lower prices, lower prices shrink the funding and wealth of other holders, and the shrinkage forces further sales [16–20]. Agent-based models of financial markets generate such episodes from the interaction of heterogeneous participants rather than from a return process prescribed in advance [21–29]. Fundamentalist and chartist populations, minority games, and adaptive strategies in these models reproduce the heavy tails and volatility clustering of the empirical record [23, 30–32]. Herding and imitation models make the amplification mechanism explicit, since a local change in behavior spreads through neighbors or through a shared price signal until it moves the aggregate [30, 32–39]. Statistical-physics models represent the same coordination through interacting spins or opinions and show how coupling strength, noise, or population composition can organize a market-wide transition [40–49], and three-state variants add an inactive state between buying and selling [50, 51]. Interaction structure shapes the route to coordination in all of these models, because global coupling, random graphs, and local lattices propagate a behavioral change differently [22, 52, 53]. A sweep over a coupling, noise, or population parameter then identifies the settings that produce a transition. The sweeps nevertheless initialize each parameter value separately [30, 42–44], so they can show that two settings differ but cannot show that a new behavior destabilized a market that was stable immediately before the behavior appeared. That causal claim requires an intervention on a running path and a matched path in which the behavior never changes. The same logic defines a treatment effect in potential-outcome methods [54–56], motivates common random numbers in paired comparisons of stochastic simulations [57–59], and underlies damage-spreading studies that follow a perturbed and an unperturbed copy of the same lattice [60–65]. The first question of this study was therefore whether activating one local rule in the middle of a run could destabilize a market model that was stable immediately before the change, which is a narrower question than whether the model reproduces the full set of empirical market statistics [8].

Execution constraints pose a second gap, because the path from a desired trade to a price change passes through stages that most interaction models compress into one. A trader may want to buy but lack cash, or want to sell but lack shares, and constraints on short selling of this kind are known to bias prices and slow their adjustment to bad news [66, 67]. Even an eligible order requires a counterparty. Market-microstructure theory makes that counterparty explicit as a dealer whose inventory and capital limit how much unmatched flow the dealer absorbs before quotes move or liquidity is withdrawn [68–73]. In the classical treatment of immediacy, liquidity depends on the number of market makers standing ready to absorb an imbalance, and the immediate price move after a liquidity shock scales inversely with that capacity [74]. Later work ties liquidity to market-maker inventories and to the capital and balance-sheet capacity of intermediaries [75–80]. Price-impact studies map signed executed flow to returns [81–87], and stylized models of this kind often represent the counterparty as an unlimited reservoir that fills every order and converts signed flow directly into a price change [81, 86]. This convention combines two distinct constraints: trader resources determine whether an order can enter the market, whereas dealer resources determine whether unmatched flow can execute. The distinction matters whenever prices feed back into behavior, because a counterparty that stops absorbing flow, an intermediary whose balance sheet is full, or a trading halt can each amplify or interrupt the loop from executed orders to prices and from prices to later decisions [12, 18, 88–91]. Agent-based markets have carried finite cash and share endowments [92], order-book clearing [93], leverage limits [90], and dealers with balance-sheet limits [94], and lattice interaction models have included trade frictions [95]. Whether a limited dealer changes a spatial instability itself, or only its effect on prices and rejected orders, was the second question of this study.

We separated these stages in a deliberately small model. Agents occupied a square lattice, held Buy, Hold, or Sell states, and updated synchronously from their four nearest neighbors according to a three-state Wolfram totalistic rule, in which the next state depends only on the sum of the neighborhood

states [96–102]. Because a cellular automaton is specified by a lookup table rather than by an energy function and a temperature, it carries no equilibrium distribution, and its finite rule space can be enumerated so that one table entry can be changed while every other entry stays fixed [96, 103–110]. Cellular automata have modeled markets before, and local majority or heterogeneous rules in those models produced abnormal fluctuations and regime changes [111–115]. The paired mid-run switch used here differs from that work by following a switched path and an unswitched path from the same realized state. Here the baseline and activated rules differed at only one of nine neighborhood sums, so the intervention changed a single entry of the lookup table. Agent states generated unit orders, cash and shares determined whether each order was eligible, an explicit dealer absorbed unmatched flow up to a capacity set per trader and rejected the rest, and executed imbalance changed the price. The price re-entered the lattice through traders on the lattice perimeter, whose states tracked the quote against an adaptive band, so that a price move could recruit new Buy or Sell states at the edge. Holding those perimeter traders in the Hold state removed the feedback without changing the interior rules or the financial bookkeeping. Switching a random fraction of sites and asking whether the whole lattice converts is structurally the question that threshold models of collective behavior and bootstrap percolation pose [116–122]. In those models a sharp threshold in the seeded fraction separates local from global spread, and the threshold on a finite lattice drifts with its size [123–125]. The lattice therefore also let us measure the growth of contiguous Sell regions directly, a quantity that phase-ordering studies track on lattices [126, 127]. It further let us follow how the curve of transition probability against switched fraction changed with system size, which is the standard finite-size test for separating a sharpening crossover from an equilibrium critical point [53, 128–132].

In this study, we asked whether a mid-run switch in one local rule could destabilize a market model that was stable immediately before the switch, and how financial constraints would alter the outcome. Each paired experiment ran two paths under the baseline rule for 1,000 steps and then switched a preselected fraction of interior sites in one path without resetting any dynamic state, while the matched path kept every site under the baseline rule. We repeated the design under a held boundary that removed price feedback, across lattice sizes from 32 to 128, across dealer capacities from zero to unlimited, and under separate cash and inventory interventions. None of the matched no-switch controls reached a Sell majority, whereas switching 20% or 25% of sites under price feedback and an unlimited dealer drove all 50 runs to a Sell majority and switching 10% or 15% drove none. The transitions persisted when the boundary was held, so the instability originated in the local rule dynamics rather than in price feedback. As the lattice grew, the switched fraction associated with a 50% transition probability fell from 0.218 to 0.166 and the transition width narrowed from 0.0275 to 0.0082, a sharpening with size that did not by itself establish an equilibrium critical point. Dealer capacity had a larger effect on drawdown and rejected orders than on the spatial transition: at 20% activation, median drawdown rose from 0.40% with a tenth of a unit of capacity per trader to 3.92% with one unit and 25.3% with unlimited capacity. These results separate the origin of the spatial instability, which lay in the lattice, from the financial constraints that determined its effect on price. They also show that a paired mid-run switch can attribute a market-wide transition to one local behavioral change in a way that separately initialized sweeps cannot.

2 Methods

2.1 Three-state cellular automaton

Agents occupied an $L \times L$ square lattice and carried one of three states: Buy ($s = 0$), Hold ($s = 1$), or Sell ($s = 2$). We modeled the synchronous local update as a cellular automaton in which only the $(L-2)^2$ strictly interior cells followed the rule. For each interior cell i, with ∂i its four von Neumann neighbors, the neighborhood sum was:

$$c_i(t) = \sum_{j \in \partial i} s_j(t) \in \{0, 1, \ldots, 8\}. \tag{1}$$

A totalistic rule table $\mathcal{R}$ maps each possible sum to one of the three states. The nine outputs, written as ternary digits $a_j = \mathcal{R}(j)$, define the Wolfram rule index [97, 100]:

$$R = \sum_{j=0}^{8} a_j 3^j. \tag{2}$$

The numerical state labels therefore entered the dynamics through both the neighborhood sum and the rule table.

We used Rule 19280, $(222110002)_3$, as the baseline and Rule 19286, $(222110022)_3$, as the activated rule. The tables differed only at $c = 1$, so a neighborhood with three Buy cells and one Hold cell returned Buy under Rule 19280 and Sell under Rule 19286. For each run, we selected a fraction f of interior sites uniformly without replacement and kept this target set fixed, so that its sites changed to Rule 19286 only at the intervention step. The initial interior contained $\mathrm{round}(0.05(L-2)^2)$ Buy sites and as many Sell sites placed without replacement, with Hold elsewhere. Randomness entered through the initial states, target sites, active boundary sample, boundary band widths, and finite-dealer rationing. Once these draws were fixed by the seed, the subsequent dynamics were deterministic.

2.2 Orders, trader books, and dealer clearing

Every interior cell and every active boundary cell represented a trader. A Buy state submitted one unit buy order, a Sell state submitted one unit sell order, and a Hold state submitted no order. Trader i carried cash w_i, integer inventory q_i, and average acquisition cost $\bar{c}_i$. A buy was eligible when $w_i \geq P_{t-1}$, and a sell was eligible when $q_i \geq 1$. We disabled the optional cost-basis restriction, which would have permitted a sale only above the trader's average acquisition cost in the manner of the disposition effect [133, 134], so average cost did not prevent a sale in the main experiments. An order that failed an eligibility test was *agent-suppressed*, which prevented execution without changing the trader's CA state.

We let B_t and S_t denote the eligible buy and sell counts at the standing quote P_{t-1}. Opposing flow of size $\min(B_t, S_t)$ cleared without using dealer resources, and an explicit dealer handled only the unmatched residual. We expressed dealer capacity as $h \geq 0$ net units per trader. For N_{tr} traders, total capacity was $H = \mathrm{round}(hN_{\mathrm{tr}})$, so larger h supplied more dealer cash and inventory. A finite dealer began with cash $W_0 = P_0 H$ at the initial quote P_0 and inventory $Q_0 = H$, could not borrow, and satisfied $0 \leq Q_t \leq 2H$. When buys exceeded sells, the executed buy and sell counts n_t^B and n_t^S were:

$$n_t^S = S_t, \qquad n_t^B = S_t + \min(B_t - S_t, Q_{t-1}). \tag{3}$$

When sells exceeded buys, the corresponding counts were:

$$n_t^B = B_t, \qquad n_t^S = B_t + \min\left(S_t - B_t, 2H - Q_{t-1}, \left\lfloor \frac{W_{t-1}}{P_{t-1}} \right\rfloor\right). \tag{4}$$

Eligible units beyond these bounds were *dealer-rejected*. When rationing was required, a counter-based score keyed by seed, step, trader identity, and side gave each eligible trader a reproducible priority, and including the step in the key prevented one trader from retaining priority over time. Filled orders updated trader books at P_{t-1}, and the opposite transfers updated the dealer ledger:

$$Q_t = Q_{t-1} + n_t^S - n_t^B, \qquad W_t = W_{t-1} + P_{t-1}(n_t^B - n_t^S). \tag{5}$$

Total cash and shares over traders plus dealer were conserved when no balance-sheet intervention occurred. The $h = \infty$ limit started from a zero dealer ledger, permitted borrowing and shorting, and filled every eligible order, whereas $h = 0$ allowed only opposing trader orders to clear. Both were limits of the same clearing rule (Algorithm S2).

2.3 Price feedback and the held-boundary control

Only executed orders moved the price. We defined normalized executed imbalance as:

$$\phi_t = \frac{n_t^B - n_t^S}{N_{\mathrm{tr}}}. \tag{6}$$

Positive imbalance represented net buying and raised the quote, and negative imbalance represented net selling and lowered it. The quote update, with price-impact coefficient λ, was:

$$P_t = P_{t-1} \exp(\lambda \phi_t). \tag{7}$$

The main experiments used $P_0 = 100$ and $\lambda = 0.02$. Active boundary traders compared the new quote with an exponential moving center μ_t and an exponential moving absolute deviation σ_t, which were updated with smoothing parameter $\alpha = 1/60$ as:

$$\mu_t = \mu_{t-1} + \alpha(P_t - \mu_{t-1}), \qquad \sigma_t = \sigma_{t-1} + \alpha(|P_t - \mu_t| - \sigma_{t-1}), \tag{8}$$

Boundary trader j had a fixed half-width $k_j \sim U[k(1-\delta), k(1+\delta)]$, with $k = 1$ and $\delta = 0.5$, and entered Buy below $\mu_t - k_j\sigma_t$, Sell above $\mu_t + k_j\sigma_t$, and Hold inside the band. One quarter of perimeter cells were active, and the rest stayed at Hold.

The simulation wrote the new boundary state before updating the interior. The complete feedback sequence was:

$$\text{CA state} \to \text{orders} \to \text{dealer fills} \to P_t \to \text{boundary state} \to \text{CA state}.$$

The implementation applied rule activation, any balance-sheet intervention, eligibility tests, clearing, settlement, price formation, boundary classification, and the synchronous CA update in that order (Algorithms S1 and S2). The matched *held-boundary* control retained the same initial lattice, target sites, active boundary coordinates, band draws, and initial books and dealer, and it used the same order, clearing, and price rules. Before each interior update, however, every active boundary state was set to Hold. Financial variables continued to evolve but could not affect the lattice, so this control isolated the lattice dynamics from price feedback.

2.4 Paired rule-switch experiments

Every paired experiment began with Rule 19280 at every interior site and evolved for $T_{\text{pre}} = 1{,}000$ steps before any rule changed. At the top of step $T_{\text{pre}} + 1$, the switched path activated Rule 19286 at the preselected fraction f, whereas the matched no-switch path left those target sites on Rule 19280, and neither path reset the lattice, quote, adaptive bands, trader books, or dealer. Because the two paths were identical through the full lattice state $X_{T_{\text{pre}}}$ and the next quote came from that common state, the first possible difference occurred during the following interior update. We verified these pre-switch equalities bit for bit for every pair.

We ran four production campaigns (Table S1), and every campaign used 3,000 post-activation steps, 50 seeds, and traders that began with cash 1,000 and 10 shares. The selected paired-switch campaign used $L = 64$, an unlimited dealer, an endogenous boundary, and $f \in \{0.10, 0.15, 0.20, 0.25\}$. The dealer-capacity screen kept $L = 64$, sampled the fraction grid 0.10, 0.15, 0.175, 0.185, 0.190, 0.195, 0.200, 0.205, 0.210, 0.225, and 0.250, and took dealer capacity values 0, 0.1, 0.5, 1, 2, 5, 10, and ∞ under endogenous and held boundaries. The finite-size experiment used $L \in \{32, 48, 64, 96, 128\}$, sampled 12 fractions from 0.140 to 0.250, and ran three conditions: a held boundary with $h = \infty$, an endogenous boundary with $h = \infty$, and an endogenous boundary with $h = 1$. Because h was defined per trader, absolute dealer capacity scaled with lattice size. The balance-sheet experiment changed trader funding and sell inventory at the same top-of-step event as rule activation by multiplying current cash and current shares independently by $g_w, g_q \in \{0.1, 1, 10\}$ and rounding shares to integers. The full 3×3 design was crossed with $f \in \{0.170, 0.175, 0.180, 0.185\}$, $h \in \{0.1, 1, \infty\}$, and both boundary conditions. All nine interventions within a $(f, h, \text{boundary}, \text{seed})$ cell shared an identical pre-intervention path, and under the held boundary their full lattice paths also had to remain identical after the intervention. That equality supplied an exact null test for a direct effect of financial variables on the interior rule.

All simulations ran in Julia 1.12 with the package environment pinned by the repository manifest, using 12 simulation threads and one garbage-collection thread. A seeded Xoshiro generator [135] produced the setup draws, the finite-size and balance-sheet campaigns wrote an atomic checkpoint per cell and resumed from it, and aggregation reconstructed a fixed row order, so thread scheduling and restart order did not alter the results. Unit tests covered parameter validation, finite-dealer limits, time-varying priorities, conservation, delayed rule activation, the held-boundary control, and the unlimited-dealer limit, in which every eligible order executes. Step-by-step tests also compared the optimized CA kernel with the general totalistic solver under homogeneous and heterogeneous rule fields. The code, environment manifest, derived data, and figure scripts are available at `https://github.com/varnerlab/Wolfram-markets-abm-paper`.

2.5 Outcomes, statistical analysis, and prespecified panels

The order parameter $m(t)$ was the strictly interior Sell fraction, so higher $m(t)$ meant that a larger share of the interior occupied the Sell state. We called a run transitioned when $m(t) \geq 0.5$ at least once after activation and defined onset as the first such crossing. We reported transition probability, onset conditional on crossing, the final 500-step mean of $m(t)$, maximum post-activation drawdown, agent-suppression fractions, and the fraction of eligible orders rejected by the dealer. At fixed

(L, f, h), the ensemble susceptibility, a standard finite-size quantity [129, 132], measured between-seed fluctuations in the final-window mean $\bar{m}$, with angle brackets denoting the average over seeds, and was computed as:

$$\chi_L = (L-2)^2 \left[\langle \bar{m}^2 \rangle - \langle \bar{m} \rangle^2\right]. \tag{9}$$

We linearly interpolated the sampled transition curve to estimate the finite-horizon midpoint $f_{50}(L)$, where transition probability reached 0.5, and the width $f_{75}(L) - f_{25}(L)$ between the fractions at which it reached 0.75 and 0.25, where a smaller width indicated a sharper crossover. These quantities described the sampled finite-time response and were not treated as equilibrium critical parameters. Binomial proportions used 95% Wilson intervals [136], and continuous summaries used interquartile ranges unless noted otherwise.

We chose every run displayed in the time-sliced market panels from these ensemble summaries before inspecting any lattice image. The activation-fraction panels included a nontransitioning $f = 0.15$ run, both outcome groups at $f = 0.175$, and transitioning runs at $f = 0.185$ and 0.20 under $h = \infty$. Within each transitioning group we chose the run whose onset was nearest the subgroup median, whereas within each nontransitioning group we chose the run whose final Sell fraction was nearest it. The dealer-capacity panels required one seed that transitioned at $h = 0.1$, 1, and ∞ for $f = 0.20$, and among eligible seeds we minimized the summed distance from the three capacity-specific median onsets. Both sets of panels displayed the lattice at offsets 0, 500, 1,500, and 3,000 after activation, so every row used the same clock.

3 Results

To determine whether Rule 19286 caused the Sell-dominated state, we compared each switched path with a matched path that never changed rules and found that only switched paths reached a Sell majority (Fig. 1). The two paths were bit-identical during the 1,000-step baseline and through the first quote after activation. None of the 50 no-switch controls reached a Sell majority in the next 3,000 steps. Switched runs at $f = 0.10$ and 0.15 also remained mixed, whereas all 50 runs crossed the Sell-majority threshold at $f = 0.20$ and 0.25 (Fig. 1A,B). Median time from activation to crossing fell from 1,044.5 steps at $f = 0.20$ to 311.5 steps at $f = 0.25$. These results showed that activating the local response could destabilize an ongoing market path and identified the interval that required finer sampling. The representative $f = 0.20$ pair, selected by the prespecified median-effect criterion, shared one lattice at activation, after which the no-switch path remained mixed while the switched path formed expanding Sell-rich regions and ended in a Sell-dominated state (Fig. 1C). Because the intervention changed only the rule at the selected sites, the later divergence was caused by Rule 19286 within this model. This comparison established causation within the model but did not show whether price feedback or dealer resources created the transition.

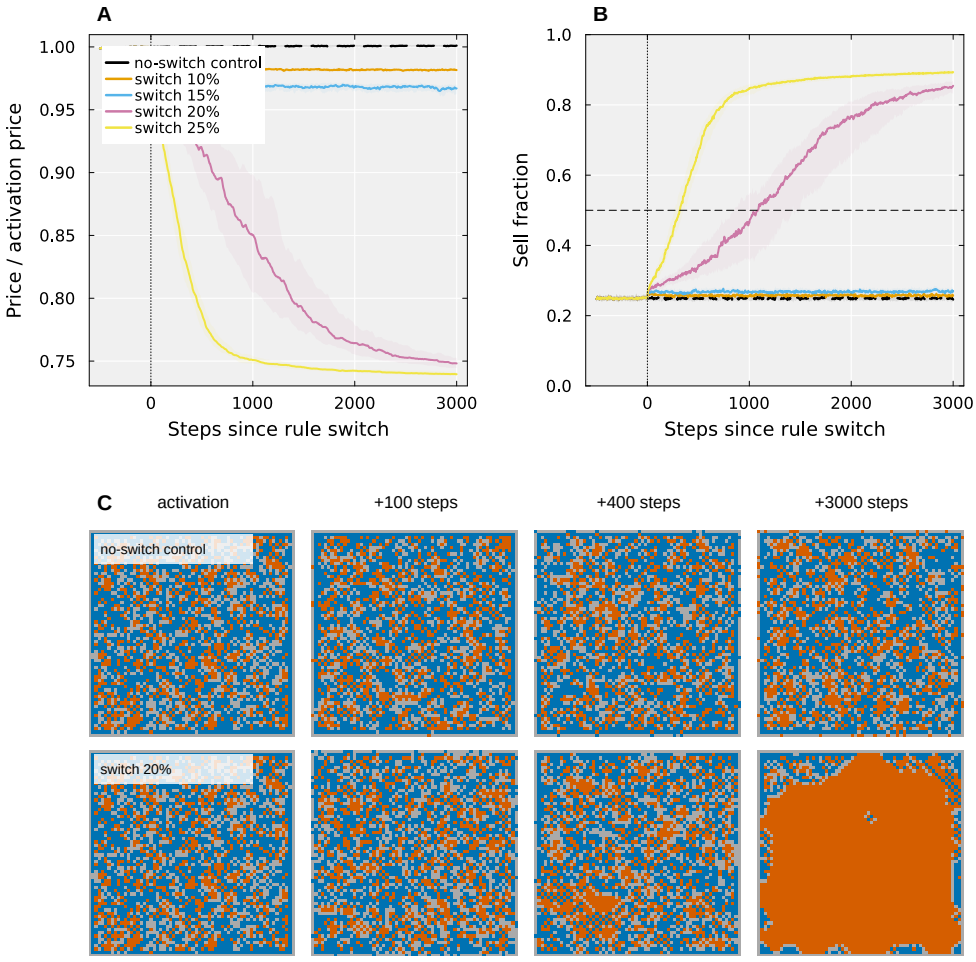


Figure 1: A paired mid-run rule switch separated initially identical market paths. Fifty seeds were run on a 64×64 lattice for 1,000 baseline steps and 3,000 post-activation steps. The no-switch path retained the same target sites but kept their baseline rules. **(A)** Median price divided by its value at activation; bands mark the interquartile range. **(B)** Median interior Sell fraction; the horizontal line marks the majority threshold. **(C)** Time-sliced lattices for seed 12, the $f = 0.20$ pair nearest the median switched-control difference in peak Sell fraction. Blue, gray, and orange cells denote Buy, Hold, and Sell. The rows show the no-switch and switched paths from their common activation state through 3,000 steps. All paired paths coincided before the vertical activation line.

To separate the roles of price feedback and dealer resources, we repeated the switch on a finer fraction grid across eight dealer capacities under both the endogenous and the held boundary and found that the transition survived the removal of price feedback (Fig. 2). Under an unlimited dealer and endogenous boundary, the finer screen found 14 of 50 crossings at $f = 0.175$, 43 at $f = 0.185$, 46 at $f = 0.190$, and 50 at every sampled fraction from $f = 0.195$ upward (Fig. 2A). Under the held-boundary control, which kept every active boundary trader in the Hold state and so removed price feedback from the lattice, transition probability rose from 0.34 at $f = 0.175$ to 0.70 at 0.185, 0.90 at 0.190, 0.98 at 0.195, and 1.00 at 0.20, so the transition did not require price feedback. For each (f, seed) pair, the held-boundary lattice path was exactly invariant to dealer capacity even though executions, prices, and dealer ledgers differed. Restoring the endogenous boundary, by contrast, shifted outcomes only within the same narrow activation window: at $f = 0.175$, transition probability ranged from 0.28 with an unlimited dealer to 0.66 at $h = 0.1$, compared with 0.34 under the held boundary, and at $f = 0.185$ the endogenous range was 0.84–0.98, compared with 0.70 under the held boundary (Fig. 2A,B). The response was not monotone in dealer capacity, yet every capacity

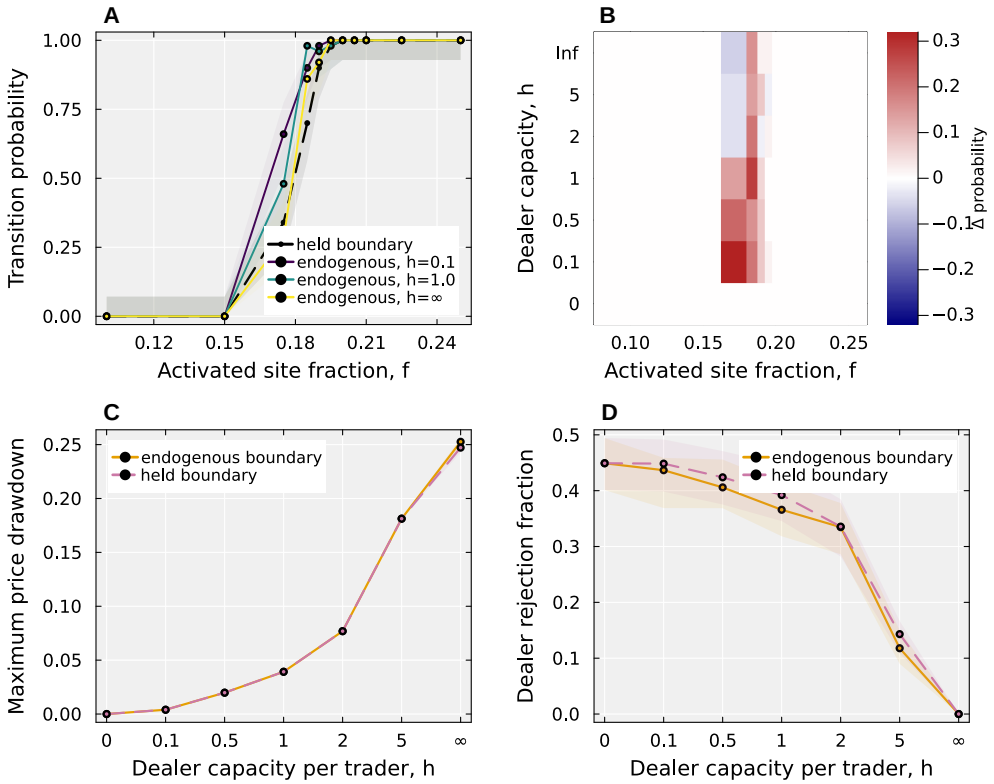


Figure 2: Rule activation across dealer capacities. **(A)** Transition probability under the held boundary and three endogenous dealer conditions; bands are 95% Wilson intervals over 50 seeds. The held boundary kept active boundary traders in the Hold state. **(B)** Difference between endogenous and held-boundary transition probabilities across the complete capacity range. **(C)** Median maximum drawdown at $f = 0.20$; bands mark the interquartile range. **(D)** Median fraction of agent-eligible orders rejected by the dealer at $f = 0.20$. Capacity $h = 10$ is omitted because it coincided exactly with $h = \infty$ in this campaign.

produced no transitions at $f \leq 0.15$ and transitions in all 50 seeds at $f \geq 0.20$. Price feedback shifted marginal outcomes but did not create the activation window. Dealer capacity, by contrast, had a larger effect on price than on the transition, and the effect was ordered in h. At $f = 0.20$, median maximum drawdown increased from zero at $h = 0$ to 0.40%, 1.98%, 3.92%, 7.69%, and 18.1% at $h = 0.1, 0.5, 1, 2$, and 5, respectively, before reaching 25.3% at $h = \infty$ (Fig. 2C). Median dealer rejection moved in the opposite direction, from 0.449 at $h = 0$ to 0.118 at $h = 5$ and zero at $h = \infty$ (Fig. 2D). Capacity $h = 10$ reproduced the unlimited path for the stated endowments and horizon. At $h = 0$, only matched flow executed, the price remained fixed, and the endogenous-boundary lattice reproduced the held-boundary lattice. Dealer resources therefore determined how much net flow reached the price but were not the source of the spatial instability.

Because the transition survived the held-boundary control, we next tested whether it also survived changes in system size and found that the crossover sharpened as the lattice grew (Fig. 3). None of the no-switch controls transitioned at any sampled size. Under the held boundary, the midpoint moved from $f_{50} = 0.2181$ at $L = 32$ through 0.1926, 0.1777, and 0.1681 to 0.1665 at $L = 128$ (Fig. 3A,C). The width $f_{75} - f_{25}$ decreased from 0.0275 to 0.0082, and peak ensemble susceptibility increased from 21.2 to 232.7 while moving from $f = 0.225$ to $f = 0.170$ (Fig. 3B,D). Over the sampled sizes and fixed time horizon, the crossover sharpened rather than disappeared. The two endogenous-boundary conditions followed the same size trend, with midpoint estimates at $L = 32$ of 0.2132 for $h = \infty$ and 0.2100 for $h = 1$ against 0.2181 under the held boundary. At $L = 128$, the corresponding estimates were 0.1655, 0.1650, and 0.1665, and their widths ranged from 0.0063 to 0.0082 (Fig. 3C). Susceptibility peaks increased in every condition, so dealer and boundary effects were largest at the smaller lattices while all three conditions had similar midpoints at the largest

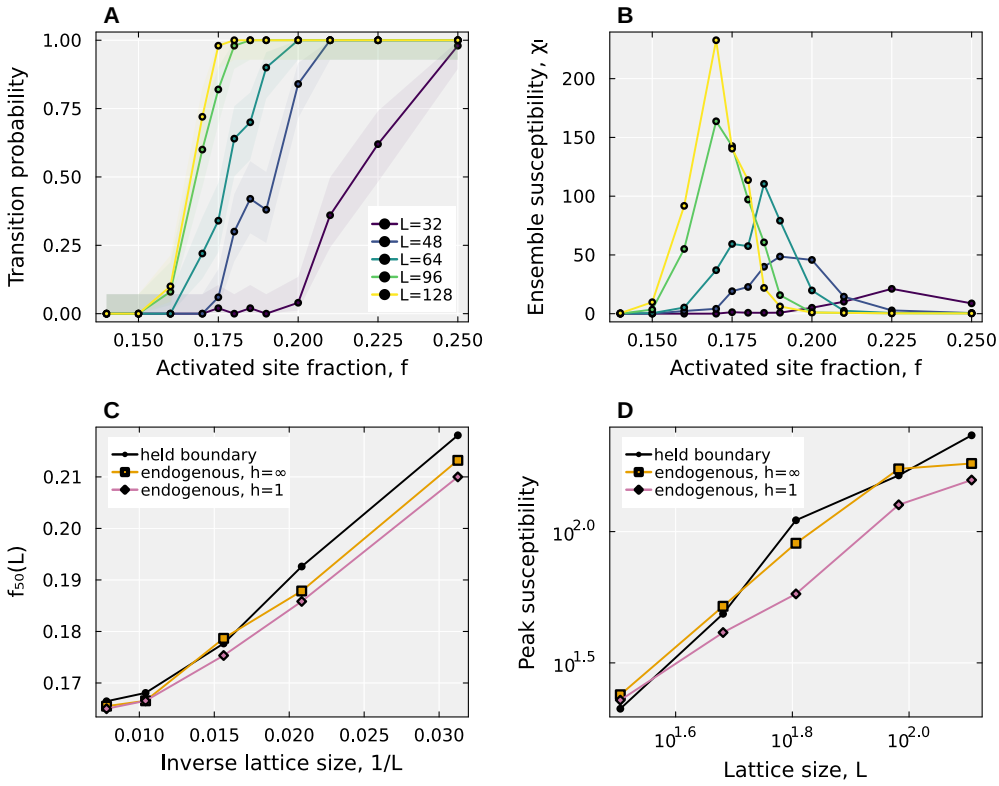


Figure 3: Finite-size response of the paired rule switch over 50 seeds. **(A)** Held-boundary transition probability; bands are 95% Wilson intervals. **(B)** Held-boundary ensemble susceptibility of the final 500-step Sell fraction. **(C)** Interpolated finite-horizon midpoint $f_{50}(L)$ for the held control and two endogenous dealer conditions. **(D)** Peak susceptibility versus lattice size on logarithmic axes. Lines connect sampled values. No asymptotic scaling law was fitted.

size tested. Five sizes and a fixed 3,000-step horizon could not establish an asymptotic transition or support a critical-exponent fit.

Having located the instability in the lattice rather than in price feedback, we next asked how individual lattices reached the Sell-dominated state and found that every selected crossing proceeded by domain growth (Fig. 4). Aggregate transition probabilities did not distinguish gradual spatial change from a near-simultaneous switch of states. At $f = 0.15$, the prespecified noncrossing run retained a fine mixed pattern for all 3,000 steps. Two runs at $f = 0.175$ followed different paths: one developed a Sell-rich region but stayed below one-half, while the other formed a large connected Sell domain and crossed after 2,186 steps. The selected $f = 0.185$ run crossed after 1,779 steps through several connected regions. The $f = 0.20$ run crossed after 1,030 steps and ended with one dominant Sell domain (Fig. 4A). Price declined as the Sell fraction rose in each crossing run, with maximum drawdowns of 17.6%, 23.7%, and 25.6% for the crossing runs at $f = 0.175$, 0.185, and 0.20 against 5.3% and 9.9% for the two noncrossing runs (Fig. 4B,C). The aggregate transition probability therefore summarized a distribution of spatial paths, not a synchronous relabeling of the lattice. Across these selected runs the common pattern was domain growth rather than classical nucleation: Sell-rich regions enlarged and merged in every crossing, but their number, shape, and timing varied across seeds, and the noncrossing $f = 0.175$ run also formed a visible domain, so the presence of a domain did not guarantee a finite-horizon majority.

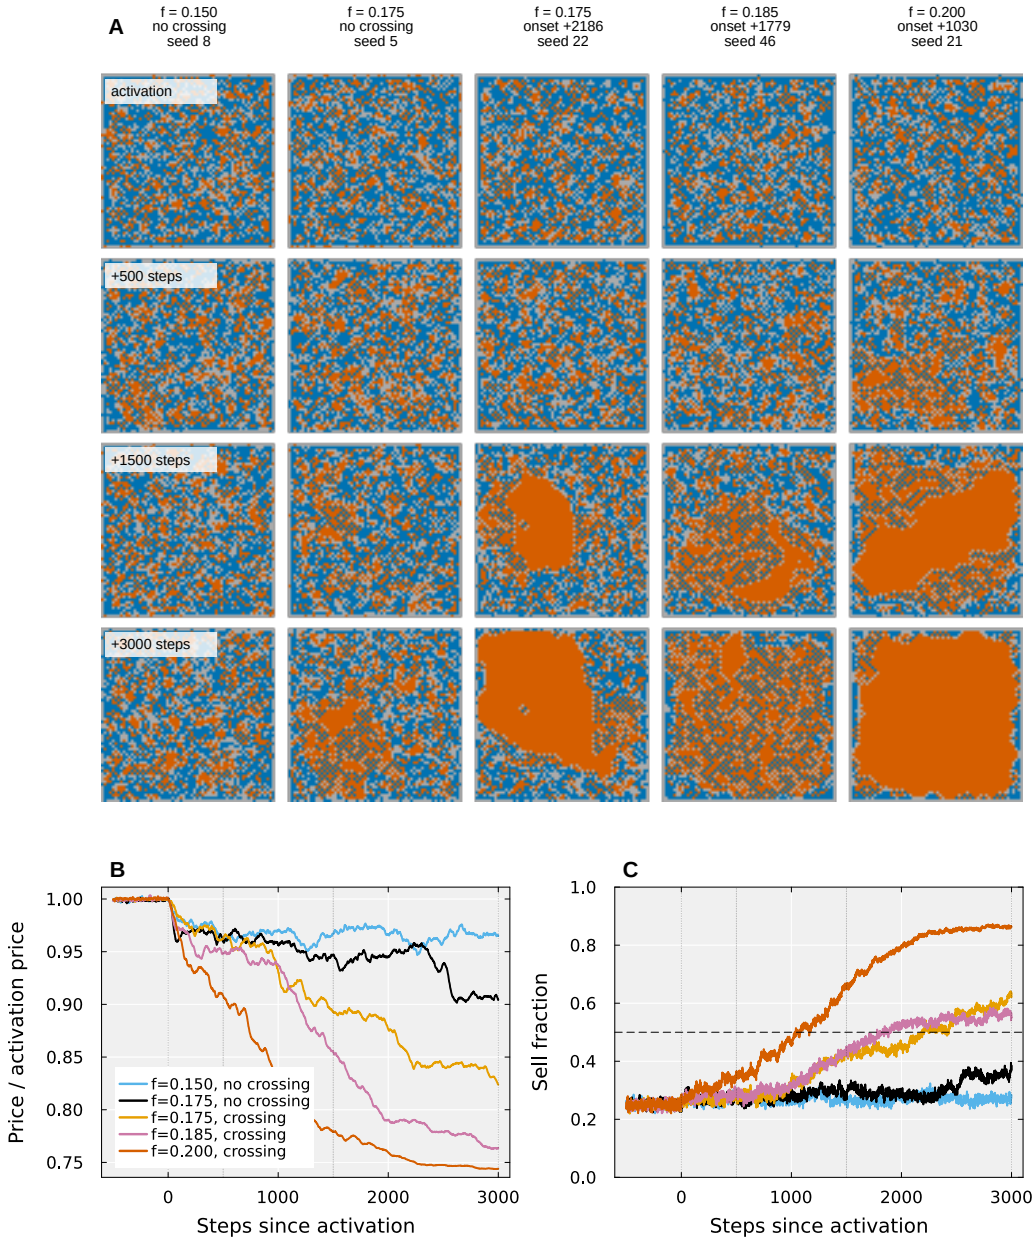


Figure 4: Prespecified time-sliced markets across the activation-fraction crossover under an endogenous boundary and unlimited dealer. Blue, gray, and orange cells denote Buy, Hold, and Sell. **(A)** Lattices for five runs (columns) at activation and after 500, 1,500, and 3,000 steps (rows). From left to right, the runs include a noncrossing $f = 0.15$ case, both outcome groups at $f = 0.175$, and crossing cases at $f = 0.185$ and 0.20. Crossing cases were nearest their subgroup median onset. Noncrossing cases were nearest their subgroup median final Sell fraction. Seeds were 8, 5, 22, 46, and 21 from left to right. **(B)** Price divided by its activation value. **(C)** Interior Sell fraction; the horizontal line marks the majority threshold. Vertical guides mark the four snapshots.

To compare spatial and price outcomes on one rule landscape, we followed a single seed that crossed at three dealer capacities and traced its lattice, price, and Sell fraction under each capacity and under the held boundary, and found that dealer capacity changed the price path but left the Sell-dominated outcome intact (Fig. 5). Seed 25 was selected jointly from runs that crossed at $h = 0.1$, 1, and ∞ by the prespecified median-onset criterion. All three endogenous paths formed expanding Sell domains and crossed within 88 steps of one another, at 1,045, 1,110, and 1,022 steps, and the held-boundary path crossed at 1,191 steps (Fig. 5A). Dealer capacity changed the detailed geometry and timing, but all four paths reached similar Sell-dominated final states. The corresponding price paths, by contrast, differed: maximum drawdown was 0.40% at $h = 0.1$, 3.92% at $h = 1$, and 25.8% with the unlimited dealer (Fig. 5B). Despite these differences, all four Sell fractions crossed one-half within 169 steps of one another and ended with final-window means between 0.81 and 0.87 (Fig. 5C). The held-boundary condition still calculated executions and price and removed only the return path from price to the lattice, so a constrained dealer could reduce drawdown at $f = 0.20$ to 0.40% without preventing Sell-domain growth.

We next asked whether trader and dealer balance sheets could change the lattice without price feedback and found that the held-boundary lattice was invariant to every intervention (Fig. 6). Within every held-boundary (h, f, seed) cell, all nine cash and share multipliers produced bit-identical lattice paths. Transition probability at $f = 0.18$ was therefore 0.64 in every held-boundary financial cell. Under the endogenous boundary, the same probability ranged from 0.50 to 0.94 across the three-by-three cash–share grid and three dealer capacities (Fig. 6A–C). The minimum occurred at $h = 0.1$ with cash multiplied by 0.1 and shares by 10, and the maximum occurred at $h = 1$ in the three cells that reduced shares tenfold or raised cash tenfold without raising shares. We reported $f = 0.18$, the sampled fraction whose held-boundary transition probability lay nearest one-half, and the other three fractions showed the same structure. The held-boundary probability was 0.22, 0.34, and 0.70 at $f = 0.170$, 0.175, and 0.185, and the endogenous range at those fractions was 0.08–0.52, 0.24–0.82, and 0.64–0.98. Financial variables affected the lattice only through executed flow, price, and the boundary response.

We then examined how the same interventions changed drawdown and found that trader inventory raised drawdown only when the dealer could absorb the resulting sell flow (Fig. 6). At baseline cash and $h = \infty$, share multipliers 0.1, 1, and 10 produced median maximum drawdowns of 2.87%, 20.1%, and 83.5% (Fig. 6D). The corresponding values were 0.040%, 3.92%, and 3.92% at $h = 1$, and 0.025%, 0.399%, and 0.399% at $h = 0.1$ (Fig. 6E,F). Once dealer inventory bound, additional eligible sells were rejected instead of increasing executed imbalance. Trader inventory set potential sell flow, while dealer resources capped the unmatched flow that reached the price. Transition probability, however, did not increase monotonically with drawdown. At $h = 1$ and baseline cash, the share-0.1 condition transitioned in 0.94 of seeds with only 0.040% median drawdown, while the share-10 condition transitioned in 0.60 of seeds with 3.92% drawdown. At baseline shares and $h = 1$, multiplying cash by 10 reduced median drawdown from 3.92% to 0.089% but increased transition probability from 0.78 to 0.94. Trader cash supplied matching buy flow and also changed the boundary signal. The crossed intervention therefore separated three outcomes that a single crash label would combine: Sell-state dominance, price drawdown, and dealer rejection.

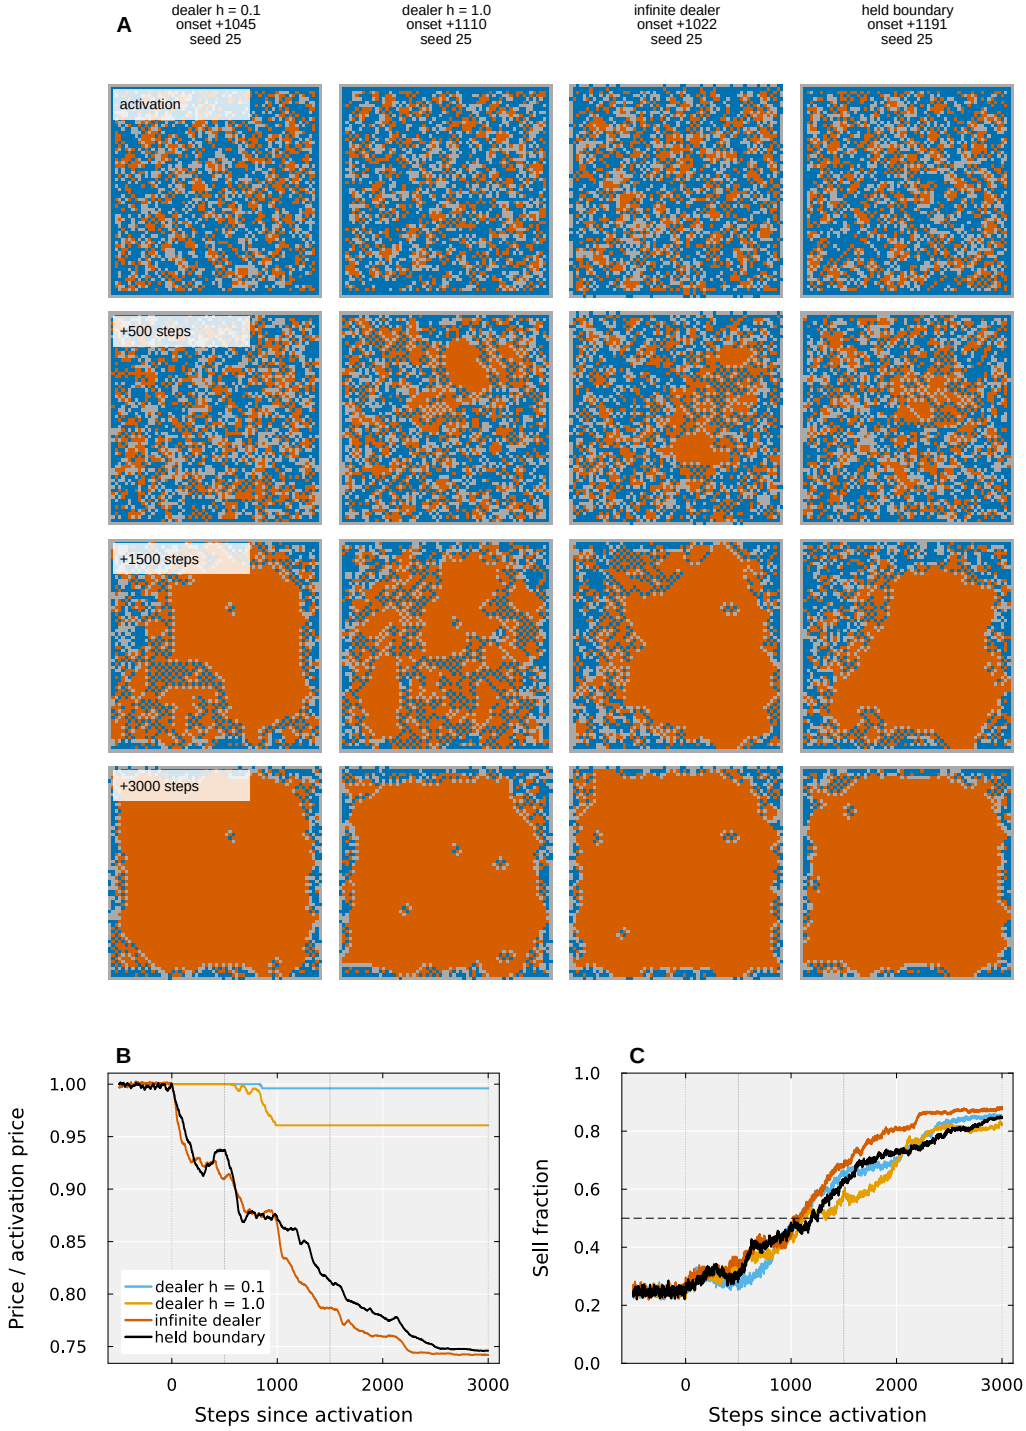


Figure 5: Time-sliced markets for one matched seed at $f = 0.20$. **(A)** Lattices for four conditions (columns) at activation and after 500, 1,500, and 3,000 steps (rows). The columns contain endogenous dealers with $h = 0.1$, $h = 1$, and $h = \infty$, followed by the held-boundary control with $h = \infty$. Seed 25 minimized the summed distance from the three endogenous capacity-specific median onsets among seeds that crossed in all three conditions. Rows use common elapsed times. **(B)** Price divided by its activation value. **(C)** Interior Sell fraction. The held boundary prevented price from changing the lattice but did not prevent price from being calculated.

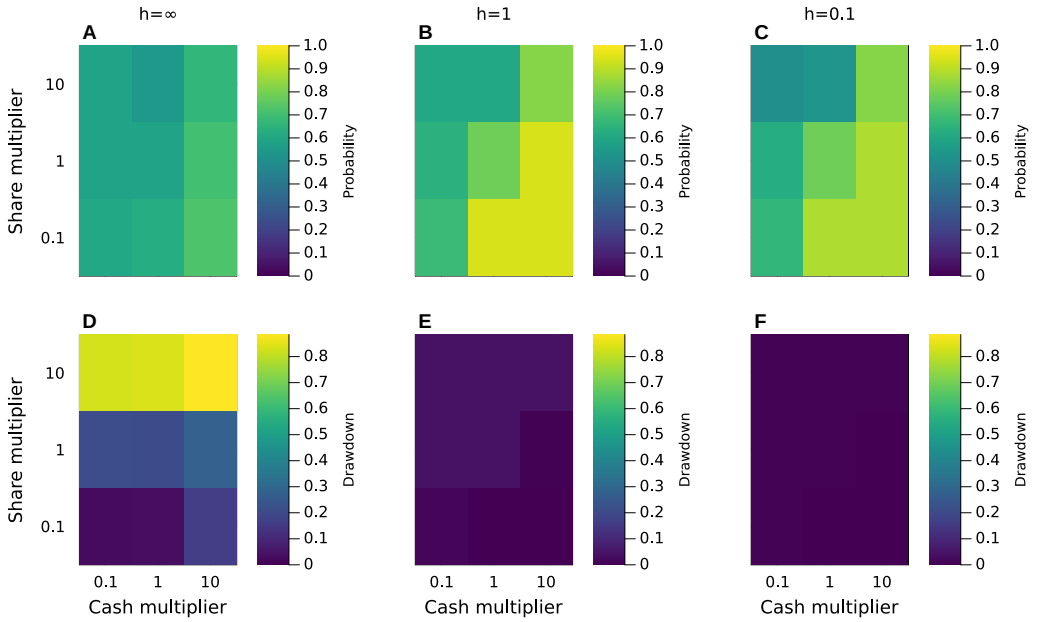


Figure 6: Crossed trader cash and share interventions at activation under the endogenous boundary, with $f = 0.18$ and 50 seeds. Columns correspond to dealer capacities $h = \infty$, 1, and 0.1. Cash multipliers run horizontally and share multipliers vertically. **(A–C)** Transition probability. **(D–F)** Median maximum drawdown on a common color scale. The matched held-boundary lattice was invariant across all nine interventions.

4 Discussion

The paired mid-run intervention supports a stronger within-model claim than a standard parameter sweep. Agent-based market models have long shown that simple local rules can produce complex aggregate dynamics [22, 23, 30, 32, 36]. Separately initialized regimes, however, cannot show that adding one behavior destabilized a path that was stable immediately beforehand, for the same reason that a treatment effect requires a matched counterfactual and a paired simulation comparison requires common random numbers [54, 55, 57, 58]. In our experiment, the switched and no-switch paths shared the realized lattice, trader books, dealer, price, adaptive bands, and target sites until activation, so changing the rule at those sites caused the later divergence. None of the matched no-switch controls reached a Sell majority, whereas switching 20% or 25% of interior sites under an unlimited dealer drove all 50 runs to a Sell majority and switching 10% or 15% drove none. The transitions persisted when the boundary was held, and as the lattice grew from 32 to 128 the midpoint of the transition curve fell from 0.218 to 0.166 while its width fell from 0.0275 to 0.0082. Median drawdown at 20% activation rose from 0.40% at a tenth of a unit of dealer capacity per trader to 25.3% with an unlimited dealer. This result establishes causation inside the model, not a calibrated claim that the same fraction of real traders would destabilize a market.

The held-boundary and finite-size results located the main instability in the heterogeneous cellular automaton. Sell-majority transitions continued when price could not affect the lattice, and the finite-time crossover narrowed from a width of 0.0275 at $L = 32$ to 0.0082 at $L = 128$. Sell-rich regions also formed, enlarged, and merged across activation fractions and dealer conditions. Persistence under the held boundary, size-dependent sharpening, and domain growth support a nonequilibrium description of the transition [65]. The size-dependent activation window also resembles the threshold in seeded fraction that separates local from global spread in threshold-contagion and bootstrap-percolation models, where the finite-lattice threshold likewise drifts with system size [116, 120, 123, 124]. This evidence does not establish classical nucleation in the sense of Johnson–Mehl–Avrami–Kolmogorov growth kinetics or Kramers-type thermal activation [137–141]: the analysis supplied no free-energy barrier, thermal activation rate, surface tension, or critical droplet. Nor did five lattice sizes observed for 3,000 post-activation steps establish an equilibrium critical point. That size dependence remains a finite-time result. Establishing an asymptotic transition would require lattices

larger than 128, horizons that scale with system size, denser sampling near the transition, and a scaling model specified before fitting [129, 132]. Such an analysis would also need to distinguish a diverging correlation length [128] from long transient coarsening near an absorbing state [65, 126, 127].

Dealer capacity controlled the price consequences of that spatial instability. Order-flow models connect executed imbalance to price changes [81, 84, 86], dealer-inventory models bound the imbalance a market maker absorbs before its quotes move [68–71, 74, 78], and market-microstructure research emphasizes the constraints imposed by counterparties, balance sheets, and trading halts [12, 18, 89]. Our dealer separated these mechanisms directly, since opposing trader orders cleared without dealer resources and the dealer handled only the residual, as in an earlier agent-based market whose dealers carried balance-sheet limits [94]. With zero capacity, executed buys and sells balanced and the quote stayed fixed, whereas with unlimited capacity every eligible order executed and the quote absorbed the full imbalance. Similar Sell domains could therefore accompany either a 0.40% or a 25.8% drawdown for the same seed at 20% activation. The dealer was useful for this controlled comparison but was not a realistic model of market depth or of a limit-order book [87, 93]. The cash and inventory interventions extended this separation to the trader side, since trader inventory determined the potential sell flow, trader cash supplied matching buy flow, and dealer cash and inventory limited the unmatched residual. At 18% activation with baseline cash, multiplying trader shares by ten produced an 83.5% median drawdown against an unlimited dealer but drawdowns of only 3.92% at one unit of capacity per trader and 0.40% at a tenth of a unit, because the dealer rejected the unmatched sells once its inventory bound. The exact held-boundary null showed that these financial variables changed the lattice only through price feedback. Transition probability, price drawdown, and dealer rejection should therefore be reported as distinct outcomes rather than combined in one crash indicator.

Several modeling choices limit the scope of the conclusions. The cellular automaton was deterministic after its random setup, updated synchronously, and used a regular lattice and one selected rule pair. Orders had unit size and a common standing quote, and the model omitted an order book, bid–ask spread, cancellation, strategic dealer, risk penalty, and endogenous replenishment of dealer capacity beyond reversing flow. Boundary traders followed one adaptive rule, and uniform balance-sheet multipliers represented causal model interventions rather than policy experiments. Finally, the price-impact coefficient, time step, and lattice geometry were dimensionless and uncalibrated. The results identify mechanisms in a minimal dynamical system and do not forecast market crashes. Several extensions follow from the same paired design. The totalistic rule space is finite [97], so every rule that differs from the baseline at a single neighborhood sum can be screened with the same matched no-switch control, which would show whether the activation window found here is specific to Rule 19286 or common to a class of single-entry switches. Asynchronous updating, a limit-order book in place of the dealer, and horizons that grow with lattice size would test whether the separation between the spatial instability and its price consequences survives more realistic execution and longer observation.

5 Conclusion

Switching one local response destabilized market-model paths that remained stable under the baseline rule. The matched no-switch paths established this causal effect within the model, while the held boundary and finite-size experiments showed that the spatial transition did not depend on price feedback. Dealer capacity instead determined how strongly the resulting order imbalance changed the price or produced rejected orders. Trader cash, trader inventory, and dealer resources constrained different stages of execution. Sell-state dominance, price drawdown, and dealer exhaustion can therefore move separately, and an unlimited counterparty should be treated as a consequential modeling assumption rather than a neutral execution convention.

References

[1] Benoit Mandelbrot. The variation of certain speculative prices. *The Journal of Business*, 36 (4):394–419, 1963.

[2] Eugene F. Fama. The behavior of stock-market prices. *The Journal of Business*, 38(1):34–105, 1965. doi: 10.1086/294743.

[3] Robert F. Engle. Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation. *Econometrica*, 50(4):987–1007, 1982. doi: 10.2307/1912773.

[4] Tim Bollerslev. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3):307–327, 1986. doi: 10.1016/0304-4076(86)90063-1.

[5] G. William Schwert. Why does stock market volatility change over time? *The Journal of Finance*, 44(5):1115–1153, 1989.

[6] Rosario N. Mantegna and H. Eugene Stanley. Scaling behaviour in the dynamics of an economic index. *Nature*, 376(6535):46–49, 1995. doi: 10.1038/376046a0.

[7] Parameswaran Gopikrishnan, Vasiliki Plerou, Luís A. Nunes Amaral, Martin Meyer, and H. Eugene Stanley. Scaling of the distribution of fluctuations of financial market indices. *Physical Review E*, 60(5):5305–5316, 1999. doi: 10.1103/PhysRevE.60.5305.

[8] Rama Cont. Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001.

[9] Robert J. Shiller. Portfolio insurance and other investor fashions as factors in the 1987 stock market crash. *NBER Macroeconomics Annual*, 3:287–297, 1988. doi: 10.1086/654091.

[10] Sanford J. Grossman. An analysis of the implications for stock and futures price volatility of program trading and dynamic hedging strategies. *The Journal of Business*, 61(3):275–298, 1988. doi: 10.1086/296433.

[11] Amir E. Khandani and Andrew W. Lo. What happened to the quants in August 2007? Evidence from factors and transactions data. *Journal of Financial Markets*, 14(1):1–46, 2011. doi: 10.1016/j.finmar.2010.07.005.

[12] Andrei Kirilenko, Albert S. Kyle, Mehrdad Samadi, and Tugkan Tuzun. The flash crash: High-frequency trading in an electronic market. *The Journal of Finance*, 72(3):967–998, 2017.

[13] David Easley, Marcos M. López de Prado, and Maureen O'Hara. The microstructure of the 'Flash Crash': flow toxicity, liquidity crashes, and the probability of informed trading. *The Journal of Portfolio Management*, 37(2):118–128, 2011. doi: 10.3905/jpm.2011.37.2.118.

[14] Albert J. Menkveld and Bart Zhou Yueshen. The Flash Crash: a cautionary tale about highly fragmented markets. *Management Science*, 65(10):4470–4488, 2019. doi: 10.1287/mnsc.2018.3040.

[15] Lasse Heje Pedersen. Game on: social networks and markets. *Journal of Financial Economics*, 146(3):1097–1119, 2022. doi: 10.1016/j.jfineco.2022.05.002.

[16] Joshua Coval and Erik Stafford. Asset fire sales (and purchases) in equity markets. *Journal of Financial Economics*, 86(2):479–512, 2007. doi: 10.1016/j.jfineco.2006.09.007.

[17] Albert S. Kyle and Wei Xiong. Contagion as a wealth effect. *The Journal of Finance*, 56(4):1401–1440, 2001. doi: 10.1111/0022-1082.00373.

[18] Markus K. Brunnermeier and Lasse Heje Pedersen. Market liquidity and funding liquidity. *The Review of Financial Studies*, 22(6):2201–2238, 2009.

[19] Tobias Adrian and Hyun Song Shin. Liquidity and leverage. *Journal of Financial Intermediation*, 19(3):418–437, 2010. doi: 10.1016/j.jfi.2008.12.002.

[20] Andrei Shleifer and Robert Vishny. Fire sales in finance and macroeconomics. *Journal of Economic Perspectives*, 25(1):29–48, 2011. doi: 10.1257/jep.25.1.29.

[21] Blake LeBaron. Agent-based computational finance: suggested readings and early research. *Journal of Economic Dynamics and Control*, 24:679–702, 2000.

[22] Blake LeBaron. Agent-based computational finance. In *Handbook of Computational Economics*, volume 2, pages 1187–1233. Elsevier, 2006.

[23] Cars H Hommes. Heterogeneous agent models in economics and finance. *Handbook of Computational Economics*, 2:1109–1186, 2006.

[24] E. Samanidou, E. Zschischang, D. Stauffer, and T. Lux. Agent-based models of financial markets. *Reports on Progress in Physics*, 70(3):409–450, 2007. doi: 10.1088/0034-4885/70/3/r03.

[25] J. Doyne Farmer and Duncan Foley. The economy needs agent-based modelling. *Nature*, 460 (7256):685–686, 2009. doi: 10.1038/460685a.

[26] Richard Bookstaber. Agent-based models for financial crises. *Annual Review of Financial Economics*, 9(1):85–100, 2017. doi: 10.1146/annurev-financial-110716-032556.

[27] W. Brian Arthur. Foundations of complexity economics. *Nature Reviews Physics*, 3:136–145, 2021.

[28] J. Doyne Farmer. Quantitative agent-based models: a promising alternative for macroeconomics. *Oxford Review of Economic Policy*, 41(2):571–590, 2025.

[29] Robert L. Axtell and J. Doyne Farmer. Agent-based modeling in economics and finance: past, present, and future. *Journal of Economic Literature*, 63(1):197–287, 2025. doi: 10.1257/jel.20221319.

[30] Thomas Lux and Michele Marchesi. Scaling and criticality in a stochastic multi-agent model of a financial market. *Nature*, 397:498–500, 1999.

[31] D. Challet and Y.-C. Zhang. Emergence of cooperation and organization in an evolutionary game. *Physica A: Statistical Mechanics and its Applications*, 246:407–418, 1997.

[32] J. Doyne Farmer and Shareen Joshi. The price dynamics of common trading strategies. *Journal of Economic Behavior & Organization*, 49(2):149–171, 2002.

[33] A. Kirman. Ants, rationality, and recruitment. *The Quarterly Journal of Economics*, 108(1): 137–156, 1993. doi: 10.2307/2118498.

[34] William A. Brock and Cars H. Hommes. Heterogeneous beliefs and routes to chaos in a simple asset pricing model. *Journal of Economic Dynamics and Control*, 22(8-9):1235–1274, 1998. doi: 10.1016/s0165-1889(98)00011-6.

[35] Blake LeBaron, W. Brian Arthur, and Richard Palmer. Time series properties of an artificial stock market. *Journal of Economic Dynamics and Control*, 23(9-10):1487–1516, 1999. doi: 10.1016/s0165-1889(98)00081-5.

[36] Rama Cont and Jean-Philippe Bouchaud. Herd behavior and aggregate fluctuations in financial markets. *Macroeconomic Dynamics*, 4(2):170–196, 2000.

[37] Thomas Lux and Michele Marchesi. Volatility clustering in financial markets: a microsimulation of interacting agents. *International Journal of Theoretical and Applied Finance*, 3(4): 675–702, 2000. doi: 10.1142/s0219024900000826.

[38] Georges Harras and Didier Sornette. How to grow a bubble: a model of myopic adapting agents. *Journal of Economic Behavior & Organization*, 80:137–152, 2012.

[39] J. J. Chen, L. Tan, and B. Zheng. Agent-based model with multi-level herding for complex financial systems. *Scientific Reports*, 5:8399, 2015.

[40] Per Bak, Maya Paczuski, and Martin Shubik. Price variations in a stock market with many agents. *Physica A: Statistical Mechanics and its Applications*, 246:430–453, 1997.

[41] Debashish Chowdhury and Dietrich Stauffer. A generalized spin model of financial markets. *The European Physical Journal B*, 8:477–482, 1999.

[42] Stefan Bornholdt. Expectation bubbles in a spin model of markets: intermittency from frustration across scales. *International Journal of Modern Physics C*, 12(5):667–674, 2001.

[43] A. Krawiecki, J. A. Hołyst, and D. Helbing. Volatility clustering and scaling for financial time series due to attractor bubbling. *Physical Review Letters*, 89:158701, 2002.

[44] Taisei Kaizoji. Speculative bubbles and crashes in stock markets: an interacting-agent model of speculative activity. *Physica A: Statistical Mechanics and its Applications*, 287:493–506, 2000.

[45] Katarzyna Sznajd-Weron and Rafal Weron. A simple model of price formation. *International Journal of Modern Physics C*, 13(1):115–123, 2002.

[46] Didier Sornette and Wei-Xing Zhou. Importance of positive feedbacks and overconfidence in a self-fulfilling ising model of financial markets. *Physica A: Statistical Mechanics and its Applications*, 370(2):704–726, 2006. doi: 10.1016/j.physa.2006.02.022.

[47] W.-X. Zhou and D. Sornette. Self-organizing Ising model of financial markets. *European Physical Journal B*, 55(2):175–181, 2007. doi: 10.1140/epjb/e2006-00391-6.

[48] Claudio Castellano, Santo Fortunato, and Vittorio Loreto. Statistical physics of social dynamics. *Reviews of Modern Physics*, 81(2):591–646, 2009. doi: 10.1103/RevModPhys.81.591.

[49] Didier Sornette. Physics and financial economics (1776–2014): puzzles, Ising and agent-based models. *Reports on Progress in Physics*, 77(6):062001, 2014. doi: 10.1088/0034-4885/77/6/062001.

[50] Taisei Kaizoji, Stefan Bornholdt, and Yoshi Fujiwara. Dynamics of price and trading volume in a spin model of stock markets with heterogeneous agents. *Physica A: Statistical Mechanics and its Applications*, 316(1-4):441–452, 2002. doi: 10.1016/s0378-4371(02)01216-5.

[51] Tetsuya Takaishi. Simulations of financial markets in a Potts-like model. *International Journal of Modern Physics C*, 16(08):1311–1317, 2005. doi: 10.1142/s0129183105007923.

[52] Stefano Battiston, J Doyne Farmer, Andreas Flache, Diego Garlaschelli, Andrew G Haldane, Hans Heesterbeek, Cars Hommes, Carlo Jaeger, Robert May, and Marten Scheffer. Complexity theory and financial regulation. *Science*, 351(6275):818–819, 2016.

[53] Simone Alfarano and Mishael Milaković. Network structure and N-dependence in agent-based herding models. *Journal of Economic Dynamics and Control*, 33(1):78–92, 2009. doi: 10.1016/j.jedc.2008.05.003.

[54] Donald B. Rubin. Estimating causal effects of treatments in randomized and nonrandomized studies. *Journal of Educational Psychology*, 66(5):688–701, 1974. doi: 10.1037/h0037350.

[55] Paul W. Holland. Statistics and causal inference. *Journal of the American Statistical Association*, 81(396):945–960, 1986. doi: 10.1080/01621459.1986.10478354.

[56] Judea Pearl. *Causality: models, reasoning, and inference*. Cambridge University Press, Cambridge, 2 edition, 2009. ISBN 9780521895606. doi: 10.1017/CBO9780511803161.

[57] Jack P. C. Kleijnen. Antithetic variates, common random numbers and optimal computer time allocation in simulation. *Management Science*, 21(10):1176–1185, 1975. doi: 10.1287/mnsc.21.10.1176.

[58] Paul Glasserman and David D. Yao. Some guidelines and guarantees for common random numbers. *Management Science*, 38(6):884–908, 1992. doi: 10.1287/mnsc.38.6.884.

[59] Averill M. Law. *Simulation modeling and analysis*. McGraw-Hill Education, New York, 5 edition, 2015. ISBN 9780073401324.

[60] S. A. Kauffman. Metabolic stability and epigenesis in randomly constructed genetic nets. *Journal of Theoretical Biology*, 22(3):437–467, 1969. doi: 10.1016/0022-5193(69)90015-0.

[61] B. Derrida and D. Stauffer. Phase transitions in two-dimensional Kauffman cellular automata. *Europhysics Letters*, 2(10):739–745, 1986. doi: 10.1209/0295-5075/2/10/001.

[62] B. Derrida and G. Weisbuch. Dynamical phase transitions in 3-dimensional spin glasses. *Europhysics Letters*, 4(6):657–662, 1987. doi: 10.1209/0295-5075/4/6/004.

[63] H. Eugene Stanley, Dietrich Stauffer, János Kertész, and Hans J. Herrmann. Dynamics of spreading phenomena in two-dimensional Ising models. *Physical Review Letters*, 59(20): 2326–2328, 1987. doi: 10.1103/PhysRevLett.59.2326.

[64] Peter Grassberger. Are damage spreading transitions generically in the universality class of directed percolation? *Journal of Statistical Physics*, 79(1-2):13–23, 1995. doi: 10.1007/BF02179381.

[65] Haye Hinrichsen. Non-equilibrium critical phenomena and phase transitions into absorbing states. *Advances in Physics*, 49(7):815–958, 2000. doi: 10.1080/00018730050198152.

[66] Edward M. Miller. Risk, uncertainty, and divergence of opinion. *The Journal of Finance*, 32 (4):1151–1168, 1977. doi: 10.1111/j.1540-6261.1977.tb03317.x.

[67] Douglas W. Diamond and Robert E. Verrecchia. Constraints on short-selling and asset price adjustment to private information. *Journal of Financial Economics*, 18(2):277–311, 1987. doi: 10.1016/0304-405x(87)90042-0.

[68] Mark B. Garman. Market microstructure. *Journal of Financial Economics*, 3(3):257–275, 1976. doi: 10.1016/0304-405X(76)90006-4.

[69] Hans R. Stoll. The supply of dealer services in securities markets. *The Journal of Finance*, 33 (4):1133–1151, 1978. doi: 10.1111/j.1540-6261.1978.tb02053.x.

[70] Yakov Amihud and Haim Mendelson. Dealership market: Market-making with inventory. *Journal of Financial Economics*, 8(1):31–53, 1980. doi: 10.1016/0304-405X(80)90020-3.

[71] Thomas Ho and Hans R. Stoll. Optimal dealer pricing under transactions and return uncertainty. *Journal of Financial Economics*, 9(1):47–73, 1981. doi: 10.1016/0304-405X(81)90020-9.

[72] Lawrence R. Glosten and Paul R. Milgrom. Bid, ask and transaction prices in a specialist market with heterogeneously informed traders. *Journal of Financial Economics*, 14(1):71–100, 1985. doi: 10.1016/0304-405x(85)90044-3.

[73] Ananth Madhavan. Market microstructure: a survey. *Journal of Financial Markets*, 3(3): 205–258, 2000. doi: 10.1016/S1386-4181(00)00007-0.

[74] Sanford J. Grossman and Merton H. Miller. Liquidity and market structure. *The Journal of Finance*, 43(3):617–633, 1988. doi: 10.1111/j.1540-6261.1988.tb04594.x.

[75] Carole Comerton-Forde, Terrence Hendershott, Charles M. Jones, Pamela C. Moulton, and Mark S. Seasholes. Time variation in liquidity: The role of market-maker inventories and revenues. *The Journal of Finance*, 65(1):295–331, 2010. doi: 10.1111/j.1540-6261.2009.01530.x.

[76] Zhiguo He and Arvind Krishnamurthy. Intermediary asset pricing. *American Economic Review*, 103(2):732–770, 2013. doi: 10.1257/aer.103.2.732.

[77] Terrence Hendershott and Albert J. Menkveld. Price pressures. *Journal of Financial Economics*, 114(3):405–423, 2014. doi: 10.1016/j.jfineco.2014.08.001.

[78] Tobias Adrian, Nina Boyarchenko, and Or Shachar. Dealer balance sheets and bond liquidity provision. *Journal of Monetary Economics*, 89:92–109, 2017. doi: 10.1016/j.jmoneco.2017.03.011.

[79] Jack Bao, Maureen O'Hara, and Xing (Alex) Zhou. The volcker rule and corporate bond market making in times of stress. *Journal of Financial Economics*, 130(1):95–113, 2018. doi: 10.1016/j.jfineco.2018.06.001.

[80] Annette Vissing-Jorgensen. The Treasury market in spring 2020 and the response of the Federal Reserve. *Journal of Monetary Economics*, 124:19–47, 2021. doi: 10.1016/j.jmoneco.2021.10.007.

[81] Albert S. Kyle. Continuous auctions and insider trading. *Econometrica*, 53(6):1315–1335, 1985.

[82] Joel Hasbrouck. Measuring the information content of stock trades. *The Journal of Finance*, 46(1):179–207, 1991. doi: 10.1111/j.1540-6261.1991.tb03749.x.

[83] Fabrizio Lillo, J. Doyne Farmer, and Rosario N. Mantegna. Master curve for price-impact function. *Nature*, 421(6919):129–130, 2003. doi: 10.1038/421129a.

[84] Xavier Gabaix, Parameswaran Gopikrishnan, Vasiliki Plerou, and H Eugene Stanley. A theory of power-law distributions in financial market fluctuations. *Nature*, 423:267–270, 2003.

[85] Jean-Philippe Bouchaud, Yuval Gefen, Marc Potters, and Matthieu Wyart. Fluctuations and response in financial markets: the subtle nature of 'random' price changes. *Quantitative Finance*, 4(2):176–190, 2004. doi: 10.1080/14697680400000022.

[86] J Doyne Farmer, Paolo Patelli, and Ilija I Zovko. The predictive power of zero intelligence in financial markets. *Proceedings of the National Academy of Sciences*, 102(6):2254–2259, 2005.

[87] Jean-Philippe Bouchaud, J. Doyne Farmer, and Fabrizio Lillo. How markets slowly digest changes in supply and demand. In Thorsten Hens and Klaus Reiner Schenk-Hoppé, editors, *Handbook of Financial Markets: Dynamics and Evolution*, pages 57–160. North-Holland, Amsterdam, 2009. ISBN 9780123742582. doi: 10.1016/B978-012374258-2.50006-3.

[88] Bruce C. Greenwald and Jeremy C. Stein. Transactional risk, market crashes, and the role of circuit breakers. *The Journal of Business*, 64(4):443–462, 1991. doi: 10.1086/296547.

[89] Avanidhar Subrahmanyam. Circuit breakers and market volatility: A theoretical perspective. *The Journal of Finance*, 49(1):237–254, 1994.

[90] Stefan Thurner, J. Doyne Farmer, and John Geanakoplos. Leverage causes fat tails and clustered volatility. *Quantitative Finance*, 12(5):695–707, 2012. doi: 10.1080/14697688.2012.674301.

[91] Hui Chen, Anton Petukhov, Jiang Wang, and Hao Xing. The dark side of circuit breakers. *The Journal of Finance*, 79(2):1405–1455, 2024. doi: 10.1111/jofi.13310.

[92] Marco Raberto, Silvano Cincotti, Sergio M. Focardi, and Michele Marchesi. Agent-based simulation of a financial market. *Physica A: Statistical Mechanics and its Applications*, 299 (1-2):319–327, 2001. doi: 10.1016/s0378-4371(01)00312-0.

[93] Carl Chiarella and Giulia Iori. A simulation analysis of the microstructure of double auction markets. *Quantitative Finance*, 2(5):346–353, 2002. doi: 10.1088/1469-7688/2/5/303.

[94] Karen Braun-Munzinger, Zijun Liu, and Arthur E. Turrell. An agent-based model of corporate bond trading. *Quantitative Finance*, 18(4):591–608, 2018. doi: 10.1080/14697688.2017.1380310.

[95] Giulia Iori. A microsimulation of traders activity in the stock market: the role of heterogeneity, agents' interactions and trade frictions. *Journal of Economic Behavior & Organization*, 49(2): 269–285, 2002. doi: 10.1016/s0167-2681(01)00164-0.

[96] Stephen Wolfram. Statistical mechanics of cellular automata. *Reviews of Modern Physics*, 55 (3):601–644, 1983. doi: 10.1103/RevModPhys.55.601.

[97] Stephen Wolfram. Universality and complexity in cellular automata. *Physica D: Nonlinear Phenomena*, 10:1–35, 1984.

[98] Stephen Wolfram. Cellular automata as models of complexity. *Nature*, 311(5985):419–424, 1984. doi: 10.1038/311419a0.

[99] Gérard Y. Vichniac. Simulating physics with cellular automata. *Physica D: Nonlinear Phenomena*, 10(1-2):96–116, 1984. doi: 10.1016/0167-2789(84)90253-7.

[100] Norman H. Packard and Stephen Wolfram. Two-dimensional cellular automata. *Journal of Statistical Physics*, 38:901–946, 1985.

[101] Stephen Wolfram. *A New Kind of Science*. Wolfram Media, 2002.

[102] Michiel Rollier, Kallil M. C. Zielinski, Aisling J. Daly, Odemir M. Bruno, and Jan M. Baetens. A comprehensive taxonomy of cellular automata. *Communications in Nonlinear Science and Numerical Simulation*, 140:108362, 2025. doi: 10.1016/j.cnsns.2024.108362.

[103] Eytan Domany and Wolfgang Kinzel. Equivalence of cellular automata to Ising models and directed percolation. *Physical Review Letters*, 53(4):311–314, 1984. doi: 10.1103/PhysRevLett.53.311.

[104] Tommaso Toffoli and Norman Margolus. *Cellular automata machines: a new environment for modeling*. MIT Press, Cambridge, MA, 1987. ISBN 9780262291019. doi: 10.7551/mitpress/1763.001.0001.

[105] Howard A. Gutowitz, Jonathan D. Victor, and Bruce W. Knight. Local structure theory for cellular automata. *Physica D: Nonlinear Phenomena*, 28(1–2):18–48, 1987.

[106] Chris G. Langton. Computation at the edge of chaos: phase transitions and emergent computation. *Physica D: Nonlinear Phenomena*, 42(1-3):12–37, 1990. doi: 10.1016/0167-2789(90)90064-V.

[107] Bastien Chopard and Michel Droz. *Cellular automata modeling of physical systems*. Cambridge University Press, Cambridge, 1998. ISBN 9780521461689. doi: 10.1017/CBO9780511549755.

[108] Andrew Ilachinski. *Cellular automata: a discrete universe*. World Scientific, Singapore, 2001. ISBN 9789810246235. doi: 10.1142/4702.

[109] Alberto Dennunzio, Enrico Formenti, and Julien Provillard. Non-uniform cellular automata: classes, dynamics, and decidability. *Information and Computation*, 215:32–46, 2012. doi: 10.1016/j.ic.2012.02.008.

[110] Kamalika Bhattacharjee, Nazma Naskar, Souvik Roy, and Sukanta Das. A survey of cellular automata: types, dynamics, non-uniformity and applications. *Natural Computing*, 19(2):433–461, 2020. doi: 10.1007/s11047-018-9696-8.

[111] Yi-ming Wei, Shang-jun Ying, Ying Fan, and Bing-Hong Wang. The cellular automaton model of investment behavior in the stock market. *Physica A: Statistical Mechanics and its Applications*, 325(3-4):507–516, 2003. doi: 10.1016/S0378-4371(03)00144-4.

[112] M. Bartolozzi and A. W. Thomas. Stochastic cellular automata model for stock market dynamics. *Physical Review E*, 69(4):046112, 2004. doi: 10.1103/PhysRevE.69.046112.

[113] G. Qiu, D. Kandhai, and P. M. A. Sloot. Understanding the complex dynamics of stock markets through cellular automata. *Physical Review E*, 75(4):046116, 2007. doi: 10.1103/PhysRevE.75.046116.

[114] Ying Fan, Shang-Jun Ying, Bing-Hong Wang, and Yi-Ming Wei. The effect of investor psychology on the complexity of stock market: an analysis based on cellular automaton model. *Computers & Industrial Engineering*, 56(1):63–69, 2009. doi: 10.1016/j.cie.2008.03.015.

[115] A. P. F. Atman and Bruna Amin Gonçalves. Influence of the investor's behavior on the complexity of the stock market. *Brazilian Journal of Physics*, 42(1-2):137–145, 2012. doi: 10.1007/s13538-011-0057-7.

[116] Mark Granovetter. Threshold models of collective behavior. *American Journal of Sociology*, 83(6):1420–1443, 1978. doi: 10.1086/226707.

[117] J. Chalupa, P. L. Leath, and G. R. Reich. Bootstrap percolation on a Bethe lattice. *Journal of Physics C: Solid State Physics*, 12(1):L31–L35, 1979. doi: 10.1088/0022-3719/12/1/008.

[118] Joan Adler. Bootstrap percolation. *Physica A: Statistical Mechanics and its Applications*, 171(3):453–470, 1991. doi: 10.1016/0378-4371(91)90295-n.

[119] Stephen Morris. Contagion. *Review of Economic Studies*, 67(1):57–78, 2000. doi: 10.1111/1467-937x.00121.

[120] Duncan J. Watts. A simple model of global cascades on random networks. *Proceedings of the National Academy of Sciences*, 99(9):5766–5771, 2002. doi: 10.1073/pnas.082090499.

[121] Damon Centola and Michael Macy. Complex contagions and the weakness of long ties. *American Journal of Sociology*, 113(3):702–734, 2007. doi: 10.1086/521848.

[122] Robert Morris. Bootstrap percolation, and other automata. *European Journal of Combinatorics*, 66:250–263, 2017. doi: 10.1016/j.ejc.2017.06.024.

[123] M. Aizenman and J. L. Lebowitz. Metastability effects in bootstrap percolation. *Journal of Physics A: Mathematical and General*, 21(19):3801–3813, 1988. doi: 10.1088/0305-4470/21/19/017.

[124] Alexander E. Holroyd. Sharp metastability threshold for two-dimensional bootstrap percolation. *Probability Theory and Related Fields*, 125(2):195–224, 2003. doi: 10.1007/s00440-002-0239-x.

[125] Béla Bollobás, Paul Smith, and Andrew Uzzell. Monotone cellular automata in a random environment. *Combinatorics, Probability and Computing*, 24(4):687–722, 2015. doi: 10.1017/S0963548315000012.

[126] A. J. Bray. Theory of phase-ordering kinetics. *Advances in Physics*, 43(3):357–459, 1994. doi: 10.1080/00018739400101505.

[127] V. Spirin, P. L. Krapivsky, and S. Redner. Fate of zero-temperature Ising ferromagnets. *Physical Review E*, 63(3):036118, 2001. doi: 10.1103/PhysRevE.63.036118.

[128] Michael E. Fisher and Michael N. Barber. Scaling theory for finite-size effects in the critical region. *Physical Review Letters*, 28(23):1516–1519, 1972. doi: 10.1103/PhysRevLett.28.1516.

[129] K. Binder. Finite size scaling analysis of Ising model block distribution functions. *Zeitschrift für Physik B Condensed Matter*, 43(2):119–140, 1981. doi: 10.1007/BF01293604.

[130] E. Egenter, T. Lux, and D. Stauffer. Finite-size effects in Monte Carlo simulations of two stock market models. *Physica A: Statistical Mechanics and its Applications*, 268(1-2):250–256, 1999. doi: 10.1016/s0378-4371(99)00059-x.

[131] V. Alfi, M. Cristelli, L. Pietronero, and A. Zaccaria. Mechanisms of self-organization and finite size effects in a minimal agent based model. *Journal of Statistical Mechanics: Theory and Experiment*, 2009(03):P03016, 2009. doi: 10.1088/1742-5468/2009/03/p03016.

[132] David P. Landau and Kurt Binder. *A guide to Monte Carlo simulations in statistical physics*. Cambridge University Press, Cambridge, 5 edition, 2021. ISBN 9781108490146. doi: 10.1017/9781108780346.

[133] Hersh Shefrin and Meir Statman. The disposition to sell winners too early and ride losers too long: Theory and evidence. *The Journal of Finance*, 40(3):777–790, 1985.

[134] Terrance Odean. Are investors reluctant to realize their losses? *The Journal of Finance*, 53(5):1775–1798, 1998.

[135] David Blackman and Sebastiano Vigna. Scrambled linear pseudorandom number generators. *ACM Transactions on Mathematical Software*, 47(4):1–32, 2021. doi: 10.1145/3460772.

[136] Edwin B. Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927. doi: 10.1080/01621459.1927.10502953.

[137] Andrei N. Kolmogorov. On the statistical theory of the crystallization of metals. *Izvestiya Akademii Nauk SSSR, Seriya Matematicheskaya*, 1(3):355–359, 1937.

[138] William A. Johnson and Robert F. Mehl. Reaction kinetics in processes of nucleation and growth. *Transactions of the American Institute of Mining and Metallurgical Engineers*, 135: 416–442, 1939.

[139] Melvin Avrami. Kinetics of phase change. i. general theory. *The Journal of Chemical Physics*, 7(12):1103–1112, 1939.

[140] Hendrik A. Kramers. Brownian motion in a field of force and the diffusion model of chemical reactions. *Physica*, 7(4):284–304, 1940.

[141] Peter Hänggi, Peter Talkner, and Michal Borkovec. Reaction-rate theory: Fifty years after Kramers. *Reviews of Modern Physics*, 62(2):251–341, 1990.

A Supplementary methods and validation

A.1 One complete simulation step

Every campaign used the complete update in Algorithm S1, with order clearing and settlement defined in Algorithm S2. Orders at step t came from the lattice X_{t-1} and the boundary states $e_{i,t-1}$, settled at the standing quote P_{t-1}, and generated the new quote P_t. The new quote classified the boundary before the cellular automaton produced X_t. Stating the full order of operations was necessary because changing that order would define a different model.

Algorithm S1 One endogenous-market step, $t-1 \to t$, in simulation order. Orders came from the old interior and boundary states, settled at the old quote, and generated the new quote. The new quote set the boundary before the interior cellular automaton update.

Require: lattice X_{t-1}; boundary states $e_i(t-1)$; trader books; dealer ledger; quote and band state $(P_{t-1}, \mu_{t-1}, \sigma_{t-1})$; rule field R_i; parameters; simulation seed
Ensure: lattice X_t; boundary states $e_i(t)$; settled books and dealer; quote and band state (P_t, μ_t, σ_t); all order counts

if t is the scheduled rule-switch step **then**
  $R_i \leftarrow 19286$ for every preselected target site i {preserve every dynamic state variable}
end if
if t is a scheduled balance-sheet intervention **then**
  for all traders i **do**
    $w_i \leftarrow g_w w_i$; $q_i \leftarrow \mathrm{round}(g_q q_i)$
  end for
end if
call Algorithm S2 using states at $t-1$ and P_{t-1}
receive $(n_t^B, n_t^S, u_t^B, u_t^S, r_t^B, r_t^S)$ and settled ledgers
$\phi_t \leftarrow (n_t^B - n_t^S)/N_{\mathrm{tr}}$
$P_t \leftarrow P_{t-1}\exp(\lambda\phi_t)$
$\mu_t \leftarrow \mu_{t-1} + \alpha(P_t - \mu_{t-1})$
$\sigma_t \leftarrow \sigma_{t-1} + \alpha(|P_t - \mu_t| - \sigma_{t-1})$
if boundary mode is endogenous **then**
  for all active boundary traders i **do**
    $e_i(t) \leftarrow 0$ if $P_t < \mu_t - k_i\sigma_t$; 2 if $P_t > \mu_t + k_i\sigma_t$; 1 otherwise
  end for
else
  $e_i(t) \leftarrow 1$ for every active boundary trader i
end if
write $e_i(t)$ to active boundary sites; leave passive sites at Hold
for all strictly interior sites i, simultaneously **do**
  $c_i(t) \leftarrow \sum_{j\in\partial i} s_j$ using old interior states and new boundary states
  $s_i(t) \leftarrow \mathcal{R}_{R_i}(c_i(t))$
end for

Algorithm S2 Order eligibility, dealer clearing, rationing, and settlement at P_{t-1}. The finite dealer filled only the unmatched residual. The $h = \infty$ branch filled every eligible order.

Require: old interior and boundary states; books $(w_i, q_i, \bar{c}_i)$; dealer (W_{t-1}, Q_{t-1}) and bounds $(Q_{\min}, Q_{\max}, C)$; P_{t-1}; h; cost-basis toggle; seed and step t

Ensure: settled books and dealer; executed counts (n_t^B, n_t^S); agent suppression (u_t^B, u_t^S); dealer rejection (r_t^B, r_t^S)

$\mathcal{B}, \mathcal{S} \leftarrow \varnothing$; $u_t^B, u_t^S \leftarrow 0$
for all interior traders, followed by active-boundary traders i **do**
  $o_i \leftarrow +1$ for Buy, -1 for Sell, and 0 for Hold {an edge trader reads its stored boundary state}
  if $o_i = +1$ and $w_i \geq P_{t-1}$ **then**
    append i to $\mathcal{B}$
  else if $o_i = -1$, $q_i \geq 1$, and the cost-basis test passes **then**
    append i to $\mathcal{S}$
  else if $o_i = +1$ **then**
    $u_t^B \leftarrow u_t^B + 1$
  else if $o_i = -1$ **then**
    $u_t^S \leftarrow u_t^S + 1$
  end if
end for
$B_t \leftarrow |\mathcal{B}|$; $S_t \leftarrow |\mathcal{S}|$; $n_t^B \leftarrow B_t$; $n_t^S \leftarrow S_t$
if $h < \infty$ and $B_t > S_t$ **then**
  $K_t^B \leftarrow Q_{t-1} - Q_{\min}$ {shares that the dealer could sell}
  $n_t^B \leftarrow S_t + \min(B_t - S_t, K_t^B)$
else if $h < \infty$ and $S_t > B_t$ **then**
  $K_t^S \leftarrow Q_{\max} - Q_{t-1}$ {inventory space}
  $K_t^S \leftarrow \min\big(K_t^S, \lfloor (W_{t-1} + C)/P_{t-1} \rfloor\big)$ {cash limit; $C = 0$ in this study}
  $n_t^S \leftarrow B_t + \min(S_t - B_t, K_t^S)$
end if
$r_t^B \leftarrow B_t - n_t^B$; $r_t^S \leftarrow S_t - n_t^S$
if $r_t^B > 0$ **then**
  sort $\mathcal{B}$ by $z(\text{seed}, t, i, +1)$, then by trader identity
end if
if $r_t^S > 0$ **then**
  sort $\mathcal{S}$ by $z(\text{seed}, t, i, -1)$, then by trader identity
end if
for all the first n_t^B traders i in $\mathcal{B}$ **do**
  $\bar{c}_i \leftarrow (\bar{c}_i q_i + P_{t-1})/(q_i + 1)$; $q_i \leftarrow q_i + 1$; $w_i \leftarrow w_i - P_{t-1}$
end for
for all the first n_t^S traders i in $\mathcal{S}$ **do**
  $q_i \leftarrow q_i - 1$; $w_i \leftarrow w_i + P_{t-1}$ {$\bar{c}_i$ was unchanged on a sale}
end for
$W_t \leftarrow W_{t-1} + P_{t-1}(n_t^B - n_t^S)$
$Q_t \leftarrow Q_{t-1} + n_t^S - n_t^B$
return $(n_t^B, n_t^S, u_t^B, u_t^S, r_t^B, r_t^S)$ and settled ledgers

The implementation kept two forms of nonexecution separate. Agent suppression meant that the trader lacked cash or shares before matching. Dealer rejection meant that the trader was eligible but the finite dealer could not absorb the unmatched residual. Neither event altered the trader's CA state. This distinction prevented a financial constraint from silently rewriting the behavioral state that generated the order.

A.2 Rationing and paired reproducibility

When dealer capacity was binding, only the excess side required rationing. Each eligible trader on that side received a 64-bit score from a SplitMix64 finalizer applied to the tuple $(\text{seed}, t, \text{trader identity}, \text{side})$. We sorted orders by score and used trader identity as a deterministic tie breaker. Priorities changed with t, so a trader who submitted the same order repeatedly did

not retain the same position. Paired paths also used the same score for a given (t, i, side) after their eligible sets diverged. The score did not consume draws from the random-number generator used to construct the lattice and quenched disorder. The resulting rationing was reproducible without giving the same traders permanent priority.

For every rule-switch pair, we asserted exact equality of the frame at T_{pre}, the Sell-fraction history through that frame, the edge coordinates and band widths, and the price history through the quote formed from the common pre-switch state. For every held-boundary balance-sheet cell we also asserted equality of the complete Sell-fraction history and final frame across all nine cash–share interventions. The assertions tested that paired experiments differed only at the stated intervention.

A.3 Conservation and limiting cases

With no intervention, agent-plus-dealer cash and shares are invariants:

$$\sum_i w_i(t) + W_t = \sum_i w_i(0) + W_0, \tag{10}$$

$$\sum_i q_i(t) + Q_t = \sum_i q_i(0) + Q_0. \tag{11}$$

Across the production reservoir and finite-size campaigns, the maximum absolute cash error was 2.53×10^{-6} and the share error was zero. In the balance-sheet campaign, conservation was checked separately before and after the intentional intervention jump. Its maximum cash error was 9.47×10^{-7} and its share error was zero.

Three exact limits supplied regression tests. With $h = \infty$, every agent-eligible order filled, and the complete path reproduced the earlier fill-all implementation. With $h = 0$, only opposing agent orders cleared, executed imbalance was zero, the price stayed fixed, and the endogenous-boundary lattice matched the held-boundary lattice. In the reservoir campaign, $h = 10$ never rejected an order and reproduced the $h = \infty$ paths exactly. The $h = 10$ equality held for the stated endowments and horizon and was not a general identity between finite and unlimited dealers.

A.4 Campaign inventory

The production inventory contained 30,550 simulations across the four campaigns (Table S1). The reservoir campaign reused one no-switch control within each $(h, \text{boundary}, \text{seed})$ cell, so its comma-separated value file contained one treated summary row per target fraction without duplicating the control. The other campaigns reconstructed a deterministic row order from independently checkpointed cells after threaded computation. This organization made a resumed threaded run reproduce the same tables as an uninterrupted run.

Table S1: Production campaign design. "Endogenous/held" indicates that both matched boundary conditions were run.

Campaign	Lattice sizes	Dealer h	Boundary	Seeds	Simulations
Selected paired switch	64	∞	endogenous	50	400
Reservoir screen	64	0–$10, \infty$	endogenous/held	50	9,600
Finite-size switch	32–128	$1, \infty$	three conditions	50	9,750
Cash–share factorial	64	$0.1, 1, \infty$	endogenous/held	50	10,800

The repository retained the earlier external-price classroom implementation and its archived Invesco QQQ exchange-traded fund input for provenance. No result from that model entered the principal analysis or figures.