Agent-Based Models of Financial Markets: A Wolfram Perspective*

Jeffrey D. Varner
R.F. Smith School of Chemical and Biomolecular Engineering
Cornell University, Ithaca, NY 14850
jdv27@cornell.edu

Abstract

We study crash-like transitions in a deterministic, three-state cellular automaton on a two-dimensional grid. Interior agents update from their four nearest neighbors using one of two totalistic rules, and their mixture fraction f is the main control parameter. In Model A, QQQ minute data drive a subset of boundary cells. The model shows a sharp finite-horizon crossover from a mixed state to a sell-dominated state. Sell domains form locally and spread across the grid as coherent fronts. Mean-field and pair approximations do not reproduce this spatial behavior. Across random initial conditions and rule placements, onset times follow an inverse-distance empirical curve consistent with a nucleation interpretation. Model B replaces the external price with an endogenous feedback loop and filters desired orders through finite cash and share inventories. In the present experiments, budgets shift onset times less than twofold and instead alter post-onset domain growth. A paired cash intervention supports a price-to-boundary feedback mechanism, although the result is specific to the model's execution and market-maker assumptions. The finite rule space makes the same analysis possible for every totalistic rule in this model class.

*Repository: https://github.com/varnerlab/Wolfram-markets-abm-paper.git.

Preprint.

1 Introduction

Large price declines in financial markets are not always traceable to news of comparable size. The crash of October 1987 removed roughly a fifth of the value of the United States equity market in a single session, and the flash crash of May 2010 took about nine percent off the major indices and returned it within the hour, neither of them following an exogenous announcement of matching magnitude [1, 2]. Market returns carry statistical signatures that an account built on independent external shocks does not reproduce either, since heavy tails, volatility clustering, and long memory in absolute returns persist across assets, sampling frequencies, and decades [3, 4]. Both observations point to the same broad question, whether the interactions among traders can generate large moves on their own, and we address one narrow version of it here: can local interactions alone produce a spatially organized transition from mixed trading states to a sell-dominated state? If they can, then the interaction network matters, because it determines where a disturbance begins and how far it spreads [5]. We study this question in a deliberately minimal model and then test how finite agent budgets affect the onset and growth of the sell-dominated domain.

The ingredients for an answer already exist, though they have not been assembled. On one side, a generation of agent-based models (ABMs) has demonstrated that heterogeneous populations of chartists and fundamentalists can reproduce fat tails and clustered volatility [6, 7], that herding on random graphs is sufficient to generate power-law-tailed returns [8], and that minority-game dynamics [9] and adaptive strategy switching [10, 11] produce non-Gaussian fluctuations without any exogenous noise assumption. On the other side, a mature literature on lattice and spin models (Ising models with market coupling [12, 13], opinion dynamics on chains [14], and the three-state speculative model of Kaizoji [15]) has shown that nearest-neighbor interactions on a regular grid can produce phase transitions between ordered (crash) and disordered (efficient) regimes. Meanwhile, cellular automata (CA) have been applied to financial markets [16–18], but the rules in those studies were chosen by hand rather than drawn from an enumerated space, which leaves them resistant to systematic comparison. What is missing is a framework that combines *local* spatial dynamics on a regular lattice with a *complete, enumerable* space of behavioral rules and a *direct* connection to real market data.

In this paper, we assemble those ingredients in a spatial cellular-automaton market. Agents occupy the cells of a two-dimensional grid, hold one of three states (Buy, Hold, or Sell), and update deterministically from their four cardinal neighbors under Wolfram totalistic CA rules [19]. The interior mixes just two such rules, a majority/herding rule and a contrarian/destabilizer rule, at the minority (contrarian) fraction f that serves throughout as the single control parameter. The totalistic rule framework supplies two properties that the models cited above do not combine: strictly local information propagation, so that sell cascades must spread from neighbor to neighbor as spatially resolved wavefronts; and a finite, enumerable rule space ($3^9 = 19{,}683$ rules), which allows systematic screening of agent behaviors rather than rule design by hand. We study two complementary configurations (Figure 1). In Model A (open), boundary agents observe real QQQ ETF minute-bar data through a Bollinger Band filter, which supplies an empirical boundary forcing. In Model B (feedback-closed), the boundary responds instead to an endogenous price formed from the grid's own net order imbalance, a reflexive loop that produces crashes without any external input, and every agent holds finite cash and shares, so that budgets can suppress the very orders sentiment demands. It is in this budget-constrained feedback loop that we separate onset from post-onset growth. Model A exhibited a sharp finite-horizon crossover between $f = 0.14$ and $f = 0.15$, above which localized sell domains grew as spatially coherent fronts, whereas mean-field and pair approximations retained a stable mixed branch and therefore did not reproduce that crossover, leaving an inverse-distance fit to the observed onset times that we treat as phenomenological rather than as a derived thermal escape law. Model B showed similar spatial onset across the tested cash levels, with onset times equivalent within a factor of two across a tenfold cash range, while budgets altered subsequent front growth, and a paired cash intervention linked that growth effect to the endogenous price and the boundary response, since cash injection accelerated the front in 14 of 19 included pairs while withdrawals produced no detectable change in the tested configurations. Taken together, these results separate the onset of a sell-dominated domain, which the tested budgets did not move, from its growth, which they did, and they locate the difference in the boundary channel through which price information reaches the interior.

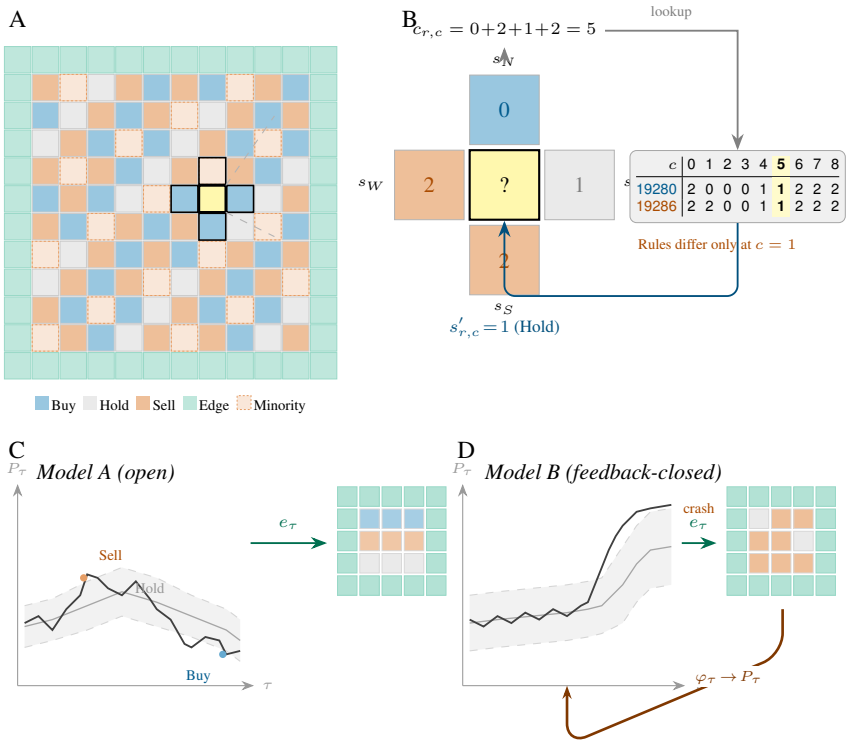


Figure 1: Model architecture. A: Agent grid ($N_r \times N_c$) with three states (Buy, Hold, Sell). Green boundary cells are edge agents; dashed-border interior cells follow the minority rule (19286); solid interior cells follow the majority rule (19280). The highlighted cell and its Von Neumann neighbors are expanded in B. B: CA update: the four neighbor states are summed to produce $c_{r,c}$, which indexes into the rule table to determine the next state $s'_{r,c}$. In this example $c_{r,c} = 5$, so both rules output Hold. The two rules differ only at $c = 1$. C: Model A (open): exogenous QQQ prices (black trace) fluctuate within a Bollinger Band (gray region). Band breakouts set the edge-agent state, which is injected into the grid boundary. D: Model B (feedback-closed): the grid's net order imbalance φ_τ feeds back (orange arrow) to update an endogenous price, closing a reflexive loop.

2 Related Work

Our work lies at the intersection of three research traditions: agent-based models of financial markets, lattice and spin models from statistical mechanics, and cellular automata theory. We review each strand below and identify where existing approaches stop short of the connection we propose.

2.1 Agent-Based Models and Stylized Facts

The quantitative study of financial market dynamics has been shaped by a tension between tractable theoretical frameworks and the empirical complexity of observed data. Mandelbrot [4] demonstrated that speculative price changes exhibit heavy tails beyond Gaussian predictions; Fama [20] showed that raw return series display negligible linear autocorrelation, consistent with the efficient markets hypothesis; yet Schwert [21] documented that the *magnitude* of returns exhibits positive autocorrelation decaying slowly over time, so that large price changes tend to be followed by large changes regardless of sign, a pattern known as *volatility clustering*. Bouchaud and Potters [22] further documented the asymmetric response of volatility to positive versus negative returns (the leverage effect). Cont [3] provides a systematic review confirming these regularities across asset classes, geographic markets, and sampling frequencies. Collectively, leptokurtosis, the absence of return predictability, and volatility clustering define the benchmark against which any model of market dynamics must be evaluated.

Agent-based models offer a bottom-up route to these stylized facts: rather than assuming the statistical structure of returns, they derive it from the simulated interaction of heterogeneous agents. The Santa Fe Artificial Stock Market [23] established the template of adaptive traders interacting through a centralized order book. Lux and Marchesi [6] showed that a two-type fundamentalist-chartist population reproduces heavy tails and volatility clustering as emergent properties, a result that established ABMs as a credible framework for studying market dynamics. Cont and Bouchaud [8] demonstrated that a simple random-graph herding structure alone is sufficient to generate fat-tailed return distributions, highlighting the role of coordination, rather than agent sophistication, in generating extreme events. The minority game of Challet and Zhang [9] introduced the paradigm of bounded-rational agents competing via strategic coordination, a framework subsequently shown to produce non-Gaussian fluctuations; Farmer and Joshi [10] demonstrated that standard trading strategies simultaneously deployed by many agents endogenously generate volatility clustering and excess kurtosis without any exogenous noise assumption; and Chen, Tan, and Zheng [24] showed that multi-level herding simultaneously reproduces sector structure and volatility clustering. Harras and Sornette [25] modeled myopic adapting agents whose local herding behavior inflates bubbles that eventually collapse, connecting the ABM literature to crash dynamics.

On the theoretical side, Hommes [11] surveys heterogeneous agent models incorporating adaptive expectations and evolutionary strategy selection; LeBaron [26] provides a systematic treatment of agent-based computational finance; Arthur [27] argued that heterogeneous agent interactions give rise to phenomena (persistent disequilibrium, emergent structure, tipping points) inaccessible to equilibrium theory; and Farmer [28] demonstrated that quantitative ABMs can produce time-series predictions competitive with mainstream macroeconomic models. Battiston et al. [5] made the case that complexity science and ABMs should inform financial regulation. These models share a limitation, however: agents interact either globally (mean-field coupling) or on ad hoc network topologies, making it difficult to disentangle the contributions of local spatial contagion from global coordination. The present work addresses this gap by placing agents on a regular lattice with strictly local CA rules, so that every cascade must propagate neighbor-by-neighbor through physical space.

2.2 Lattice Models, Spin Models, and Phase Transitions

The mapping of financial agent interactions onto lattice models draws on the statistical mechanics tradition, where the mathematical machinery for analyzing cooperative phenomena on regular grids (phase transitions, critical exponents, universality classes) is well developed. Ising-type models map agent states (buy/sell) onto spins and impose nearest-neighbor coupling on a lattice [29, 30]. These models naturally capture local herding and can exhibit phase transitions between ordered (bubble/crash) and disordered (efficient) regimes. Bornholdt [12] introduced a two-dimensional Ising model with a global coupling term representing market impact, producing intermittent bubbles and crashes from the interplay of local herding and global frustration across scales. Sznajd-Weron and Weron [14] proposed a one-dimensional opinion model for price formation in which local consensus propagates outward, demonstrating that even the simplest spatial rule can generate nontrivial collective dynamics. Krawiecki et al. [13] showed that attractor-bubbling dynamics on a lattice generate volatility clustering with realistic scaling exponents. Kaizoji [15] extended the Ising framework to three states (buy/hold/sell) to model speculative activity, which is the closest antecedent to our three-state CA formulation.

The crash literature provides the other essential ingredient. Sornette [1] argued that large crashes are the culmination of cooperative herding behavior and can be modeled as critical points analogous to phase transitions in physical systems. The self-organized criticality (SOC) framework of Bak, Tang, and Wiesenfeld [31] suggests that complex systems may naturally evolve toward a critical state where scale-free avalanches occur without parameter tuning, and Bak and Paczuski [32] applied this idea to economic systems. Gabaix et al. [33] derived power-law tails from the size distribution of large institutional trades, providing an alternative mechanism for heavy tails that does not require agent heterogeneity. Scheffer et al. [34] proposed generic early-warning signals, rising autocorrelation and variance, for systems approaching a tipping point, a framework that has been applied to ecosystems, climate, and financial markets [5]. Farmer et al. [35] demonstrated that even zero-intelligence agents can reproduce aspects of market microstructure when placed in a realistic order-book mechanism, suggesting that the structure of interaction may matter more than the sophistication of individual decision rules.

Statistical-mechanics accounts describe crashes as collective transitions with slow relaxation or endogenous amplification [36], a distinction that has been tested empirically in other complex systems by separating endogenous from exogenous shocks [37], and related market models study feedback between funding and market liquidity, the 2010 Flash Crash, circuit breakers, and anticipated bailouts [2, 38–40]. Our model is not a calibrated version of those settings, and it differs from an Ising lattice in two further ways, since we replace the single coupling constant with Wolfram totalistic CA rules, which provide a richer and fully enumerable space of $3^9 = 19{,}683$ local update functions, and we couple the lattice to an external or endogenous price signal through boundary agents. The external series supplies an empirical forcing. It does not validate the model's prices. The fixed-seed sweep shows a sharp finite-horizon crossover rather than a parameter-free SOC transition, and the narrower advantage of the construction is that the same deterministic realization can be rerun with and without a cash intervention.

2.3 Cellular Automata in Complex Systems and Finance

Wolfram [41] classified one-dimensional CA into four behavioral classes (fixed points, periodic orbits, spatial chaos, and complex/undecidable dynamics) and later provided an exhaustive treatment of the subject [19]. Packard and Wolfram [42] extended this classification to two-dimensional totalistic rules on Von Neumann neighborhoods, demonstrating rich pattern-formation dynamics including self-replication, spiral waves, and spatial chaos. The classification program reveals that even deterministic, strictly local rules can produce dynamics ranging from trivially ordered to computationally universal, a property that makes CA attractive as a modeling framework wherever the goal is to understand how microscopic rules generate macroscopic complexity.

Applications to financial markets, however, are few, and we are aware of three that apply CA directly to market dynamics. Wei et al. [16] proposed a CA model for securities markets in which agent states evolve based on local majority rules and demonstrated qualitative agreement with empirical return distributions. Qiu et al. [17] developed a CA model specifically targeting abnormal fluctuations, showing that local interaction rules can produce heavy-tailed returns. Yang et al. [18] simulated financial markets using CA with heterogeneous update rules and showed that the ratio of contrarian to trend-following agents controls the transition between calm and volatile regimes, a result qualitatively consistent with our phase diagram, though their model uses a different rule framework and does not characterize the transition as a sharp phase boundary. None of these studies uses the full Wolfram totalistic rule framework, and none couples the CA to real market data or resolves the spatial structure of crash cascades.

The present work embeds agent dynamics in the complete Wolfram totalistic rule space (3^9 rules), which allows systematic enumeration and comparison rather than rule design by hand; it couples the CA to real market data (Model A) or an endogenous reflexive price (Model B) through boundary agents, bridging the gap between abstract CA dynamics and empirical financial time series; and it identifies a discontinuous transition once the contrarian fraction crosses a threshold, resolving its spatial structure as nucleation of sell cascades from localized seeds followed by wavefront propagation, which connects the CA dynamics to the tipping-point framework of Scheffer et al. [34].

3 Methods

We first establish the formal ingredients of the model, the CA rule framework that governs interior agent dynamics, the Bollinger Band filter that translates price signals into edge-agent states, and the specific rules used in this study, and then describe the two simulation configurations and the phase-sweep protocol used to map the crash transition.

3.1 Three-State Totalistic Cellular Automata

We consider a two-dimensional grid of $N = N_r \times N_c$ agents, each in one of three states: Buy ($s = 0$), Hold ($s = 1$), or Sell ($s = 2$). At each discrete time step, every interior agent updates its state deterministically based on the states of its four cardinal (Von Neumann) neighbors. For an agent at position (r, c) with neighbor states $\{s_N, s_S, s_E, s_W\}$, the neighborhood sum is:

$$c_{r,c} = s_N + s_S + s_E + s_W \in \{0, 1, \ldots, 8\}. \tag{1}$$

Table 1: Wolfram totalistic CA rule taxonomy. Base-3 encoding reads right-to-left: digit a_i gives the next state when the neighborhood sum index equals i.

Rule	Base-3 encoding	Interpretation	Role
6158	$(022110002)_3$	Imitation / herding	Homogeneous model; stable waves
19280	$(222110002)_3$	Majority rule	80% crash model; herding
19286	$(222110022)_3$	Contrarian	20% crash model; destabilizer

A *totalistic rule* is a mapping $\mathcal{R} : \{0, 1, \ldots, 8\} \to \{0, 1, 2\}$ that assigns a next state to each of the nine possible neighborhood sums. Since each of the nine outputs can take three values, there are $3^9 = 19{,}683$ distinct rules, each identified by its decimal index:

$$\text{Rule index} = \sum_{j=0}^{8} a_j \cdot 3^j, \tag{2}$$

where $a_j = \mathcal{R}(j)$ is the next state when the neighborhood sum equals j. This sum uses the ordered numerical encoding Buy $= 0$, Hold $= 1$, and Sell $= 2$. Relabeling the states would change the rule. The ordering is therefore part of the model, not a neutral label choice. This rule space is finite, enumerable, and complete: every possible deterministic, translation-invariant, totalistic update function on the Von Neumann neighborhood is represented exactly once.

3.2 Edge Agent Decision Model

Boundary cells serve as *edge agents* that inject external information into the grid (Figure 1A, green cells). A fraction f_e of boundary cells, selected uniformly at random with a fixed seed, are designated as edge agents. The remaining boundary cells stay fixed in the Hold state. The CA update acts only on strictly interior cells.

Each edge agent classifies the current price P_τ relative to an exponential moving average (EMA) and its associated volatility band:

$$\hat{S}_\tau = \hat{S}_{\tau-1} + \alpha\big(P_\tau - \hat{S}_{\tau-1}\big), \tag{3}$$

$$\hat{\sigma}_\tau = \hat{\sigma}_{\tau-1} + \alpha\big(|P_\tau - \hat{S}_\tau| - \hat{\sigma}_{\tau-1}\big), \tag{4}$$

where $\alpha = 1/60$ corresponds to a 60-minute smoothing window. The decision rule is a three-way classifier:

$$e_\tau = \begin{cases} 0 \text{ (Buy)} & \text{if } P_\tau < \hat{S}_\tau - k\,\hat{\sigma}_\tau, \\ 1 \text{ (Hold)} & \text{if } \hat{S}_\tau - k\,\hat{\sigma}_\tau \le P_\tau \le \hat{S}_\tau + k\,\hat{\sigma}_\tau, \\ 2 \text{ (Sell)} & \text{if } P_\tau > \hat{S}_\tau + k\,\hat{\sigma}_\tau, \end{cases} \tag{5}$$

where k controls the width of the hold band. When the price is within k standard deviations of its moving average, the edge agent holds; when the price breaks out above (below) the band, it sells (buys). This is a standard mean-reversion signal widely used in technical analysis, and its simplicity is deliberate: the edge agents provide a minimal, interpretable channel through which real market information enters the grid.

3.3 Rule Taxonomy

Three rules are used in this study (Table 1). Rule 19280 is a majority/herding rule: agents tend to adopt the dominant state in their neighborhood, promoting coordination and spatial coherence. Rule 19286 is a contrarian rule that differs from 19280 at a single neighborhood sum, $c = 1$, the case with three Buy neighbors and one Hold neighbor. A Rule 19280 agent responds to this configuration by *buying* ($a_1 = 0$), reinforcing the local majority; a Rule 19286 agent instead *sells* ($a_1 = 2$), acting against the prevailing state. This single-digit difference in the base-3 encoding is the minimal perturbation that converts a stabilizing agent into a destabilizing one, and it is this minimal difference that drives the phase transition studied below.

3.4 Model A: Open (Boundary-Driven) Simulation

In Model A (Figure 1C), edge agents observe real market data, namely one-minute OHLC bars for the QQQ ETF (Invesco QQQ Trust, January 2 through March 28, 2024, $T = 23{,}790$ time steps). The vendor feed is no longer publicly available, so an archival copy travels with the repository alongside its SHA-256 checksum and a record of the bar accounting, 61 trading days of exactly 390 regular-session bars each. At each step τ the typical price $P_\tau = (H_\tau + L_\tau + C_\tau)/3$ is computed from the OHLC bar, each edge agent then evaluates Eq. (5) using the online EMA and volatility from Eqs. (3)–(4), and all interior agents update simultaneously via their assigned CA rule. The grid is initialized to all-Hold ($s = 1$), so that any Buy or Sell activity must originate from the boundary and propagate inward.

For the heterogeneous (crash) model, a fraction f of interior agents are assigned Rule 19286 (contrarian) uniformly at random, while the remaining $1 - f$ agents use Rule 19280 (majority), and the minority agent positions are fixed for the duration of the simulation (quenched disorder), so that the only control parameter is the mixture fraction f. The default grid is 256×256 ($N = 65{,}536$ agents), edge agents comprise 25% of boundary cells ($f_e = 0.25$, yielding 255 active edge agents), selected with a fixed random seed for reproducibility, the EMA smoothing parameter is $\alpha = 1/60$ (a 60-minute window), and the band width is $k = 1.0$ standard deviations.

3.5 Model B: Feedback-Closed (Endogenous Price) Simulation

In Model B (Figure 1D), there is no external price feed. The price is generated endogenously from the grid's own collective state, and every trading cell, each interior agent together with each active edge agent, now carries a *book* that gates whether its sentiment can be acted upon: a cash balance initialized to `initial_cash` (the sweep variable, default 1,000), a share inventory initialized to 10 shares acquired at the reference cost basis $P_0 = 100$, and a running average cost basis that tracks the price paid across subsequent fills. Sentiment still flows by the same totalistic rule as in Model A, which is what keeps the two models' interiors directly comparable, but a cell now expresses that sentiment through orders: a cell in the Buy state submits a one-share buy order every step it remains in that state, a cell in the Sell state submits a one-share sell order, and a Hold cell submits nothing.

Orders do not clear against one another. They fill at the standing quote $P_{\tau-1}$ against an unmodeled infinite-inventory market maker in the spirit of the price-impact accounting of Kyle [43] and the market-clearing conventions of Farmer and collaborators [35]. A buy fills if and only if the cell holds cash $\geq P$; a sell fills if and only if the cell holds at least one share and, when the disposition toggle is on, the standing quote exceeds the cell's average cost basis, $P > \bar{c}_i$. This cost-basis sell rule is the *disposition effect* of Shefrin and Statman [44], the reluctance to realize losses that Odean later documented in the trading records of individual investors [45]. We treat it not as a modeling assumption but as a behavioral bias and hence a sweep factor, toggled on and off across campaigns. Every order that the book cannot honor is tallied as a *suppressed* order rather than executed: budget constraints suppress order *flow*, never the underlying CA states. The interior uses the same update rule as Model A, while cash and inventory throttle its market footprint. Because the market maker lies outside the modeled population, cash and shares are not conserved among grid agents. Model B is closed with respect to price feedback, not economically closed.

Because fills, not states, now drive the price, the imbalance is measured on executed flow rather than on the raw sentiment census:

$$\varphi_\tau = \frac{(\text{executed buys})_\tau - (\text{executed sells})_\tau}{N_{\text{traders}}}, \qquad N_{\text{traders}} = N_{\text{interior}} + N_{\text{active edge}}, \tag{6}$$

and the price updates multiplicatively from this executed imbalance,

$$P_\tau = P_{\tau-1} \cdot \exp\big(\lambda\, \varphi_\tau\big), \tag{7}$$

where $\lambda > 0$ is the price-impact parameter controlling how strongly the aggregate order flow moves the quote. Edge agents are full participants in φ, since their fills enter the numerator on the same footing as interior fills, and, to break the all-or-none lockstep that a single shared band would impose, each edge cell draws its own band half-width $k_i \sim U\big(k(1-\delta),\ k(1+\delta)\big)$, where δ is the band-spread parameter; the edge cells then evaluate Eq. (5) against the newly generated price P_τ, with the online EMA and volatility recursions of Eqs. (3)–(4) observing that same new price. The result is a fully

self-contained loop: the grid's executed orders generate the price, the price drives the edge agents, and the edge agents inject signals back into the grid.

To probe that loop causally we expose a deterministic intervention hook, a function of the current step and the full collection of books, (t, books), invoked at the top of step t before any orders execute, through which we implement mid-run cash shocks by scaling every cell's cash by a common factor. These are uniform wealth shocks within the model, not conventional order-book liquidity operations. The hook takes no randomness of its own, so the same seed still yields an identical trajectory up to the moment of intervention. This same-seed, identical-trajectory contract is what the paired counterfactuals rely on to attribute any downstream divergence to the liquidity shock alone.

3.6 Phase Sweep Protocol

To map the finite-horizon crossover, we swept the minority fraction f from 0.0 to 0.50 in two passes. A coarse sweep with $\Delta f = 0.05$ covered the full range; a fine-grained sweep with $\Delta f = 0.01$ resolved the transition region $f \in [0.08, 0.22]$. For each value of f, the full T-step price series drove a 128×128 grid (chosen for computational efficiency), and we recorded the sell fraction at every time step. Three summary statistics characterize each run: mean sell fraction, peak sell fraction, and crash frequency (the fraction of time steps at which the sell fraction exceeds 0.50). Each value of f used one fixed disorder realization. The sweep therefore brackets a finite-horizon crossover; it does not estimate a thermodynamic critical point.

3.7 Model B Nucleation Campaigns

The endogenous-price machinery of Model B lets us treat a market crash as a nucleation event and interrogate it with a family of campaigns. Unless a campaign states otherwise, all of them share one configuration, a 64×64 grid, price impact $\lambda = 0.02$, nominal band half-width $k = 1$, band spread $\delta = 0.5$, the disposition rule switched off, an initial quote $P_0 = 100$, and deterministic seeds, and they differ only in the minority fraction f, the initial cash, and the horizon over which we let the loop run, which is $1{,}000$ steps for the seed sweep, $1{,}500$ steps for the budget-toggle and boundary-versus-bulk campaigns, and $2{,}000$ steps for the escape-time campaign. Two definitions, fixed once and reused by every campaign, let the campaigns be read together: we label sell clusters by breadth-first search over the 4-connected Sell cells of the grid, and we define the crash onset t_{on} as the first step at which the largest sell cluster reaches 5% of the grid. Cluster labeling runs on the full grid in every campaign but one, since the boundary-versus-bulk birth analysis instead labels on the interior alone, so that a cluster attached to the active-edge ring can be told apart from one born in the bulk, and because onset and cluster size are measured identically everywhere else, front velocities and onset times reported by different analyses can be read on a common footing.

The seed sweep compares variation across configurations with variation across disorder realizations. We ran three configurations, a *stable* regime at $f = 0.10$ with cash $1{,}000$, a *freeze* regime at $f = 0.20$ with cash 500, and a *collapse* regime at $f = 0.20$ with cash $5{,}000$, across 20 seeds each. The budget toggle then turns the deterministic hook into a controlled experiment. For each seed, we used the control trajectory to identify t_{on}. In the paired counterfactual, every cell's cash is scaled at the start of that step by $\times 10$ (an injection) or by $\times 0.1$ or $\times 0.01$ (two withdrawal doses). Histories are bit-identical through step $t_{\text{on}} - 1$, so the intervention immediately precedes the control's first detectable onset frame. We then measured outcomes in an onset-aligned window of $W = 400$ steps and took the front velocity to be the endpoint difference quotient, the change in largest-cluster fraction between onset and onset$+W$ divided by W, rather than a fitted slope. A pair is excluded from an arm when its control run had no onset at all or when t_{on} fell too late to leave a full window before the horizon ($t_{\text{on}} >$ horizon $- W$).

The boundary-versus-bulk campaign asks where the winning cluster is born: we tracked the founding cell of the onset cluster and compared its location against the analytic null for spatially uniform births, under which the fraction of births falling within distance 2 of the boundary is 0.125, and we complemented this observational arm with an interventional one that stamps a 3×3 sell patch onto the grid at boundary-adjacent, mid, and center placements, each paired against an unstamped control.

The onset-time campaign measures how the waiting time changes with f at three cash levels, cash $\in \{500, 1{,}000, 5{,}000\}$. A pilot pass swept $f \in \{0.12, 0.15, 0.18, 0.20, 0.25, 0.30\}$ at 20 seeds, and having localized the onset cliff to the interval between 0.15 and 0.18 at all three cash levels,

a refinement pass added $f \in \{0.156, 0.162, 0.168, 0.174, 0.177\}$ at 40 seeds. We took the escape time to be $\tau = t_{\text{on}}$, right-censored at the 2,000-step horizon and never discarded, so runs without onset still inform the Kaplan–Meier estimates. We fitted the empirical form $\ln \tilde{\tau} = A + B/(f - f_0)$ separately at each cash level using the finite Kaplan–Meier medians, with the pseudo-threshold f_0 left free, and obtained 95% confidence intervals on (A, B, f_0) from 2,000 bootstrap resamples of the seeds within each f.

A companion velocity–suppression dose-response holds $f = 0.20$ and extends the cash ladder with the levels 250 and 2,000 so that five doses are probed at 20 seeds each. For every run we formed, over the window $[t_{\text{on}} + 50,\ t_{\text{on}} + 400]$, a front velocity as the OLS slope of the front extent and a suppression load as the mean *absolute* number of suppressed buy orders per step, and we tested their association by Spearman rank correlation with a 10,000-draw permutation p-value, excluding and separately counting any run that is censored before onset or whose window endpoint $t_{\text{on}} + 400$ would exceed the horizon.

Finally, to place these targeted campaigns against the model's full parameter dependence, we drew on a large factorial sweep whose crash statistics anchor the Model B regime discussion: the full crossing of $f \in \{0,\ 0.1,\ 0.15,\ 0.2,\ 0.3\}$, $\lambda \in \{0.005,\ 0.01,\ 0.02,\ 0.05\}$, cash $\in \{500,\ 1{,}000,\ 5{,}000\}$, disposition $\in \{\text{off},\ \text{on}\}$, band spread $\delta \in \{0,\ 0.5,\ 0.9\}$, and band half-width $k \in \{0.5,\ 1,\ 2\}$ yields 1,080 configurations which, replicated across 10 seeds of 2,000 steps, comprise 10,800 runs. The full design and its variance decomposition are deferred to Appendix B.

4 Results

Model A provides an unconstrained spatial baseline. Prices arrive from outside the model, and CA states are not filtered through budgets. We first inspected representative grid trajectories and then located the finite-horizon crossover in a fixed-seed sweep.

We simulated the 256×256 agent grid under Model A for three values of the minority fraction, $f = 0$ (homogeneous), $f = 0.15$ (near-critical), and $f = 0.20$ (supercritical), and recorded grid snapshots at four logarithmically spaced time points to capture both the early transient and the late-time steady state (Figure 2). In the homogeneous model ($f = 0$, top row), the grid reached a dynamic steady state dominated by Buy (blue) and Sell (orange) agents in a fine-grained spatial mixture. The pattern fluctuated locally but never exhibited large-scale coordination, and the sell fraction remained near its equilibrium value of $\sim$0.26 throughout the simulation (cf. Figure 3). At $f = 0.15$ (middle row), the first sampled value above the crossover bracket, the system initially resembled the homogeneous case, but at $t = 8{,}236$ we observed a sell cascade that nucleated from a localized seed and propagated outward as a spatially coherent wavefront. The orange region grew concentrically until it engulfed the grid by $t = 23{,}790$, which suggests a contagion mechanism qualitatively different from the instantaneous global crashes produced by mean-field models: the crash had a *spatial history*, starting at one location, spreading through a physical medium, and arriving at different parts of the grid at different times (Figure 2, middle row). That sequence, a localized seed, a coherent front, and a spatial history, is nucleation-and-growth in the classical sense. At $f = 0.20$ (bottom row), the cascade occurred earlier and more rapidly, and the grid was already predominantly sell-dominated by $t = 2{,}851$, consistent with a system deep in the supercritical regime.

Having characterized the two regimes, we next located the crossover between them. We swept the minority fraction f from 0.0 to 0.50 and computed three summary statistics, the mean sell fraction, the peak sell fraction, and the crash frequency, for each run (Figure 3). All three summaries changed sharply between the sampled values $f = 0.14$ and $f = 0.15$. The dashed line at $f = 0.145$ marks the midpoint of this bracket, and it is not an independently fitted critical point. At and below $f = 0.14$, the system remained stable: the mean sell fraction hovered near 0.26 (close to the $1/3$ expected for a uniform random mixture of three states), the peak sell fraction never exceeded 0.30, and no crashes occurred. At and above $f = 0.15$, the system entered a sell-dominated regime: the mean sell fraction jumped to $\sim$0.50 at $f = 0.15$ and saturated near 0.91 by $f = 0.20$, and crash frequency exceeded 95% (Figure 3C). The fixed-seed scan therefore showed a sharp finite-size, finite-horizon crossover, and whether that crossover sits at a fixed value of f calls for repeating the measurement across five lattice sizes with replicated disorder.

These results have an economic interpretation: a small minority of contrarian agents was harmless, because Rule 19280 (majority) agents damped local sell signals through herding toward the dominant

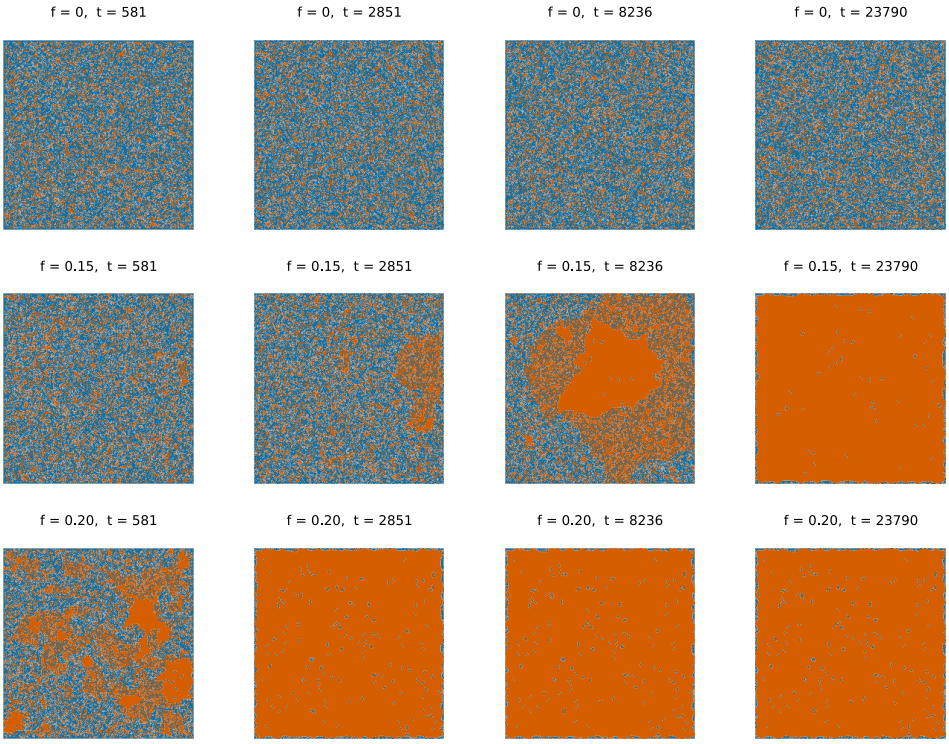


Figure 2: Spatiotemporal dynamics of the 256×256 agent grid under Model A (Rules 19280/19286, QQQ 2024 Q1 data). Columns share logarithmically spaced time points; rows vary the minority fraction f. Blue = Buy, gray = Hold, orange = Sell. At $f = 0.15$ (middle row, $t = 8{,}236$), a sell cascade nucleates from a localized seed and propagates as a coherent wavefront before engulfing the grid.

state. But once the contrarian fraction entered the crossover bracket, Rule 19286 agents were numerous enough to seed sell clusters that the majority-rule agents could no longer absorb; instead, the herding mechanism *amplified* the cascade, because the local majority had tipped from buy to sell. The result was a coordination failure triggered by the interaction between two deterministic rules, without any stochastic component, learning, or adaptive strategy switching. Below the bracket, sell clusters still formed, but each died before reaching the critical size at which the herding mechanism turns from damper to amplifier. The crash was an emergent property of the rule mixture, not of any individual rule.

The crossover was not located at a fixed value of f (Figure 4). We repeated the sweep on lattices from 32×32 to 512×512, with twenty disorder seeds at each size, on periodic boundaries and with no price forcing. Larger lattices crashed at lower contrarian fractions. At $f = 0.14$, no run out of twenty converted on the 64×64 lattice, half converted at 128×128, and every run converted at 256×256 and above (Figure 4A). Certain conversion required $f \approx 0.155$ at 64×64 but only $f \approx 0.14$ at 512×512. The apparent threshold is therefore a property of the lattice and the observation window, not of the rule mixture alone.

The direction of that shift runs opposite to the usual intuition. One might expect a larger population to absorb contrarian pressure better, since idiosyncratic behavior averages out as the number of agents grows. That reasoning applies to quantities that are averages. Whether the lattice crashes is not an average: it is decided by the most favorable local region anywhere on the grid, since a single sell domain that survives to critical size is enough to convert the whole system. Enlarging the lattice adds candidate regions, so it raises the chance that at least one is favorable. The mechanism is the size effect familiar from brittle fracture, where a larger specimen fails at lower stress than a small one of the same material because it is more likely to contain a critical flaw. Size does not make a stable configuration unstable, since above the spreading threshold every lattice converts eventually, because

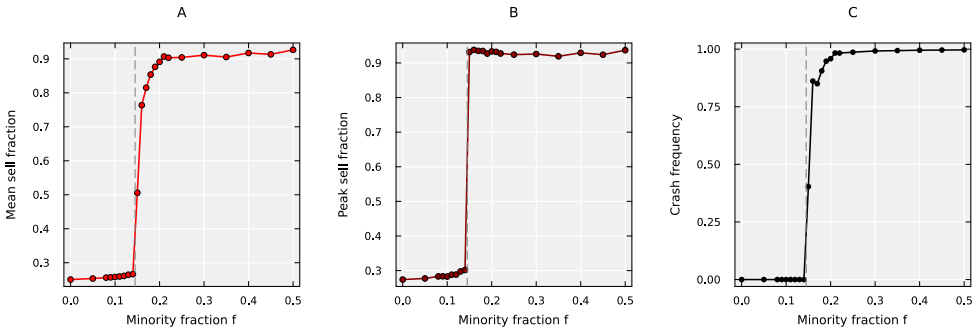


Figure 3: Phase diagram for Model A (Rules 19280/19286, 128×128 grid). A: Mean sell fraction vs. minority fraction f. B: Peak sell fraction. C: Crash frequency (fraction of time steps with sell fraction > 0.50). The dashed vertical line at $f = 0.145$ marks the midpoint of the observed $[0.14, 0.15]$ crossover bracket, not a fitted thermodynamic critical point.

the sell state is absorbing and the waiting time is finite, but it does decide what has already happened inside a fixed window.

The averaging intuition is nonetheless correct about a different quantity (Figure 4). As the lattice grew, conversion times became sharply more predictable: the coefficient of variation across seeds fell from 0.66 at 32×32 to 0.035 at 512×512 at $f = 0.20$, halving on each doubling of the linear size and so tracking $N^{-1/2}$ in the agent count N (Figure 4B). Larger systems also stopped taking longer to convert. Above $f \approx 0.17$ the median conversion time was independent of size, within 9% across a 16-fold range in area (Figure 4C), which is the signature of a conversion carried by many domains growing and merging at once rather than by one domain sweeping the lattice. Size averages away the uncertainty in when the crash arrives without removing the crash.

What sets that common clock is visible in the shape of the conversion curve (Figure 4). Writing the converted fraction as $X(t) = 1 - \exp[-(Kt)^n]$, the exponent n distinguishes two mechanisms: $n = d + 1 = 3$ in two dimensions if new domains keep appearing throughout the transformation, and $n = d = 2$ if the domains are all present at the outset and the transformation is growth alone. Fitted per run, n ran from 2.54 at $f = 0.16$ to 1.91 at $f = 0.20$ on the largest lattice ($R^2 = 0.994$ to 0.999), approaching 2 as f rose. The rate constant K did not depend on system size: read as a characteristic time, $1/K$ was 1445, 1363, 1367, and 1358 steps across $4{,}096$ to $262{,}144$ agents at $f = 0.18$, a spread of 6%. Rescaling each run's time by its own fitted K collapsed all 320 runs, across four sizes and four values of f, onto a single curve (Figure 4D).

An exponent near 2 says that the vulnerable regions were fixed before the clock started. That is what the construction implies: the contrarian sites are assigned once, at setup, and never move, so the favorable regions are a property of the quenched disorder rather than something that arises during the run. What differs between seeds is how quickly the transformation spreads from regions that were present all along. The distinction matters for how the onset-time fit should be read. Thermal activation would mean ignition at random times; this transformation is measurably not that, which is why the inverse-distance form is a compact empirical summary rather than a derived law.

Two limits bound these measurements. The 32×32 lattice behaves anomalously because a domain of the size used to mark local clustering already spans a tenth of it, so local and system-wide events are not separable at that size, and we plotted it (Figure 4) but excluded it from the conversion-curve fits. The fixed cluster-size criterion used to flag local clustering is also not a scale-free detector of nucleation, since the mixed state produces sell clusters of 30 to 60 cells unprompted and the largest such cluster grows with the number of sites. Conclusions about rates here rest on the conversion curve, which requires no cluster identification.

The fixed-seed scan showed a sharp finite-horizon crossover that spread through coherent fronts. We next compared the lattice simulations with two low-order approximations and summarized the observed onset times. Neither mean-field nor pair closure reproduced the crossover, and across the sampled range the onset times were well fit by $\ln \bar{\tau} = 2.78 + 0.389/(f - 0.101)$ ($R^2 = 0.991$), a

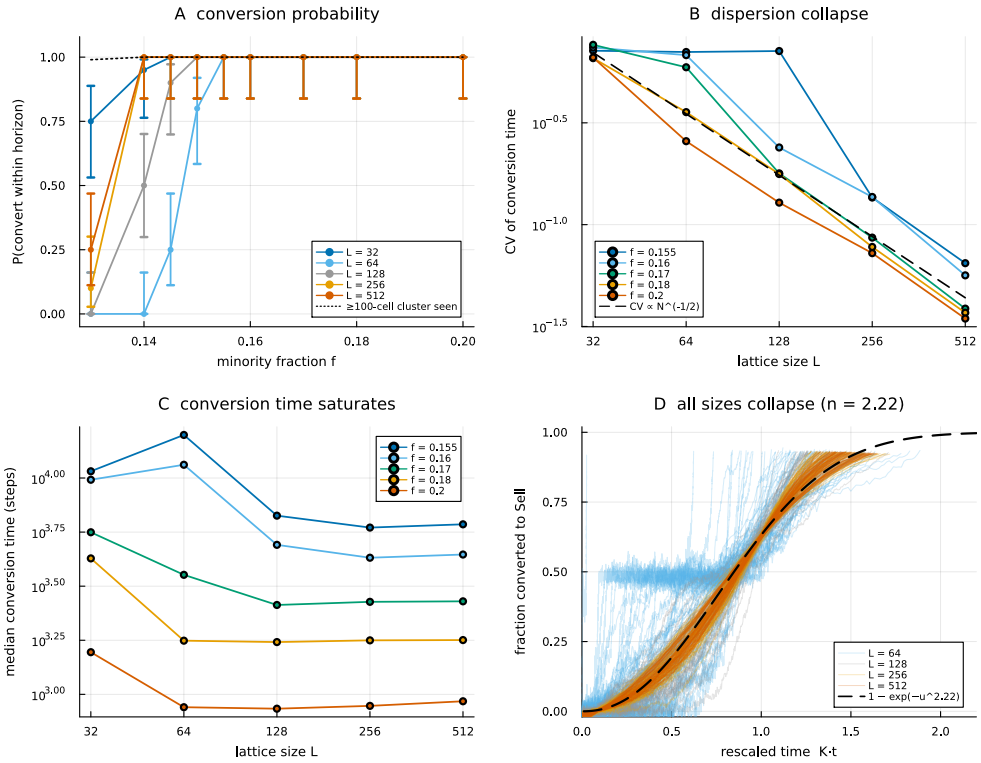


Figure 4: System-size dependence of the Model A crossover (Rules 19280/19286, periodic boundaries, no price forcing, twenty disorder seeds per cell). A: Fraction of seeds converting within the 100,000-step horizon against f, one curve per lattice size, with Wilson score intervals. Larger lattices convert at lower f. The dotted curve marks where a 100-cell sell cluster was seen at any size; such clusters appear even where nothing converts, so what moves with size is spreading rather than clustering. This is a descriptive threshold and not a certified critical nucleus. B: Coefficient of variation of the conversion time against lattice size, for cells in which every seed converted, against the $N^{-1/2}$ self-averaging reference. C: Median conversion time against lattice size, saturating above $f \approx 0.17$. D: All 320 conversion curves plotted against time rescaled by each run's own fitted rate constant, with $1 - \exp(-u^n)$ overlaid. Sizes lie on top of one another; the band tightens with size, and the smallest lattice visibly fails to collapse.

curve we use as a phenomenological description of the finite data rather than as a derived thermal escape law.

To measure local ordering, we used the nearest-neighbor alignment statistic associated with a three-state Potts model:

$$e \;=\; -\frac{1}{N_{\text{bonds}}} \sum_{\langle i,j \rangle} \delta(\sigma_i,\, \sigma_j), \tag{8}$$

where the sum runs over all distinct nearest-neighbor pairs on the lattice and δ is the Kronecker delta. In a perfectly aligned (single-state) configuration $e = -1$; in a maximally disordered configuration $e \approx -1/3$. This statistic measures order, but it is not a Hamiltonian that generates the CA dynamics. Rule 19280 acts on the numerical sum of state labels rather than on counts of matching neighbors, and a synchronous update need not decrease e. For example, three Hold neighbors and one Buy neighbor have $c = 3$, for which the rule outputs Buy rather than the local majority state. A further anomaly occurs at $c = 0$ (all neighbors Buy), which outputs Sell rather than Buy. This rule entry means that the all-Buy configuration is *unstable*: every interior cell flips to Sell in a single time step, and the all-Sell absorbing state ($c = 8 \to$ Sell) has maximal alignment, $e = -1$. The "buy-dominated" phase observed in simulations is therefore not a static equilibrium but a *dynamic mixed-state attractor* with fractions $p_0 \approx 0.49$, $p_1 \approx 0.24$, $p_2 \approx 0.26$ and $e \approx -0.37$ (Figure 5A).

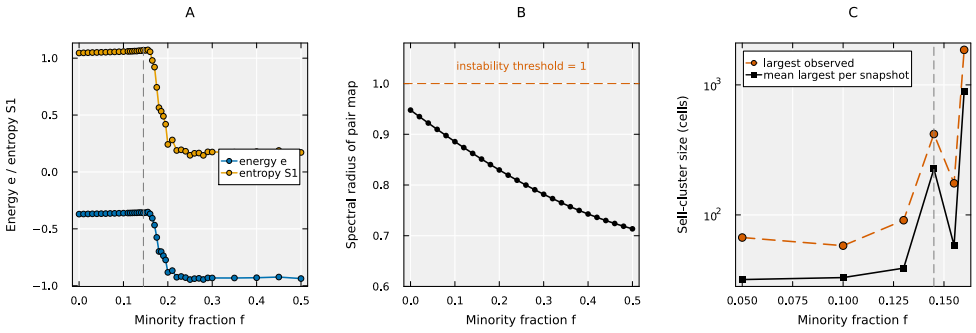


Figure 5: Theory overview for Model A (Rules 19280/19286). A: Potts energy e and site entropy S_1 vs. minority fraction f (128×128 lattice, periodic boundaries, time-averaged over 2,000 steps). The mixed attractor persists at $e \approx -0.37$ and $S_1 \approx 1.06$; above the transition the system collapses to the all-Sell state at $e \approx -0.93$ and $S_1 \approx 0.18$, a sharp finite-size drop ($\Delta e \approx 0.56$, $\Delta S_1 \approx 0.88$); near the crossover the 2,000-step averages retain the mixed state because metastable lifetimes exceed the window, so the crossover completes to the right of the marked f_c. B: Spectral radius of the pair-approximation Jacobian at the uncorrelated fixed point, declining monotonically from 0.95 at $f = 0$ to 0.71 at $f = 0.5$, strictly below the instability threshold (dashed line at 1). This tested pair closure does not capture the crossover. C: Sell-cluster sizes in the mixed state (log scale): largest cluster ever observed (orange) and mean largest per snapshot (black), over 100 snapshots of a 128×128 lattice. The mean largest cluster (black) holds at ∼30–60 cells in the subcritical mixed state, while the largest observed cluster (orange) reaches ∼400 cells near $f = 0.145$ and exceeds 1,800 cells by $f = 0.16$ (the dip at $f = 0.155$ is a single-realization fluctuation of the fixed-seed protocol). Dashed vertical lines mark the midpoint $f = 0.145$ of the observed $[0.14, 0.15]$ bracket.

The contrarian rule (19286), which differs only at $c = 1$ by outputting Sell instead of Buy, provides quenched rule disorder: a fraction f of sites carry a fixed sell-biased response to the three-Buy, one-Hold neighborhood.

We measured the Potts energy and the single-site Shannon entropy $S_1 = -\sum_s \rho_s \ln \rho_s$ across the phase sweep on a 128×128 lattice with periodic boundaries. Below f_c the mixed attractor persisted at $e \approx -0.37$ and $S_1 \approx 1.06$ (near the maximum $\ln 3 \approx 1.10$); above the crossover the system approached the all-Sell absorbing state at $e \approx -0.93$ and $S_1 \approx 0.18$, a sharp finite-size drop consistent with a discontinuous transition (Figure 5A).

To isolate the role of spatial correlations, we derived a mean-field self-consistency map for the mixed 19280/19286 population. The mean-field approximation replaces the lattice with a well-mixed population in which each agent's four neighbors are drawn independently from the global state distribution (p_0, p_1, p_2), where p_s denotes the fraction of agents in state s ($p_0 + p_1 + p_2 = 1$). Because each neighbor contributes independently to the neighborhood sum $c = s_N + s_S + s_E + s_W$, the probability $\pi_k = \Pr(c = k)$ follows from the multinomial convolution of four i.i.d. draws from $\{0, 1, 2\}$ with weights (p_0, p_1, p_2):

$$\pi_k = \sum_{\substack{n_0+n_1+n_2=4 \\ n_1+2n_2=k}} \binom{4}{n_0,\, n_1,\, n_2} p_0^{n_0}\, p_1^{n_1}\, p_2^{n_2}, \qquad k = 0, 1, \ldots, 8. \tag{9}$$

The verified rule tables (Table 1) assign next states as follows: $c \in \{2, 3\} \to$ Buy; $c \in \{4, 5\} \to$ Hold; $c \in \{0, 6, 7, 8\} \to$ Sell (both rules agree at these values); and $c = 1$, where the two rules diverge, with Rule 19280 outputting Buy and Rule 19286 Sell. In a population with minority fraction f, the mean-field map is therefore:

$$p_0' = (1 - f)\,\pi_1 + \pi_2 + \pi_3, \tag{10}$$

$$p_1' = \pi_4 + \pi_5, \tag{11}$$

$$p_2' = \pi_0 + f\,\pi_1 + \pi_6 + \pi_7 + \pi_8, \tag{12}$$

where primes denote the next-step fractions and π_k depends on the current (p_0, p_2) through Eq. (9).

We iterated Eqs. (10)–(12) to convergence ($\|\mathbf{p}' - \mathbf{p}\|_1 < 10^{-14}$) from multiple initial conditions and swept f from 0 to 1. For every value of f the map admitted two stable fixed points, a *buy-dominated*

equilibrium with $p_0^* \approx 0.63$, $p_1^* \approx 0.18$, $p_2^* \approx 0.19$ at $f = 0$, shifting smoothly to $p_0^* \approx 0.45$, $p_2^* \approx 0.23$ at $f = 1$, and the all-Sell absorbing state $p_2^* = 1$. The Jacobian at the buy-dominated fixed point had complex eigenvalues with modulus $|\lambda| \approx 0.95$ (a stable spiral), confirming local stability. At no value of f did the buy-dominated branch disappear or cross the sell-majority threshold $p_2 > 0.50$. The mean-field map predicted no phase transition: the sell fraction increased by only $\sim$0.04 across the entire range of f, a smooth monotone drift with no discontinuity.

We next extended the mean-field analysis to the pair approximation, which tracks nearest-neighbor pair densities $\rho_{ab} = \Pr(\sigma_i = a,\, \sigma_j = b)$ for adjacent sites i, j rather than single-site marginals alone. In the pair approximation closure [46], the three neighbors of a focal site *other* than its paired partner are drawn independently from the conditional distribution $q(s \mid a) = \rho_{as}/\rho_a$, where $\rho_a = \sum_b \rho_{ab}$ is the marginal density of state a. For a synchronous update, the new pair density is

$$\rho'_{bb'} = \sum_{a,a'} \rho_{aa'}\; T_b(a, a')\; T_{b'}(a', a), \tag{13}$$

where $T_b(a, a')$ is the probability that a site in state a with paired neighbor in state a' transitions to state b; it is computed from the effective rule table and the trinomial distribution of the remaining three neighbors' sum (Appendix A).

The pair approximation also predicted no phase transition. We computed the 6×6 Jacobian of the pair map (13) evaluated at the uncorrelated (mean-field) fixed point $\rho^*_{ab} = \rho^*_a \rho^*_b$ and found that its spectral radius was strictly less than unity at every f we sampled ($f \in [0, 0.5]$), declining monotonically from $|\lambda_{\max}| \approx 0.95$ at $f = 0$ to $|\lambda_{\max}| \approx 0.71$ at $f = 0.5$ (Figure 5B). A systematic search over correlated initial conditions (same-state correlation ε ranging from 0.1 to 0.5, both buy- and sell-biased) confirmed that every trajectory converged to one of only two fixed points: the uncorrelated buy-dominated state (identical to mean-field) or the all-Sell absorbing state. Spatial mutual information $I = 2S_1 - S_2$ decayed to zero in all cases.

The pair approximation gave the same negative result as mean field: within this particular nearest-neighbor closure, the mixed branch remained locally stable. That does not rule out every finite-order closure, but it shows that single-site marginals and the tested pair correlations are insufficient. The explicit lattice simulations instead showed localized clusters and propagating fronts.

Those localized clusters and seed-dependent waiting times motivated a *nucleation* description, in which a small sell domain appears, survives local fluctuations, and grows into a macroscopic domain, and we characterized that interpretation through three measurements. We first measured the metastable lifetime. We initialized 64×64 lattices from random states, equilibrated for 200 steps to the mixed attractor, then recorded the time τ at which the sell fraction first exceeded 0.90 (collapse to the absorbing state). For each value of f we ran 20 independent realizations with distinct random seeds for both initial conditions and minority-agent placement. Below $f \approx 0.14$ the mixed state survived in all realizations for 10^5 steps; above $f \approx 0.155$ all realizations collapsed, with the mean lifetime decreasing from $\bar{\tau} \approx 24{,}000$ at $f = 0.155$ to $\bar{\tau} \approx 86$ at $f = 0.30$.

We next fit the observed onset times. Motivated by the inverse-distance form used in Kramers-type escape problems [47, 48], we fit

$$\ln \bar{\tau} \;=\; A + \frac{B}{f - f_0}, \tag{14}$$

with best-fit parameters $A = 2.78$, $B = 0.389$, and $f_0 = 0.101$ ($R^2 = 0.991$; $n = 12$ data points with all seeds collapsed; collected, with the Model B fits, in Table 2). Because the fit excludes values at which not all seeds collapsed within the horizon and is not derived from the CA dynamics, f_0 is an extrapolated pseudo-threshold rather than proof of absolute stability. The observed change near $f = 0.15$ is best read as an observation-time-dependent crossover: the fitted waiting time enters the simulation window over a narrow range of f.

As an extrapolation, Eq. (14) gives an effective crossover at a chosen observation window: for example, $f_c(T = 10{,}000) \approx 0.162$ and $f_c(T = 1{,}000) \approx 0.195$, consistent with the steepening of the phase boundary observed in the simulations.

We then measured the spatial influence of an isolated contrarian agent embedded in an otherwise all-majority lattice. The sell density at the contrarian's own site was elevated ($\rho_2 \approx 0.42$ vs. the background $\rho_2 \approx 0.26$), but this excess decayed to zero within one to two lattice spacings. Individual contrarians therefore act as point-like sell seeds with negligible spatial range; the transition requires

the *collective* effect of many such seeds. Sell cluster analysis confirmed this: below f_c the largest sell cluster in the mixed attractor contained $\sim$30–60 cells; at $f = 0.145$ transient clusters of $\sim$400 cells appeared; and at $f = 0.16$ macroscopic clusters exceeding $1{,}800$ cells formed and triggered system-wide collapse (Figure 5C).

The measurements support a disorder-driven nucleation interpretation: the quenched random placement of contrarian agents at density f creates local density fluctuations; at $f > f_0$, rare regions with an anomalously high local contrarian density seed sell clusters that exceed the critical size; the majority/herding rule then amplifies these seeds into coherent wavefronts that engulf the lattice. The system-size dependence of f_c was weak ($f_c \approx 0.17$ at $L = 32$, $f_c \approx 0.16$ at $L = 128$), consistent with a critical nucleus that is small relative to the lattice dimensions at all tested sizes.

Taken together, the sequence of approximations narrows the plausible description of the lattice crossover. Mean-field theory, by erasing all spatial structure, missed the transition. The pair approximation, which restores nearest-neighbor correlations, fared no better: the uncorrelated fixed point remained linearly stable at all f, and the spatial mutual information vanished. The lattice simulations added the missing ingredients: localized clusters, broad seed-dependent waiting times, and propagating fronts. The Potts-inspired statistic summarized the accompanying increase in alignment, with a finite-size drop $\Delta e \approx 0.56$ from the mixed attractor ($e \approx -0.37$) toward the all-Sell state ($e \approx -0.93$). These observations are consistent with a nucleation-like discontinuous crossover but do not establish a thermodynamic free-energy barrier (Figure 5A).

Model B provides a direct comparison. Retaining the same two rules and quenched disorder but adding endogenous price feedback and per-agent cash and inventory constraints, it let us ask whether the same empirical onset-time form described each cash level, whether the spatial crossover survived budgets, whether cash was associated with onset, and how budgets affected post-onset growth. The experiments are designed to detect large shifts, and overlapping uncertainty intervals cannot exclude smaller budget effects.

We first asked whether the crossover survived the budget constraint, drawing on a factorial sweep of $10{,}800$ runs across minority fraction, price impact, cash, the disposition switch, and two order-band parameters (Appendix B). A sell-dominated regime appeared between the sampled values $f = 0.15$ and $f = 0.20$ at every tested price impact. A variance decomposition attributed 94–97% of each summary metric's variance to differences among the swept configurations and at most 5.4% to disorder seeds within a configuration. That decomposition is descriptive, and it does not identify which individual parameter causes the difference. The classifier assigned every run to one of two realized regimes and left the other three candidate classes empty, with $6{,}589$ runs (61%) called oscillatory and $4{,}211$ (39%) crash-prone. "Crash-prone" is a grid-state label: the Sell fraction is at least 0.5 on at least 1% of steps. It does not imply a large price drawdown. That classification also had a price signature. The modal crash-prone outcome was a liquidity *freeze*, in which the interior locked near a Sell fraction of 0.88 while suppressed orders limited price movement. The median maximum drawdown in this class was about 5%, and its median minimum price was 100.00. Price collapses formed a tail, with 897 of the $4{,}211$ crash-prone runs falling below 90. Across matched design cells, enabling the disposition rule increased the return standard deviation by a factor of 2.40 on average. One tail trajectory sat beside an oscillatory run (Figure 6), and it was not the modal case.

We next compared the largest Sell clusters in stable, freeze, and collapse configurations, using 20 disorder seeds per configuration (Figure 7). In the stable regime the at-peak ratio of the largest cluster to the total Sell population had a median of 0.022 (IQR 0.018–0.027) across every seed, since the Sell cells present were scattered into roughly 500 dispersed fragments, the largest of which never exceeded about 0.5% of the grid, and all 20 trajectories stayed flat, whereas the freeze and collapse regimes gathered their Sell cells into a single dominant domain, reaching median at-peak ratios of 0.26 and 0.44 respectively with zero overlap against the stable band. That gap was the observed signature of localized domain growth rather than a synchronous grid-wide flip, and the above-threshold trajectories that produced it formed time-shifted sigmoids whose onset varied more than their eventual shape. That pattern motivates the onset-aligned front velocities, budget toggles, and onset-time fits that follow. It is consistent with variable waiting times, but does not by itself establish a Kramers escape process.

Having established that a nucleus exists, we asked where on the grid it was born, and in particular whether the active-edge ring that injects the budget-constrained orders is also the cradle of the winning cluster (Figure 8). The boundary-versus-bulk campaign tracked the founding cell of the

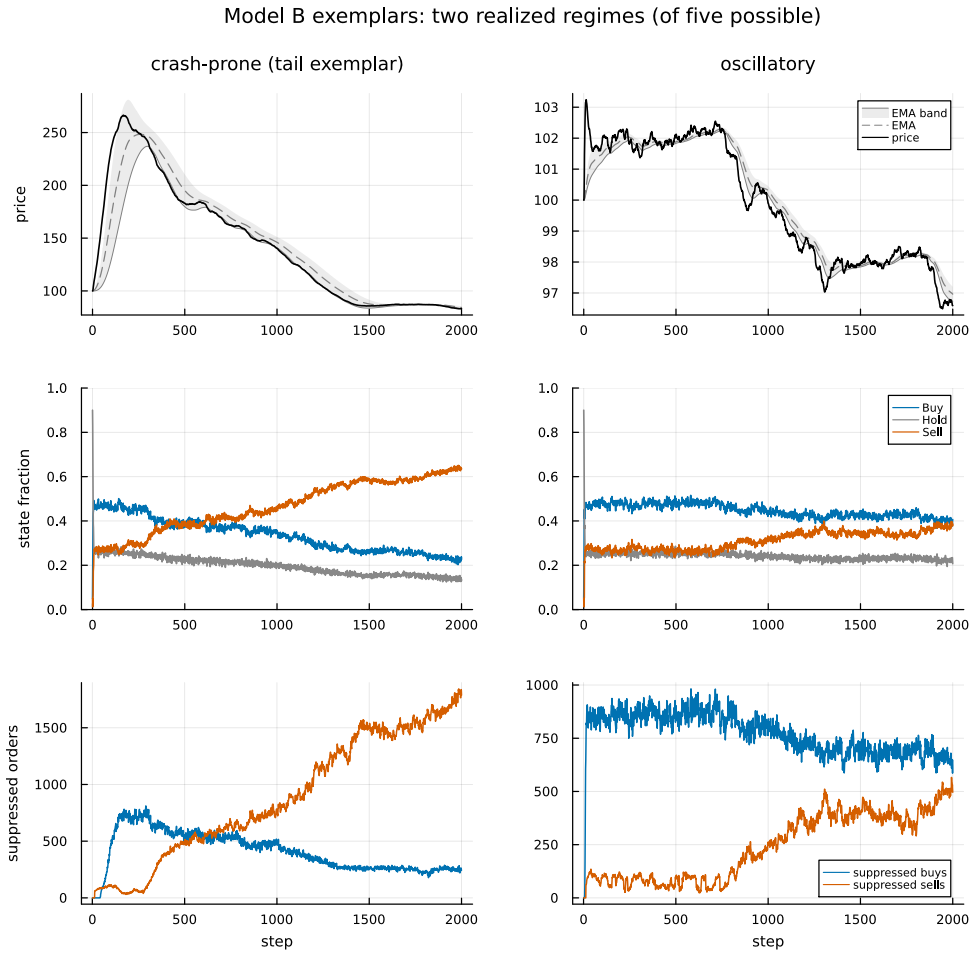


Figure 6: Model B regime exemplars: price, interior state fractions, and suppressed orders for a crash-prone run (left, run 9652) and an oscillatory run (right, run 5916). The crash-prone exemplar is the dramatic tail of its regime, a bubble-and-collapse trajectory ($100 \to 266 \to 83$), shown because it makes the mechanism visible; the *typical* crash-prone run is a liquidity freeze in which the grid locks near a Sell fraction of 0.88 while the price barely moves. "Crash-prone" throughout denotes majority-sell lock-in (Sell fraction ≥ 0.5 on at least 1% of steps), not a price discontinuity.

onset cluster across 40 runs, of which 39 nucleated within the horizon, and compared its position against the analytic null for spatially uniform births under which a fraction 0.125 of foundings fall within distance 2 of the boundary. Births were, if anything, biased away from the edge rather than toward it: the pooled fraction of onset clusters founded in the boundary band was 0.077, statistically indistinguishable from and slightly below the uniform expectation (binomial $p = 0.47$; the freeze configuration taken alone gave $p = 1.0$); the median founding cell sat a distance 12 into the deep bulk; and the typical founding was a single Sell cell (median birth size 1 cell), with a median birth step of 256 and individual onsets spanning 75 to $1{,}115$. The winning nucleus, in short, was overwhelmingly a bulk object, an interior single-cell fluctuation that happened to survive and grow, rather than a structure templated by the boundary. The one place the edge did assert itself was *after* birth: pre-onset attachment of the growing cluster to the active-edge ring occurred only in the collapse configuration (6 of 19 runs) and never in the freeze (0 of 20), a first hint that the boundary governs how a nucleus grows rather than where it is born, a distinction we sharpen when we ask where in the market the budget constraint actually bites.

A spatial, bulk-born nucleus raises a direct question: is cash associated with its onset time? Across a tenfold cash range, onset times proved equivalent within a factor of two (Table 2). We fitted the same

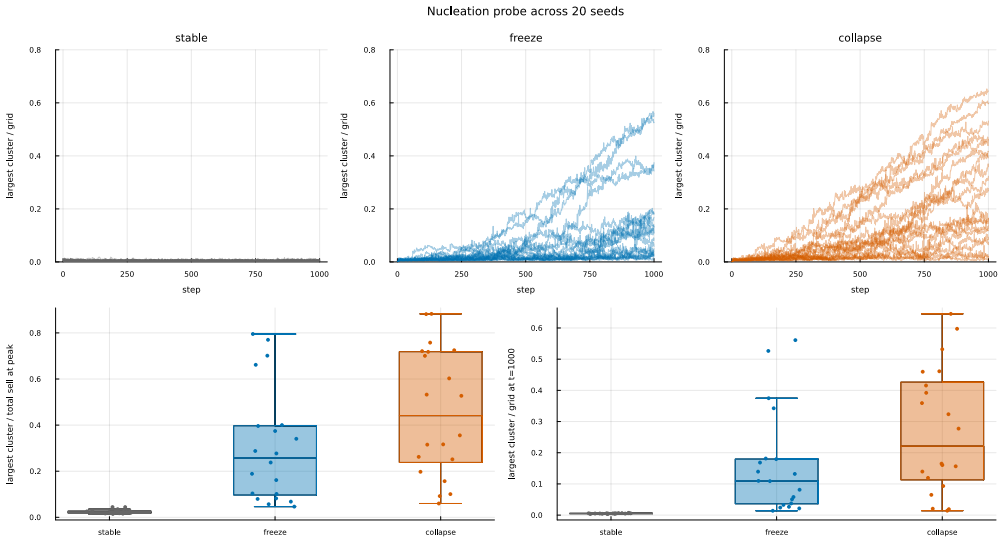


Figure 7: Model B seed sweep: largest-sell-cluster trajectories per seed. (Top row) For each of the three budget configurations (stable at $f = 0.10$, cash 1,000; freeze at $f = 0.20$, cash 500; and collapse at $f = 0.20$, cash 5,000), the largest sell cluster, expressed as a fraction of the grid, is drawn step by step for every one of 20 disorder seeds. (Bottom left) Per-seed ratio of the largest cluster to the total sell population at its peak: the stable configuration sits at a median ratio of 0.022, dispersed fluctuations with no dominant domain, while the freeze and collapse configurations concentrate into a single growing domain at median ratios of 0.26 and 0.44; the stable band shows zero overlap against either above-threshold configuration, the signature of a spatial nucleus above threshold and its absence below. (Bottom right) Final largest-cluster fraction of the grid at the end of the run. The time-shifted-sigmoid fan within each above-threshold configuration reflects stochastic nucleation waiting times rather than a spread of eventual outcomes.

empirical form, $\ln \tilde{\tau} = A + B/(f - f_0)$, to Kaplan–Meier median onset times at each cash level. The fits achieved $R^2 = 0.978$, 0.990, and 0.980 for cash 500, 1,000, and 5,000, using only 5, 6, and 6 finite median points, respectively (Table 2). Cells with fewer than half their seeds reaching onset have undefined medians and do not enter these fits. The small number of fitted points and the three free parameters make the curves descriptive rather than confirmatory.

Between the tightest and loosest budgets, the bootstrap confidence intervals overlapped on both fitted shape parameters, with the curves drawn over the medians (Table 2, Figure 9): for the pseudo-threshold f_0, cash 500 gave $(0.000, 0.153)$ against $(0.077, 0.157)$ at cash 5,000; for B, cash 500 gave $(0.070, 1.135)$ against $(0.050, 0.415)$; and the onset cliff itself sat between $f = 0.15$ and $f = 0.18$ at every cash level. We state the accompanying power caveat just as plainly, because an overlap of intervals is a bound and not a demonstration of equality: the B point estimates differ by $2.5\times$ across the cash ladder, and the cash-500 lower bound of 0.000 on f_0 is the physical boundary of the fit domain rather than a data quantile. Interval overlap is not a test of equality.

We therefore tested the difference directly. The three cash arms were run on identical disorder seeds at every f, so the bootstrap can be paired: each of 4,000 draws resampled seed positions once per f, applied the same resampled positions to all three arms, refit each arm, and differenced the fits within the draw, which removes the shared disorder variance from the difference. Because the three fitted parameters trade off against one another, we compared the quantity they jointly determine, the predicted log median onset time at a given f, whose difference between two arms is the log ratio of their predicted onset times, at every f retaining a defined median in all three arms. The margin was fixed in advance at a factor of two, $|\Delta \ln \tilde{\tau}| < \ln 2$, and assessed by two one-sided tests at $\alpha = 0.05$.

Every pairwise comparison met the margin at every shared f: budgets spanning a tenfold range changed the predicted median onset time by less than a factor of two. The largest point difference anywhere in the comparison was 0.354 in $\ln \tilde{\tau}$, a factor of 1.42, between cash 500 and 5,000 at $f = 0.20$, with a 95% interval of $(-0.098, 0.718)$. The individual parameters remain poorly

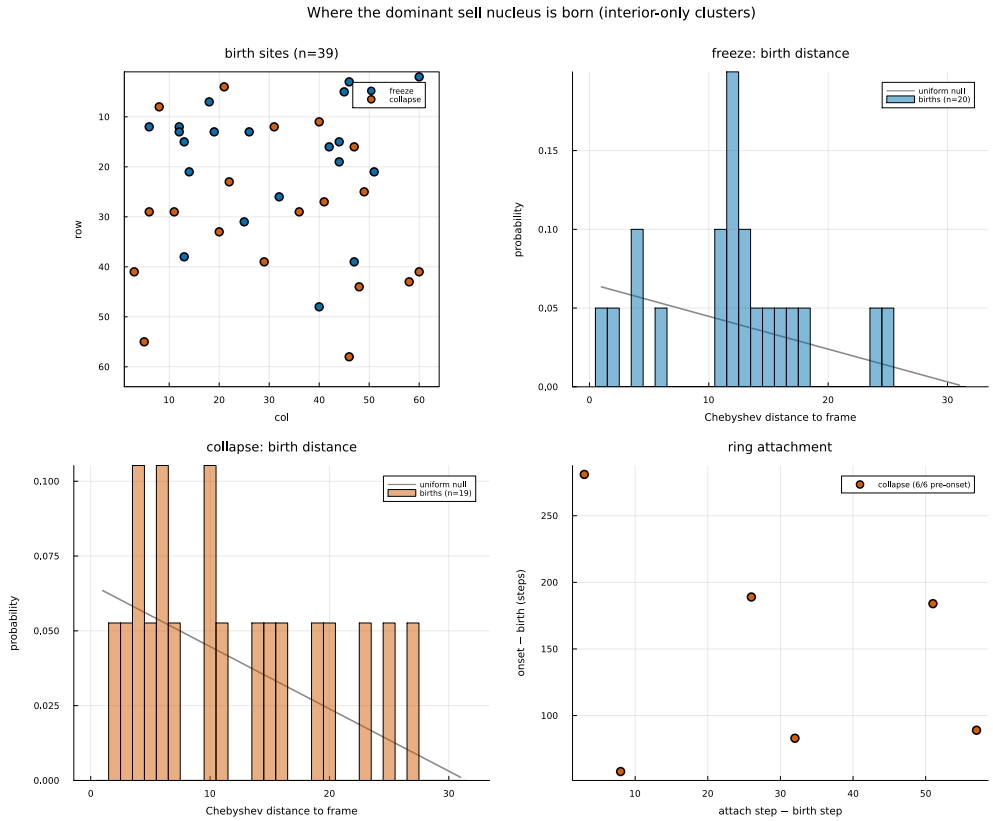


Figure 8: Birth locations of the onset cluster. (Top left) Map of the founding cell of the largest sell cluster at crash onset across the 39 boundary-versus-bulk runs that nucleated, labeled on the interior alone so that a bulk birth can be told apart from attachment to the active-edge ring. (Top right and bottom left) Histograms of the founding cell's distance from the nearest boundary, for the 20 freeze and 19 collapse runs separately, each against the uniform-null expectation. Births were bulk-dominated and consistent with a uniform distribution: the pooled fraction founded within distance 2 of the boundary was 0.077, below rather than above the 0.125 uniform null and statistically indistinguishable from it (binomial $p = 0.47$), with a median founding distance of 12 cells into the deep bulk. (Bottom right) Timing of ring attachment relative to birth and onset for the runs whose growing cluster reached the active-edge ring before onset (6 of 19 collapse runs and 0 of 20 freeze runs), the edge asserting itself only after birth. The winning nucleus is born in the bulk, not templated by the boundary.

determined by comparison, since the difference in B between cash 500 and $1{,}000$ has a 95% interval of $(-0.298,\ 0.965)$, which is why an argument from overlapping parameter intervals carried so little information. One comparison is worth flagging without promoting: between cash $1{,}000$ and $5{,}000$ near the threshold, the intervals exclude zero at $f = 0.177$ and $f = 0.18$, $(0.014,\ 0.433)$ and $(0.021,\ 0.398)$, placing onset roughly 19% later under the tighter budget. Those are adjacent points on the same pair of fitted curves and so are one signal rather than two, and neither survives correction for the fifteen comparisons in the family. We read it as a lead, not a result. What the onset-time analysis supports is therefore bounded rather than null: across a tenfold cash range, onset times were equivalent within a factor of two, and any residual effect is smaller than that.

The present experiments did not resolve a large cash-associated difference in founding sites or fitted onset-time parameters. They therefore direct attention to post-onset growth, without proving that cash has no effect on onset. One exploratory observation, which we report uncorrected for multiple comparisons and weigh accordingly, already points the same way: near the cliff the onset rate was monotone in cash, and at $f = 0.174$ the fraction of seeds nucleating within the horizon rose from 0.45 to 0.62 to 0.70 from tight to loose cash, while at $f = 0.177$ the Kaplan–Meier medians fell in the same order, $1{,}598$, $1{,}427$, and $1{,}208$ steps (Appendix C). Because the onset trigger is itself a growth

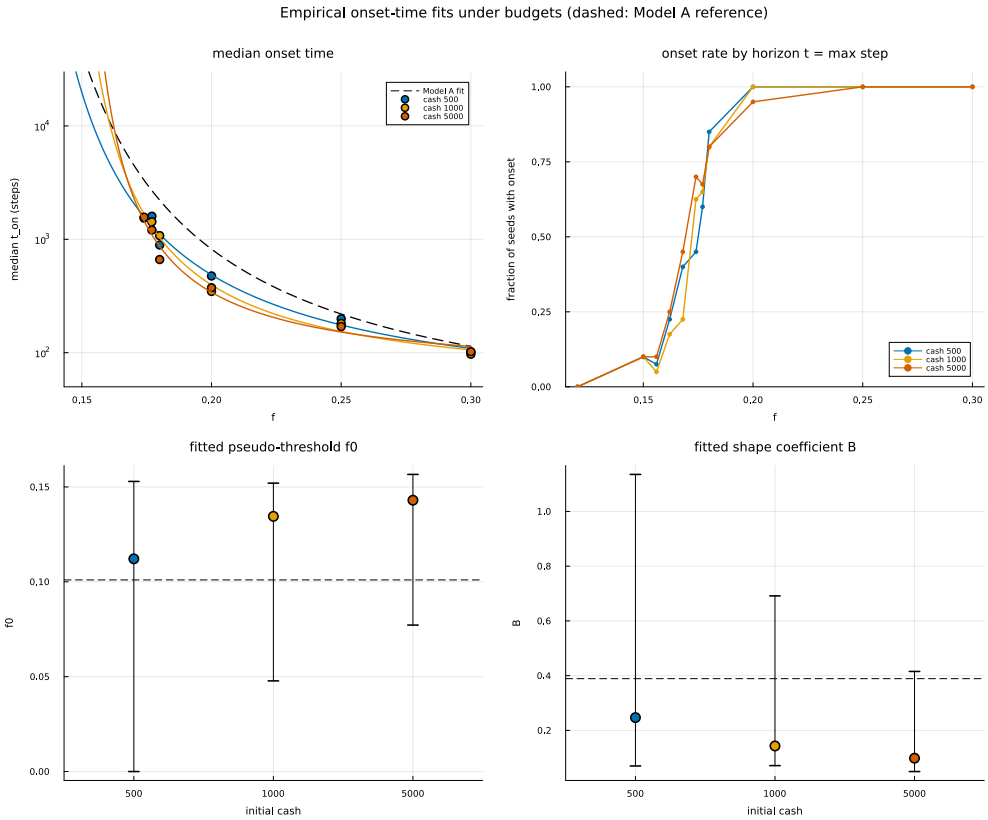


Figure 9: Empirical onset-time fits under budgets. (Top left) Kaplan–Meier median onset times (points) against the minority fraction f at cash 500, 1,000, and 5,000, each with its fitted divergence $\ln \tilde{\tau} = A + B/(f - f_0)$ (solid curves) and the Model A fit of Equation (14) dashed for reference; the fitted curves are drawn only up to a display ceiling so that the medians are not crushed by the divergence near f_0. (Top right) Fraction of seeds that nucleated within the 2,000-step horizon: the onset cliff sits between $f = 0.15$ and $f = 0.18$ at every cash level. (Bottom row) Fitted pseudo-threshold f_0 (left) and shape coefficient B (right) per cash level, with 95% bootstrap confidence intervals; the dashed lines mark the Model A values $f_0 = 0.101$ and $B = 0.389$. The intervals overlap between tight and loose cash on both parameters. That overlap is not itself a test; the paired equivalence analysis reported in the text is, and it bounds the between-budget difference in predicted median onset time to within a factor of two. Censoring at the horizon and its handling in the medians are detailed in Appendix C.

milestone, in that t_{on} marks the step at which the largest sell cluster first reaches 5% of the grid, a faster-growing front reaches the trigger sooner, and a monotone-in-cash onset rate is the signature one expects if budgets gate growth rather than birth. It is that gating we set out to demonstrate directly.

If budgets act on growth rather than birth, two predictions follow, one negative and one positive, and the nucleation campaigns let us test both (Figure 10). The negative prediction is that handing the interior a ready-made nucleus should buy nothing, because the interior was never short of nuclei; the interventional arm of the boundary-versus-bulk campaign therefore stamped a 3×3 sell patch onto the grid at boundary-adjacent (edge), mid, and center placements in both the freeze and the collapse configurations, each stamped run paired against an unstamped control, and the patch never founded the winning domain: in all six cells of the configuration-by-placement design the dominance rate was $0/20$, the domain standing largest at onset never once descended from the stamped patch, the patch's forward lineages grew only to 9–15 cells before dying out, never alive at the moment of onset, and no placement advanced the onset itself. The interior, evidently, does not lack seeds. The single-cell fluctuations from which onset clusters are born supply them constantly and at no cost, and what is scarce above threshold is not a nucleus but the *growth* that would make one matter.

Table 2: Empirical fits $\ln \tilde{\tau} = A + B/(f - f_0)$ to Kaplan–Meier median onset times, per cash level, with the Model A fit for reference. The Model A row is the mean-lifetime fit of Equation (14); the Model B rows fit Kaplan–Meier medians. Brackets are 95% bootstrap CIs (2,000 resamples of seeds within each f). The cash-500 f_0 lower bound of 0.000 is the physical boundary of the fit domain, not a data quantile. CIs overlap between tight (500) and loose (5000) cash on both B and f_0, but this overlap does not establish parameter equality.

Configuration	n_f	A	B	f_0	R^2
Model A (reference)	12	2.78	0.389	0.101	0.991
Model B, cash 500	5	3.38	0.247 [0.070, 1.135]	0.112 [0.000, 0.153]	0.978
Model B, cash 1000	6	3.80	0.143 [0.071, 0.691]	0.135 [0.048, 0.152]	0.990
Model B, cash 5000	6	4.10	0.098 [0.050, 0.415]	0.143 [0.077, 0.157]	0.980

The paired onsets, the dominance rates, and exemplar growth trajectories also contain one nominal placement effect, which ran, if anywhere, in the wrong direction for seeding: the freeze-configuration mid-grid patch *delayed* onset (median paired onset advance -257 steps, sign-flip $p = 0.0238$), an observation we hold as exploratory, since it is uncorrected across a six-test family, fails the Bonferroni threshold of 0.0083, and rests on a sign-flip test of the mean, and we accordingly report it as a lead, not a result. Even taken at face value, though, it points the same way as the null it accompanies: a stamped patch could, at best, perturb the interior into nucleating *later*, and nothing we planted made the market crash sooner.

The positive prediction is a dose-response: if suppressed buy orders are the brake on the front, the front should slow in proportion to how hard the brake is pressed, whatever the nominal cash endowment that produced the pressing (Figure 11). Across the five budget doses at $f = 0.20$ (cash 250–5,000; $n = 97$ runs included, with 3 excluded by the window rule), the front velocity fell as the *absolute* suppression rate rose, at a pooled Spearman $\rho = -0.372$ with permutation $p = 0.0003$, and the relationship was not a between-dose confound, for it held within every dose separately, with ρ between -0.38 and -0.52, and it survived the removal of any single dose entirely (leave-one-dose-out $p \leq 0.0014$). We state the shape of this result as carefully as its strength. The median suppression load varied little across the ladder (517–595 orders per step), and the median velocity was U-shaped in cash, highest at the tightest dose, lowest in the middle of the ladder, and rising again at the loosest, so the claim the data support is the pooled, and within-dose, negative velocity–suppression relationship, not a monotone-in-cash dose ladder: initial cash was a weak proxy for the suppression a given front actually experienced, and it was the realized suppression, wherever it landed, that tracked the slowing of the front.

Together, these experiments located the budget effect after onset. The winning domain was bulk-born, and planted nine-cell patches did not found it. The onset-time fits bounded any cash-associated difference to within a factor of two, while front velocity covaried with suppressed buy orders.

We next used the budget toggle as a paired intervention to test how a cash shock reaches the front through the price-mediated boundary (Figure 12). Scaling cash $\times 10$ in the tight freeze configuration (cash 500) *accelerated* the front in 14 of 19 seeds (sign test $p = 0.064$; we report the count and the p-value as they stand), since the median front velocity rose from 1.39 to 3.55×10^{-4} grid fractions per step, the median front extent at $t_{\text{on}} + 400$ from 0.112 to 0.192, and the suppressed buy orders fell by 29% (the post-toggle window rule admitted 19 of the 20 seeds in every arm). Handing the modeled population's cash tenfold just before detected onset made the sell domain spread *faster*. This is a causal result within the paired simulation, not evidence about real-market liquidity policy.

An acceleration under injection runs opposite to the naive mechanism, which is what makes it diagnostic. Injected cash cannot feed the front directly, since the front is a *sell* wave, sells need shares rather than cash, and the interior CA never sees the books at all, because budgets suppress order flow and never sentiment, so the only conduit by which a cash shock can reach interior sentiment is the loop's own: executed flow moves the price, the price crosses the edge bands, and the band crossings set the state of the active-edge ring whose sentiment the boundary feeds into the grid. The edge-ring trace recorded that conduit directly (Figure 13). At the toggle, injection flipped the ring *instantly*, from the control's ~ 0.5 Buy and 0 Sell to 0 Buy and ~ 1.0 Sell: the injected cash let the buy orders the books had been suppressing execute at last, the supported price, 113 against the control's 93 at $t_{\text{on}} + 400$, broke above the upper bands, and the entire ring, contrarian in the value trader's sense of

buying below its band and selling above it, dumped at once, the flipped state decaying back over the following 100–200 steps as the EMA bands caught up with the new price level. Under tight budgets the boundary bought into the falling price and braked the front, whereas under injection it sold into the supported price and fed the front instead. The boundary switched from brake to accelerant, a sign reversal.

Withdrawal, which by symmetry ought to have braked the crash, did nothing of the kind at either dose (Figure 12), and the two doses matter because together they rule out mere under-dosing. The mild cut ($\times 0.1$, applied in the loose collapse configuration, cash $5{,}000$) left the front statistically untouched (sign-flip permutation $p = 0.67$) and barely disturbed the suppression load, suppressed buys moving $+7\%$. The harsh cut ($\times 0.01$) *bound*, in that it dropped every cell to ~ 50 in cash, below the standing price, so buys died outright and the suppressed-buy load jumped from 524 to 680 orders per step, and yet the median front velocity moved only from 2.23 to 2.38×10^{-4} grid fractions per step (11 of 19 seeds in the predicted direction; sign-flip permutation $p = 0.29$), while the price barely registered the cut at all, 119.5 against the control's 123.6 at $t_{\text{on}} + 400$. The null has an explanation in the loose configuration (Figure 13): the control ring already sat near 0.46 Buy after onset, a brake the crash itself had engaged, the sell wave having pressed the price beneath nearly half the ring's bands, and the cash cut could only strengthen it transiently, lifting the ring's Buy fraction to ~ 0.72–0.74 before it relaxed back to the control level within ~ 200 steps: the same sign, a modest margin, and a decay.

Why the asymmetry? Two features of the model's economy combine to produce it. The first is a matter of sign and ceiling: only injection can *reverse* the boundary's state, because propping the price above the upper bands flips the ring from buying to selling, whereas withdrawal can move the ring only *further* in the direction it already points, and only transiently, an engaged brake pressed harder rather than a new force. The second follows from the execution convention: every cell the front recruits can sell up to its 10-share endowment to the unmodeled market maker, receiving roughly $1{,}000$ in cash at prices near $P \approx 100$–120. A one-time cash reduction can therefore be offset by later sales to the external counterparty. Together these mechanisms produce a cash-intervention asymmetry within Model B: injections prop the price, flip the contrarian boundary to selling, and accelerate the sentiment crash, while withdrawals are near-inert. What this asymmetry means for the design of liquidity interventions requires a market model with explicit counterparties and calibrated institutions, which is outside the scope of this study.

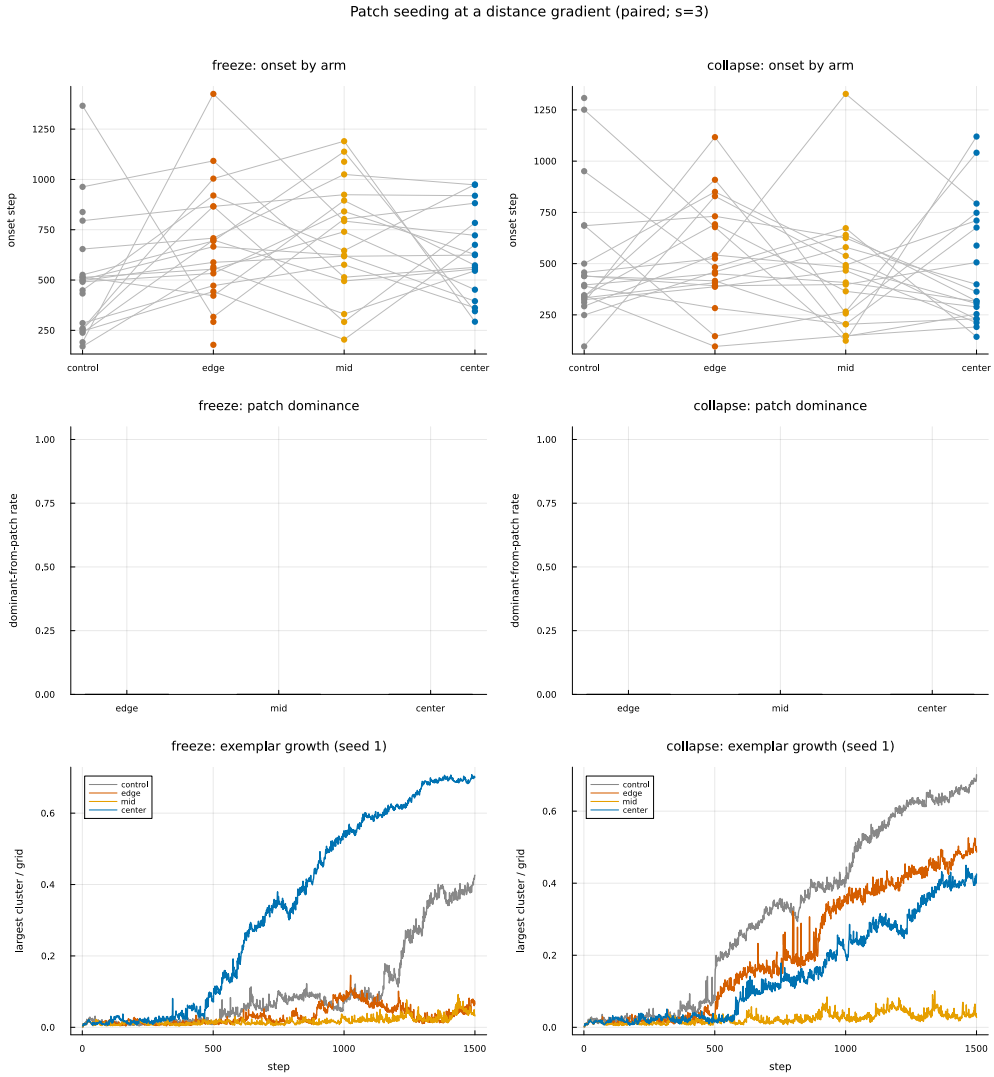


Figure 10: Stamped-patch seeding outcomes across placements and configurations. (Top row) Onset step per seed under each arm (an unstamped control and a 3×3 sell patch stamped at edge, mid, and center placements) for the freeze (left) and collapse (right) configurations, with gray lines pairing the same seed across arms; the pairings cross in both directions and no placement advances onset. The apparent delay under the freeze mid placement (median paired onset advance -257 steps, sign-flip $p = 0.0238$) is exploratory, uncorrected across the six-test family and failing the Bonferroni threshold of 0.0083, a lead rather than a result. (Middle row) Dominance rate per placement: the fraction of runs in which the cluster standing largest at onset descends from the stamped patch is $0/20$ in every one of the six configuration-by-placement cells. (Bottom row) Largest-cluster growth trajectories for a single exemplar seed under each arm: which arm reaches domain scale first differs between the two configurations rather than tracking the stamped placement, and the domain that grows is in every run one founded away from the patch.

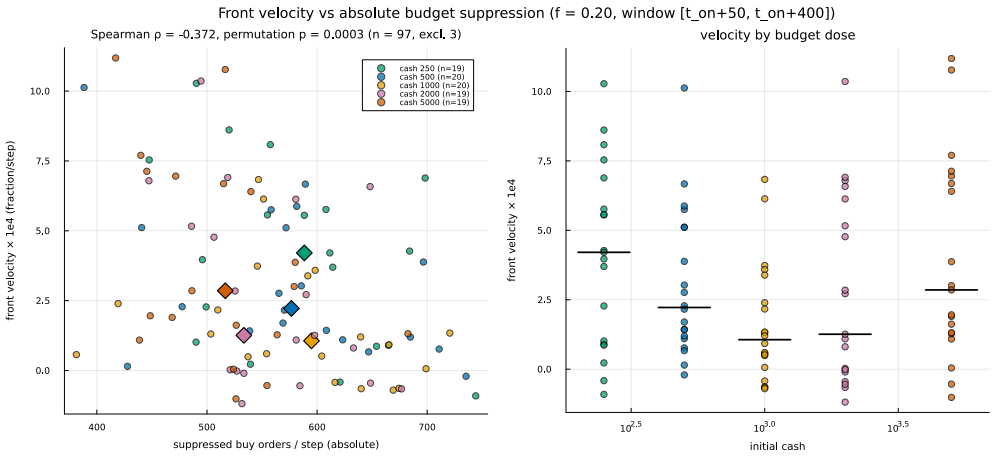


Figure 11: Front velocity against absolute budget suppression at $f = 0.20$. (Left) For the 97 runs included across the five cash doses (250–5,000; 3 runs excluded by the window rule), the front velocity, the OLS slope of the front extent over the onset-aligned window $[t_{\text{on}} + 50,\ t_{\text{on}} + 400]$, is plotted against the mean absolute number of suppressed buy orders per step over the same window, colored by dose with the per-cash medians marked as diamonds; velocity falls as suppression rises, Spearman $\rho = -0.372$, permutation $p = 0.0003$. (Right) The same velocities arranged by initial cash on a logarithmic axis, with a bar at each per-dose median: the medians are U-shaped in cash, highest at cash 250, lowest near 1,000–2,000, and rising again at 5,000, while the per-dose median suppression spans only 517–595 orders per step, so the supported claim is the pooled and within-dose velocity–suppression relationship, not a monotone ordering by dose.

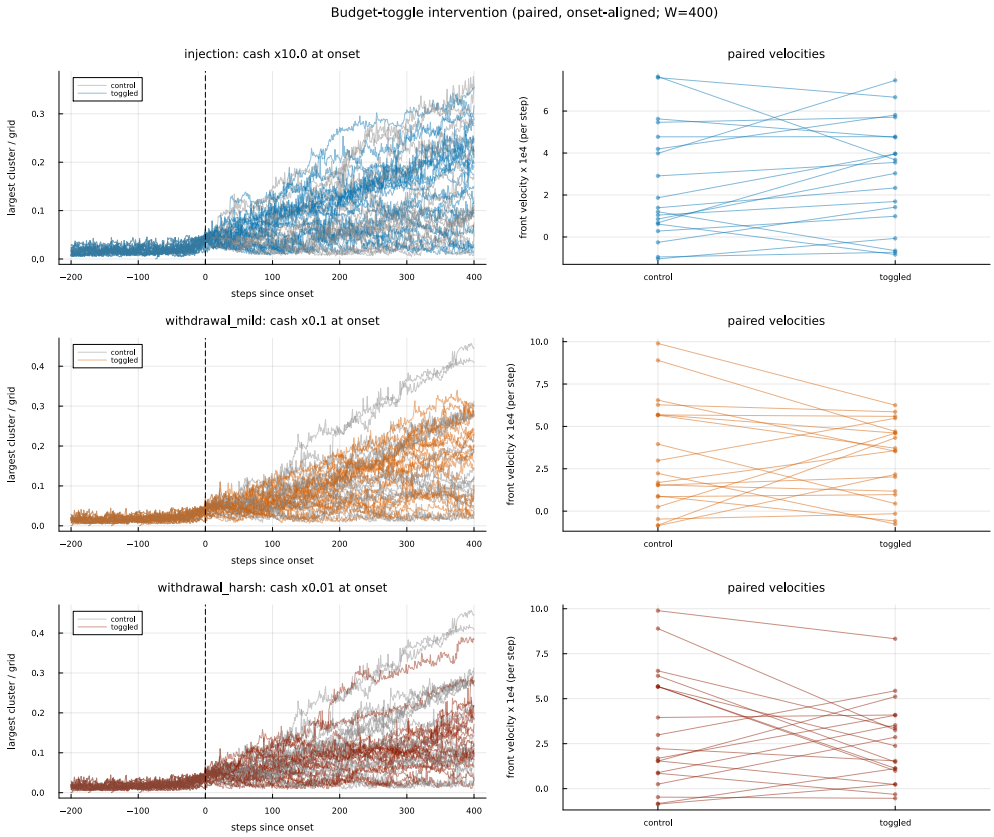


Figure 12: Budget-toggle intervention: paired, onset-aligned counterfactuals at $f = 0.20$ with a $W = 400$ post-toggle window. (Left column) Largest sell cluster as a fraction of the grid against steps since onset, for the three arms (injection at cash $\times 10$ on the tight cash-500 base, and the mild ($\times 0.1$) and harsh ($\times 0.01$) withdrawals, both on the loose cash-5,000 base), with each toggled trajectory (color) drawn over its control (gray) for the 19 included seed pairs per arm; the dashed line marks the toggle at t_{on}, and the pre-toggle trajectories coincide by construction. (Right column) Paired front velocities over the window, one line per seed pair, control against toggled: injection accelerates the front in 14 of 19 pairs (sign test $p = 0.064$), while withdrawal is null at both doses (sign-flip permutation $p = 0.67$ mild, $p = 0.29$ harsh).

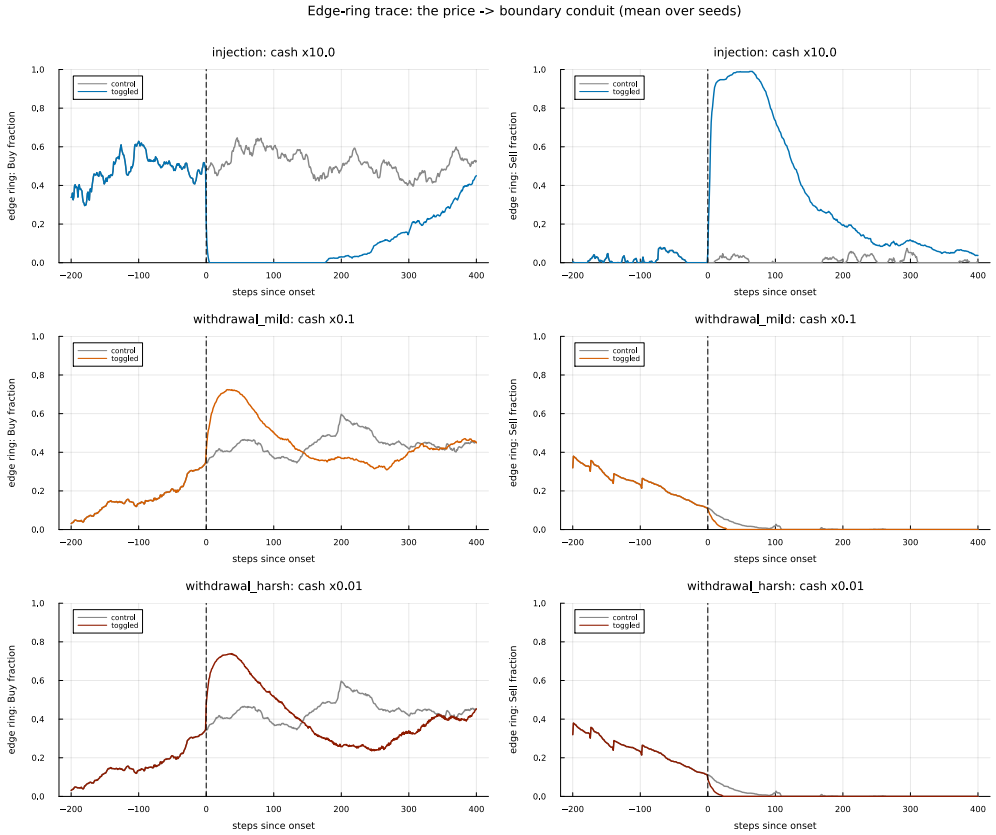


Figure 13: Edge-ring trace of the price-to-boundary conduit around the toggle. Each panel draws an edge-ring state fraction, averaged over the 19 included seed pairs and aligned at the per-seed onset, for the toggled arm (color) against its control (gray), with the dashed line marking the toggle at t_{on}: the ring's Buy fraction occupies the left column and its Sell fraction the right, one row per arm (injection $\times 10$; mild withdrawal $\times 0.1$; harsh withdrawal $\times 0.01$), and the pre-toggle traces coincide by construction. (Top row) Injection flips the ring at the toggle, from ~ 0.5 Buy and 0 Sell in the control to 0 Buy and ~ 1.0 Sell, and the flipped state decays back over the following 100–200 steps as the EMA bands catch up with the supported price. (Middle and bottom rows) Withdrawal transiently deepens an already-engaged brake: the control ring sits near 0.46 Buy after onset, either cut lifts the Buy fraction to ~ 0.72–0.74 before the transient excess dissipates by ~ 200 steps, and the Sell fraction, which in this loose configuration declines from ~ 0.37 well before onset to zero just after, remains indistinguishable between control and toggled throughout.

5 Discussion

The lattice produced a sharp finite-horizon crossover between $f = 0.14$ and $f = 0.15$, above which localized sell domains grew as spatially coherent fronts, and neither the mean-field map nor the pair closure reproduced that behavior. Against that baseline, the paired intervention identified a model-specific feedback in which added cash lets suppressed buys execute, raises the price, and moves the contrarian boundary from Buy to Sell, so that the front accelerated in 14 of 19 included pairs (two-sided sign test $p = 0.064$) while withdrawals produced no detectable change at either dose. Later sales to the unmodeled market maker replenish agent cash, which helps explain why a one-time reduction has little effect. This mechanism is structurally related to feedback between funding and market liquidity [38], but it is not a calibrated account of the 2010 Flash Crash or a policy model [2], and in particular the agents do not form expectations, so the mechanism is distinct from ex-ante collective moral hazard [40]. That mechanism also sharpens the study's central separation: the experiments distinguished onset from growth, since birth sites were consistent with uniform interior placement and the onset-time fits bounded any cash-associated difference to within a factor of two, whereas post-onset growth covaried with suppressed buy orders and responded to the paired cash shock through the price-to-boundary loop. That separation is a model-generated hypothesis, not policy advice, and circuit breakers and other real interventions involve strategic responses absent here [39].

The limitations of the study bound how far these readings travel. Model A uses one boundary-forcing series, QQQ one-minute bars from January 2 through March 28, 2024, and does not constitute empirical price validation. Robustness across instruments and market regimes remains to be established. The Model B nucleation campaigns all ran on a single 64×64 grid trading a single asset against an infinite-inventory market maker, and while the weak system-size dependence of the Model A transition makes grid-size robustness plausible, the budget results have not been rechecked at other scales and the model is silent on cross-asset contagion. The agents themselves are deterministic, since the CA rules incorporate no stochastic noise, no learning, and no adaptive strategy switching, features of real markets whose inclusion would enrich the dynamics at the cost of breaking the enumerability of the rule space [10, 11]. On the theory side, the analytical hierarchy treats the contrarian disorder as annealed, averaging over rule types in the mean-field and pair maps, whereas the simulations fixed each contrarian's placement for the duration of a run, quenched disorder in the proper sense, and a calculation that carried that quenching through would sharpen the prediction of f_0 and connect the critical nucleus size to the combinatorics of contrarian clustering. And the Model B statistics carry flags that interpretation must not launder: the monotone-in-cash onset rate near the cliff and the freeze mid-patch onset delay are exploratory, uncorrected observations; the injection effect stands at 14 of 19 seeds with a sign-test $p = 0.064$, and it is the paired design, not the count, that carries the causal reading; and the onset-fit comparison is power-bounded. Overlapping confidence intervals do not establish parameter equality or exclude a small effect.

A last question applies to any two-rule study: how much hangs on the two rules? The totalistic framework makes the question answerable, which models with continuous behavioral parameters cannot match because their parameter space is not enumerable, since the majority/contrarian pair studied here, two rules that differ at a single neighborhood sum, is one point in a finite and enumerable rule space. Whether the same sequence, a metastable attractor, a disorder-seeded nucleus, and budget-gated growth, recurs across other rule pairings, and whether the intervention asymmetry survives boundary rules other than the contrarian band, are open questions. A screen over rule pairs that differ at a single neighborhood sum, the same minimal perturbation studied here, would test the first directly, and swapping the band classifier for a momentum or threshold rule at the boundary would test the second.

6 Conclusion

We studied a deterministic three-state cellular automaton with two totalistic rules on a two-dimensional lattice, and in Model A a fixed-seed sweep showed a sharp finite-horizon crossover between $f = 0.14$ and $f = 0.15$, above which localized sell domains grew through coherent fronts that neither mean-field nor the tested pair closure reproduced. Onset times were well summarized by an inverse-distance empirical curve, a fit consistent with a nucleation interpretation but neither a derived Kramers law

nor evidence of an infinite thermodynamic barrier, and the Potts-inspired alignment statistic similarly measures spatial order without generating the CA dynamics.

Model B added endogenous price feedback and finite cash and share inventories. Birth sites were consistent with uniform interior placement, and the present experiments bounded any cash-associated difference in onset time to within a factor of two, whereas budgets were associated with post-onset growth. In paired simulations, a tenfold cash injection just before the control's detected onset accelerated the front in 14 of 19 included pairs by changing the price-mediated boundary state. This result is specific to the model's unmodeled infinite-inventory market maker and uniform cash-shock intervention.

The finite set of $3^9 = 19{,}683$ totalistic rules makes systematic extension possible. Future work should recheck the budget results at other grid sizes, compare alternative onset-time models under censoring, include explicit counterparties, and test whether the same spatial mechanism appears for other rule pairs.

References

[1] Didier Sornette. *Why Stock Markets Crash: Critical Events in Complex Financial Systems*. Princeton University Press, 2017.

[2] Andrei Kirilenko, Albert S. Kyle, Mehrdad Samadi, and Tugkan Tuzun. The flash crash: High-frequency trading in an electronic market. *The Journal of Finance*, 72(3):967–998, 2017.

[3] Rama Cont. Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001.

[4] Benoit Mandelbrot. The variation of certain speculative prices. *The Journal of Business*, 36(4): 394–419, 1963.

[5] Stefano Battiston, J Doyne Farmer, Andreas Flache, Diego Garlaschelli, Andrew G Haldane, Hans Heesterbeek, Cars Hommes, Carlo Jaeger, Robert May, and Marten Scheffer. Complexity theory and financial regulation. *Science*, 351(6275):818–819, 2016.

[6] Thomas Lux and Michele Marchesi. Scaling and criticality in a stochastic multi-agent model of a financial market. *Nature*, 397:498–500, 1999.

[7] Per Bak, Maya Paczuski, and Martin Shubik. Price variations in a stock market with many agents. *Physica A: Statistical Mechanics and its Applications*, 246:430–453, 1997.

[8] Rama Cont and Jean-Philippe Bouchaud. Herd behavior and aggregate fluctuations in financial markets. *Macroeconomic Dynamics*, 4(2):170–196, 2000.

[9] D. Challet and Y.-C. Zhang. Emergence of cooperation and organization in an evolutionary game. *Physica A: Statistical Mechanics and its Applications*, 246:407–418, 1997.

[10] J. Doyne Farmer and Shareen Joshi. The price dynamics of common trading strategies. *Journal of Economic Behavior & Organization*, 49(2):149–171, 2002.

[11] Cars H Hommes. Heterogeneous agent models in economics and finance. *Handbook of Computational Economics*, 2:1109–1186, 2006.

[12] Stefan Bornholdt. Expectation bubbles in a spin model of markets: intermittency from frustration across scales. *International Journal of Modern Physics C*, 12(5):667–674, 2001.

[13] A. Krawiecki, J. A. Hołyst, and D. Helbing. Volatility clustering and scaling for financial time series due to attractor bubbling. *Physical Review Letters*, 89:158701, 2002.

[14] Katarzyna Sznajd-Weron and Rafal Weron. A simple model of price formation. *International Journal of Modern Physics C*, 13(1):115–123, 2002.

[15] Taisei Kaizoji. Speculative bubbles and crashes in stock markets: an interacting-agent model of speculative activity. *Physica A: Statistical Mechanics and its Applications*, 287:493–506, 2000.

[16] Yu-Cheng Wei, Shao-Hua Wang, and Fang-Mei Yang. A cellular automaton model for the securities market. *Physica A: Statistical Mechanics and its Applications*, 325:285–290, 2003.

[17] Guoping Qiu, Drona Kandhai, and Peter M. A. Sloot. A cellular automata model for the abnormal fluctuations in stock markets. *Advances in Complex Systems*, 10:241–255, 2007.

[18] Guang Yang, Yun Chen, and Ji-Ping Huang. Financial market simulation based on cellular automaton. *Physica A: Statistical Mechanics and its Applications*, 521:586–594, 2019.

[19] Stephen Wolfram. *A New Kind of Science*. Wolfram Media, 2002.

[20] Eugene F. Fama. Efficient capital markets: a review of theory and empirical work. *The Journal of Finance*, 25(2):383–417, 1970.

[21] G. William Schwert. Why does stock market volatility change over time? *The Journal of Finance*, 44(5):1115–1153, 1989.

[22] Jean-Philippe Bouchaud and Marc Potters. An apparent stylized fact about the dynamics of financial markets. *Quantitative Finance*, 1:105–112, 2001.

[23] Blake LeBaron. Agent-based computational finance: suggested readings and early research. *Journal of Economic Dynamics and Control*, 24:679–702, 2000.

[24] J. J. Chen, L. Tan, and B. Zheng. Agent-based model with multi-level herding for complex financial systems. *Scientific Reports*, 5:8399, 2015.

[25] Georges Harras and Didier Sornette. How to grow a bubble: a model of myopic adapting agents. *Journal of Economic Behavior & Organization*, 80:137–152, 2012.

[26] Blake LeBaron. Agent-based computational finance. In *Handbook of Computational Economics*, volume 2, pages 1187–1233. Elsevier, 2006.

[27] W. Brian Arthur. Foundations of complexity economics. *Nature Reviews Physics*, 3:136–145, 2021.

[28] J. Doyne Farmer. Quantitative agent-based models: a promising alternative for macroeconomics. *Oxford Review of Economic Policy*, 41(2):571–590, 2025.

[29] Dietrich Stauffer. Generalization of a two-dimensional Ising-like model to arbitrary dimension. *Computing in Science & Engineering*, 2(3):69–72, 2000.

[30] Debashish Chowdhury and Dietrich Stauffer. A generalized spin model of financial markets. *The European Physical Journal B*, 8:477–482, 1999.

[31] Per Bak, Chao Tang, and Kurt Wiesenfeld. Self-organized criticality: an explanation of $1/f$ noise. *Physical Review Letters*, 59:381–384, 1987.

[32] Per Bak and Maya Paczuski. Complexity, contingency, and criticality. *Proceedings of the National Academy of Sciences*, 92:6689–6696, 1995.

[33] Xavier Gabaix, Parameswaran Gopikrishnan, Vasiliki Plerou, and H Eugene Stanley. A theory of power-law distributions in financial market fluctuations. *Nature*, 423:267–270, 2003.

[34] Marten Scheffer, Jordi Bascompte, William A Brock, Victor Brovkin, Stephen R Carpenter, Vasilis Dakos, Hermann Held, Egbert H Van Nes, Max Rietkerk, and George Sugihara. Early-warning signals for critical transitions. *Nature*, 461:53–59, 2009.

[35] J Doyne Farmer, Paolo Patelli, and Ilija I Zovko. The predictive power of zero intelligence in financial markets. *Proceedings of the National Academy of Sciences*, 102(6):2254–2259, 2005.

[36] Jean-Philippe Bouchaud. Crises and collective socio-economic phenomena: Simple models and challenges. *Journal of Statistical Physics*, 151(3–4):567–606, 2013.

[37] Didier Sornette, Fabrice Deschatres, Thomas Gilbert, and Yann Ageon. Endogenous versus exogenous shocks in complex networks: an empirical test using book sale rankings. *Physical Review Letters*, 93:228701, 2004.

[38] Markus K. Brunnermeier and Lasse Heje Pedersen. Market liquidity and funding liquidity. *The Review of Financial Studies*, 22(6):2201–2238, 2009.

[39] Avanidhar Subrahmanyam. Circuit breakers and market volatility: A theoretical perspective. *The Journal of Finance*, 49(1):237–254, 1994.

[40] Emmanuel Farhi and Jean Tirole. Collective moral hazard, maturity mismatch, and systemic bailouts. *American Economic Review*, 102(1):60–93, 2012.

[41] Stephen Wolfram. Universality and complexity in cellular automata. *Physica D: Nonlinear Phenomena*, 10:1–35, 1984.

[42] Norman H. Packard and Stephen Wolfram. Two-dimensional cellular automata. *Journal of Statistical Physics*, 38:901–946, 1985.

[43] Albert S. Kyle. Continuous auctions and insider trading. *Econometrica*, 53(6):1315–1335, 1985.

[44] Hersh Shefrin and Meir Statman. The disposition to sell winners too early and ride losers too long: Theory and evidence. *The Journal of Finance*, 40(3):777–790, 1985.

[45] Terrance Odean. Are investors reluctant to realize their losses? *The Journal of Finance*, 53(5): 1775–1798, 1998.

[46] Howard A. Gutowitz, Jonathan D. Victor, and Bruce W. Knight. Local structure theory for cellular automata. *Physica D: Nonlinear Phenomena*, 28(1–2):18–48, 1987.

[47] Hendrik A. Kramers. Brownian motion in a field of force and the diffusion model of chemical reactions. *Physica*, 7(4):284–304, 1940.

[48] Peter Hänggi, Peter Talkner, and Michal Borkovec. Reaction-rate theory: Fifty years after Kramers. *Reviews of Modern Physics*, 62(2):251–341, 1990.

A Pair-Approximation Derivation

This appendix records the pair-approximation calculation in the form in which it was implemented. The quenched rule mixture enters in its annealed form: Rules 19280 and 19286 share every entry of their totalistic tables except at neighborhood sum $c = 1$, where the majority rule outputs Buy and the contrarian rule outputs Sell, so a site whose rule assignment is drawn afresh at minority fraction f follows the effective stochastic rule

$$R_f(c \to b) \;=\; (1-f)\,\mathbf{1}[\mathcal{R}_{19280}(c) = b] \;+\; f\,\mathbf{1}[\mathcal{R}_{19286}(c) = b]\,, \tag{15}$$

deterministic at every $c \neq 1$ and splitting $c = 1$ between Buy, with probability $1 - f$, and Sell, with probability f.

The state of the approximation is the symmetric matrix of nearest-neighbor pair densities ρ_{ab}, $a, b \in \{0, 1, 2\}$, normalized by $\sum_{ab} \rho_{ab} = 1$, with site marginals $\rho_a = \sum_b \rho_{ab}$ and conditional neighbor distribution $q(s \mid a) = \rho_{as}/\rho_a$. The closure is the local-structure ansatz of Gutowitz, Victor, and Knight [46]: given a focal site in state a, its three neighbors *other* than the paired partner are drawn independently from $q(\cdot \mid a)$. Their sum m therefore follows the trinomial convolution

$$\Pr(m = k \mid a) \;=\; \sum_{\substack{n_0+n_1+n_2=3 \\ n_1+2n_2=k}} \binom{3}{n_0,\, n_1,\, n_2}\, q(0 \mid a)^{n_0}\, q(1 \mid a)^{n_1}\, q(2 \mid a)^{n_2}, \qquad k = 0, \ldots, 6, \tag{16}$$

and the focal site's full neighborhood sum is the partner's state plus that draw, $c = a' + m$, which never exceeds 8. Averaging the effective rule over the trinomial weights gives the transition kernel

$$T_b(a, a') \;=\; \sum_{k=0}^{6} \Pr(m = k \mid a)\; R_f(a' + k \to b), \tag{17}$$

the probability that a site in state a with paired neighbor in state a' lands in state b after one synchronous step. Substituting this kernel for both members of a pair (each site conditioning on its own partner) yields the synchronous pair update $\rho'_{bb'} = \sum_{a,a'} \rho_{aa'}\, T_b(a, a')\, T_{b'}(a', a)$ quoted as Equation (13) of the main text. The implementation symmetrizes and renormalizes ρ' after every step to guard against numerical drift.

We iterated the map to its fixed points, declaring convergence when the largest componentwise change $\max_{ab} |\rho'_{ab} - \rho_{ab}|$ fell below 10^{-12}, and every trajectory reported here converged within a few hundred iterations. Stability of the uncorrelated fixed point $\rho^*_{ab} = \rho^*_a \rho^*_b$, the product state built from the mean-field marginals, was assessed through the Jacobian of the map restricted to the six independent components $(\rho_{00}, \rho_{01}, \rho_{02}, \rho_{11}, \rho_{12}, \rho_{22})$ that survive symmetry and normalization, computed by finite differences with step $\delta = 10^{-7}$, each perturbed matrix projected back onto the probability simplex before the map was applied. The spectral radius sat strictly below unity at every value of f we sampled (twelve values from 0 to 0.5), declining monotonically from 0.95 at $f = 0$ to 0.71 at $f = 0.5$ (Figure 5B).

The search for correlated fixed points used initial conditions

$$\rho_{ab} \;=\; p_a p_b + \varepsilon\,(\delta_{ab}\, p_a - p_a p_b)\,, \tag{18}$$

which interpolate between the uncorrelated product state at $\varepsilon = 0$ and perfect same-state correlation at $\varepsilon = 1$. We ran $\varepsilon \in \{0.1, 0.3, 0.5\}$ on buy-biased marginals $(0.50, 0.30, 0.20)$, $\varepsilon \in \{0.1, 0.3\}$ on sell-biased marginals $(0.20, 0.10, 0.70)$, and $\varepsilon = 0.5$ on $(0.10, 0.05, 0.85)$, together with an uncorrelated uniform control, at thirteen values of f between 0 and 0.5. Every trajectory converged either to the uncorrelated buy-dominated fixed point or to the all-Sell absorbing state, with the spatial mutual information $I = 2S_1 - S_2$ decaying to zero in all cases: the pair approximation supports no correlated fixed point, which is the negative result on which the nucleation argument rests.

B Model B Stage 2 Factorial

The Stage 2 factorial sweep crossed six factors in full, giving the design from which every Model B summary statistic in this study is drawn (Table 3).

Table 3: Model B Stage 2 factorial design. The six factors are crossed in full: $5 \times 4 \times 3 \times 2 \times 3 \times 3 = 1{,}080$ configurations, each run at 10 seeds for 2,000 steps (10,800 runs).

Factor	Levels	Count
Minority fraction f	0, 0.1, 0.15, 0.2, 0.3	5
Price impact λ	0.005, 0.01, 0.02, 0.05	4
Initial cash	500, 1,000, 5,000	3
Disposition rule	off, on	2
Band spread δ	0, 0.5, 0.9	3
Band half-width k	0.5, 1, 2	3

Every run was assigned to one of five candidate regimes by a classifier evaluated in strict order, each test applying only to runs that fail all earlier ones. A run with zero return standard deviation is *inert*. Failing that, a run is *crash-prone* when its crash frequency, the fraction of steps at which the grid's Sell fraction stands at or above 0.5, reaches at least 0.01. Failing that, a run is *oscillatory* when the net order imbalance φ changes sign at least 20 times per 2,000 steps. That threshold is a rate, one flip per 100 steps, stated here against the sweep's horizon so that it transfers to other run lengths. Failing that, a run whose total price range stays below $0.10\,P_0$ is a *plateau*, and any remaining run is a *drift*. Of the 10,800 runs, 6,589 (61%) were oscillatory and 4,211 (39%) crash-prone; no run was ever inert, plateau, or drift. The sell-dominated (crash-prone) regime appeared abruptly between $f = 0.15$ and $f = 0.20$ at every price impact λ in the design.

Grouping runs by their parameter configuration split each summary metric's variance into a between-configuration share, attributable to the swept parameters jointly, and a within-configuration share, attributable to the disorder seed. The parameters accounted for 94–97% of the variance across the four summary metrics, at between-configuration shares of 0.946 for crash frequency, 0.952 for return kurtosis, 0.959 for return standard deviation, and 0.974 for the φ sign-flip count, leaving the seed at most 5.4%. Across matched design cells, enabling the disposition rule increased return standard deviation by a factor of 2.40 on average.

The crash-prone class divided into the modal freeze and a collapsing tail, the decomposition on which the transition analysis builds: the median crash-prone run had a maximum drawdown of about 5% and a price minimum of 100.00, so half the class never traded below the initial quote at all, while 897 of the 4,211 runs drove the price below 90 and 451 below 80. The deepest collapses, drawdowns of 40–69% arriving after bubbles of 2.5–2.7× the initial price, were concentrated at $f = 0.2$, $\lambda = 0.05$, cash 5,000, disposition off, the loosest budget regime in the design.

C Escape-Time Censoring and Survival Detail

The escape-time campaign ran each seed to a horizon of 2,000 steps, and any run that has not nucleated by then is right-censored: its onset time $\tau = t_{\text{on}}$ is treated as exceeding the horizon, and the run enters the Kaplan–Meier estimator carrying that censoring mark, since censored runs are never dropped. The resulting counts are recorded for every (f, cash) cell of the campaign (Table 4); rows at 20 seeds belong to the pilot pass and rows at 40 seeds to the refinement pass. The companion velocity–suppression dose-response added two further cells at $f = 0.20$, cash 250 and cash 2,000, each with 20 seeds, one censored run, and an onset rate of 0.95. Those cells feed the dose-response analysis, not the barrier fits.

Because censoring occurs only at the horizon, the Kaplan–Meier survival curve of a cell steps down at each observed onset and its median is finite exactly when the onset rate reaches one half. The fits therefore rest on five points at cash 500, at $f \in \{0.177, 0.18, 0.20, 0.25, 0.30\}$, and on six points at cash 1,000 and 5,000, which add $f = 0.174$ (Table 2). Cash 500 has one fewer because its onset rate at $f = 0.174$ was 0.45, just below the one-half a finite median requires. It was this gap that prompted extending the refinement grid to $f = 0.177$, so that the tightest budget still contributed a point hard against the cliff. The reported confidence intervals (Table 2, Figure 9) come from 2,000 bootstrap resamples drawn over seeds within each f: each resample recomputes the Kaplan–Meier medians and refits the barrier law, and the intervals are the 2.5th and 97.5th percentiles of the resampled parameters.

Table 4: Escape-time campaign: run count n, censored count c, and onset rate $(n - c)/n$ for every (f, cash) cell. Rows at $n = 20$ come from the pilot pass, rows at $n = 40$ from the refinement pass. A Kaplan–Meier median is finite only where the onset rate reaches 0.5. The qualifying cells are exactly the $n_f = 5/6/6$ fit points of Table 2.

	Cash 500			Cash 1,000			Cash 5,000		
f	n	c	rate	n	c	rate	n	c	rate
0.12	20	20	0.00	20	20	0.00	20	20	0.00
0.15	20	18	0.10	20	18	0.10	20	18	0.10
0.156	40	37	0.07	40	38	0.05	40	36	0.10
0.162	40	31	0.23	40	33	0.17	40	30	0.25
0.168	40	24	0.40	40	31	0.23	40	22	0.45
0.174	40	22	0.45	40	15	0.62	40	12	0.70
0.177	40	16	0.60	40	14	0.65	40	13	0.68
0.18	20	3	0.85	20	4	0.80	20	4	0.80
0.20	20	0	1.00	20	0	1.00	20	1	0.95
0.25	20	0	1.00	20	0	1.00	20	0	1.00
0.30	20	0	1.00	20	0	1.00	20	0	1.00

Finally, the onset-rate observation near the cliff is visible directly in Table 4, and it carries the same label as in the main text: exploratory, reported uncorrected for multiple comparisons, and weighed accordingly. At $f = 0.174$ the onset rate rose monotonically in cash, 0.45, 0.62, 0.70; at $f = 0.177$ the Kaplan–Meier medians fell in the same order, $1{,}598$, $1{,}427$, and $1{,}208$ steps.