Local rule activation and dealer capacity in a cellular automaton market*

Jeffrey D. Varner
R.F. Smith School of Chemical and Biomolecular Engineering
Cornell University, Ithaca, NY 14850
jdv27@cornell.edu

Abstract

A change in how traders respond to their neighbors can spread through a market, but its price consequence depends on whether the resulting orders execute. We developed a three-state cellular automaton market to examine this connection with explicit trader holdings and dealer capacity. Paired trajectories shared the same realized market state until a fraction of interior traders changed one entry of their local response rule. Persistent activation produced Sell majorities above a window of activation fractions. A one-time Sell pulse at the same sites produced no majority crossings at the tested fractions, while several other single-entry rule changes did. Holding boundary traders inactive preserved the transition, showing that local interactions could sustain collective selling without price feedback. Larger lattices had a sharper response on refined fraction grids, although the midpoint continued to shift with the observation horizon. Dealer capacity changed majority probability non-monotonically when prices fed back through boundary traders. Greater capacity increased drawdown at the main activation fraction across the tested price-impact coefficients. Cash and share interventions traced the price consequence to order eligibility and execution: trader holdings limited available flow, dealer resources bounded unmatched fills, and executed imbalance changed the quote. Similar Sell majorities could therefore accompany price declines that differed by more than an order of magnitude. An unlimited dealer allowed all eligible selling to execute, while finite capacity could limit the decline despite persistent Sell-state dominance.

*Repository: https://github.com/varnerlab/Wolfram-markets-abm-paper.git.

Preprint.

1 Introduction

During the October 1987 market decline, portfolio insurance concentrated selling across participating traders [1, 2]. The forced unwind of leveraged funds transmitted losses across quantitative equity strategies in August 2007 [3]. Such episodes motivate a common modeling question: how can a change in the behavior of some traders become a market-wide event? Agent-based models address collective behavior through imitation, strategy adaptation, and responses to a shared price [4–7]. Herding models describe how individual decisions produce collective activity [8–10], while lattice models restrict influence to neighboring traders [11–13]. These interactions can generate large fluctuations and clustered volatility, two features of financial returns [14]. However, reproducing aggregate fluctuations does not by itself explain how an established pattern of trading responds when part of the population changes its behavior. The size of the changed group matters, but so does the nature of the change. A temporary increase in selling and a persistent change in how traders respond to their neighbors need not produce the same collective outcome.

A cellular automaton makes that distinction concrete, because it gives each trader an explicit local decision rule. Traders occupy Buy, Hold, or Sell states on a square lattice and update from their four nearest neighbors, and a totalistic rule assigns the next state from the sum of those neighbor states [15–17]. Local interactions of this kind have been used to generate market fluctuations and regime changes [12, 18–20]. The lookup table makes a persistent behavioral change precise: one response can change while the remaining entries stay fixed. A one-time Sell pulse changes only the current configuration, whereas changing the rule changes the response whenever the affected neighborhood recurs, and the altered response remains in place even after a trader leaves the Sell state. Following matched trajectories from the same realized lattice allows the effect of that persistent change to be measured against continued baseline behavior [21, 22]. The question is then how many traders must adopt the changed rule before local selling grows into a Sell majority, and whether that majority can develop without feedback from the market price.

Whether a Sell majority forms is one question. What it does to the price is a second question, and the answer depends on the resources available to execute the resulting orders. A trader needs shares to sell or cash to buy, and an eligible order needs a counterparty. Dealer inventory and capital therefore constrain the unmatched flow that can execute [23–27]. Price-impact models connect that executed flow to changes in the quote [28–31]. Agent-based markets have incorporated finite endowments, order books, leverage limits, and dealers with constrained balance sheets [32–35]. For a lattice market, separating desired trades from executed orders allows a Sell-dominated state to coexist with a limited price decline. The execution constraint also matters when traders respond to the price, because restricting unmatched flow changes the quote and therefore changes the signal used in subsequent decisions. Dealer capacity can then influence both the spread of selling and its financial consequences. An unlimited dealer can absorb all eligible unmatched flow, whereas a finite dealer may reach its cash or inventory limit while many traders remain in Sell.

In this study, we used a three-state cellular automaton market to determine how a persistent change in one local response produced collective selling and how dealer capacity shaped the resulting price decline. We followed paired trajectories from a common market history and changed one rule entry at a selected fraction of interior sites, and we repeated the design under a held boundary that kept perimeter traders inactive, which removed price feedback while preserving order execution and financial accounts. Persistent activation at 20% or 25% of sites produced a Sell majority in all 50 seeds within 3,000 steps under an unlimited dealer with price feedback, whereas activation at 10% or 15% produced none (Supplementary Table S2). A matched one-time Sell pulse did not reproduce the transition, while several other single-entry rule changes did. The transition survived removal of price feedback, so the instability lived in the lattice, although its measured location depended on lattice size and observation time (Supplementary Table S3). With feedback active, dealer capacity changed majority probability non-monotonically near the activation window. Greater capacity increased drawdown at 20% activation, and separate cash and inventory interventions identified the resources that limited execution. Similar Sell-dominated lattices therefore accompanied median drawdowns from below 1% to above 25%, depending on how much of the desired selling could execute (Fig. 2). Collective selling and a large price decline were separate outcomes in this model, and the second required a dealer able to keep buying.

2 Methods

2.1 Three-state cellular automaton

We represented local trader interactions on an $L \times L$ square lattice, where interior traders responded to their immediate neighbors. Agents occupied Buy ($s = 0$), Hold ($s = 1$), or Sell ($s = 2$). Only the $(L-2)^2$ interior cells followed the cellular-automaton rule. Each interior cell updated synchronously from its four nearest neighbors, denoted by ∂i. Their state sum $c_i(t)$ defined the local input. A totalistic table $\mathcal{R}$ assigned one output $a_j = \mathcal{R}(j)$ to each possible sum. The sum represents a scalar score of selling pressure, with Hold between Buy and Sell, but discards neighborhood composition. Two Holds and two Buys receive the same response as one Sell and three Buys. Changing one table entry therefore altered the response to a specified neighbor sum while preserving the other eight responses. Neighborhoods with the same sum remained equivalent under the altered rule. The local input and the Wolfram index R encoding the nine outputs were defined by the expressions [15, 16]:

$$c_i(t) = \sum_{j \in \partial i} s_j(t) \in \{0, 1, \ldots, 8\}, \tag{1}$$

$$R = \sum_{j=0}^{8} a_j 3^j. \tag{2}$$

Rule 19280, $(222110002)_3$, supplied the baseline response. The activated Rule 19286, $(222110022)_3$, differed only at $c = 1$. A neighborhood containing three Buys and one Hold therefore returned Buy under the baseline and Sell under activation. Before each run, we selected a fraction f of interior sites uniformly without replacement as the target set for the later switch. The initial interior contained $\mathrm{round}(0.05(L-2)^2)$ Buy sites and the same number of Sell sites, with Hold elsewhere. Initial states were placed without replacement. A seed fixed the initial placements, target sites, active boundary sample, and boundary band multipliers. It also fixed the priorities used when eligible orders exceeded dealer capacity. Once these choices were made, the market evolved deterministically. Delaying activation allowed the market to develop under the baseline rule before the response changed. At the intervention, each targeted trader retained its current state and holdings but used the altered rule in subsequent updates.

2.2 Orders, trader books, and dealer clearing

A trader needed sufficient cash or shares to submit an eligible order, and that order needed a counterparty to execute. Every interior cell and active boundary cell represented a trader with cash w_i, integer inventory q_i, and average acquisition cost $\bar{c}_i$. Buy and Sell states requested one unit on their respective sides of the market; Hold requested no trade. Eligibility depended on the resources needed to settle at the standing quote P_{t-1}. A buy required $w_i \geq P_{t-1}$, while a sell required $q_i \geq 1$. The main experiments disabled the optional cost-basis restriction, which permits sales only above average acquisition cost in the manner of the disposition effect [36, 37]. An ineligible order was recorded as *agent-suppressed*. Suppression prevented execution without changing the trader's cellular-automaton state. A trader could therefore remain in Sell after exhausting its shares, or in Buy without enough cash to purchase a share. Its neighbors still responded to that state during the lattice update. We passed eligible orders to the dealer clearing procedure, which determined how many could fill against opposing trader orders and available dealer resources (Algorithm S2).

Eligible orders first cleared against opposing trader flow. With B_t eligible buys and S_t eligible sells, $\min(B_t, S_t)$ units on each side executed without using dealer resources. The dealer handled the remaining imbalance up to its cash and inventory limits. Capacity $h \geq 0$ specified net units per trader, giving total capacity $H = \mathrm{round}(hN_{\mathrm{tr}})$ for N_{tr} traders. A finite dealer began with cash $W_0 = P_0 H$ and inventory $Q_0 = H$. It could not borrow and had to satisfy $0 \leq Q_t \leq 2H$. When buys exceeded sells, executed counts were:

$$n_t^S = S_t, \qquad n_t^B = S_t + \min(B_t - S_t, Q_{t-1}). \tag{3}$$

When sells exceeded buys, execution also depended on the dealer's remaining inventory space and ability to pay for shares:

$$n_t^B = B_t, \qquad n_t^S = B_t + \min\left(S_t - B_t, 2H - Q_{t-1}, \left\lfloor \frac{W_{t-1}}{P_{t-1}} \right\rfloor\right). \tag{4}$$

Eligible orders beyond these limits were *dealer-rejected*. A counter-based score keyed by seed, step, trader identity, and side set the priority for rationing. Including the step allowed priority to change over time while keeping it reproducible across paired runs. Settlement at P_{t-1} updated trader holdings and the opposing dealer transfers:

$$Q_t = Q_{t-1} + n_t^S - n_t^B, \qquad W_t = W_{t-1} + P_{t-1}(n_t^B - n_t^S). \tag{5}$$

These transfers conserved total cash and shares across traders and dealer, except at an intentional balance-sheet intervention. The unlimited limit $h = \infty$ started from a zero dealer ledger and permitted borrowing and shorting to fill every eligible order. At $h = 0$, only opposing trader orders could clear. Both limits used the same clearing procedure (Algorithm S2).

2.3 Price feedback and the held-boundary control

Executed imbalance determined the price signal received by boundary traders. We normalized the difference between filled buys and sells by the fixed trader count N_{tr} to obtain imbalance ϕ_t. Positive imbalance raised the quote, negative imbalance lowered it, and balanced execution left it unchanged. The exponential update used price-impact coefficient $\lambda = 0.02$ and initial quote $P_0 = 100$. Active boundary traders compared each new quote with a moving center μ_t and moving absolute deviation σ_t. Each trader had a fixed multiplier $k_j \sim U[k(1-\delta), k(1+\delta)]$, with $k = 1$ and $\delta = 0.5$. Its decision band extended from $\mu_t - k_j\sigma_t$ to $\mu_t + k_j\sigma_t$. The trader entered Buy below the band, Sell above it, and Hold inside it. One quarter of perimeter cells were active; the others stayed at Hold. Boundary traders responded to departures from a moving price reference. Their new states entered the update of adjacent interior cells, providing the route by which executions could affect local behavior. Imbalance, price, and the moving band statistics followed these updates, with smoothing parameter $\alpha = 1/60$:

$$\phi_t = \frac{n_t^B - n_t^S}{N_{\mathrm{tr}}}, \tag{6}$$

$$P_t = P_{t-1}\exp(\lambda\phi_t), \tag{7}$$

$$\mu_t = \mu_{t-1} + \alpha(P_t - \mu_{t-1}), \qquad \sigma_t = \sigma_{t-1} + \alpha(|P_t - \mu_t| - \sigma_{t-1}). \tag{8}$$

Orders came from the current trader states, while the resulting price influenced the next interior update (Algorithm S1). At each step, rule activation and any financial intervention preceded eligibility tests, clearing, and settlement. Executed orders then generated the new quote. Boundary traders responded to the quote before the interior cells updated synchronously. The complete step followed the sequence:

$$\text{lattice state} \to \text{orders} \to \text{dealer fills} \to P_t \to \text{boundary state} \to \text{lattice state}.$$

The *held-boundary* control removed the price-to-behavior connection by setting every active boundary trader to Hold before each interior update. It retained the same initial lattice, target sites, active coordinates, band draws, trader holdings, and dealer. Orders still passed through the same eligibility, clearing, and price rules. Cash, inventory, and price could still change, but none could alter the held-boundary lattice. The interior then evolved through neighbor responses alone, allowing us to determine whether the rule change could spread to a majority without a price signal (Algorithms S1 and S2).

2.4 Paired rule-switch experiments

We introduced the rule change into paired market paths with a common history. Both paths used Rule 19280 at every interior site for $T_{\mathrm{pre}} = 1{,}000$ baseline steps. At the start of step $T_{\mathrm{pre}}+1$, one path activated Rule 19286 at the preselected fraction f. The matched control left those sites on Rule 19280. Neither path reset the lattice, quote, adaptive bands, trader holdings, or dealer. Their common state $X_{T_{\mathrm{pre}}}$ also generated the first quote after activation, so the earliest possible difference occurred in the following interior update. We checked the shared pre-switch histories and first post-switch quote bit for bit in every pair. Each switched path thus began with the same trader resources and neighbor configurations as its control. The four production campaigns used 50 seeds and 3,000 post-activation steps, with initial trader cash of 1,000 and inventory of ten shares. Additional controls extended this protocol to other rules, a matched Sell pulse, frozen boundary bands, different impact coefficients,

and longer horizons. The supplementary protocols specify which part of the intervention or feedback was changed in each control (Appendix A.5).

We varied the activation fraction, lattice size, and financial resources to determine when the rule change spread and how much selling executed (Supplementary Table S1). The selected switch experiment used $L = 64$, an unlimited dealer, and an endogenous boundary at four activation fractions. A finer screen crossed the fraction with eight dealer capacities under both boundary conditions. The finite-size campaign used $L \in \{32, 48, 64, 96, 128\}$ under a held boundary with $h = \infty$ and endogenous boundaries with $h = 1$ or ∞. Because capacity was defined per trader, absolute dealer resources increased with lattice size. The financial campaign multiplied current trader cash and shares independently by $g_w, g_q \in \{0.1, 1, 10\}$ at activation, rounding shares to integers. Its nine interventions shared the same pre-intervention path within each $(f, h, \text{boundary}, \text{seed})$ cell. Under the held boundary, they also had to retain identical lattice trajectories afterward because neither cash nor shares entered the interior update directly. A discrepancy would therefore indicate an implementation error. The appendix gives the full parameter grids and the checks used to verify these paired and held-boundary relationships.

2.5 Outcomes, statistical analysis, and prespecified panels

We used the strictly interior Sell fraction $m(t)$ to measure collective selling. Larger values indicated that more interior traders occupied Sell. A run transitioned when $m(t) \geq 0.5$ at least once after activation; onset was the first crossing time, and transition probability was the fraction of seeds that crossed. We summarized onset among crossing seeds and used the final 500-step mean Sell fraction to describe each path near the end of the observation window. Between-seed variation in the final-window mean $\bar{m}$ defined the ensemble susceptibility, a standard finite-size quantity [38, 39]:

$$\chi_L = (L-2)^2 \left[\langle \bar{m}^2 \rangle - \langle \bar{m} \rangle^2\right]. \tag{9}$$

Here, $(L-2)^2$ is the number of interior sites and angle brackets denote the average over seeds. At a fixed size, higher susceptibility indicates greater variation in the final-window mean across seeds. Linear interpolation gave $f_q(L)$, the activation fraction at which the transition probability reached $q/100$. The midpoint was $f_{50}(L)$, and the width $f_{75}(L) - f_{25}(L)$ measured how much additional activation separated probabilities of 0.25 and 0.75. A smaller width therefore indicated a sharper response. Interpolation used the first upward crossing without forcing the sampled curve to be monotone. Binomial proportions used 95% Wilson intervals [40]; continuous summaries used interquartile ranges unless otherwise specified.

A high Sell fraction need not produce a large price decline if the requested orders cannot execute. We measured the decline by maximum drawdown, the largest proportional fall from a running price peak during the post-activation window. We recorded unfilled orders at the stage where resources prevented execution. Agent suppression counted desired orders excluded by insufficient trader cash or shares. Dealer rejection counted eligible orders left unfilled after matching because the dealer lacked cash or inventory capacity. In the production campaigns, the reported rejection fraction used all eligible orders over the full run, including the baseline steps, as its denominator, whereas the supplementary controls used post-activation orders only. Higher values indicated that the dealer prevented a larger share of available flow from executing. Financial summaries used medians and interquartile ranges across seeds. Drawdown measured the price consequence of filled imbalance, while suppression and rejection identified the resource constraints that prevented desired orders from contributing to that imbalance.

A coarse activation grid could miss the narrow range over which majority probability changed. We refined each curve using a common resolution criterion. Intervals intersecting the empirical 5–95% response were bisected until every size and boundary/capacity condition had at least five fractions strictly inside that probability range. Each added fraction used the original 50 seeds. We then resampled whole seeds across all fractions 2,000 times to estimate uncertainty in the crossings, width, and sampled susceptibility maximum. These percentile intervals preserved dependence between fractions and included reselection of the susceptibility peak. A monotone logistic fit supplied a separate check on the assumed shape:

$$p_L(f) = \{1 + \exp[-a_L - b_L(f - 0.18)/0.01]\}^{-1}, \qquad b_L > 0.$$

We fitted this curve by binomial likelihood and used the same seed blocks to bootstrap its midpoint and width. Fitted standard errors were bootstrap standard deviations. Both analyses condition on the

chosen grid and horizon; neither estimates interpolation bias, adaptive-grid selection uncertainty, or an infinite-time limit.

We selected lattice illustrations from ensemble summaries so that crossing time or final Sell fraction determined which run was displayed. Selection preceded inspection of the images. The activation-fraction panels included a nontransitioning $f = 0.15$ run, both outcome groups at $f = 0.175$, and transitioning runs at $f = 0.185$ and 0.20 under $h = \infty$. Within a transitioning group, selection minimized the distance from the subgroup median onset. Within a nontransitioning group, it minimized the distance from the subgroup median final Sell fraction. The dealer panels required a single seed that transitioned at $h = 0.1$, 1, and ∞ for $f = 0.20$. Among eligible seeds, selection minimized the summed distance from the three capacity-specific median onsets. Both panel sets showed snapshots at activation and at offsets 500, 1,500, and 3,000, allowing the spatial trajectories to be compared at common elapsed times. The resulting panels followed typical onsets within the selected crossing groups and typical final Sell fractions within the noncrossing groups. Individual trajectories illustrated how the spatial patterns developed; ensemble probabilities measured how often a majority occurred.

3 Results

To determine whether a lasting change in one local response could move a stable market path to a Sell majority, we activated Rule 19286 at four fractions of interior sites and compared each switched path with a matched no-switch control, and we found that only the switched paths crossed (Fig. 1). The selected campaign used a 64×64 lattice, an unlimited dealer, and an endogenous boundary (Supplementary Table S1). Each switched path shared its 1,000-step history and its first post-activation quote with its control, so the two paths differed only in how the targeted traders responded to a neighbor sum of one. All 50 controls remained below a majority for 3,000 post-activation steps at every fraction, as did the switched paths at $f = 0.10$ and 0.15, whereas activation at $f = 0.20$ or 0.25 produced a crossing in every seed, with median onset falling from 1,044.5 to 311.5 steps between the two larger fractions (Fig. 1B; Supplementary Table S2). In the representative $f = 0.20$ pair, Sell-rich regions expanded through the switched lattice while the control remained mixed (Fig. 1C). Two controls asked which part of the intervention carried this effect. A matched one-time Sell pulse, which placed the same target sites in Sell once and left their rule unchanged, produced no majority crossings at $f = 0.15$, 0.20, or 0.25 under either boundary, so the lasting response change produced an outcome that the transient state change did not (Supplementary Table S5). An exploratory screen of all 18 single-entry neighbors of Rule 19280 found majority transitions for six of them at $f = 0.20$ and none for the other twelve (Supplementary Table S6). Within the tested interventions, collective selling therefore depended on the persistence of the response change, on which response changed, and on how many traders changed it.

A Sell majority could spread through neighbor responses alone or through the price signal that boundary traders fed back into the lattice. To separate the two routes, we repeated the switch on a finer fraction grid across eight dealer capacities under both the endogenous boundary and a held boundary that kept every active boundary trader in Hold, and we found that the transition survived removal of the price signal (Fig. 2). Under the held boundary, majority probability rose from 0.34 at $f = 0.175$ to 0.70 at 0.185 and 1.00 at 0.20 (Fig. 2A). Changing dealer capacity left every held-boundary (f, seed) lattice trajectory exactly unchanged, despite differences in executions, prices, and dealer holdings, because nothing the dealer did could reach the interior once the boundary stopped responding to price. Local neighbor responses were therefore sufficient to sustain collective selling. Restoring the boundary response to price changed the probability only near the activation window (Fig. 2A,B). With an endogenous boundary and unlimited dealer, 14 of 50 seeds crossed at $f = 0.175$, 43 crossed at 0.185, and all 50 crossed at every sampled fraction from 0.195 upward, and across all dealer capacities the endogenous runs produced no transitions at $f \leq 0.15$ and transitions in every seed at $f \geq 0.20$. Price feedback thus shifted how often the majority formed inside the activation window but was not required for it to develop. The instability lived in the lattice, which left two questions: how the lattice transition depended on size and time, and what determined its price consequence.

Because the majority grew out of finite Sell-rich regions, its location on the activation axis could depend on the size of the lattice and on how long each run was observed. We repeated the held-boundary switch at five lattice sizes on refined fraction grids and found that larger lattices reached a majority at lower activation fractions and over a narrower crossover (Fig. 3). At 3,000 steps under the

held boundary, the midpoint fell from $f_{50} = 0.216$ at $L = 32$ to 0.166 at $L = 128$, and the width decreased from 0.023 [0.018, 0.027] to 0.008 [0.004, 0.010], where brackets give 95% seed-bootstrap intervals (Fig. 3A,C; Supplementary Table S3). Peak susceptibility, the seed-to-seed variance of the final-window Sell fraction scaled by the number of interior sites, increased from 21.2 to 241.2 over the same range (Fig. 3B,D). Refined grids resolved each crossover, and logistic fits under a monotone curve assumption placed each midpoint within 0.004 of its interpolated value, while the fitted widths were narrower under the held boundary (Supplementary Table S10). Observation time also mattered: extending the $L = 64$ window from 3,000 to 9,000 steps lowered the majority midpoint from about 0.178 to 0.167, and size-scaled windows and alternative Sell-fraction criteria retained the endpoint size ordering while longer runs continued to shift the midpoint (Supplementary Table S9). The sharper response on larger lattices therefore described a finite-time crossover whose location depended on the observation horizon, and we did not treat it as an asymptotic threshold.

Ensemble probabilities did not show how an individual lattice reached a majority, or whether it did so by gradual spatial growth or by a near-simultaneous change of many states. We examined the prespecified time-sliced runs across the activation crossover, which were selected from ensemble summaries before any image was inspected, and found that every crossing proceeded by growth and merger of Sell-rich regions (Fig. 4). At $f = 0.15$, the noncrossing run retained a fine mixed pattern for all 3,000 steps. The two $f = 0.175$ runs followed different paths: one formed a Sell-rich region but stayed below one-half, while the other developed a large connected domain and crossed after 2,186 steps. The $f = 0.185$ run crossed after 1,779 steps through several connected regions, and the $f = 0.20$ run crossed after 1,030 steps and ended with one dominant Sell domain (Fig. 4A). Price declined as Sell occupancy rose in each crossing run, with maximum drawdowns of 17.6%, 23.7%, and 25.6% at $f = 0.175$, 0.185, and 0.20, against 5.3% and 9.9% for the two noncrossing runs (Fig. 4B,C). The number, shape, and timing of the domains varied across seeds, and the noncrossing $f = 0.175$ run also formed a visible domain, so the appearance of a Sell-rich region did not by itself guarantee a majority within the horizon. Each displayed transition therefore proceeded by growth and merger of Sell-rich domains rather than by a synchronous relabeling of the lattice, and the ensemble probability counted how often a majority was reached within the horizon.

Having located the instability in the lattice, we asked what determined its price consequence. We compared majority probability, maximum drawdown, and dealer rejection across the eight dealer capacities under the endogenous boundary and found that capacity ordered the price decline while changing majority probability non-monotonically (Fig. 2). Near the activation window, at $f = 0.175$, majority probability was 0.66 with $h = 0.1$ and 0.28 with unlimited capacity, a 38-percentage-point advantage for the smaller dealer (Fig. 2A,B). Freezing each path's boundary decision width at its pre-switch value changed this contrast from 0.38 to -0.02, which implicated the adaptation of that width to the moving quote in the advantage while leaving other non-monotone differences in place (Supplementary Table S7). At $f = 0.20$, every seed transitioned under every capacity, but greater capacity allowed more eligible orders to execute and so increased the drawdown: median maximum drawdown rose from zero at $h = 0$ to 0.40% at $h = 0.1$, 3.92% at $h = 1$, and 25.3% with unlimited capacity, while median dealer rejection fell from 0.449 to zero (Fig. 2C,D). Zero capacity admitted only balanced trader flow and left the quote fixed, and $h = 10$ rejected no orders and matched unlimited capacity for these endowments and horizon. The ordering of drawdowns by capacity survived a decade-wide sweep of the price-impact coefficient, although the magnitudes changed by a factor of seven to ten across the sweep (Supplementary Table S8). The same separation appeared within one seed. At $f = 0.20$, seed 25 reached a majority at $h = 0.1$, 1, and ∞ within 88 steps of one another, and its final-window mean Sell fraction lay between 0.81 and 0.87 in all three endogenous paths and in the held-boundary path, yet its maximum drawdown rose from 0.40% at $h = 0.1$ to 3.92% at $h = 1$ and 25.8% with unlimited capacity (Fig. 5). A dealer could therefore limit the price decline after most interior traders had entered Sell, while price feedback changed how often the majority formed near the window.

Dealer capacity limited which eligible orders executed, but trader holdings determined which orders were eligible in the first place. To separate the two resources, we multiplied trader cash and shares independently by 0.1, 1, or 10 at activation under three dealer capacities and both boundaries, and we found that holdings changed majority probability only when prices fed back to the lattice, whereas their effect on drawdown depended on the dealer (Fig. 6). At $f = 0.18$, all nine cash–share interventions produced identical held-boundary lattice paths and a majority probability of 0.64 across dealer capacities, while endogenous probabilities ranged from 0.50 to 0.94, and the other sampled

fractions showed the same boundary dependence (Fig. 6A; Supplementary Table S4). Holdings altered order eligibility, but their effect on local behavior required the resulting price signal. With baseline cash and an unlimited dealer, increasing the share multiplier from 0.1 to 1 to 10 raised median drawdown from 2.87% to 20.1% to 83.5%, because every eligible sell executed and the dealer absorbed the whole unmatched residual (Fig. 6B). Finite capacity capped that increase: baseline and tenfold shares both produced 3.92% at $h = 1$ and 0.399% at $h = 0.1$, because the dealer rejected the unmatched sells that exceeded its capacity (Fig. 6C). Majority probability did not follow drawdown (Fig. 6D). At $h = 1$ and baseline shares, tenfold cash reduced median drawdown from 3.92% to 0.089% while raising majority probability from 0.78 to 0.94, because the extra cash made more buy orders eligible, and those buys matched sells before they reached the dealer while the resulting quote also changed the boundary signal. Trader and dealer resources therefore determined how much collective selling translated into a price decline, and a Sell majority and a large drawdown were separate outcomes of the same intervention.

4 Discussion

Two results organize this study. A persistent change in how traders responded to their neighbors was sufficient to move the modeled market from mixed behavior to a Sell majority, and what that majority did to the price was decided at execution rather than on the lattice. The altered rule remained in place when a targeted trader left Sell, allowing the new response to act whenever its neighborhood input recurred, whereas a matched one-time Sell pulse changed only the current state and produced no majority crossings in the tested 50-seed controls. Persistent Rule 19286 activation produced crossings in every seed at the two larger fractions, and several other single-entry changes also produced majorities, although most screened rules did not. Collective selling therefore depended on the response being changed as well as on the fraction of traders adopting it. Because paired paths shared the same realized market state, their later divergence followed the intervention by construction. The measured transition curve supplied the additional result: it described how often that local change spread to a majority and how the outcome depended on the activation fraction.

The local rule dynamics could sustain the transition without a price signal. Holding boundary traders inactive left the interior free to evolve under neighbor interactions alone, and Sell majorities still developed. In the displayed crossing trajectories, Sell-rich regions formed, enlarged, and merged, which connected the local responses to the aggregate Sell fraction in the way that evolving spatial patterns do in other nonequilibrium lattice systems [41, 42]. The dependence on activation fraction resembled the spread found in threshold-contagion and bootstrap-percolation models [43–46], where a seeded fraction likewise separates local from global spread. Because the majority grew out of finite domains, its location depended on the size of the lattice and on the time allowed. The response narrowed between the smallest and largest lattices on the refined grids, yet its midpoint continued to fall as the observation window increased, so the measured activation fraction was a property of both size and time. Establishing an asymptotic transition would require horizon convergence and a specified scaling model [38, 39]. The spatial evidence also stops short of classical nucleation, since no free-energy barrier, activation rate, surface tension, or critical droplet was identified. The observed mechanism was growth and merger of Sell-rich domains over finite times.

A Sell state generated a desired order, but trader and dealer resources determined whether that order could move the price. Trader holdings first limited order eligibility, opposing eligible orders then cleared against one another, and the dealer absorbed only the unmatched flow that remained, up to its cash and inventory bounds. Those bounds connect this model to inventory-based accounts of intermediation [23–25, 35], and only the imbalance in filled orders entered the price-impact relation [28, 30, 31]. At zero dealer capacity, balanced execution left the quote fixed, whereas unlimited capacity allowed every eligible order to execute. The same illustrated seed consequently reached a Sell majority with either a 0.40% or a 25.8% drawdown. At 18% activation with baseline cash, tenfold trader shares produced an 83.5% median drawdown against the unlimited dealer, but finite capacity limited the corresponding drawdowns to 3.92% or 0.399%. More shares made additional selling eligible, but additional unmatched selling lowered the quote only while the dealer could continue buying. A lattice full of sellers therefore lowered the price only as far as someone was able to buy.

Price feedback allowed the execution constraint to influence the spread of selling as well as its financial consequence. Near the activation window, capacity $h = 0.1$ raised majority probability

by 38 percentage points relative to the unlimited dealer. The quote determined boundary trader states through a moving center and an adaptive decision width, and freezing each path's width at its pre-switch value removed the low-capacity advantage in the point estimates, linking that advantage to adaptation after the switch. Frozen widths still differed across seeds and capacities, the moving center and financial accounts continued to evolve, and other non-monotone capacity differences remained, so width adaptation did not explain the entire capacity response. The impact sweep placed a further bound on the price interpretation: drawdowns remained ordered in capacity across a decade of impact coefficients, but their magnitudes changed by a factor of seven to ten. Thus, limiting execution could reduce drawdown while price feedback increased the probability of a Sell majority. The two outcomes measured different consequences of the same resource constraint, and a single crash indicator would have combined them.

Several modeling choices limit how far these conclusions extend. Traders updated synchronously on a regular lattice using one baseline rule, and the totalistic score assigned the same response to distinct neighbor compositions, so the rule screen did not test arbitrary composition-based responses. Unit orders settled at a common quote, and the dealer had no strategic quoting policy, risk penalty, or source of replenishment beyond reversing flow. These assumptions made the financial constraints explicit but supplied neither an order book nor an empirical calibration of market depth, price impact, or time. Ten no-switch seeds under each boundary remained below a Sell majority for 100,000 additional steps, which extended the observed baseline persistence without proving global stability, and neither the 1,000-step warm-up nor the final 500-step summary established a stationary distribution. Asynchronous updates, other interaction rules, and an order book would allow the same question to be examined under different assumptions: whether a persistent local response can spread through the market, and how available resources limit the resulting price decline.

5 Conclusion

A lasting change in one local trading response produced Sell majorities from paired market paths that remained mixed under the baseline rule. The matched one-time Sell pulse did not reproduce those transitions, while several other single-entry rule changes did. The activation fraction needed for a majority depended on lattice size and observation time. Holding boundary traders inactive showed that local interactions could sustain the spatial transition without price feedback, so whether selling spread was settled on the lattice. Its financial consequence was settled at execution: trader cash and shares determined eligibility, and dealer cash and inventory limited the unmatched flow that could fill. Restoring price feedback also allowed these constraints to change majority probability non-monotonically, and greater dealer capacity increased drawdown at the main activation fraction even though a smaller dealer could produce a higher majority probability near the activation window. Within this model, collective selling and a large price decline were therefore separate outcomes. Many traders could remain in Sell after the dealer exhausted its ability to buy their shares.

References

[1] Robert J. Shiller. Portfolio insurance and other investor fashions as factors in the 1987 stock market crash. *NBER Macroeconomics Annual*, 3:287–297, 1988. doi: 10.1086/654091.

[2] Sanford J. Grossman. An analysis of the implications for stock and futures price volatility of program trading and dynamic hedging strategies. *The Journal of Business*, 61(3):275–298, 1988. doi: 10.1086/296433.

[3] Amir E. Khandani and Andrew W. Lo. What happened to the quants in August 2007? Evidence from factors and transactions data. *Journal of Financial Markets*, 14(1):1–46, 2011. doi: 10.1016/j.finmar.2010.07.005.

[4] Blake LeBaron. Agent-based computational finance. In *Handbook of Computational Economics*, volume 2, pages 1187–1233. Elsevier, 2006.

[5] Cars H Hommes. Heterogeneous agent models in economics and finance. *Handbook of Computational Economics*, 2:1109–1186, 2006.

[6] E. Samanidou, E. Zschischang, D. Stauffer, and T. Lux. Agent-based models of financial markets. *Reports on Progress in Physics*, 70(3):409–450, 2007. doi: 10.1088/0034-4885/70/3/r03.

[7] Robert L. Axtell and J. Doyne Farmer. Agent-based modeling in economics and finance: past, present, and future. *Journal of Economic Literature*, 63(1):197–287, 2025. doi: 10.1257/jel.20221319.

[8] A. Kirman. Ants, rationality, and recruitment. *The Quarterly Journal of Economics*, 108(1):137–156, 1993. doi: 10.2307/2118498.

[9] Rama Cont and Jean-Philippe Bouchaud. Herd behavior and aggregate fluctuations in financial markets. *Macroeconomic Dynamics*, 4(2):170–196, 2000.

[10] Thomas Lux and Michele Marchesi. Scaling and criticality in a stochastic multi-agent model of a financial market. *Nature*, 397:498–500, 1999.

[11] Stefan Bornholdt. Expectation bubbles in a spin model of markets: intermittency from frustration across scales. *International Journal of Modern Physics C*, 12(5):667–674, 2001.

[12] M. Bartolozzi and A. W. Thomas. Stochastic cellular automata model for stock market dynamics. *Physical Review E*, 69(4):046112, 2004. doi: 10.1103/PhysRevE.69.046112.

[13] G. Qiu, D. Kandhai, and P. M. A. Sloot. Understanding the complex dynamics of stock markets through cellular automata. *Physical Review E*, 75(4):046116, 2007. doi: 10.1103/PhysRevE.75.046116.

[14] Rama Cont. Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001.

[15] Stephen Wolfram. Universality and complexity in cellular automata. *Physica D: Nonlinear Phenomena*, 10:1–35, 1984.

[16] Norman H. Packard and Stephen Wolfram. Two-dimensional cellular automata. *Journal of Statistical Physics*, 38:901–946, 1985.

[17] Stephen Wolfram. *A New Kind of Science*. Wolfram Media, 2002.

[18] Yi-ming Wei, Shang-jun Ying, Ying Fan, and Bing-Hong Wang. The cellular automaton model of investment behavior in the stock market. *Physica A: Statistical Mechanics and its Applications*, 325(3-4):507–516, 2003. doi: 10.1016/S0378-4371(03)00144-4.

[19] Ying Fan, Shang-Jun Ying, Bing-Hong Wang, and Yi-Ming Wei. The effect of investor psychology on the complexity of stock market: an analysis based on cellular automaton model. *Computers & Industrial Engineering*, 56(1):63–69, 2009. doi: 10.1016/j.cie.2008.03.015.

[20] A. P. F. Atman and Bruna Amin Gonçalves. Influence of the investor's behavior on the complexity of the stock market. *Brazilian Journal of Physics*, 42(1-2):137–145, 2012. doi: 10.1007/s13538-011-0057-7.

[21] Paul Glasserman and David D. Yao. Some guidelines and guarantees for common random numbers. *Management Science*, 38(6):884–908, 1992. doi: 10.1287/mnsc.38.6.884.

[22] Averill M. Law. *Simulation modeling and analysis*. McGraw-Hill Education, New York, 5 edition, 2015. ISBN 9780073401324.

[23] Mark B. Garman. Market microstructure. *Journal of Financial Economics*, 3(3):257–275, 1976. doi: 10.1016/0304-405X(76)90006-4.

[24] Hans R. Stoll. The supply of dealer services in securities markets. *The Journal of Finance*, 33 (4):1133–1151, 1978. doi: 10.1111/j.1540-6261.1978.tb02053.x.

[25] Thomas Ho and Hans R. Stoll. Optimal dealer pricing under transactions and return uncertainty. *Journal of Financial Economics*, 9(1):47–73, 1981. doi: 10.1016/0304-405X(81)90020-9.

[26] Ananth Madhavan. Market microstructure: a survey. *Journal of Financial Markets*, 3(3): 205–258, 2000. doi: 10.1016/S1386-4181(00)00007-0.

[27] Tobias Adrian, Nina Boyarchenko, and Or Shachar. Dealer balance sheets and bond liquidity provision. *Journal of Monetary Economics*, 89:92–109, 2017. doi: 10.1016/j.jmoneco.2017.03.011.

[28] Albert S. Kyle. Continuous auctions and insider trading. *Econometrica*, 53(6):1315–1335, 1985.

[29] Joel Hasbrouck. Measuring the information content of stock trades. *The Journal of Finance*, 46 (1):179–207, 1991. doi: 10.1111/j.1540-6261.1991.tb03749.x.

[30] J Doyne Farmer, Paolo Patelli, and Ilija I Zovko. The predictive power of zero intelligence in financial markets. *Proceedings of the National Academy of Sciences*, 102(6):2254–2259, 2005.

[31] Jean-Philippe Bouchaud, J. Doyne Farmer, and Fabrizio Lillo. How markets slowly digest changes in supply and demand. In Thorsten Hens and Klaus Reiner Schenk-Hoppé, editors, *Handbook of Financial Markets: Dynamics and Evolution*, pages 57–160. North-Holland, Amsterdam, 2009. ISBN 9780123742582. doi: 10.1016/B978-012374258-2.50006-3.

[32] Marco Raberto, Silvano Cincotti, Sergio M. Focardi, and Michele Marchesi. Agent-based simulation of a financial market. *Physica A: Statistical Mechanics and its Applications*, 299 (1-2):319–327, 2001. doi: 10.1016/s0378-4371(01)00312-0.

[33] Carl Chiarella and Giulia Iori. A simulation analysis of the microstructure of double auction markets. *Quantitative Finance*, 2(5):346–353, 2002. doi: 10.1088/1469-7688/2/5/303.

[34] Stefan Thurner, J. Doyne Farmer, and John Geanakoplos. Leverage causes fat tails and clustered volatility. *Quantitative Finance*, 12(5):695–707, 2012. doi: 10.1080/14697688.2012.674301.

[35] Karen Braun-Munzinger, Zijun Liu, and Arthur E. Turrell. An agent-based model of corporate bond trading. *Quantitative Finance*, 18(4):591–608, 2018. doi: 10.1080/14697688.2017.1380310.

[36] Hersh Shefrin and Meir Statman. The disposition to sell winners too early and ride losers too long: Theory and evidence. *The Journal of Finance*, 40(3):777–790, 1985.

[37] Terrance Odean. Are investors reluctant to realize their losses? *The Journal of Finance*, 53(5): 1775–1798, 1998.

[38] K. Binder. Finite size scaling analysis of Ising model block distribution functions. *Zeitschrift für Physik B Condensed Matter*, 43(2):119–140, 1981. doi: 10.1007/BF01293604.

[39] David P. Landau and Kurt Binder. *A guide to Monte Carlo simulations in statistical physics*. Cambridge University Press, Cambridge, 5 edition, 2021. ISBN 9781108490146. doi: 10.1017/9781108780346.

[40] Edwin B. Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927. doi: 10.1080/01621459.1927.10502953.

[41] Haye Hinrichsen. Non-equilibrium critical phenomena and phase transitions into absorbing states. *Advances in Physics*, 49(7):815–958, 2000. doi: 10.1080/00018730050198152.

[42] A. J. Bray. Theory of phase-ordering kinetics. *Advances in Physics*, 43(3):357–459, 1994. doi: 10.1080/00018739400101505.

[43] Mark Granovetter. Threshold models of collective behavior. *American Journal of Sociology*, 83 (6):1420–1443, 1978. doi: 10.1086/226707.

[44] Duncan J. Watts. A simple model of global cascades on random networks. *Proceedings of the National Academy of Sciences*, 99(9):5766–5771, 2002. doi: 10.1073/pnas.082090499.

[45] M. Aizenman and J. L. Lebowitz. Metastability effects in bootstrap percolation. *Journal of Physics A: Mathematical and General*, 21(19):3801–3813, 1988. doi: 10.1088/0305-4470/21/19/017.

[46] Alexander E. Holroyd. Sharp metastability threshold for two-dimensional bootstrap percolation. *Probability Theory and Related Fields*, 125(2):195–224, 2003. doi: 10.1007/s00440-002-0239-x.

[47] David Blackman and Sebastiano Vigna. Scrambled linear pseudorandom number generators. *ACM Transactions on Mathematical Software*, 47(4):1–32, 2021. doi: 10.1145/3460772.

Tables

Table S1: Production campaign design. "Endogenous/held" indicates that both matched boundary conditions were run.

Campaign	Lattice sizes	Dealer h	Boundary	Seeds	Simulations
Selected paired switch	64	∞	endogenous	50	400
Reservoir screen	64	0–$10, \infty$	endogenous/held	50	9,600
Finite-size switch	32–128	$1, \infty$	three conditions	50	9,750
Cash–share factorial	64	$0.1, 1, \infty$	endogenous/held	50	10,800

Table S2: Selected paired-switch campaign on a 64×64 lattice with an unlimited dealer and an endogenous boundary. Transition probability is the fraction of seeds whose peak interior Sell fraction reached one half after activation. Median onset is taken over crossing seeds and is undefined where no seed crossed.

		Transition probability		
Fraction f	Seeds	No-switch	Switched	Median onset (steps)
0.10	50	0.00	0.00	Undefined
0.15	50	0.00	0.00	Undefined
0.20	50	0.00	1.00	1,044.5
0.25	50	0.00	1.00	311.5

Table S3: Refined finite-size response over 50 seeds and 3,000 post-activation steps. Brackets give 95% confidence intervals (CI) from seed resampling. n_{int} counts fractions with $0.05 < \hat{p} < 0.95$. Crossings use linear interpolation. The susceptibility maximum is over the sampled grid, and its interval includes reselection of that maximum within each bootstrap replicate. Intervals condition on the refined grid.

Boundary	h	L	n_{int}	f_{50} [95% CI]	Width [95% CI]	Peak χ_L [95% CI]
Endog.	1	32	8	0.210 [0.207, 0.218]	0.024 [0.018, 0.031]	23.6 [19.4, 30.0]
Endog.	1	48	5	0.186 [0.182, 0.188]	0.016 [0.010, 0.020]	41.3 [32.5, 56.2]
Endog.	1	64	8	0.178 [0.172, 0.178]	0.009 [0.007, 0.011]	70.1 [56.1, 91.8]
Endog.	1	96	6	0.166 [0.164, 0.168]	0.010 [0.007, 0.012]	126.5 [104.2, 163.7]
Endog.	1	128	5	0.165 [0.164, 0.166]	0.006 [0.005, 0.007]	191.4 [154.7, 261.0]
Endog.	∞	32	8	0.212 [0.208, 0.215]	0.023 [0.011, 0.030]	23.9 [18.5, 31.7]
Endog.	∞	48	6	0.188 [0.179, 0.192]	0.019 [0.016, 0.021]	52.0 [42.7, 67.5]
Endog.	∞	64	5	0.179 [0.177, 0.181]	0.009 [0.006, 0.014]	90.3 [78.2, 121.0]
Endog.	∞	96	5	0.168 [0.166, 0.169]	0.010 [0.005, 0.012]	173.4 [136.1, 230.6]
Endog.	∞	128	7	0.165 [0.164, 0.167]	0.007 [0.005, 0.008]	262.5 [226.5, 338.3]
Held	∞	32	5	0.216 [0.211, 0.224]	0.023 [0.018, 0.027]	21.2 [16.5, 27.2]
Held	∞	48	5	0.193 [0.184, 0.195]	0.019 [0.016, 0.022]	48.6 [40.3, 65.3]
Held	∞	64	5	0.178 [0.176, 0.180]	0.015 [0.008, 0.019]	110.5 [78.5, 138.5]
Held	∞	96	5	0.168 [0.167, 0.170]	0.009 [0.006, 0.012]	168.7 [137.6, 221.0]
Held	∞	128	7	0.166 [0.165, 0.167]	0.008 [0.004, 0.010]	241.2 [202.6, 326.0]

Table S4: Transition probability in the cash–share campaign over 50 seeds. Each fraction contains 27 held-boundary cells, which returned one common value, and 27 endogenous cells spanning the stated range across the three-by-three cash–share grid and three dealer capacities.

Fraction f	Held boundary	Endogenous range
0.170	0.22	0.08–0.52
0.175	0.34	0.24–0.82
0.180	0.64	0.50–0.94
0.185	0.70	0.64–0.98

Table S5: Matched one-time Sell pulse and persistent Rule 19286 activation. Each arm used the same target sites and 3,000 post-intervention steps. Entries are majority probabilities with 95% Wilson intervals; the dealer was unlimited under both boundary conditions.

Boundary	f	Seeds	Sell pulse	Persistent rule
Endogenous	0.15	50	0.00 [0.00, 0.07]	0.00 [0.00, 0.07]
Endogenous	0.20	50	0.00 [0.00, 0.07]	1.00 [0.93, 1.00]
Endogenous	0.25	50	0.00 [0.00, 0.07]	1.00 [0.93, 1.00]
Held	0.15	50	0.00 [0.00, 0.07]	0.00 [0.00, 0.07]
Held	0.20	50	0.00 [0.00, 0.07]	1.00 [0.93, 1.00]
Held	0.25	50	0.00 [0.00, 0.07]	1.00 [0.93, 1.00]

Table S6: Single-entry neighbors of Rule 19280. Each probability pair is held / endogenous boundary over 20 seeds; the intervention is persistent. The output labels are Buy $= 0$, Hold $= 1$, Sell $= 2$.

Sum c	Output change	Rule	$f = 0.15$	$f = 0.20$	$f = 0.25$
0	$2 \to 0$	19278	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
0	$2 \to 1$	19279	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
1	$0 \to 1$	19283	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
1	$0 \to 2$	19286	0.00 / 0.00	1.00 / 1.00	1.00 / 1.00
2	$0 \to 1$	19289	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
2	$0 \to 2$	19298	0.95 / 1.00	1.00 / 1.00	1.00 / 1.00
3	$0 \to 1$	19307	1.00 / 1.00	1.00 / 1.00	1.00 / 1.00
3	$0 \to 2$	19334	1.00 / 1.00	1.00 / 1.00	1.00 / 1.00
4	$1 \to 0$	19199	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
4	$1 \to 2$	19361	1.00 / 1.00	1.00 / 1.00	1.00 / 1.00
5	$1 \to 0$	19037	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
5	$1 \to 2$	19523	1.00 / 1.00	1.00 / 1.00	1.00 / 1.00
6	$2 \to 0$	17822	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
6	$2 \to 1$	18551	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
7	$2 \to 0$	14906	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
7	$2 \to 1$	17093	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
8	$2 \to 0$	6158	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00
8	$2 \to 1$	12719	0.00 / 0.00	0.00 / 0.00	0.00 / 0.00

Table S7: Band-deviation control at $f = 0.175$ over 50 seeds. Probability brackets are 95% Wilson confidence intervals (CI). Boundary Sell occupancy is the mean fraction of active boundary traders in Sell over post-switch steps and seeds.

h	Adaptive p [95% CI]	Frozen p [95% CI]	Adaptive Sell	Frozen Sell
0	0.34 [0.22, 0.48]	0.34 [0.22, 0.48]	0.000	0.000
0.1	0.66 [0.52, 0.78]	0.50 [0.37, 0.63]	0.154	0.038
0.5	0.56 [0.42, 0.69]	0.30 [0.19, 0.44]	0.158	0.045
1	0.48 [0.35, 0.61]	0.42 [0.29, 0.56]	0.142	0.058
2	0.30 [0.19, 0.44]	0.56 [0.42, 0.69]	0.096	0.107
5	0.30 [0.19, 0.44]	0.52 [0.39, 0.65]	0.136	0.271
10	0.28 [0.17, 0.42]	0.52 [0.39, 0.65]	0.113	0.189
∞	0.28 [0.17, 0.42]	0.52 [0.39, 0.65]	0.113	0.189

Table S8: Price-impact sweep at $f = 0.20$ under the endogenous boundary, 50 seeds per cell. Drawdown is a percentage; rejection is the fraction of post-switch eligible orders rejected.

h	λ	Transition p	Median drawdown (%)	Median rejection
0.1	0.005	1.00	0.10	0.507
0.1	0.020	1.00	0.40	0.506
0.1	0.050	1.00	0.99	0.520
1	0.005	1.00	1.00	0.441
1	0.020	1.00	3.92	0.412
1	0.050	1.00	9.52	0.395
∞	0.005	1.00	7.10	0.000
∞	0.020	1.00	25.26	0.000
∞	0.050	1.00	49.56	0.000

Table S9: Held-boundary horizon and threshold sensitivity, 50 seeds. Entries are interpolated f_{50} [95% seed-bootstrap interval] for each crossing criterion. The horizon grids are independent of adaptive refinement, so their interpolated values can differ slightly from Supplementary Table S3.

L	Post-switch steps	Criterion 0.4	Criterion 0.5	Criterion 0.6
32	750	0.236 [0.228, 0.238]	0.249 [0.245, 0.255]	0.266 [0.260, 0.270]
32	1,500	0.216 [0.210, 0.220]	0.230 [0.228, 0.234]	0.245 [0.238, 0.250]
32	3,000	0.207 [0.205, 0.210]	0.219 [0.210, 0.225]	0.229 [0.226, 0.232]
64	3,000	0.172 [0.169, 0.176]	0.178 [0.176, 0.180]	0.186 [0.180, 0.188]
64	9,000	0.165 [0.163, 0.167]	0.167 [0.165, 0.168]	0.171 [0.169, 0.175]
128	3,000	0.163 [0.162, 0.164]	0.167 [0.165, 0.168]	0.169 [0.168, 0.171]
128	6,000	0.154 [0.153, 0.156]	0.157 [0.156, 0.158]	0.159 [0.158, 0.160]
128	12,000	0.150 [0.148, 0.151]	0.152 [0.150, 0.153]	0.153 [0.152, 0.155]

Table S10: Logistic shape sensitivity on the refined grids. Parentheses give seed-bootstrap standard errors (SE), not binomial standard errors treating fractions as independent. The last column is the location of the empirical susceptibility maximum: it is a sampled fraction with local grid resolution, not an interpolated optimum.

Boundary	h	L	Logistic f_{50} (SE)	Logistic width (SE)	f at peak
Endog.	1	32	0.212 (0.0016)	0.020 (0.0014)	0.217
Endog.	1	48	0.184 (0.0011)	0.014 (0.0011)	0.190
Endog.	1	64	0.175 (0.0007)	0.010 (0.0008)	0.182
Endog.	1	96	0.166 (0.0007)	0.008 (0.0006)	0.170
Endog.	1	128	0.165 (0.0006)	0.006 (0.0006)	0.172
Endog.	∞	32	0.212 (0.0015)	0.019 (0.0014)	0.225
Endog.	∞	48	0.186 (0.0013)	0.014 (0.0011)	0.190
Endog.	∞	64	0.178 (0.0007)	0.009 (0.0008)	0.185
Endog.	∞	96	0.168 (0.0007)	0.007 (0.0005)	0.170
Endog.	∞	128	0.166 (0.0004)	0.005 (0.0003)	0.168
Held	∞	32	0.218 (0.0016)	0.018 (0.0014)	0.225
Held	∞	48	0.189 (0.0011)	0.013 (0.0009)	0.190
Held	∞	64	0.178 (0.0010)	0.011 (0.0008)	0.185
Held	∞	96	0.169 (0.0006)	0.007 (0.0006)	0.172
Held	∞	128	0.166 (0.0005)	0.006 (0.0004)	0.168

Figures

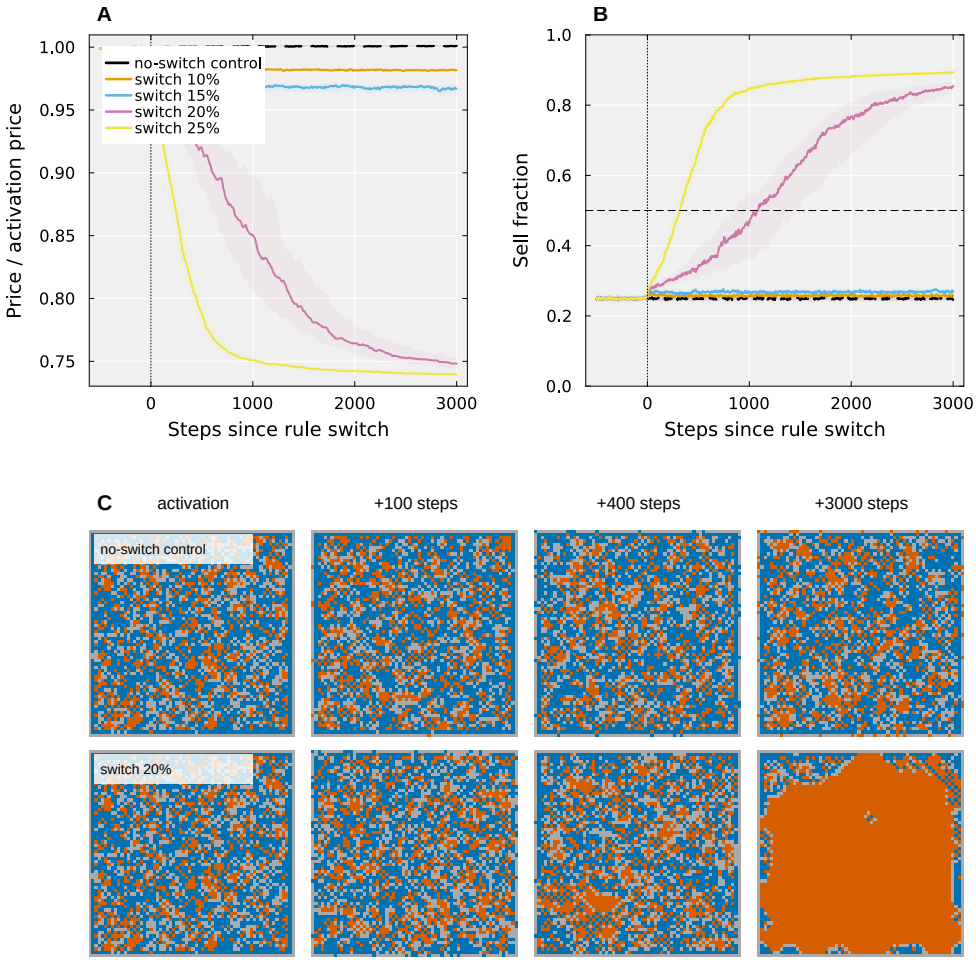


Figure 1: A paired mid-run rule switch separated initially identical market paths. Fifty seeds were run on a 64×64 lattice for 1,000 baseline steps and 3,000 post-activation steps. The no-switch path retained the same target sites but kept their baseline rules. **(A)** Median price divided by its value at activation; bands mark the interquartile range. **(B)** Median interior Sell fraction; the horizontal line marks the majority threshold. **(C)** Time-sliced lattices for seed 12, the $f = 0.20$ pair nearest the median switched-control difference in peak Sell fraction. Blue, gray, and orange cells denote Buy, Hold, and Sell. The rows show the no-switch and switched paths from their common activation state through 3,000 steps. All paired paths coincided before the vertical activation line.

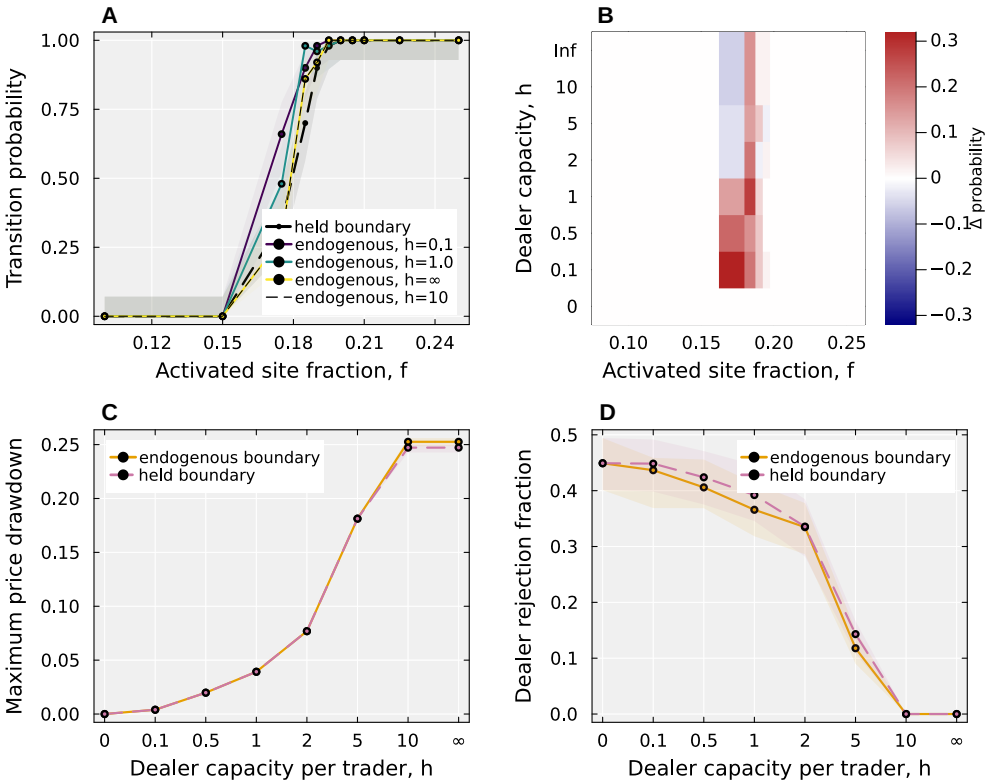


Figure 2: Rule activation across dealer capacities. **(A)** Transition probability under the held boundary and selected endogenous dealer conditions; bands are 95% Wilson intervals over 50 seeds. The held boundary kept active boundary traders in the Hold state. **(B)** Difference between endogenous and held-boundary transition probabilities across the complete capacity range. **(C)** Median maximum drawdown at $f = 0.20$; bands mark the interquartile range. **(D)** Median fraction of agent-eligible orders rejected by the dealer at $f = 0.20$. Capacity $h = 10$ is shown as a dashed overlay in (A) and as an explicit capacity level in (B–D); it coincided exactly with $h = \infty$.

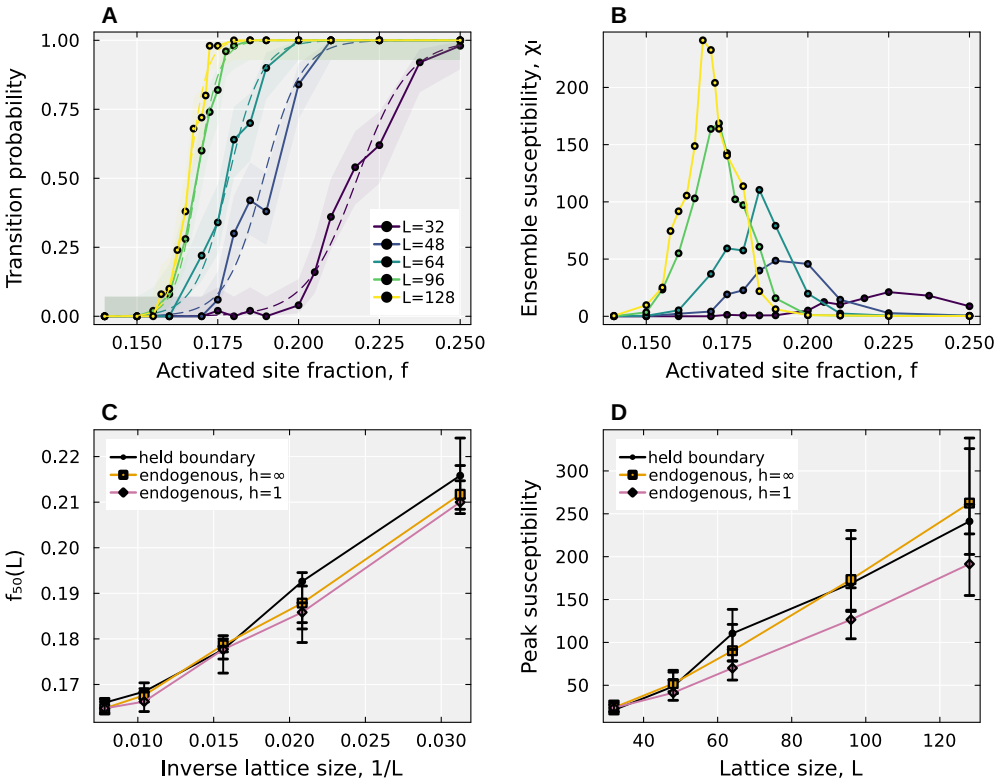


Figure 3: Finite-size response of the paired rule switch over 50 seeds. **(A)** Held-boundary transition probability; bands are 95% Wilson intervals. **(B)** Held-boundary ensemble susceptibility of the final 500-step mean Sell fraction. **(C)** Interpolated finite-horizon midpoint $f_{50}(L)$ for the held control and two endogenous dealer conditions. **(D)** Peak susceptibility versus lattice size on linear axes. Error bars in (C,D) are 95% seed-bootstrap intervals. Solid lines connect sampled values; dashed curves in (A) are logistic fits. No asymptotic scaling law was fitted.

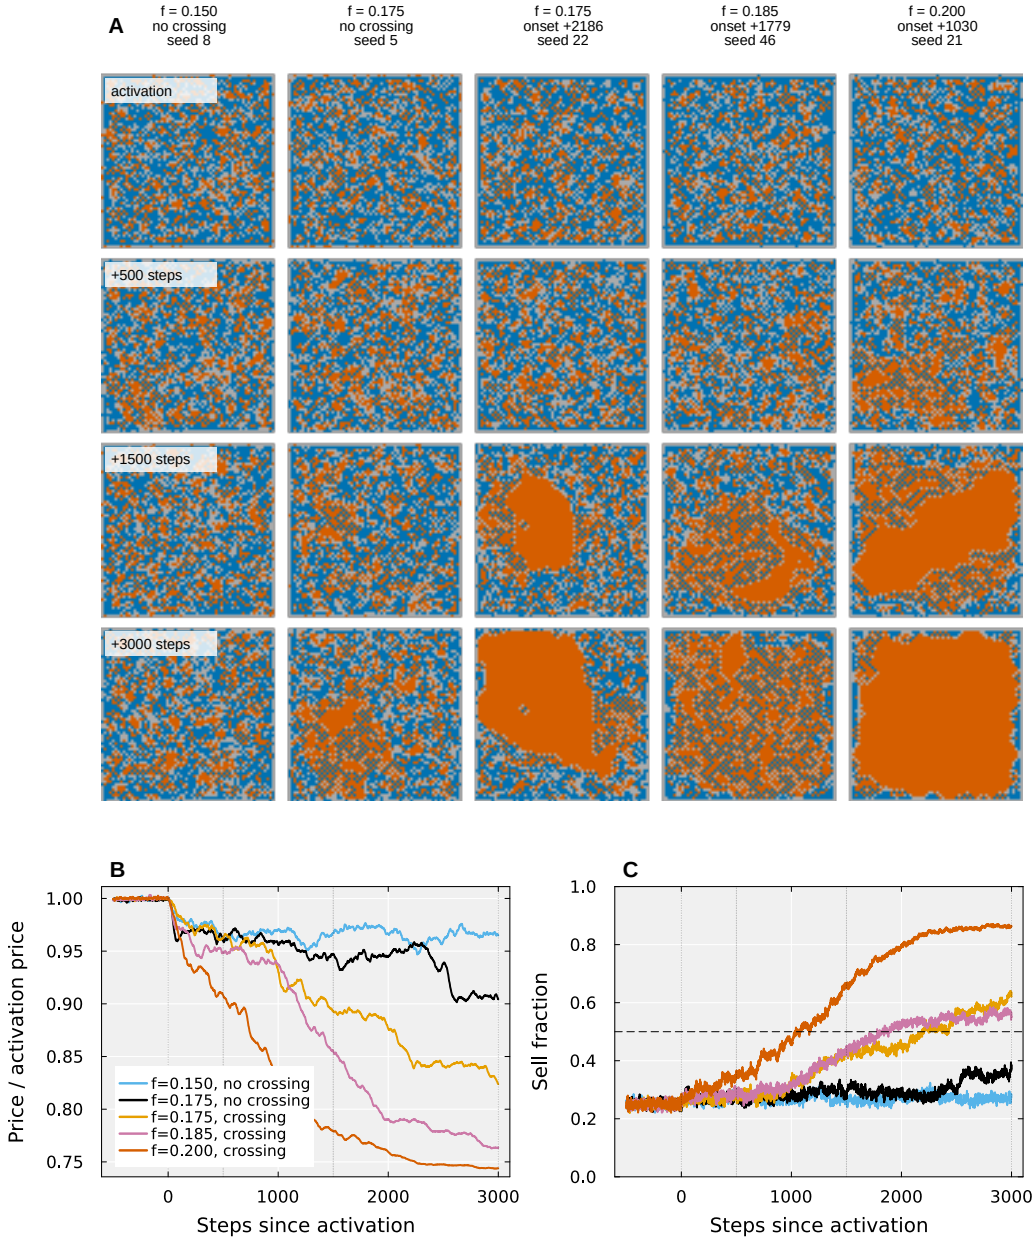


Figure 4: Prespecified time-sliced markets across the activation-fraction crossover under an endogenous boundary and unlimited dealer. Blue, gray, and orange cells denote Buy, Hold, and Sell. **(A)** Lattices for five runs (columns) at activation and after 500, 1,500, and 3,000 steps (rows). From left to right, the runs include a noncrossing $f = 0.15$ case, both outcome groups at $f = 0.175$, and crossing cases at $f = 0.185$ and 0.20. Crossing cases were nearest their subgroup median onset. Noncrossing cases were nearest their subgroup median final Sell fraction. Seeds were 8, 5, 22, 46, and 21 from left to right. **(B)** Price divided by its activation value. **(C)** Interior Sell fraction; the horizontal line marks the majority threshold. Vertical guides mark the four snapshots.

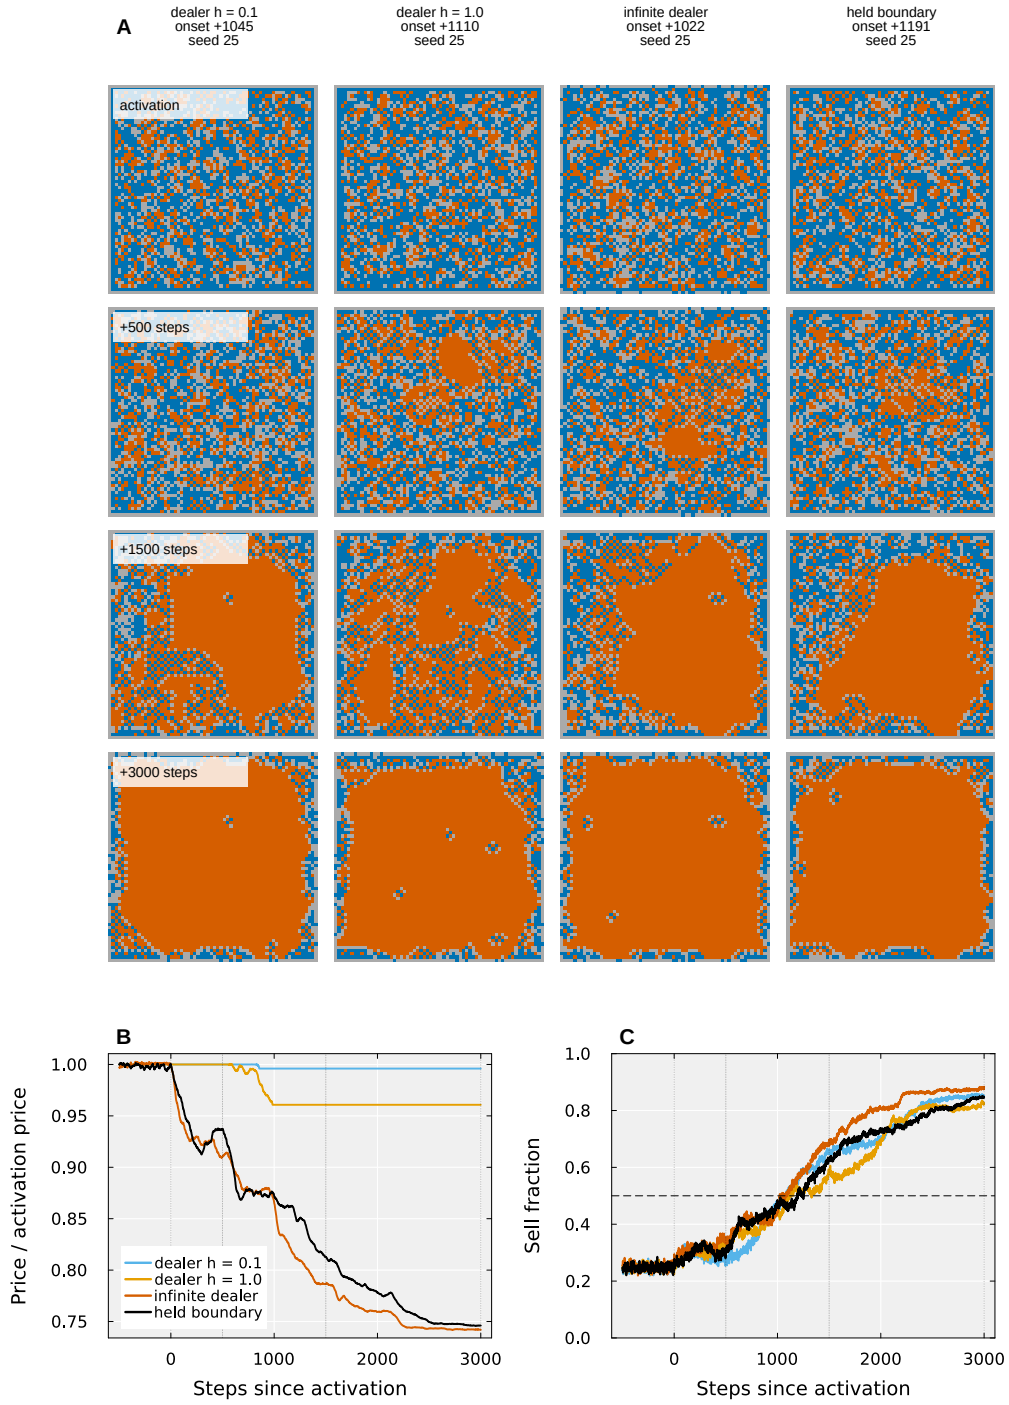


Figure 5: Time-sliced markets for one matched seed at $f = 0.20$. **(A)** Lattices for four conditions (columns) at activation and after 500, 1,500, and 3,000 steps (rows). The columns contain endogenous dealers with $h = 0.1$, $h = 1$, and $h = \infty$, followed by the held-boundary control with $h = \infty$. Seed 25 minimized the summed distance from the three endogenous capacity-specific median onsets among seeds that crossed in all three conditions. Rows use common elapsed times. **(B)** Price divided by its activation value. **(C)** Interior Sell fraction. The held boundary prevented price from changing the lattice but did not prevent price from being calculated. The single seed illustrates trajectories; the capacity comparison supplies ensemble evidence (Fig. 2).

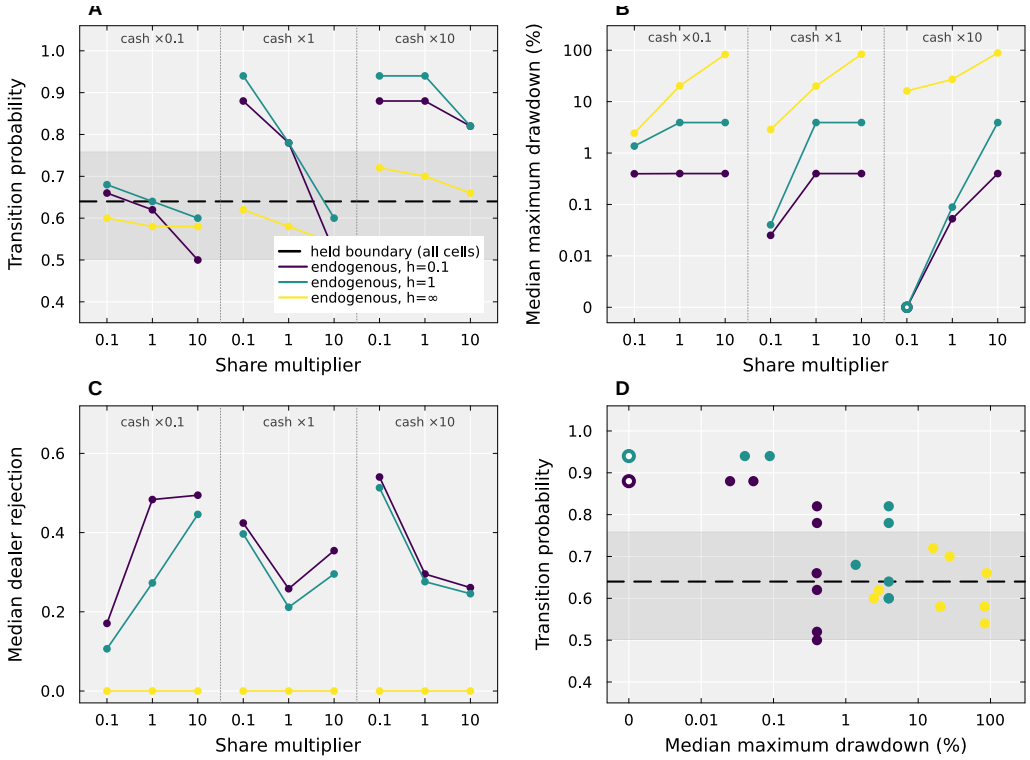


Figure 6: Crossed trader cash and share interventions at activation under the endogenous boundary, with $f = 0.18$ and 50 seeds. In panels A to C the nine intervention cells are grouped by cash multiplier with the share multiplier on the axis, and each series is one dealer capacity. **(A)** Transition probability. The dashed line and band give the held-boundary value and its 95% Wilson interval, which were identical in all nine cells because the held-boundary lattice was invariant to every intervention. **(B)** Median maximum drawdown on a logarithmic axis. Medians of exactly zero are drawn at the axis floor with open markers. **(C)** Median fraction of agent-eligible orders rejected by the dealer. **(D)** Transition probability against median maximum drawdown for all 27 endogenous cells, on the same logarithmic axis and with zero medians drawn as in B, with the held-boundary reference repeated.

A Supplementary methods and validation

A.1 One complete simulation step

The standing trader states generated orders before the new quote could change those states (Algorithm S1). Orders passed through eligibility, clearing, and settlement before generating a new quote (Algorithm S2). Orders at step t came from the lattice X_{t-1} and stored boundary states $e_{i,t-1}$. They settled at the standing quote P_{t-1} before executed imbalance generated P_t. Boundary traders then classified the new quote, setting the perimeter states used to produce X_t. Rule activation occurred before this sequence without changing the current state. A balance-sheet intervention, when scheduled, changed trader resources before eligibility was evaluated. This timing explains why a rule switch leaves the first post-switch quote unchanged, whereas a financial intervention can affect it immediately. The held-boundary control retained the same sequence but replaced boundary classification with Hold. Under either boundary, a shortage of resources could stop an order without changing the trader's cellular-automaton state. We recorded orders excluded by trader resources as agent suppression and eligible orders left unfilled by the dealer as dealer rejection. Neighboring traders continued to respond to the stored Buy, Hold, or Sell state even when the associated order did not execute.

Algorithm S1 One endogenous-market step, $t-1 \to t$, in simulation order. Orders came from the old interior and boundary states, settled at the old quote, and generated the new quote. The new quote set the boundary before the interior cellular automaton update.

Require: lattice X_{t-1}; boundary states $e_i(t-1)$; trader books; dealer ledger; quote and band state $(P_{t-1}, \mu_{t-1}, \sigma_{t-1})$; rule field R_i; parameters; simulation seed
Ensure: lattice X_t; boundary states $e_i(t)$; settled books and dealer; quote and band state (P_t, μ_t, σ_t); all order counts

if t is the scheduled rule-switch step **then**
  $R_i \leftarrow 19286$ for every preselected target site i {preserve every dynamic state variable}
end if
if t is a scheduled balance-sheet intervention **then**
  for all traders i **do**
    $w_i \leftarrow g_w w_i$; $q_i \leftarrow \mathrm{round}(g_q q_i)$
  end for
end if
call Algorithm S2 using states at $t-1$ and P_{t-1}
receive $(n_t^B, n_t^S, u_t^B, u_t^S, r_t^B, r_t^S)$ and settled ledgers
$\phi_t \leftarrow (n_t^B - n_t^S)/N_{\mathrm{tr}}$
$P_t \leftarrow P_{t-1} \exp(\lambda \phi_t)$
$\mu_t \leftarrow \mu_{t-1} + \alpha(P_t - \mu_{t-1})$
$\sigma_t \leftarrow \sigma_{t-1} + \alpha(|P_t - \mu_t| - \sigma_{t-1})$
if boundary mode is endogenous **then**
  for all active boundary traders i **do**
    $e_i(t) \leftarrow 0$ if $P_t < \mu_t - k_i\sigma_t$; 2 if $P_t > \mu_t + k_i\sigma_t$; 1 otherwise
  end for
else
  $e_i(t) \leftarrow 1$ for every active boundary trader i
end if
write $e_i(t)$ to active boundary sites; leave passive sites at Hold
for all strictly interior sites i, simultaneously **do**
  $c_i(t) \leftarrow \sum_{j \in \partial i} s_j$ using old interior states and new boundary states
  $s_i(t) \leftarrow \mathcal{R}_{R_i}(c_i(t))$
end for

Algorithm S2 Order eligibility, dealer clearing, rationing, and settlement at P_{t-1}. The finite dealer filled only the unmatched residual. The $h = \infty$ branch filled every eligible order.

Require: old interior and boundary states; books $(w_i, q_i, \bar{c}_i)$; dealer (W_{t-1}, Q_{t-1}) and bounds $(Q_{\min}, Q_{\max}, C)$; P_{t-1}; h; cost-basis toggle; seed and step t

Ensure: settled books and dealer; executed counts (n_t^B, n_t^S); agent suppression (u_t^B, u_t^S); dealer rejection (r_t^B, r_t^S)

$\mathcal{B}, \mathcal{S} \leftarrow \varnothing$; $u_t^B, u_t^S \leftarrow 0$
for all interior traders, followed by active-boundary traders i **do**
  $o_i \leftarrow +1$ for Buy, -1 for Sell, and 0 for Hold {an edge trader reads its stored boundary state}
  if $o_i = +1$ and $w_i \geq P_{t-1}$ **then**
    append i to $\mathcal{B}$
  else if $o_i = -1$, $q_i \geq 1$, and the cost-basis test passes **then**
    append i to $\mathcal{S}$
  else if $o_i = +1$ **then**
    $u_t^B \leftarrow u_t^B + 1$
  else if $o_i = -1$ **then**
    $u_t^S \leftarrow u_t^S + 1$
  end if
end for
$B_t \leftarrow |\mathcal{B}|$; $S_t \leftarrow |\mathcal{S}|$; $n_t^B \leftarrow B_t$; $n_t^S \leftarrow S_t$
if $h < \infty$ and $B_t > S_t$ **then**
  $K_t^B \leftarrow Q_{t-1} - Q_{\min}$ {shares that the dealer could sell}
  $n_t^B \leftarrow S_t + \min(B_t - S_t, K_t^B)$
else if $h < \infty$ and $S_t > B_t$ **then**
  $K_t^S \leftarrow Q_{\max} - Q_{t-1}$ {inventory space}
  $K_t^S \leftarrow \min\big(K_t^S, \lfloor (W_{t-1} + C)/P_{t-1} \rfloor\big)$ {cash limit; $C = 0$ in this study}
  $n_t^S \leftarrow B_t + \min(S_t - B_t, K_t^S)$
end if
$r_t^B \leftarrow B_t - n_t^B$; $r_t^S \leftarrow S_t - n_t^S$
if $r_t^B > 0$ **then**
  sort $\mathcal{B}$ by $z(\text{seed}, t, i, +1)$, then by trader identity
end if
if $r_t^S > 0$ **then**
  sort $\mathcal{S}$ by $z(\text{seed}, t, i, -1)$, then by trader identity
end if
for all the first n_t^B traders i in $\mathcal{B}$ **do**
  $\bar{c}_i \leftarrow (\bar{c}_i q_i + P_{t-1})/(q_i + 1)$; $q_i \leftarrow q_i + 1$; $w_i \leftarrow w_i - P_{t-1}$
end for
for all the first n_t^S traders i in $\mathcal{S}$ **do**
  $q_i \leftarrow q_i - 1$; $w_i \leftarrow w_i + P_{t-1}$ {$\bar{c}_i$ was unchanged on a sale}
end for
$W_t \leftarrow W_{t-1} + P_{t-1}(n_t^B - n_t^S)$
$Q_t \leftarrow Q_{t-1} + n_t^S - n_t^B$
return $(n_t^B, n_t^S, u_t^B, u_t^S, r_t^B, r_t^S)$ and settled ledgers

A.2 Rationing and paired reproducibility

When eligible buys and sells differed, the dealer rationed the excess side if its resources could not absorb all unmatched orders. Each eligible trader on that side received a 64-bit score from a SplitMix64 finalizer applied to $(\text{seed}, t, \text{trader identity}, \text{side})$. Orders were sorted by score, using trader identity to break ties reproducibly. The dependence on t prevented a trader from retaining the same priority across repeated orders. Paired paths used the same score for a given (t, i, side) even after their eligible sets diverged. Scoring did not consume draws from the generator used to initialize the lattice and target sites. For each pair, we checked the frame at T_{pre}, its preceding Sell history, the boundary coordinates and band multipliers, and the price history through the first post-switch quote. Held-boundary balance-sheet cells additionally required exact agreement of the full Sell history and final frame across all nine interventions. Exact pre-switch agreement ensured that a pair

reached activation with the same history. Exact post-switch agreement across held-boundary financial interventions ensured that trader resources affected execution without entering the local rule update.

A.3 Conservation and limiting cases

Settlement conserved total cash and shares across traders and dealer when no intervention occurred:

$$\sum_i w_i(t) + W_t = \sum_i w_i(0) + W_0, \tag{10}$$

$$\sum_i q_i(t) + Q_t = \sum_i q_i(0) + Q_0. \tag{11}$$

The maximum absolute cash error across the production reservoir and finite-size campaigns was 2.53×10^{-6}; the share error was zero. In the balance-sheet campaign, checks were applied separately before and after the intentional resource change. Its maximum cash error was 9.47×10^{-7}, again with zero share error. Three limiting cases also supplied regression checks. Unlimited capacity filled every eligible order and reproduced the earlier fill-all implementation. Zero capacity allowed only balanced trader flow, giving a fixed price and identical endogenous and held-boundary lattice paths. Capacity $h = 10$ never rejected an order in the reservoir campaign and matched the unlimited path. That last equality applies to the stated endowments and horizon. A different flow history, larger trader holdings, or longer observation could cause this finite dealer's cash or inventory limits to bind.

A.4 Campaign inventory

We used four production campaigns to measure activation, dealer capacity, lattice size, and trader resources, totaling 30,550 simulations (Supplementary Table S1). Each used 50 seeds and 3,000 post-activation steps as fixed campaign settings. Pilot and full modes recorded how the experiments were developed. The seed count and horizon were fixed campaign settings, without a prospective power calculation or independent preregistration of the horizon. With 50 seeds, a probability near one-half has a binomial standard error of about 0.07. The reservoir campaign reused one no-switch control within each $(h, \text{boundary}, \text{seed})$ cell. Its comma-separated value (CSV) file therefore stored the control metrics with each treated summary row rather than as a separate trajectory for every fraction. Checkpointed campaigns reconstructed a fixed row order after threaded computation. A resumed campaign therefore retained the same control comparisons and row order as an uninterrupted run. Reusing the no-switch path avoided repeating an identical simulation for each activation fraction.

The selected paired-switch campaign used $L = 64$, an unlimited dealer, and an endogenous boundary at $f = 0.10, 0.15, 0.20, 0.25$. The reservoir screen used the same lattice size at fractions 0.10, 0.15, 0.175, 0.185, 0.190, 0.195, 0.200, 0.205, 0.210, 0.225, and 0.250. It crossed these fractions with capacities $h = 0, 0.1, 0.5, 1, 2, 5, 10, \infty$ under both boundaries. The original finite-size grid used fractions 0.140, 0.150, 0.160, 0.170, 0.175, 0.180, 0.185, 0.190, 0.200, 0.210, 0.225, and 0.250 at $L = 32, 48, 64, 96, 128$. Its three conditions were a held boundary with $h = \infty$ and endogenous boundaries with $h = 1$ or ∞. The cash–share campaign crossed nine multiplier pairs with $f = 0.170, 0.175, 0.180, 0.185$ and $h = 0.1, 1, \infty$ under both boundaries. Capacity was defined per trader throughout; a fixed-absolute-capacity size series was not run. The supplementary controls below retain this initialization while varying the stated intervention, parameter, or observation window.

Reproducing a market path required both the seeded setup and the specified update order. We ran the simulations in Julia 1.12 with the environment pinned by the repository manifest. Original campaigns used 12 simulation threads and one garbage-collection thread; supplementary campaigns used four or eight simulation threads and one garbage-collection thread. A seeded Xoshiro generator [47] supplied the setup draws. Finite-size and balance-sheet campaigns wrote atomic checkpoints, then aggregated them in a deterministic order. Tests covered parameter validation, dealer limits, conservation, delayed activation, boundary controls, and time-varying priorities. Stepwise comparisons checked the optimized lattice kernel against the general totalistic solver under both homogeneous and heterogeneous rule fields. Bootstrap analysis and generated LaTeX tables used Python, with package versions recorded in `code/revision/requirements.txt`. Source code, environment manifests, derived data, and figure scripts are available at https://github.com/varnerlab/Wolfram-markets-abm-paper. The repository also retains the earlier external-price classroom implementation and its archived Invesco QQQ exchange-traded fund input for provenance. No result from that model entered the principal analysis or figures.

Majority probability and onset described whether collective selling occurred and how long it took. We used these summaries to compare activation fractions, lattice sizes, and trader resources. In the selected switch campaign, transition status was evaluated from peak Sell fraction because the stored onset used zero for a path with no crossing. Median onset was therefore undefined at the two lower fractions, where no switched seed transitioned (Supplementary Table S2). The refined size comparisons used midpoints, widths, and susceptibility maxima calculated over each sampled fraction grid (Supplementary Table S3). The susceptibility peak location remained a sampled fraction rather than an interpolated optimum (Supplementary Table S10). In the financial campaign, all 27 held-boundary conditions returned one transition probability at each fraction. The 27 endogenous conditions spanned a range of probabilities (Supplementary Table S4). The shared held-boundary probability came from local rule dynamics at the given fraction. The endogenous range included the additional influence of prices on boundary states. Each probability counted crossings within its stated observation window, and each interpolated midpoint depended on the fractions sampled around that crossing.

A.5 Additional controls and sensitivity analyses

We varied the persistence of the intervention, the boundary response, and the observation window to identify which parts of the model supported the transition. These controls were added in response to review and retained the original initialization and 1,000-step baseline. Unless specified otherwise, they used $L = 64$, $\lambda = 0.02$, $h = \infty$, and 3,000 post-intervention steps. Each new rule or pulse arm had a matched arm that retained the baseline rule. We required exact pre-intervention agreement of their lattice and Sell histories, and checked the price prehistory for financial arms. Held-boundary experiments needing only lattice observables used an allocation-free kernel with the production setup draws and rule tables. Tests matched its full Sell history and final grid to the market simulator, and 1,800 overlapping archived seed/fraction outcomes also agreed. The reduced kernel reproduced the local dynamics when prices could not feed back, but did not calculate financial observables.

Matched transient perturbation and rule screen. The matched pulse placed targeted traders in Sell once while leaving their subsequent response rule unchanged (Supplementary Table S5). At activation, we wrote Sell at the rule-switch target sites before orders executed and retained Rule 19280 everywhere. The pulse could therefore change the first quote, whereas persistent activation preserved that quote and changed later local updates. Fractions matched the number of targeted sites; sites already in Sell did not change state. For $f = 0.15, 0.20, 0.25$ under both boundaries, no pulse reached a majority in 50 seeds (0/50; Wilson upper bound 0.071). Persistent Rule 19286 activation yielded 0/50, 50/50, and 50/50, respectively. The Wilson lower bound for 50/50 was 0.929. A one-time increase in Sell occupancy at these targets was thus insufficient to reproduce the collective response to a lasting rule change. The comparison applied to this matched pulse; other choices of state perturbation could have different outcomes.

The rule screen tested all 18 ways to change one entry of Rule 19280 (Supplementary Table S6). Each of the nine neighbor sums had two alternative outputs. We activated each resulting rule at the same three fractions using 20 seeds under both boundaries. At $f = 0.20$, six rules produced majority crossings in all 20 seeds under each boundary, while the other twelve produced none. Effective changes included entries $c = 2, 3, 4, 5$ as well as the selected $c = 1$ change. Every change that replaced a Buy or Hold response with Sell at a neighbor sum from one to five produced a majority, as did the change from Buy to Hold at a sum of three; the same Buy-to-Hold change at a sum of one or two, every change that moved a response away from Sell, and every change at a sum of six or more produced none. The input being altered and its replacement output therefore mattered even when the number of changed entries was fixed. The screen was descriptive and used no adjustment for multiple comparisons. An observed 0/20 had a Wilson upper bound of 0.161, and 20/20 had a lower bound of 0.839. The screen identified other rules capable of producing majorities over the tested window without defining a universal class of activation thresholds.

Frozen band and non-monotone dealer response. Boundary decision widths changed as the quote moved, allowing the same price departure to elicit different trader responses over time (Supplementary Table S7). We removed this adaptation by freezing each path's decision deviation immediately before activation. Traders retained their k_j values and the moving center continued to update. The half-width $k_j \sigma_{1000}$ then stayed constant within a path, although the diagnostic deviation still evolved.

Each frozen arm shared its adaptive arm's exact prehistory, so widths could differ across seeds and capacities. At $f = 0.175$, the $h = 0.1$ minus $h = \infty$ majority contrast changed from 0.38 with adaptation to -0.02 with frozen widths. The paired difference of these contrasts was 0.40, with a 95% seed-bootstrap interval of [0.16, 0.64]. Removing width adaptation therefore removed the low-capacity advantage in the point estimates. Boundary Sell occupancy also changed, but the frozen capacity response remained non-monotone and individual probability intervals were wide. Width adaptation contributed to the observed advantage without isolating a unique microscopic cause of the entire capacity response.

Price-impact sensitivity. The price-impact coefficient set the quote change produced by a given executed imbalance (Supplementary Table S8). At $f = 0.20$, we varied λ over $\{0.005, 0.02, 0.05\}$ and capacity over $h \in \{0.1, 1, \infty\}$. A smaller coefficient reduced the immediate price movement from a given imbalance. Under the endogenous boundary, the changed quote also entered the decision bands that set perimeter states. All 50 seeds transitioned in every cell. Median drawdown increased with capacity at each coefficient, although its magnitude changed by a factor of seven to ten across the coefficient range. Dealer resources thus limited the price decline across a decade of impact choices, while collective selling persisted at this activation fraction throughout the sweep. These uncalibrated coefficients describe model sensitivity; they do not convert the simulated drawdowns into estimates for an empirical market. Both the amount of eligible flow the dealer could absorb and the impact assigned to filled imbalance determined the financial magnitude of the transition.

Horizon scaling and transition criterion. A longer observation window gave the altered local dynamics more time to reach the selected Sell fraction (Supplementary Table S9). Held-boundary checks used 50 seeds on $f = 0.140, 0.145, \ldots, 0.250$, extended to 0.300 at $L = 32$ to bracket the short-horizon response. We compared the original 3,000 steps with $T_{\text{post}} = 3000(L/64)^z$ for $z = 1, 2$ at the endpoint sizes. These windows were 1,500 and 750 steps at $L = 32$, and 6,000 and 12,000 at $L = 128$. Both scalings preserved the ordering of endpoint midpoints. At $L = 64$, increasing the horizon from 3,000 to 9,000 steps lowered the majority midpoint from about 0.178 to 0.167. Longer endpoint runs likewise continued to move the curve. Changing the crossing criterion to 0.4 or 0.6 retained the endpoint ordering and downward drift with time. At $L = 64$ and 3,000 steps, the three criteria gave midpoints of approximately 0.172, 0.178, and 0.186. The qualitative size response persisted, but the numerical midpoint depended on both the criterion and horizon. Neither these comparisons nor the final 500-step mean established equilibration at any size.

Baseline characterization and long no-switch paths. The pre-switch paths remained mixed, with similar mean Sell occupancy in the two windows preceding activation. Across 50 endogenous unlimited-dealer seeds, the mean interior Sell fraction was 0.2494 over steps 501–750 and 0.2503 over steps 751–1,000. In the latter window, the median within-seed return standard deviation was 0.0003, and the median within-seed 1st and 99th percentiles were -0.00068 and 0.00066. The mean zero-return fraction was 0.006. These summaries characterized the starting ensemble without establishing stationarity or empirical calibration. We extended the first ten production seeds under each boundary for 100,000 additional no-switch steps. None crossed a majority (0/10 per boundary; Wilson upper bound 0.278 for the ensemble event probability). The extension followed existing seeds and was not an independent replication. It showed that those paths remained below a majority over that longer interval, while leaving global stability and the distinction between long transients and attractors unresolved.

Fit and grid uncertainty. A monotone logistic curve supplied a second estimate of transition midpoint and width alongside empirical interpolation (Supplementary Table S10). Interpolation used the first upward crossing of each target probability and allowed local decreases in the sampled curve. The logistic fit imposed a monotone response as a working shape model. Both analyses used 2,000 bootstrap replicates that resampled whole seeds across fractions. Keeping each seed together preserved the dependence arising from shared initialization, so standard errors reflected seed variation without treating fractions as independent observations. Each susceptibility replicate reselected the maximum over the sampled grid. Its location remained a sampled fraction with finite local resolution. The resulting intervals measured uncertainty conditional on the refined grid and observation horizon. They did not account for discretization error, adaptive-grid selection uncertainty, or changes caused by extending the horizon. No susceptibility exponent was estimated from these finite-size summaries.