Agent-Based Models of Financial Markets: A Wolfram Perspective*

Jeffrey D. Varner
R.F. Smith School of Chemical and Biomolecular Engineering
Cornell University, Ithaca, NY 14850
jdv27@cornell.edu

Abstract

We study crash-like transitions in a deterministic, three-state cellular automaton on a two-dimensional grid. Interior agents update from their four nearest neighbors using one of two totalistic rules, and their mixture fraction f is the main control parameter. In Model A, QQQ minute data drive a subset of boundary cells. The model shows a sharp finite-horizon crossover from a mixed state to a sell-dominated state. Sell domains form locally and spread across the grid as coherent fronts. Mean-field and pair approximations do not reproduce this spatial behavior. Across random initial conditions and rule placements, onset times follow an inverse-distance empirical curve consistent with a nucleation interpretation. Model B replaces the external price with an endogenous feedback loop and filters desired orders through finite cash and share inventories. In the present experiments, budgets shift onset times less than twofold and instead alter post-onset domain growth. A paired cash intervention supports a price-to-boundary feedback mechanism, although the result is specific to the model's execution and market-maker assumptions. The finite rule space makes the same analysis possible for every totalistic rule in this model class.

*Repository: https://github.com/varnerlab/Wolfram-markets-abm-paper.git.

Preprint.

1 Introduction

Large price declines in financial markets are not always traceable to news of comparable size. The crash of October 1987 removed roughly a fifth of the value of the United States equity market in a single session, and the flash crash of May 2010 took about nine percent off the major indices and returned it within the hour, neither of them following an exogenous announcement of matching magnitude [1, 2]. The everyday statistics of returns raise the same difficulty in a milder form. Heavy tails, volatility clustering, and long memory in absolute returns persist across assets, sampling frequencies, and decades [3, 4], and a market driven only by independent external shocks would inherit the arrival statistics of those shocks rather than produce structure of its own. Both observations point to the same broad question, whether the interactions among traders can generate large moves on their own, and we address one narrow version of it here: can local interactions alone produce a spatially organized transition from mixed trading states to a sell-dominated state? If they can, then the interaction network matters, because it determines where a disturbance begins and how far it spreads [5]. We study this question in a deliberately minimal model and then test how finite agent budgets affect the onset and growth of the sell-dominated domain.

The ingredients for an answer already exist, though they have not been assembled. On one side, a generation of agent-based models (ABMs) has demonstrated that heterogeneous populations of chartists and fundamentalists can reproduce fat tails and clustered volatility [6, 7], that herding on random graphs is sufficient to generate power-law-tailed returns [8], and that minority-game dynamics [9] and adaptive strategy switching [10, 11] produce non-Gaussian fluctuations without any exogenous noise assumption. On the other side, a mature literature on lattice and spin models (Ising models with market coupling [12, 13], opinion dynamics on chains [14], and the three-state speculative model of Kaizoji [15]) has shown that nearest-neighbor interactions on a regular grid can produce phase transitions between ordered (crash) and disordered (efficient) regimes. Meanwhile, cellular automata (CA) have been applied to financial markets [16–18], but the rules in those studies were chosen by hand rather than drawn from an enumerated space, which leaves them resistant to systematic comparison. What is missing is a framework that combines *local* spatial dynamics on a regular lattice with a *complete, enumerable* space of behavioral rules and a *direct* connection to real market data. Wolfram totalistic CA rules [19] supply the first two of those at once. Information reaches a cell only from its immediate neighbors, so a sell cascade cannot appear everywhere at once but must travel from neighbor to neighbor as a spatially resolved wavefront whose origin and speed are both measurable. The space of such rules is also finite ($3^9 = 19{,}683$ rules), so agent behaviors can be screened by enumeration rather than designed by hand.

In this paper, we assemble those ingredients in a spatial cellular-automaton market. Agents occupy the cells of a two-dimensional grid, hold one of three states (Buy, Hold, or Sell), and update deterministically from their four cardinal neighbors under Wolfram totalistic rules (Figure 1). The interior mixes just two such rules, a majority/herding rule and a contrarian/destabilizer rule, at the minority (contrarian) fraction f that serves throughout as the single control parameter. We study two complementary configurations. In Model A (open), boundary agents observe real QQQ ETF minute-bar data through a Bollinger Band filter, which supplies an empirical boundary forcing. In Model B (feedback-closed), the boundary responds instead to an endogenous price formed from the grid's own net order imbalance, a reflexive loop that produces crashes without any external input, and every agent holds finite cash and shares, so that budgets can suppress the very orders sentiment demands. Model A showed a sharp finite-horizon crossover between $f = 0.14$ and $f = 0.15$, above which localized sell domains grew as spatially coherent fronts that neither a mean-field map nor a pair approximation reproduced. Model B kept that crossover under a budget, and across the tested cash levels the budgets moved onset by less than a factor of two while altering the domain's post-onset growth, a separation that a paired cash intervention traced to the boundary channel through which price information reaches the interior. We read the transition as nucleation and growth, a crash with a location, a start time, and a speed.

2 Related Work

Our work lies at the intersection of three research traditions: agent-based models of financial markets, lattice and spin models from statistical mechanics, and cellular automata theory. We review each strand below and identify where existing approaches stop short of the connection we propose.

2.1 Agent-Based Models and Stylized Facts

The quantitative study of financial market dynamics has been shaped by a tension between tractable theoretical frameworks and the empirical complexity of observed data. Mandelbrot [4] demonstrated that speculative price changes exhibit heavy tails beyond Gaussian predictions; Fama [20] showed that raw return series display negligible linear autocorrelation, consistent with the efficient markets hypothesis; yet Schwert [21] documented that the *magnitude* of returns exhibits positive autocorrelation decaying slowly over time, so that large price changes tend to be followed by large changes regardless of sign, a pattern known as *volatility clustering*. Bouchaud and Potters [22] further documented the asymmetric response of volatility to positive versus negative returns (the leverage effect). Cont [3] provides a systematic review confirming these regularities across asset classes, geographic markets, and sampling frequencies. Collectively, leptokurtosis, the absence of return predictability, and volatility clustering define the benchmark against which any model of market dynamics must be evaluated.

Agent-based models offer a bottom-up route to these stylized facts: rather than assuming the statistical structure of returns, they derive it from the simulated interaction of heterogeneous agents. The Santa Fe Artificial Stock Market [23] established the template of adaptive traders interacting through a centralized order book. Lux and Marchesi [6] showed that a two-type fundamentalist-chartist population reproduces heavy tails and volatility clustering as emergent properties, a result that established ABMs as a credible framework for studying market dynamics. Cont and Bouchaud [8] demonstrated that a simple random-graph herding structure alone is sufficient to generate fat-tailed return distributions, highlighting the role of coordination, rather than agent sophistication, in generating extreme events. The minority game of Challet and Zhang [9] introduced the paradigm of bounded-rational agents competing via strategic coordination, a framework subsequently shown to produce non-Gaussian fluctuations; Farmer and Joshi [10] demonstrated that standard trading strategies simultaneously deployed by many agents endogenously generate volatility clustering and excess kurtosis without any exogenous noise assumption; and Chen, Tan, and Zheng [24] showed that multi-level herding simultaneously reproduces sector structure and volatility clustering. Harras and Sornette [25] modeled myopic adapting agents whose local herding behavior inflates bubbles that eventually collapse, connecting the ABM literature to crash dynamics.

On the theoretical side, Hommes [11] surveys heterogeneous agent models incorporating adaptive expectations and evolutionary strategy selection; LeBaron [26] provides a systematic treatment of agent-based computational finance; Arthur [27] argued that heterogeneous agent interactions give rise to phenomena (persistent disequilibrium, emergent structure, tipping points) inaccessible to equilibrium theory; and Farmer [28] demonstrated that quantitative ABMs can produce time-series predictions competitive with mainstream macroeconomic models. Battiston et al. [5] made the case that complexity science and ABMs should inform financial regulation. These models share a limitation, however: agents interact either globally (mean-field coupling) or on ad hoc network topologies, making it difficult to disentangle the contributions of local spatial contagion from global coordination. The present work addresses this gap by placing agents on a regular lattice with strictly local CA rules, so that every cascade must propagate neighbor-by-neighbor through physical space.

2.2 Lattice Models, Spin Models, and Phase Transitions

The mapping of financial agent interactions onto lattice models draws on the statistical mechanics tradition, where the mathematical machinery for analyzing cooperative phenomena on regular grids (phase transitions, critical exponents, universality classes) is well developed. Ising-type models map agent states (buy/sell) onto spins and impose nearest-neighbor coupling on a lattice [29, 30]. These models naturally capture local herding and can exhibit phase transitions between ordered (bubble/crash) and disordered (efficient) regimes. Bornholdt [12] introduced a two-dimensional Ising model with a global coupling term representing market impact, producing intermittent bubbles and crashes from the interplay of local herding and global frustration across scales. Sznajd-Weron and Weron [14] proposed a one-dimensional opinion model for price formation in which local consensus propagates outward, demonstrating that even the simplest spatial rule can generate nontrivial collective dynamics. Krawiecki et al. [13] showed that attractor-bubbling dynamics on a lattice generate volatility clustering with realistic scaling exponents. Kaizoji [15] extended the Ising framework to three states (buy/hold/sell) to model speculative activity, which is the closest antecedent to our three-state CA formulation.

The crash literature provides the other essential ingredient. Sornette [1] argued that large crashes are the culmination of cooperative herding behavior and can be modeled as critical points analogous to phase transitions in physical systems. The self-organized criticality (SOC) framework of Bak, Tang, and Wiesenfeld [31] suggests that complex systems may naturally evolve toward a critical state where scale-free avalanches occur without parameter tuning, and Bak and Paczuski [32] applied this idea to economic systems. Gabaix et al. [33] derived power-law tails from the size distribution of large institutional trades, providing an alternative mechanism for heavy tails that does not require agent heterogeneity. Scheffer et al. [34] proposed generic early-warning signals, rising autocorrelation and variance, for systems approaching a tipping point, a framework that has been applied to ecosystems, climate, and financial markets [5]. Farmer et al. [35] demonstrated that even zero-intelligence agents can reproduce aspects of market microstructure when placed in a realistic order-book mechanism, suggesting that the structure of interaction may matter more than the sophistication of individual decision rules.

Statistical-mechanics accounts describe crashes as collective transitions with slow relaxation or endogenous amplification [36], a distinction that has been tested empirically in other complex systems by separating endogenous from exogenous shocks [37], and related market models study feedback between funding and market liquidity, the 2010 Flash Crash, circuit breakers, and anticipated bailouts [2, 38–40]. Our model is not a calibrated version of those settings, and it differs from an Ising lattice in two further ways, since we replace the single coupling constant with Wolfram totalistic CA rules, which provide a richer and fully enumerable space of $3^9 = 19{,}683$ local update functions, and we couple the lattice to an external or endogenous price signal through boundary agents. The external series supplies an empirical forcing. It does not validate the model's prices. The fixed-seed sweep shows a sharp finite-horizon crossover rather than a parameter-free SOC transition, and the narrower advantage of the construction is that the same deterministic realization can be rerun with and without a cash intervention.

2.3 Cellular Automata in Complex Systems and Finance

Wolfram [41] classified one-dimensional CA into four behavioral classes (fixed points, periodic orbits, spatial chaos, and complex/undecidable dynamics) and later provided an exhaustive treatment of the subject [19]. Packard and Wolfram [42] extended this classification to two-dimensional totalistic rules on Von Neumann neighborhoods, demonstrating rich pattern-formation dynamics including self-replication, spiral waves, and spatial chaos. The classification program reveals that even deterministic, strictly local rules can produce dynamics ranging from trivially ordered to computationally universal, a property that makes CA attractive as a modeling framework wherever the goal is to understand how microscopic rules generate macroscopic complexity.

Applications to financial markets, however, are few, and we are aware of three that apply CA directly to market dynamics. Wei et al. [16] proposed a CA model for securities markets in which agent states evolve based on local majority rules and demonstrated qualitative agreement with empirical return distributions. Qiu et al. [17] developed a CA model specifically targeting abnormal fluctuations, showing that local interaction rules can produce heavy-tailed returns. Yang et al. [18] simulated financial markets using CA with heterogeneous update rules and showed that the ratio of contrarian to trend-following agents controls the transition between calm and volatile regimes, a result qualitatively consistent with our phase diagram, though their model uses a different rule framework and does not characterize the transition as a sharp phase boundary. None of these studies uses the full Wolfram totalistic rule framework, and none couples the CA to real market data or resolves the spatial structure of crash cascades.

The present work embeds agent dynamics in the complete Wolfram totalistic rule space (3^9 rules), which allows systematic enumeration and comparison rather than rule design by hand; it couples the CA to real market data (Model A) or an endogenous reflexive price (Model B) through boundary agents, bridging the gap between abstract CA dynamics and empirical financial time series; and it identifies a discontinuous transition once the contrarian fraction crosses a threshold, resolving its spatial structure as nucleation of sell cascades from localized seeds followed by wavefront propagation, which connects the CA dynamics to the tipping-point framework of Scheffer et al. [34].

3 Methods

We first establish the formal ingredients of the model, the CA rule framework that governs interior agent dynamics, the Bollinger Band filter that translates price signals into edge-agent states, and the specific rules used in this study, and then describe the two simulation configurations and the phase-sweep protocol used to map the crash transition.

3.1 Three-State Totalistic Cellular Automata

We consider a two-dimensional grid of $N = N_r \times N_c$ agents, each in one of three states: Buy ($s = 0$), Hold ($s = 1$), or Sell ($s = 2$). At each discrete time step, every interior agent updates its state deterministically based on the states of its four cardinal (Von Neumann) neighbors. For an agent at position (r, c) with neighbor states $\{s_N, s_S, s_E, s_W\}$, the neighborhood sum is:

$$c_{r,c} = s_N + s_S + s_E + s_W \in \{0, 1, \ldots, 8\}. \tag{1}$$

A *totalistic rule* is a mapping $\mathcal{R} : \{0, 1, \ldots, 8\} \to \{0, 1, 2\}$ that assigns a next state to each of the nine possible neighborhood sums. Since each of the nine outputs can take three values, there are $3^9 = 19{,}683$ distinct rules, each identified by its decimal index:

$$\text{Rule index} = \sum_{j=0}^{8} a_j \cdot 3^j, \tag{2}$$

where $a_j = \mathcal{R}(j)$ is the next state when the neighborhood sum equals j. This sum uses the ordered numerical encoding Buy $= 0$, Hold $= 1$, and Sell $= 2$. Relabeling the states would change the rule. The ordering is therefore part of the model, not a neutral label choice. This rule space is finite, enumerable, and complete: every possible deterministic, translation-invariant, totalistic update function on the Von Neumann neighborhood is represented exactly once.

3.2 Edge Agent Decision Model

Boundary cells serve as *edge agents* that inject external information into the grid (Figure 1A, green cells). A fraction f_e of boundary cells, selected uniformly at random with a fixed seed, are designated as edge agents. The remaining boundary cells stay fixed in the Hold state. The CA update acts only on strictly interior cells.

Each edge agent classifies the current price P_τ relative to an exponential moving average (EMA) and its associated volatility band:

$$\hat{S}_\tau = \hat{S}_{\tau-1} + \alpha\big(P_\tau - \hat{S}_{\tau-1}\big), \tag{3}$$

$$\hat{\sigma}_\tau = \hat{\sigma}_{\tau-1} + \alpha\big(|P_\tau - \hat{S}_\tau| - \hat{\sigma}_{\tau-1}\big), \tag{4}$$

where $\alpha = 1/60$ corresponds to a 60-minute smoothing window. The decision rule is a three-way classifier:

$$e_\tau = \begin{cases} 0 \text{ (Buy)} & \text{if } P_\tau < \hat{S}_\tau - k\,\hat{\sigma}_\tau, \\ 1 \text{ (Hold)} & \text{if } \hat{S}_\tau - k\,\hat{\sigma}_\tau \le P_\tau \le \hat{S}_\tau + k\,\hat{\sigma}_\tau, \\ 2 \text{ (Sell)} & \text{if } P_\tau > \hat{S}_\tau + k\,\hat{\sigma}_\tau, \end{cases} \tag{5}$$

where k controls the width of the hold band. When the price is within k standard deviations of its moving average, the edge agent holds; when the price breaks out above (below) the band, it sells (buys). This is a standard mean-reversion signal widely used in technical analysis, and its simplicity is deliberate: the edge agents provide a minimal, interpretable channel through which real market information enters the grid.

3.3 Rule Taxonomy

Three rules are used in this study (Table 1). Rule 19280 is a majority/herding rule: agents tend to adopt the dominant state in their neighborhood, promoting coordination and spatial coherence. Rule 19286 is a contrarian rule that differs from 19280 at a single neighborhood sum, $c = 1$, the case with three Buy neighbors and one Hold neighbor. A Rule 19280 agent responds to this configuration

by *buying* ($a_1 = 0$), reinforcing the local majority; a Rule 19286 agent instead *sells* ($a_1 = 2$), acting against the prevailing state. This single-digit difference in the base-3 encoding is the minimal perturbation that converts a stabilizing agent into a destabilizing one, and it is this minimal difference that drives the phase transition studied below.

3.4 Model A: Open (Boundary-Driven) Simulation

In Model A, edge agents observe real market data, namely one-minute OHLC bars for the QQQ ETF (Invesco QQQ Trust, January 2 through March 28, 2024, $T = 23{,}790$ time steps). The vendor feed is no longer publicly available, so an archival copy travels with the repository alongside its SHA-256 checksum and a record of the bar accounting, 61 trading days of exactly 390 regular-session bars each. At each step τ the typical price $P_\tau = (H_\tau + L_\tau + C_\tau)/3$ is computed from the OHLC bar, each edge agent then evaluates Eq. (5) using the online EMA and volatility from Eqs. (3)–(4), and all interior agents update simultaneously via their assigned CA rule. The grid is initialized to all-Hold ($s = 1$), so that any Buy or Sell activity must originate from the boundary and propagate inward.

For the heterogeneous (crash) model, a fraction f of interior agents are assigned Rule 19286 (contrarian) uniformly at random, while the remaining $1 - f$ agents use Rule 19280 (majority), and the minority agent positions are fixed for the duration of the simulation (quenched disorder), so that the only control parameter is the mixture fraction f. The default grid is 256×256 ($N = 65{,}536$ agents), edge agents comprise 25% of boundary cells ($f_e = 0.25$, yielding 255 active edge agents), selected with a fixed random seed for reproducibility, the EMA smoothing parameter is $\alpha = 1/60$ (a 60-minute window), and the band width is $k = 1.0$ standard deviations.

3.5 Model B: Feedback-Closed (Endogenous Price) Simulation

In Model B, there is no external price feed. The price is generated endogenously from the grid's own collective state, and every trading cell, each interior agent together with each active edge agent, now carries a *book* that gates whether its sentiment can be acted upon: a cash balance initialized to `initial_cash` (the sweep variable, default 1,000), a share inventory initialized to 10 shares acquired at the reference cost basis $P_0 = 100$, and a running average cost basis that tracks the price paid across subsequent fills. Sentiment still flows by the same totalistic rule as in Model A, which is what keeps the two models' interiors directly comparable, but a cell now expresses that sentiment through orders: a cell in the Buy state submits a one-share buy order every step it remains in that state, a cell in the Sell state submits a one-share sell order, and a Hold cell submits nothing.

Orders do not clear against one another. They fill at the standing quote $P_{\tau-1}$ against an unmodeled infinite-inventory market maker in the spirit of the price-impact accounting of Kyle [43] and the market-clearing conventions of Farmer and collaborators [35]. A buy fills if and only if the cell holds cash $\geq P$; a sell fills if and only if the cell holds at least one share and, when the disposition toggle is on, the standing quote exceeds the cell's average cost basis, $P > \bar{c}_i$. This cost-basis sell rule is the *disposition effect* of Shefrin and Statman [44], the reluctance to realize losses that Odean later documented in the trading records of individual investors [45]. We treat it not as a modeling assumption but as a behavioral bias and hence a sweep factor, toggled on and off across campaigns. Every order that the book cannot honor is tallied as a *suppressed* order rather than executed: budget constraints suppress order *flow*, never the underlying CA states. The interior uses the same update rule as Model A, while cash and inventory throttle its market footprint. Because the market maker lies outside the modeled population, cash and shares are not conserved among grid agents. Model B is closed with respect to price feedback, not economically closed.

Because fills, not states, now drive the price, the imbalance is measured on executed flow rather than on the raw sentiment census:

$$\varphi_\tau = \frac{(\text{executed buys})_\tau - (\text{executed sells})_\tau}{N_{\text{traders}}}, \qquad N_{\text{traders}} = N_{\text{interior}} + N_{\text{active edge}}, \tag{6}$$

and the price updates multiplicatively from this executed imbalance,

$$P_\tau = P_{\tau-1} \cdot \exp\big(\lambda\, \varphi_\tau\big), \tag{7}$$

where $\lambda > 0$ is the price-impact parameter controlling how strongly the aggregate order flow moves the quote. Edge agents are full participants in φ, since their fills enter the numerator on the same footing as interior fills, and, to break the all-or-none lockstep that a single shared band would impose,

each edge cell draws its own band half-width $k_i \sim U\big(k(1-\delta),\ k(1+\delta)\big)$, where δ is the band-spread parameter; the edge cells then evaluate Eq. (5) against the newly generated price P_τ, with the online EMA and volatility recursions of Eqs. (3)–(4) observing that same new price. The result is a fully self-contained loop: the grid's executed orders generate the price, the price drives the edge agents, and the edge agents inject signals back into the grid.

To probe that loop causally we expose a deterministic intervention hook, a function of the current step and the full collection of books, (t, books), invoked at the top of step t before any orders execute, through which we implement mid-run cash shocks by scaling every cell's cash by a common factor. These are uniform wealth shocks within the model, not conventional order-book liquidity operations. The hook takes no randomness of its own, so the same seed still yields an identical trajectory up to the moment of intervention. This same-seed, identical-trajectory contract is what the paired counterfactuals rely on to attribute any downstream divergence to the liquidity shock alone.

3.6 Phase Sweep Protocol

To map the finite-horizon crossover, we swept the minority fraction f from 0.0 to 0.50 in two passes. A coarse sweep with $\Delta f = 0.05$ covered the full range; a fine-grained sweep with $\Delta f = 0.01$ resolved the transition region $f \in [0.08, 0.22]$. For each value of f, the full T-step price series drove a 128×128 grid (chosen for computational efficiency), and we recorded the sell fraction at every time step. Three summary statistics characterize each run: mean sell fraction, peak sell fraction, and crash frequency (the fraction of time steps at which the sell fraction exceeds 0.50). Each value of f used one fixed disorder realization. The sweep therefore brackets a finite-horizon crossover; it does not estimate a thermodynamic critical point.

3.7 Model B Nucleation Campaigns

The endogenous-price machinery of Model B lets us treat a market crash as a nucleation event and interrogate it with a family of campaigns. Unless a campaign states otherwise, all of them share one configuration, a 64×64 grid, price impact $\lambda = 0.02$, nominal band half-width $k = 1$, band spread $\delta = 0.5$, the disposition rule switched off, an initial quote $P_0 = 100$, and deterministic seeds, and they differ only in the minority fraction f, the initial cash, and the horizon over which we let the loop run, which is $1{,}000$ steps for the seed sweep, $1{,}500$ steps for the budget-toggle and boundary-versus-bulk campaigns, and $2{,}000$ steps for the escape-time campaign. Two definitions, fixed once and reused by every campaign, let the campaigns be read together: we label sell clusters by breadth-first search over the 4-connected Sell cells of the grid, and we define the crash onset t_{on} as the first step at which the largest sell cluster reaches 5% of the grid. Cluster labeling runs on the full grid in every campaign but one, since the boundary-versus-bulk birth analysis instead labels on the interior alone, so that a cluster attached to the active-edge ring can be told apart from one born in the bulk, and because onset and cluster size are measured identically everywhere else, front velocities and onset times reported by different analyses can be read on a common footing.

The seed sweep compares variation across configurations with variation across disorder realizations. We ran three configurations, a *stable* regime at $f = 0.10$ with cash $1{,}000$, a *freeze* regime at $f = 0.20$ with cash 500, and a *collapse* regime at $f = 0.20$ with cash $5{,}000$, across 20 seeds each. The budget toggle then turns the deterministic hook into a controlled experiment. For each seed, we used the control trajectory to identify t_{on}. In the paired counterfactual, every cell's cash is scaled at the start of that step by $\times 10$ (an injection) or by $\times 0.1$ or $\times 0.01$ (two withdrawal doses). Histories are bit-identical through step $t_{\text{on}} - 1$, so the intervention immediately precedes the control's first detectable onset frame. We then measured outcomes in an onset-aligned window of $W = 400$ steps and took the front velocity to be the endpoint difference quotient, the change in largest-cluster fraction between onset and onset$+W$ divided by W, rather than a fitted slope. A pair is excluded from an arm when its control run had no onset at all or when t_{on} fell too late to leave a full window before the horizon ($t_{\text{on}} >$ horizon $- W$).

The boundary-versus-bulk campaign asks where the winning cluster is born: we tracked the founding cell of the onset cluster and compared its location against the analytic null for spatially uniform births, under which the fraction of births falling within distance 2 of the boundary is 0.125, and we complemented this observational arm with an interventional one that stamps a 3×3 sell patch onto the grid at boundary-adjacent, mid, and center placements, each paired against an unstamped control.

The onset-time campaign measures how the waiting time changes with f at three cash levels, cash $\in \{500, 1{,}000, 5{,}000\}$. A pilot pass swept $f \in \{0.12, 0.15, 0.18, 0.20, 0.25, 0.30\}$ at 20 seeds, and having localized the onset cliff to the interval between 0.15 and 0.18 at all three cash levels, a refinement pass added $f \in \{0.156, 0.162, 0.168, 0.174, 0.177\}$ at 40 seeds. We took the escape time to be $\tau = t_{\mathrm{on}}$, right-censored at the 2,000-step horizon and never discarded, so runs without onset still inform the Kaplan–Meier estimates. We fitted the empirical form $\ln \tilde{\tau} = A + B/(f - f_0)$ separately at each cash level using the finite Kaplan–Meier medians, with the pseudo-threshold f_0 left free, and obtained 95% confidence intervals on (A, B, f_0) from 2,000 bootstrap resamples of the seeds within each f.

A companion velocity–suppression dose-response holds $f = 0.20$ and extends the cash ladder with the levels 250 and 2,000 so that five doses are probed at 20 seeds each. For every run we formed, over the window $[t_{\mathrm{on}} + 50,\ t_{\mathrm{on}} + 400]$, a front velocity as the OLS slope of the front extent and a suppression load as the mean *absolute* number of suppressed buy orders per step, and we tested their association by Spearman rank correlation with a 10,000-draw permutation p-value, excluding and separately counting any run that is censored before onset or whose window endpoint $t_{\mathrm{on}} + 400$ would exceed the horizon.

Finally, to place these targeted campaigns against the model's full parameter dependence, we drew on a large factorial sweep whose crash statistics anchor the Model B regime discussion: the full crossing of $f \in \{0, 0.1, 0.15, 0.2, 0.3\}$, $\lambda \in \{0.005, 0.01, 0.02, 0.05\}$, cash $\in \{500, 1{,}000, 5{,}000\}$, disposition $\in \{\text{off}, \text{on}\}$, band spread $\delta \in \{0, 0.5, 0.9\}$, and band half-width $k \in \{0.5, 1, 2\}$ yields 1,080 configurations which, replicated across 10 seeds of 2,000 steps, comprise 10,800 runs. The full design and its variance decomposition are deferred to Appendix B.

4 Results

Model A provides an unconstrained spatial baseline. Prices arrive from outside the model, and CA states are not filtered through budgets. We first inspected representative grid trajectories and then located the finite-horizon crossover in a fixed-seed sweep.

We simulated the 256×256 agent grid under Model A across the minority fraction and observed a sharp finite-horizon crossover with spatially coherent front growth above it (Figure 2). Three values of f fix the picture, each sampled at four logarithmically spaced times. In the homogeneous model, $f = 0$, the grid settled into a fine-grained mixture of Buy and Sell agents that fluctuated locally, never coordinated over large distances, and held the sell fraction near 0.26 throughout. At $f = 0.15$, the first sampled value above the crossover bracket, the grid at first looked the same, but at $t = 8{,}236$ a sell cascade nucleated from a localized seed and spread outward as a coherent wavefront, growing concentrically until it engulfed the grid by $t = 23{,}790$. The crash had a spatial history, beginning at one location, moving through a physical medium, and reaching different parts of the grid at different times, which is nucleation and growth in the classical sense rather than the instantaneous global collapse of a mean-field model. At $f = 0.20$ the same sequence ran faster and the grid was sell-dominated by $t = 2{,}851$. Sweeping f from 0.0 to 0.50 at a fixed seed on a 128×128 lattice placed the change between the sampled values $f = 0.14$ and $f = 0.15$. At and below $f = 0.14$ the system was stable, with a mean sell fraction near 0.26, close to the $1/3$ expected of a uniform random mixture of three states, and no crashes. At $f = 0.15$ the mean sell fraction jumped to $\sim$0.50, saturating near 0.91 by $f = 0.20$, where crash frequency rose above 95%. The dashed marker at $f = 0.145$ is the midpoint of that bracket and not an independently fitted critical point. What the sweep describes is a coordination failure built from two deterministic rules and nothing else, with no noise, no learning, and no strategy switching. A small contrarian minority was harmless, because the majority agents (Rule 19280) damped local sell signals by herding toward the dominant state, and the sell clusters that did form died before reaching the size at which herding turns from damper into amplifier. Inside the bracket the contrarians (Rule 19286) grew numerous enough to seed clusters the majority could no longer absorb, the local majority tipped from buy to sell, and the same herding mechanism amplified the cascade it had been suppressing. One fixed seed on one lattice cannot say whether that bracket belongs to the rule mixture or to the grid we happened to simulate, so we repeated the measurement across five lattice sizes with replicated disorder.

Having located the crossover, we next asked whether it sits at a fixed value of f (Figure 3). It does not. We repeated the sweep on lattices from 32×32 to 512×512, twenty disorder seeds at each size,

with periodic boundaries and no price forcing, and larger lattices crashed at lower contrarian fractions. At $f = 0.14$ no run out of twenty converted at 64×64, half converted at 128×128, and all converted at 256×256 and above; certain conversion required $f \approx 0.155$ at 64×64 but only $f \approx 0.14$ at 512×512. That direction is the opposite of the usual intuition, because whether the lattice crashes is not an average but is settled by the most favorable local region anywhere on the grid: one sell domain surviving to critical size converts the whole system, and a larger lattice offers more candidate regions. It is the size effect of brittle fracture, where a bigger specimen fails at lower stress than a small one of the same material because it more likely contains a critical flaw. Size does not make a stable configuration unstable, since the sell state is absorbing and every lattice above the spreading threshold converts eventually; it decides only what has happened inside a fixed window. At one absolute time, $t = 17{,}235$ steps at $f = 0.15$, the 64×64 lattice was still mixed while the 512×512 lattice had already converted (Figure 3A). The averaging intuition holds instead for the timing. The coefficient of variation of the conversion time fell from 0.66 at 32×32 to 0.035 at 512×512 at $f = 0.20$, halving on each doubling of the linear size and so tracking $N^{-1/2}$ in the agent count N, and above $f \approx 0.17$ the median conversion time was size-independent to within 9% over a 16-fold range in area (Figure S7). Writing the converted fraction as $X(t) = 1 - \exp[-(Kt)^n]$ identifies that clock: in two dimensions $n = d + 1 = 3$ if new domains keep appearing and $n = d = 2$ if all are present at the outset and the transformation is growth alone. Fitted per run, n ran from 2.54 at $f = 0.16$ to 1.91 at $f = 0.20$ on the largest lattice ($R^2 = 0.994$ to 0.999), and the rate constant did not depend on size: read as a characteristic time, $1/K$ was 1445, 1363, 1367, and 1358 steps from $4{,}096$ to $262{,}144$ agents at $f = 0.18$, a spread of 6%, and rescaling each run by its own fitted K collapsed all 320 runs onto one curve. An exponent near 2 says the vulnerable regions were fixed before the clock started, as the construction implies: contrarian sites are assigned once at setup and never move. Two limits bound the instrument. The 32×32 lattice is anomalous, because a 100-cell domain, the size used to mark local clustering, already spans a tenth of it, so we kept it out of the conversion-curve fits; and that criterion is not a scale-free detector of nucleation, since the mixed state threw up clusters of the same size unprompted, its median largest-ever sell cluster reaching 98 cells at $f = 0.145$ (Figure 4C). Rate conclusions therefore rest on the conversion curve, which needs no cluster identification. A threshold that moves with size implicates spatial structure, so we next removed the space.

We next compared the lattice against two low-order closures and found that neither reproduced the crossover (Figure 4). Local order was tracked with the nearest-neighbor alignment statistic e of a three-state Potts model, the negated fraction of neighbor bonds whose two sites share a state, so that a single-state lattice reads -1 and a maximally disordered one about $-1/3$. It measures order without generating the dynamics, since the rules act on the numerical sum of neighbor labels rather than on counts of matching neighbors, and one entry is decisive: a cell with four Buy neighbors outputs Sell, so the all-Buy configuration is *unstable* and the "buy-dominated" phase is not a static equilibrium but a dynamic mixed-state attractor at $p_0 \approx 0.49$, $p_1 \approx 0.24$, $p_2 \approx 0.26$. On a 128×128 lattice it held at $e \approx -0.37$ and site entropy $S_1 \approx 1.06$ below the crossover and fell to -0.93 and 0.18 above. Replacing the lattice with a well-mixed population, each agent's four neighbors drawn independently from the global state distribution, gave a map (Appendix A) whose only stable fixed points at every f were the all-Sell absorbing state and a buy-dominated state, at $p_0^* \approx 0.63$ and $p_2^* \approx 0.19$ when $f = 0$ and locally stable as a spiral with $|\lambda| \approx 0.95$. That branch never crossed the sell-majority threshold and its sell fraction drifted up by only ~ 0.04 over the whole range of f, so mean field predicted no transition. Restoring nearest-neighbor correlations did not help. The pair approximation [46] tracks pair densities ρ_{ab} under the closure $q(s \mid a) = \rho_{as}/\rho_a$, giving the synchronous update

$$\rho'_{bb'} = \sum_{a,a'} \rho_{aa'}\, T_b(a, a')\, T_{b'}(a', a), \tag{8}$$

with kernel $T_b(a, a')$ derived in Appendix A. Its Jacobian at the uncorrelated fixed point had spectral radius strictly below unity at every sampled f, falling from 0.95 at $f = 0$ to 0.71 at $f = 0.5$, and correlated initial conditions all converged to that fixed point or to all-Sell, with the spatial mutual information $I = 2S_1 - S_2$ decaying to zero. Other finite-order closures are not excluded, but single-site marginals and the tested pair correlations are insufficient, because they erase exactly what the lattice supplies. An isolated contrarian lifted the sell density at its own site to $\rho_2 \approx 0.42$ against a background of 0.26, an excess gone within one to two lattice spacings, so the seeds are point-like and the transition needs many together. On 64×64 lattices the mixed state survived 10^5 steps in all twenty realizations below $f \approx 0.14$, while above $f \approx 0.155$ all collapsed past a sell fraction of 0.90,

the mean lifetime falling from $\bar{\tau} \approx 24{,}000$ at $f = 0.155$ to $\bar{\tau} \approx 86$ at $f = 0.30$. Those onset times follow the inverse-distance form of Kramers-type escape problems [47, 48],

$$\ln \bar{\tau} \;=\; A + \frac{B}{f - f_0}, \tag{9}$$

with $A = 2.78$, $B = 0.389$, and $f_0 = 0.101$ over twelve points ($R^2 = 0.991$; collected with the Model B fits in Table 2). The fit drops values at which not every seed collapsed within the horizon and is not derived from the CA dynamics, so f_0 is an extrapolated pseudo-threshold, not proof of absolute stability, and the curve is phenomenological, not a thermal escape law. Read as an extrapolation it places the crossover at $f_c \approx 0.162$ for a 10,000-step window and $f_c \approx 0.195$ for a 1,000-step one. Model B keeps these rules and this quenched disorder but adds endogenous price feedback with per-agent cash and inventory limits, letting us ask whether the crossover, the onset-time form, and the front's growth survive a budget; those experiments are sized to detect large shifts, so overlapping intervals cannot exclude smaller budget effects.

We next asked whether the crossover survived finite budgets and where on the grid the winning domain was born (Figure 5). The 10,800-run factorial sweep (Appendix B) placed a sell-dominated regime between the sampled values $f = 0.15$ and $f = 0.20$ at every tested price impact, so a budget did not remove the crossover. A variance decomposition attributed 94–97% of each summary metric's variance to differences among the swept configurations and at most 5.4% to disorder seeds within a configuration; that decomposition is descriptive, and it does not identify which individual parameter causes the difference. The classifier assigned every run to one of two realized regimes and left the other three candidate classes empty, 6,589 runs (61%) called oscillatory and 4,211 (39%) crash-prone, where "crash-prone" is a grid-state label, the Sell fraction standing at or above 0.5 on at least 1% of steps, and not a price drawdown. The modal crash-prone outcome was a liquidity *freeze* in which the interior locked near a Sell fraction of 0.88 while suppressed orders limited price movement, its median maximum drawdown about 5% and its median minimum price 100.00, while the price collapses formed a tail of the class rather than its center (Figure S1). Above threshold the Sell cells gathered into one dominant domain. Across 20 disorder seeds per configuration the at-peak ratio of the largest cluster to the total Sell population held a median of 0.022 (IQR 0.018–0.027) in the stable regime, whose Sell cells were scattered into roughly 500 dispersed fragments the largest of which never exceeded about 0.5% of the grid, against median ratios of 0.26 in the freeze and 0.44 in the collapse with zero overlap against the stable band, and the above-threshold trajectories formed time-shifted sigmoids whose onset varied more than their eventual shape (Figure S2). That gap is the observed signature of localized domain growth rather than a synchronous grid-wide flip, and it is consistent with variable waiting times without by itself establishing a Kramers escape process. The domain was born in the bulk. Across 40 boundary-versus-bulk runs, 39 of which nucleated within the horizon, the founding cell of the onset cluster fell in the boundary band at a pooled fraction of 0.077, slightly below and statistically indistinguishable from the analytic null of 0.125 for spatially uniform births within distance 2 of the edge (binomial $p = 0.47$; the freeze configuration taken alone gave $p = 1.0$), the median founding cell sat a distance 12 into the deep bulk, and the typical founding was a single Sell cell (median birth size 1 cell), with a median birth step of 256 and individual onsets spanning 75 to 1,115. One collapse run traced that whole history, from a single founding cell 15 cells deep in the interior at $t = 203$, through 208 cells at onset, to a 973-cell domain by $t = 528$ (Figure 5A). The winning nucleus was an interior single-cell fluctuation that happened to survive and grow, not a structure templated by the boundary, and the one place the active-edge ring asserted itself was *after* birth: pre-onset attachment of the growing cluster to the ring occurred in 6 of 19 collapse runs and never in the 20 freeze runs. A boundary that touches growth but not birth sharpens the question of which of the two a budget acts on.

Having established a bulk-born nucleus, we next tested whether budgets act on its birth or on its growth (Figure 6). Onset times proved equivalent within a factor of two across a tenfold cash range. We fitted the same empirical form, $\ln \tilde{\tau} = A + B/(f - f_0)$, to Kaplan–Meier median onset times at cash 500, 1,000, and 5,000, on the 5, 6, and 6 cells whose onset rate reached one half (Table 2; Appendix C); the few points and the three free parameters make the curves descriptive rather than confirmatory. The bootstrap intervals on both fitted shape parameters overlapped between the tightest and the loosest budget, and the onset cliff sat between $f = 0.15$ and $f = 0.18$ at every cash level (Figure S4), but an overlap of intervals is a bound and not a demonstration of equality, the B point estimates differing by 2.5× across the cash ladder and the cash-500 lower bound of 0.000 on f_0 being the physical boundary of the fit domain rather than a data quantile. We therefore tested the difference

directly, with a 4,000-draw bootstrap paired across the three arms on the disorder seeds they share at every f, comparing the predicted log median onset time, which the three fitted parameters jointly determine, against a margin fixed in advance at a factor of two, $|\Delta \ln \tilde{\tau}| < \ln 2$, by two one-sided tests at $\alpha = 0.05$. Every pairwise comparison met that margin at every shared f. The largest point difference anywhere was 0.354 in $\ln \tilde{\tau}$, a factor of 1.42, between cash 500 and $5{,}000$ at $f = 0.20$, with a 95% interval of $(-0.098,\ 0.718)$, while the 95% interval on the difference in B between cash 500 and $1{,}000$ spanned $(-0.298,\ 0.965)$, so the individual parameters remain poorly determined by comparison. One comparison is worth flagging without promoting: between cash $1{,}000$ and $5{,}000$ near the threshold the intervals exclude zero at $f = 0.177$ and $f = 0.18$, $(0.014,\ 0.433)$ and $(0.021,\ 0.398)$, placing onset roughly 19% later under the tighter budget, but those are adjacent points on the same pair of fitted curves and so are one signal rather than two, and neither survives correction for the fifteen comparisons in the family. Birth was equally indifferent to a nucleus handed to it. The stamped 3×3 sell patch, set at boundary-adjacent, mid, and center placements in both the freeze and the collapse configurations, never founded the winning domain: the dominance rate was $0/20$ in all six cells of that design, the patch's forward lineages grew only to 9–15 cells before dying out and were never alive at onset, no placement advanced the onset, and the one nominal placement effect ran, if anywhere, in the direction of delay (Figure S5). The interior does not lack seeds, and what is scarce above threshold is the *growth* that would make one matter. Growth did track the budget. One seed run at both budgets put the dominant domain at 873 cells 400 steps past its own onset under the loose budget against 516 cells under the tight one (Figure 6A). Across the five doses at $f = 0.20$ (cash 250–5,000; $n = 97$ runs included, 3 excluded by the window rule) the front velocity fell as the absolute suppression rate rose, at a pooled Spearman $\rho = -0.372$ with permutation $p = 0.0003$, and the relationship was not a between-dose confound, holding within every dose separately, with ρ between -0.38 and -0.52, and surviving the removal of any single dose (leave-one-dose-out $p \leq 0.0014$). The median suppression load varied little across the ladder (517–595 orders per step) and the median velocity was U-shaped in cash (Figure S4), so initial cash was a weak proxy for the suppression a given front actually experienced, and it was the realized suppression that tracked the slowing of the front. A brake that acts on growth is best tested by releasing it.

We finally used the budget toggle as a paired intervention to test how a cash shock reaches the front (Figure 7). Releasing the brake made the crash faster, and tightening it did nothing at all. Scaling cash $\times 10$ in the tight freeze configuration (cash 500) just before detected onset *accelerated* the front in 14 of 19 seeds (sign test $p = 0.064$; we report the count and the p-value as they stand), the median front velocity rising from 1.39 to 3.55×10^{-4} grid fractions per step, the median front extent at $t_{\mathrm{on}} + 400$ from 0.112 to 0.192, and the suppressed buy orders falling by 29%, with 19 of the 20 seed pairs meeting the window rule in every arm. In one such pair the two arms held the same 190-cell domain one step before the toggle, after which the injected arm grew to 198 cells at $t_{\mathrm{on}} + 400$ while the control's largest cluster broke up to 38 (Figure 7A). That is a causal result within the paired simulation, not evidence about real-market liquidity policy, and it runs opposite to the naive mechanism, which makes it diagnostic. Injected cash cannot feed the front directly, since the front is a *sell* wave, sells need shares rather than cash, and budgets suppress order flow and never sentiment, so the only conduit to interior sentiment is the loop's own: executed flow moves the price, the price crosses the edge bands, and the band crossings set the state of the active-edge ring the boundary feeds into the grid. At the toggle, injection flipped the ring *instantly*, from the control's ~ 0.5 Buy and 0 Sell to 0 Buy and ~ 1.0 Sell (Figure 7B): the injected cash let the buy orders the books had been suppressing execute at last, the supported price, 113 against the control's 93 at $t_{\mathrm{on}} + 400$, broke above the upper bands, and the entire ring, contrarian in the value trader's sense of buying below its band and selling above it, dumped at once, the flipped state decaying back over the following 100–200 steps as the EMA bands caught up. A boundary that under tight budgets had bought into the falling price and braked the front now sold into the supported price and fed it, a sign reversal from brake to accelerant. Withdrawal, which by symmetry ought to have braked the crash, did nothing of the kind at either dose, and the two doses matter because together they rule out mere under-dosing. The mild cut ($\times 0.1$, applied in the loose collapse configuration, cash $5{,}000$) left the front statistically untouched (sign-flip permutation $p = 0.67$) and barely disturbed the suppression load, suppressed buys moving $+7\%$; the harsh cut ($\times 0.01$) *bound*, in that it dropped every cell to ~ 50 in cash, below the standing price, so buys died outright and the suppressed-buy load jumped from 524 to 680 orders per step, and yet the median front velocity moved only from 2.23 to 2.38×10^{-4} grid fractions per step (11 of 19 seeds in the predicted direction; sign-flip permutation $p = 0.29$), while the price barely registered the cut, 119.5 against the control's 123.6 at $t_{\mathrm{on}} + 400$ (Figure S6). Two features of the model's economy

combine to produce that asymmetry. The first is a matter of sign and ceiling: only injection can *reverse* the boundary's state, because propping the price above the upper bands flips the ring from buying to selling, whereas withdrawal can move the ring only *further* in the direction it already points, and only transiently, the control ring already sitting near 0.46 Buy after onset, a brake the crash itself had engaged, and either cut lifting its Buy fraction to $\sim$0.72–0.74 before it relaxed back to the control level within $\sim$200 steps (Figure S6). The second follows from the execution convention: every cell the front recruits can sell up to its 10-share endowment to the unmodeled market maker, receiving roughly $1{,}000$ in cash at prices near $P \approx 100$–120, so a one-time cash reduction can be offset by later sales to the external counterparty. Within Model B, then, injections prop the price, flip the contrarian boundary to selling, and accelerate the sentiment crash, while withdrawals are near-inert; what that asymmetry would mean for the design of liquidity interventions requires a market model with explicit counterparties and calibrated institutions, which is outside the scope of this study.

5 Discussion

We asked whether local interaction alone could carry a market from a mixed trading state into a sell-dominated one, and the lattice answered yes, with more structure than the question required. The crossover between $f = 0.14$ and $f = 0.15$ is sharp, it needs no noise, no learning, and no strategy switching, and above it the sell state does not arrive everywhere at once. It arrives somewhere first, and then it spreads. That last property is what makes the model worth building on a lattice at all. A transition with a location, a start time, and a speed is a different object from one that a population-level description can hold, and the two closures we tested make the difference concrete, since the first retains only how many agents are selling and the second adds the correlation between neighbors, and neither retains the geometry a domain needs in order to have an edge. Once a transition has a location and a speed, its onset and its growth become separable, and there is no reason for the same thing to limit both. Model B is the configuration that separates them. It makes every agent pay for the orders its sentiment demands, so a budget can bear on one of the two without bearing on the other, and the question becomes which one.

That question came back asymmetric, and the asymmetry is the mechanism this study proposes. Birth is cheap, and two measurements say so independently. In the unconstrained model the conversion curves carry a growth exponent near two, the value for a transformation whose sites were all present when the clock started rather than appearing as it ran; in the budgeted model the founding cell of the winning domain fell wherever the quenched disorder favored it, well into the interior and at a rate indistinguishable from uniform placement. Seeds are not the scarce resource, and the experiment that settled the point is a null with a working lever, since the stamped patches did grow for a while before dying, so the design could have shown a placement effect and showed none. Growth is where the budget acts, and it acts through exactly one channel. The interior update never reads a book, since sentiment propagates by the totalistic rule whether or not a cell can afford to act on it and budgets suppress orders rather than states, so no cash balance enters the automaton directly. It enters only where the model closes its loop, at the ring of edge cells that classify the price and feed that classification back into the grid, and that ring is one cell deep. A channel of that width is what allows the same boundary to be irrelevant to birth and decisive for growth, which otherwise reads as a contradiction. A nucleus is born out of a fluctuation among four interior neighbors, which the ring is too far away to touch, whereas a front is a moving surface that eventually reaches the ring and from then on counts it as a neighbor. Distance to the boundary is the wrong variable for birth and the right one for growth, and the model draws that line cleanly only because the economics has just the one way in. The dose-response locates the middle of that channel and carries a warning with it. We set the initial cash, but what the front answered to was the suppression it actually met, and the same endowment buys different amounts of suppression depending on where the price has already gone. The controlled variable is a weak proxy for the causal one, which is why a velocity ordering read off the cash ladder alone would have looked like no effect, and also why the onset comparison should be read as what it is, a comparison across the proxy rather than across the quantity the front responds to. Stated this way the mechanism makes predictions that could break it. A budget that gated sentiment rather than order flow should move the onset, where the present one moved it by less than a factor of two across a tenfold range in cash. An intervention that moved the quote and the volatility bands together, changing the price by as much but leaving every edge classification where it was, should do nothing, since on this account it is the band crossings and not the price that reach the interior. And a

boundary frozen through onset should leave the birth statistics untouched while removing the gating on growth entirely.

The arm of the intervention that did nothing is the more informative of the two, and its two explanations do not travel equally far. Symmetry suggests that if adding cash speeds the front then taking cash away should slow it, and taking it away did not, at either dose. One explanation is about sign, in that only injection can reverse the boundary, since propping the quote above the bands turns a ring that had been buying a falling price into a ring selling a supported one, while withdrawal can only deepen a brake the crash had already engaged. The other is about bookkeeping, in that the front recruits sellers, selling converts inventory into cash at a counterparty that always takes the other side, and a crash therefore finances the very dip buyers a levy on cash was meant to remove. The first explanation would survive any change of counterparty convention, because it needs only a boundary rule with a sign that the price can flip. The second would not, being a direct consequence of the infinite-inventory market maker, and in an economy where cash was conserved among the agents a withdrawal would have nothing to be undone by. That difference is testable. Capping the counterparty, or conserving cash within the population, should restore a withdrawal effect while leaving the injection asymmetry in place, and a version in which neither returned would put the whole asymmetry on the sign explanation. We take the harsh dose as an operational check rather than as a stronger version of the mild one, since it bound in the model's own accounting and the front still did not respond, and we report the outcome as no detectable effect at either dose, which is a weaker claim than no effect and the only one this design supports.

Feedback between funding and market liquidity has a familiar destabilizing form, in which tighter funding forces sales, the sales move prices, and the price move tightens funding again [38]. Model B closes a loop of similar shape and returns the opposite sign, which is instructive about where the sign comes from. In the funding spiral the constraint reaches the side of the book on which the forced action is selling, so a withdrawal is what destabilizes; in Model B the cash constraint reaches only buying, since a cell needs cash to buy but only shares to sell, so a withdrawal removes buyers and cannot force a sale out of anyone. The direction in which a liquidity intervention destabilizes thus depends on which side of the book the constraint reaches, a dependence this model illustrates rather than establishes, and one that a calibrated model with explicit counterparties could test properly. The circuit-breaker literature arrives at a comparable warning by a route unavailable to us, since a breaker can be counterproductive because traders anticipate it and pull their orders forward [39], and anticipation is precisely what deterministic automaton agents cannot do, just as they cannot take correlated risk in expectation of a rescue [40]. What they supply instead is a mechanical version of the same hazard, in which an intervention alters the state of the rule that had been damping the cascade. Work on the 2010 Flash Crash makes intermediation capacity and the speed at which it is withdrawn the object of study [2], which is part of why a spatially resolved model is worth building, although nothing here is a calibrated account of that episode. Where the result does sit comfortably is on the endogenous side of the endogenous-versus-exogenous distinction [36, 37], since the feedback-closed configuration has no external input to assign the crash to. One consequence is worth flagging as a hypothesis rather than a finding. If what limits a crash is growth rather than birth, then early-warning statistics computed on an aggregate price series [34] may be watching a quantity that need not move, since the modal crash-prone outcome here locked the grid into a sell-dominated state while the price stayed close to where it began, and the observable that separated the regimes without overlap was spatial, the concentration of sell activity into a single domain. That is the kind of structural quantity that the case for complexity-informed financial regulation [5] would place on the monitoring side, and it is not recoverable from a price series alone.

The limitations of the study bound how far these readings travel. Model A uses one boundary-forcing series, QQQ one-minute bars from January 2 through March 28, 2024, and does not constitute empirical price validation. Robustness across instruments and market regimes remains to be established. The Model B nucleation campaigns all ran on a single 64×64 grid trading a single asset against an infinite-inventory market maker, and while the Model A sweep found a conversion-curve rate constant that barely depended on lattice size, it also found the threshold itself moving with size, so grid-size robustness of the budget results is an assumption and not a check, and the model is silent on cross-asset contagion. The agents themselves are deterministic, since the CA rules incorporate no stochastic noise, no learning, and no adaptive strategy switching, features of real markets whose inclusion would enrich the dynamics at the cost of breaking the enumerability of the rule space [10, 11]. On the theory side, the analytical hierarchy treats the contrarian disorder as annealed, averaging over

rule types in the mean-field and pair maps, whereas the simulations fixed each contrarian's placement for the duration of a run, quenched disorder in the proper sense, and a calculation that carried that quenching through would sharpen the prediction of f_0 and connect the critical nucleus size to the combinatorics of contrarian clustering. And the Model B statistics carry flags that interpretation must not launder: the monotone-in-cash onset rate near the cliff and the freeze mid-patch onset delay are exploratory, uncorrected observations; the injection effect stands at 14 of 19 seeds with a sign-test $p = 0.064$, and it is the paired design, not the count, that carries the causal reading; and the onset-fit comparison is power-bounded, so that overlapping confidence intervals do not establish parameter equality or exclude a small effect. A last limitation applies to any two-rule study, which is how much hangs on the two rules. The majority and contrarian pair studied here, differing at a single neighborhood sum, is one point in a finite and enumerable rule space, and that enumerability is what makes the question answerable at all, where models with continuous behavioral parameters cannot pose it in the same form. Whether the same sequence, a metastable attractor, a disorder-seeded nucleus, and budget-gated growth, recurs across other rule pairings, and whether the intervention asymmetry survives boundary rules other than the contrarian band, are open questions. A screen over rule pairs that differ at a single neighborhood sum, the same minimal perturbation studied here, would test the first directly, and swapping the band classifier for a momentum or threshold rule at the boundary would test the second.

6 Conclusion

We studied a deterministic three-state cellular automaton with two totalistic rules on a two-dimensional lattice, and in Model A a fixed-seed sweep showed a sharp finite-horizon crossover between $f = 0.14$ and $f = 0.15$, above which localized sell domains grew through coherent fronts that neither mean-field nor the tested pair closure reproduced. Onset times were well summarized by an inverse-distance empirical curve, a fit consistent with a nucleation interpretation but neither a derived Kramers law nor evidence of an infinite thermodynamic barrier, and the Potts-inspired alignment statistic similarly measures spatial order without generating the CA dynamics.

Model B added endogenous price feedback and finite cash and share inventories. Birth sites were consistent with uniform interior placement, and the present experiments bounded any cash-associated difference in onset time to within a factor of two, whereas budgets were associated with post-onset growth. In paired simulations, a tenfold cash injection just before the control's detected onset accelerated the front in 14 of 19 included pairs by changing the price-mediated boundary state. This result is specific to the model's unmodeled infinite-inventory market maker and uniform cash-shock intervention.

The finite set of $3^9 = 19{,}683$ totalistic rules makes systematic extension possible. Future work should recheck the budget results at other grid sizes, compare alternative onset-time models under censoring, include explicit counterparties, and test whether the same spatial mechanism appears for other rule pairs.

References

[1] Didier Sornette. *Why Stock Markets Crash: Critical Events in Complex Financial Systems*. Princeton University Press, 2017.

[2] Andrei Kirilenko, Albert S. Kyle, Mehrdad Samadi, and Tugkan Tuzun. The flash crash: High-frequency trading in an electronic market. *The Journal of Finance*, 72(3):967–998, 2017.

[3] Rama Cont. Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001.

[4] Benoit Mandelbrot. The variation of certain speculative prices. *The Journal of Business*, 36(4): 394–419, 1963.

[5] Stefano Battiston, J Doyne Farmer, Andreas Flache, Diego Garlaschelli, Andrew G Haldane, Hans Heesterbeek, Cars Hommes, Carlo Jaeger, Robert May, and Marten Scheffer. Complexity theory and financial regulation. *Science*, 351(6275):818–819, 2016.

[6] Thomas Lux and Michele Marchesi. Scaling and criticality in a stochastic multi-agent model of a financial market. *Nature*, 397:498–500, 1999.

[7] Per Bak, Maya Paczuski, and Martin Shubik. Price variations in a stock market with many agents. *Physica A: Statistical Mechanics and its Applications*, 246:430–453, 1997.

[8] Rama Cont and Jean-Philippe Bouchaud. Herd behavior and aggregate fluctuations in financial markets. *Macroeconomic Dynamics*, 4(2):170–196, 2000.

[9] D. Challet and Y.-C. Zhang. Emergence of cooperation and organization in an evolutionary game. *Physica A: Statistical Mechanics and its Applications*, 246:407–418, 1997.

[10] J. Doyne Farmer and Shareen Joshi. The price dynamics of common trading strategies. *Journal of Economic Behavior & Organization*, 49(2):149–171, 2002.

[11] Cars H Hommes. Heterogeneous agent models in economics and finance. *Handbook of Computational Economics*, 2:1109–1186, 2006.

[12] Stefan Bornholdt. Expectation bubbles in a spin model of markets: intermittency from frustration across scales. *International Journal of Modern Physics C*, 12(5):667–674, 2001.

[13] A. Krawiecki, J. A. Hołyst, and D. Helbing. Volatility clustering and scaling for financial time series due to attractor bubbling. *Physical Review Letters*, 89:158701, 2002.

[14] Katarzyna Sznajd-Weron and Rafal Weron. A simple model of price formation. *International Journal of Modern Physics C*, 13(1):115–123, 2002.

[15] Taisei Kaizoji. Speculative bubbles and crashes in stock markets: an interacting-agent model of speculative activity. *Physica A: Statistical Mechanics and its Applications*, 287:493–506, 2000.

[16] Yu-Cheng Wei, Shao-Hua Wang, and Fang-Mei Yang. A cellular automaton model for the securities market. *Physica A: Statistical Mechanics and its Applications*, 325:285–290, 2003.

[17] Guoping Qiu, Drona Kandhai, and Peter M. A. Sloot. A cellular automata model for the abnormal fluctuations in stock markets. *Advances in Complex Systems*, 10:241–255, 2007.

[18] Guang Yang, Yun Chen, and Ji-Ping Huang. Financial market simulation based on cellular automaton. *Physica A: Statistical Mechanics and its Applications*, 521:586–594, 2019.

[19] Stephen Wolfram. *A New Kind of Science*. Wolfram Media, 2002.

[20] Eugene F. Fama. Efficient capital markets: a review of theory and empirical work. *The Journal of Finance*, 25(2):383–417, 1970.

[21] G. William Schwert. Why does stock market volatility change over time? *The Journal of Finance*, 44(5):1115–1153, 1989.

[22] Jean-Philippe Bouchaud and Marc Potters. An apparent stylized fact about the dynamics of financial markets. *Quantitative Finance*, 1:105–112, 2001.

[23] Blake LeBaron. Agent-based computational finance: suggested readings and early research. *Journal of Economic Dynamics and Control*, 24:679–702, 2000.

[24] J. J. Chen, L. Tan, and B. Zheng. Agent-based model with multi-level herding for complex financial systems. *Scientific Reports*, 5:8399, 2015.

[25] Georges Harras and Didier Sornette. How to grow a bubble: a model of myopic adapting agents. *Journal of Economic Behavior & Organization*, 80:137–152, 2012.

[26] Blake LeBaron. Agent-based computational finance. In *Handbook of Computational Economics*, volume 2, pages 1187–1233. Elsevier, 2006.

[27] W. Brian Arthur. Foundations of complexity economics. *Nature Reviews Physics*, 3:136–145, 2021.

[28] J. Doyne Farmer. Quantitative agent-based models: a promising alternative for macroeconomics. *Oxford Review of Economic Policy*, 41(2):571–590, 2025.

[29] Dietrich Stauffer. Generalization of a two-dimensional Ising-like model to arbitrary dimension. *Computing in Science & Engineering*, 2(3):69–72, 2000.

[30] Debashish Chowdhury and Dietrich Stauffer. A generalized spin model of financial markets. *The European Physical Journal B*, 8:477–482, 1999.

[31] Per Bak, Chao Tang, and Kurt Wiesenfeld. Self-organized criticality: an explanation of $1/f$ noise. *Physical Review Letters*, 59:381–384, 1987.

[32] Per Bak and Maya Paczuski. Complexity, contingency, and criticality. *Proceedings of the National Academy of Sciences*, 92:6689–6696, 1995.

[33] Xavier Gabaix, Parameswaran Gopikrishnan, Vasiliki Plerou, and H Eugene Stanley. A theory of power-law distributions in financial market fluctuations. *Nature*, 423:267–270, 2003.

[34] Marten Scheffer, Jordi Bascompte, William A Brock, Victor Brovkin, Stephen R Carpenter, Vasilis Dakos, Hermann Held, Egbert H Van Nes, Max Rietkerk, and George Sugihara. Early-warning signals for critical transitions. *Nature*, 461:53–59, 2009.

[35] J Doyne Farmer, Paolo Patelli, and Ilija I Zovko. The predictive power of zero intelligence in financial markets. *Proceedings of the National Academy of Sciences*, 102(6):2254–2259, 2005.

[36] Jean-Philippe Bouchaud. Crises and collective socio-economic phenomena: Simple models and challenges. *Journal of Statistical Physics*, 151(3–4):567–606, 2013.

[37] Didier Sornette, Fabrice Deschatres, Thomas Gilbert, and Yann Ageon. Endogenous versus exogenous shocks in complex networks: an empirical test using book sale rankings. *Physical Review Letters*, 93:228701, 2004.

[38] Markus K. Brunnermeier and Lasse Heje Pedersen. Market liquidity and funding liquidity. *The Review of Financial Studies*, 22(6):2201–2238, 2009.

[39] Avanidhar Subrahmanyam. Circuit breakers and market volatility: A theoretical perspective. *The Journal of Finance*, 49(1):237–254, 1994.

[40] Emmanuel Farhi and Jean Tirole. Collective moral hazard, maturity mismatch, and systemic bailouts. *American Economic Review*, 102(1):60–93, 2012.

[41] Stephen Wolfram. Universality and complexity in cellular automata. *Physica D: Nonlinear Phenomena*, 10:1–35, 1984.

[42] Norman H. Packard and Stephen Wolfram. Two-dimensional cellular automata. *Journal of Statistical Physics*, 38:901–946, 1985.

[43] Albert S. Kyle. Continuous auctions and insider trading. *Econometrica*, 53(6):1315–1335, 1985.

[44] Hersh Shefrin and Meir Statman. The disposition to sell winners too early and ride losers too long: Theory and evidence. *The Journal of Finance*, 40(3):777–790, 1985.

[45] Terrance Odean. Are investors reluctant to realize their losses? *The Journal of Finance*, 53(5): 1775–1798, 1998.

[46] Howard A. Gutowitz, Jonathan D. Victor, and Bruce W. Knight. Local structure theory for cellular automata. *Physica D: Nonlinear Phenomena*, 28(1–2):18–48, 1987.

[47] Hendrik A. Kramers. Brownian motion in a field of force and the diffusion model of chemical reactions. *Physica*, 7(4):284–304, 1940.

[48] Peter Hänggi, Peter Talkner, and Michal Borkovec. Reaction-rate theory: Fifty years after Kramers. *Reviews of Modern Physics*, 62(2):251–341, 1990.

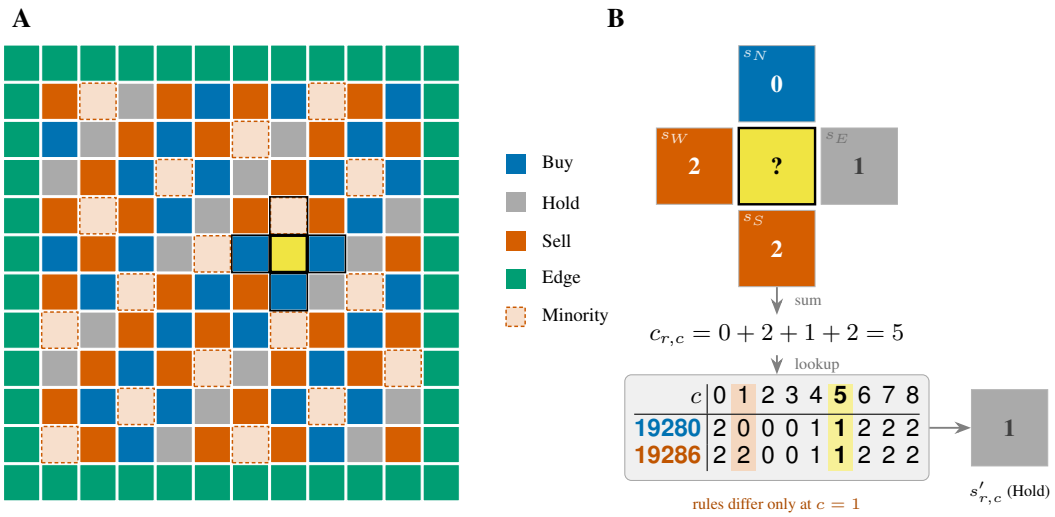


Figure 1: Model architecture. **(A)** Agent grid ($N_r \times N_c$) with three states (Buy, Hold, Sell). Green boundary cells are edge agents; dashed-outline interior cells follow the minority rule (19286); solid interior cells follow the majority rule (19280). The outlined cell and its Von Neumann neighbors are expanded in B. **(B)** CA update for the outlined cell: the four neighbor states are summed to $c_{r,c} = 5$, the sum indexes the rule table, and both rules output Hold, so the next state is $s'_{r,c} = 1$. The two rules differ only at $c = 1$ (shaded column).

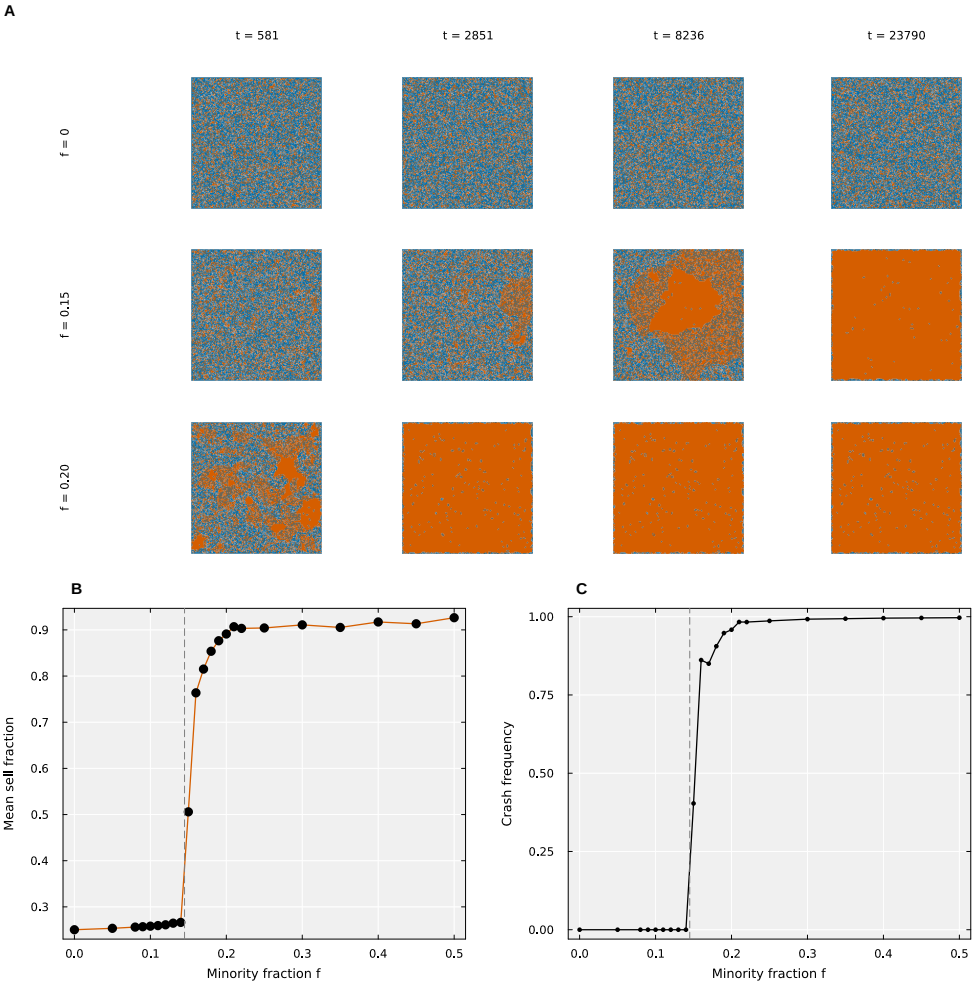


Figure 2: The Model A crossover (Rules 19280/19286, QQQ 2024 Q1 data). **(A)** Spatiotemporal snapshots of the 256×256 agent grid at four logarithmically spaced times ($t = 581$, 2,851, 8,236, and 23,790) and three minority fractions ($f = 0$, $f = 0.15$, and $f = 0.20$), with blue for Buy, gray for Hold, and orange for Sell. At $f = 0.15$ and $t = 8{,}236$ a sell cascade nucleated from a localized seed and propagated as a coherent wavefront before engulfing the grid. **(B)** Mean sell fraction against the minority fraction f over a fixed-seed sweep on a 128×128 grid. **(C)** Crash frequency, the fraction of time steps whose sell fraction exceeds 0.50, over the same sweep. The dashed vertical line at $f = 0.145$ marks the midpoint of the observed $[0.14, 0.15]$ crossover bracket, not a fitted thermodynamic critical point. Peak sell fraction over the same sweep is given in Figure S3.

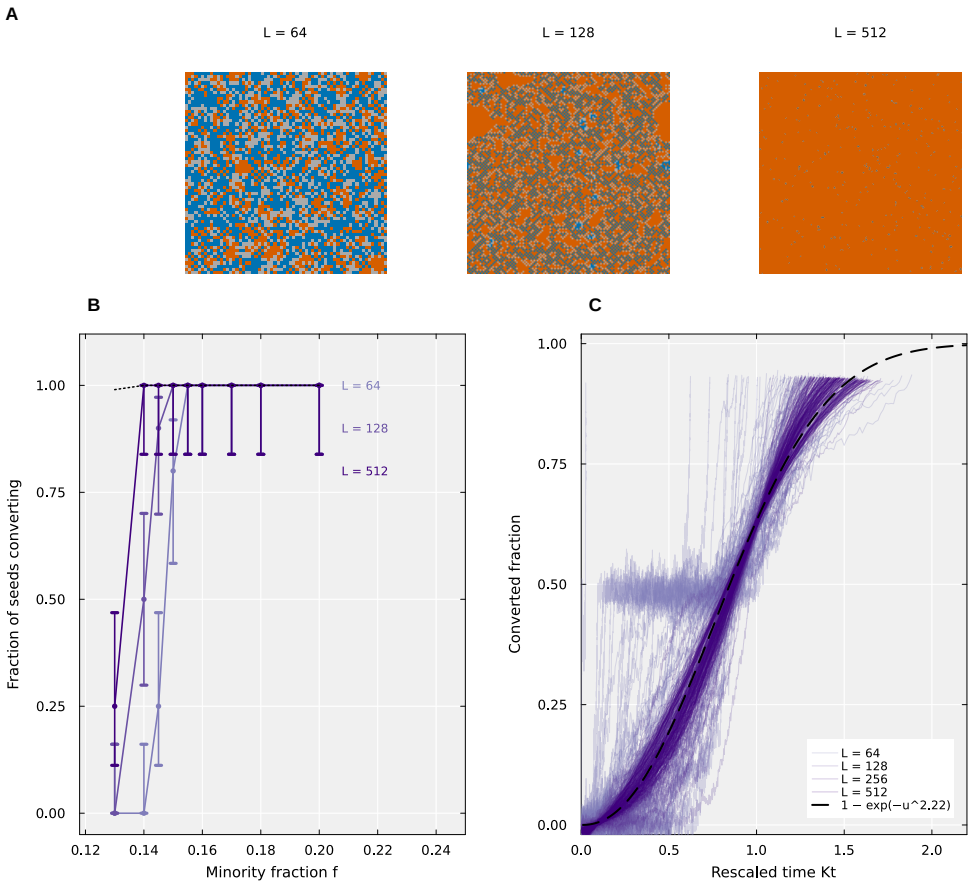


Figure 3: System-size dependence of the Model A crossover (Rules 19280/19286, periodic boundaries, no price forcing, twenty disorder seeds per cell). **(A)** One grid per lattice size, $L = 64$, 128, and 512, each at $f = 0.15$, seed 1, and the same absolute time $t = 17{,}235$ steps, with blue for Buy, gray for Hold, and orange for Sell. At that one clock reading the smallest lattice was still mixed, at a sell fraction of 0.28, the middle one was part converted at 0.65, and the largest had gone over to Sell at 0.99, so what a fixed window records depends on the size of the market it watches. **(B)** Fraction of seeds converting within the $100{,}000$-step horizon against f for those same three sizes, with Wilson score intervals and each curve labeled at its right end. Larger lattices convert at lower f. The dotted curve is the fraction of runs, pooled over all five sizes, in which a 100-cell sell cluster was ever seen; such clusters appear even where nothing converts, so what moves with size is spreading rather than clustering. This is a descriptive threshold and not a certified critical nucleus. **(C)** All 320 conversion curves plotted against time rescaled by each run's own fitted rate constant, with $1 - \exp(-u^n)$ overlaid at $n = 2.22$, the mean fitted exponent on the largest lattice. Sizes lie on top of one another; the band tightens with size, and 64×64, the smallest lattice fitted, visibly fails to collapse. The dispersion and timing of the conversion times themselves are given in Figure S7.

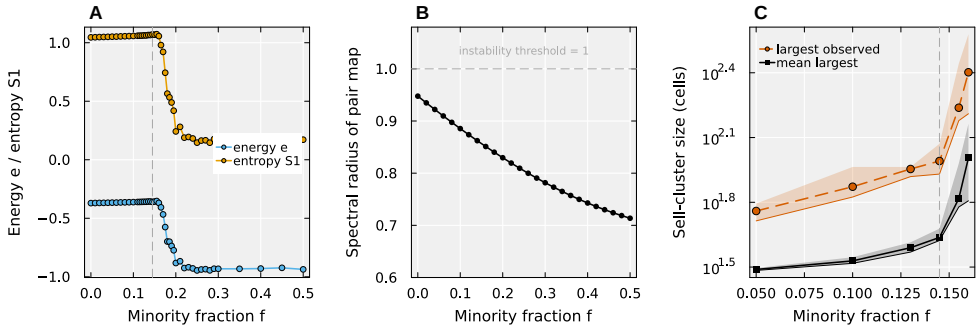


Figure 4: Theory overview for Model A (Rules 19280/19286). **(A)** Potts energy e and site entropy S_1 vs. minority fraction f (128×128 lattice, periodic boundaries, time-averaged over $2{,}000$ steps). The mixed attractor persists at $e \approx -0.37$ and $S_1 \approx 1.06$; above the transition the system collapses to the all-Sell state at $e \approx -0.93$ and $S_1 \approx 0.18$, a sharp finite-size drop ($\Delta e \approx 0.56$, $\Delta S_1 \approx 0.88$); near the crossover the $2{,}000$-step averages retain the mixed state because metastable lifetimes exceed the window, so the crossover completes to the right of the marked f_c. **(B)** Spectral radius of the pair-approximation Jacobian at the uncorrelated fixed point, declining monotonically from 0.95 at $f = 0$ to 0.71 at $f = 0.5$, strictly below the instability threshold (dashed line at 1). This tested pair closure does not capture the crossover. **(C)** Sell-cluster sizes in the mixed state (log scale): largest cluster ever observed (orange, dashed) and mean largest per snapshot (black), each the median over 20 disorder seeds with an interquartile band, taken over 100 snapshots of a 128×128 lattice. The mean largest cluster held at 31 cells at $f = 0.05$ and 43 cells at $f = 0.145$, while the largest observed cluster rose from 58 to 98 cells over the same range and to 253 cells at $f = 0.16$. The band widens above the bracket, where seeds disagree: the middle half of them spanned 162 to 379 cells at $f = 0.16$. Dashed vertical lines mark the midpoint $f = 0.145$ of the observed $[0.14, 0.15]$ bracket.

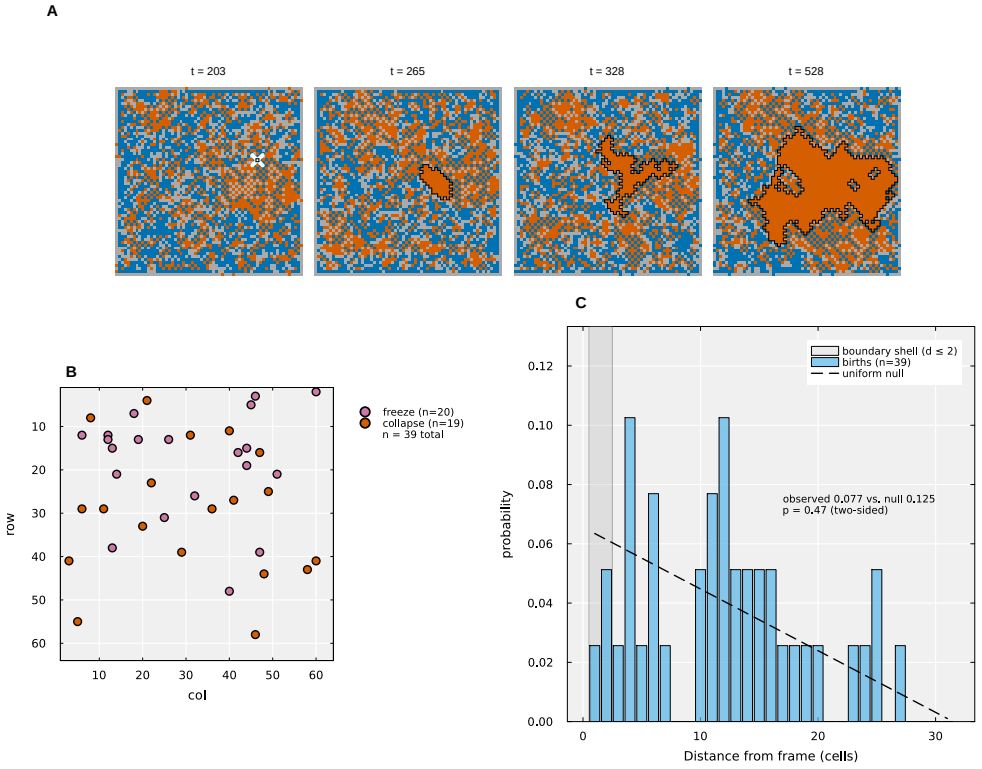


Figure 5: The winning sell domain is a single dominant nucleus born in the bulk. **(A)** Four snapshots of the 64×64 grid in one collapse run ($f = 0.20$, cash 5,000, seed 20), whose onset at $t = 328$ sits near the median of 374 steps over the 19 collapse runs that nucleated, with blue for Buy, gray for Hold, and orange for Sell. The cross marks the founding cell at its birth step $t = 203$, 15 cells from the nearest frame, and the black outline traces that one cell's forward lineage: 75 cells at $t = 265$, 208 cells at onset, and 973 cells by $t = 528$, grown without ever reaching the frame. **(B)** Map of the founding cell of the largest sell cluster at crash onset across the 39 boundary-versus-bulk runs that nucleated, plotted on the interior alone so that a bulk birth can be told apart from attachment to the active-edge ring. **(C)** Distance of that founding cell from the nearest boundary, pooled over the 20 freeze and 19 collapse runs, against the uniform-null expectation and with the boundary shell ($d \leq 2$) shaded. Births were bulk-dominated and consistent with a uniform distribution: the pooled fraction founded within distance 2 of the boundary was 0.077, below rather than above the 0.125 uniform null and statistically indistinguishable from it (binomial $p = 0.47$), with a median founding distance of 12 cells into the deep bulk. Per-seed ratios of the largest cluster to the total sell population at its peak, and the growth trajectories behind them, are given in Figure S2.

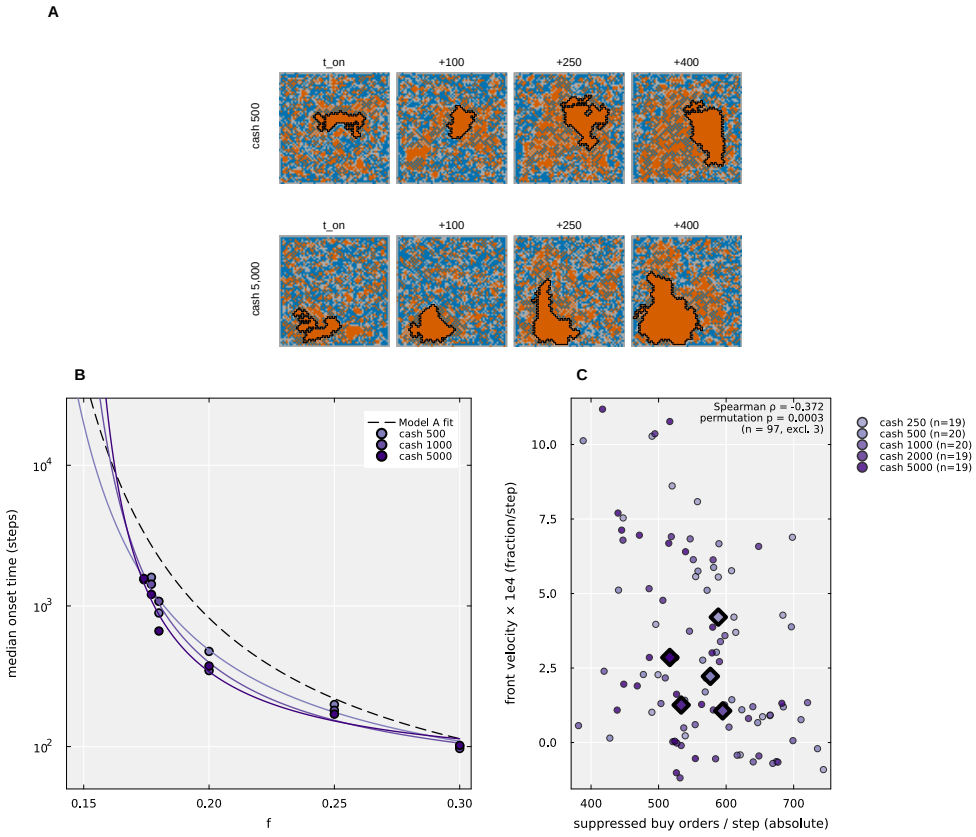


Figure 6: Budgets gate the growth of the nucleus rather than its birth. **(A)** Four snapshots of the 64×64 grid under a tight budget (cash 500, top row) and a loose one (cash 5,000, bottom row) at $f = 0.20$, both from seed 17, with blue for Buy, gray for Hold, and orange for Sell. Each row is cut at its own onset step, $t_{\mathrm{on}} = 350$ above and 374 below, and then at 100, 250, and 400 steps after it, so the two arms are compared on their own clocks rather than on one shared one. The black outline traces the largest sell cluster in each frame. Both arms enter the first tile just past the 205-cell onset threshold, at 210 and 225 cells, and by 400 steps the loose budget's domain has reached 873 cells against the tight budget's 516. **(B)** Kaplan–Meier median onset times (points) against the minority fraction f at cash 500, 1,000, and 5,000, each with its fitted divergence $\ln \tilde{\tau} = A + B/(f - f_0)$ (solid curves) and the Model A fit of Equation (9) dashed for reference; the fitted curves are drawn only up to a display ceiling so that the medians are not crushed by the divergence near f_0. Censoring at the horizon and its handling in the medians are detailed in Appendix C, and the fitted parameters with their bootstrap intervals in Figure S4. **(C)** Front velocity against absolute budget suppression at $f = 0.20$ for the 97 runs included across the five cash doses (250–5,000; 3 runs excluded by the window rule). The velocity is the OLS slope of the front extent over the onset-aligned window $[t_{\mathrm{on}} + 50,\ t_{\mathrm{on}} + 400]$ and the suppression is the mean absolute number of suppressed buy orders per step over the same window, colored by dose with the per-dose medians marked as diamonds; velocity fell as suppression rose, Spearman $\rho = -0.372$, permutation $p = 0.0003$. The same velocities arranged by initial cash are given in Figure S4, where the per-dose medians are U-shaped rather than monotone, and the stamped-patch design behind the $0/20$ dominance rate in Figure S5.

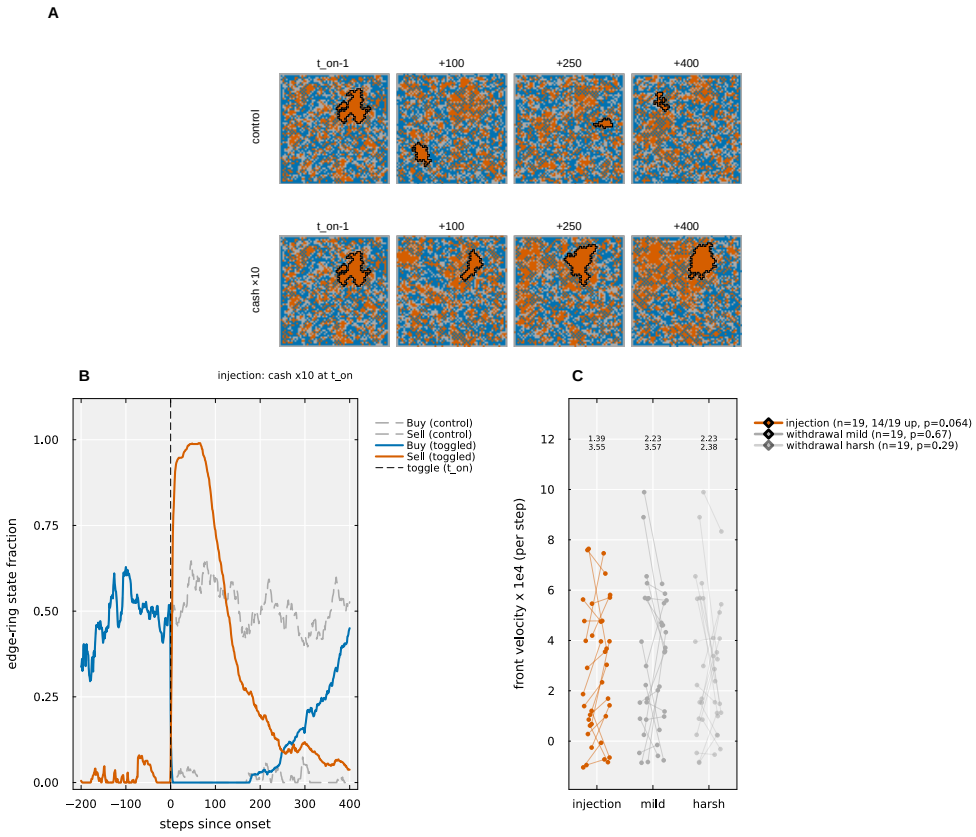


Figure 7: Liquidity intervention is asymmetric. Paired, onset-aligned counterfactuals at $f = 0.20$ with a $W = 400$ post-toggle window and 19 included seed pairs per arm; the pre-toggle trajectories coincide by construction. **(A)** Four snapshots of the 64×64 grid in one injection pair, control above and cash $\times 10$ below (seed 16, $t_{\text{on}} = 268$, the pair whose velocity gain sits at the median of the 14 that accelerated), with blue for Buy, gray for Hold, and orange for Sell. The first column is one step before the toggle, where the two rows are identical cell for cell, and the remaining three are 100, 250, and 400 steps after it. The black outline traces the largest sell cluster in each frame: the shared 190-cell domain grew to 198 cells under injection while the control's broke up to 38. **(B)** Edge-ring Buy and Sell fractions under injection, averaged over the 19 seed pairs and aligned at the per-seed onset, for the toggled arm against its control, with the dashed line marking the toggle at t_{on}. Injection flipped the ring at the toggle, from the control's ~ 0.5 Buy and 0 Sell to 0 Buy and ~ 1.0 Sell, and the flipped state decayed back over the following 100–200 steps as the EMA bands caught up with the supported price. **(C)** Front velocity over the window for every seed pair, one control-to-toggled segment per pair, with injection at left in orange and the mild ($\times 0.1$) and harsh ($\times 0.01$) withdrawals beside it in gray, the harsh dose the paler of the two; injection sits on the tight cash-500 base and both withdrawals on the loose cash-5,000 base. Each arm carries its control and toggled medians as text above it, in that order. Injection accelerated the front in 14 of 19 pairs (sign test $p = 0.064$), while withdrawal was null at both doses (sign-flip permutation $p = 0.67$ mild, $p = 0.29$ harsh). The mild arm's two medians rise even so, 2.23 to 3.57, although 11 of its 19 pairs individually fell, so its marginal medians and the paired difference the test uses disagree in sign. The cluster trajectories and the withdrawal ring traces for all three arms are given in Figure S6.

Table 1: Wolfram totalistic CA rule taxonomy. Base-3 encoding reads right-to-left: digit a_i gives the next state when the neighborhood sum index equals i.

Rule	Base-3 encoding	Interpretation	Role
6158	$(022110002)_3$	Imitation / herding	Homogeneous model; stable waves
19280	$(222110002)_3$	Majority rule	80% crash model; herding
19286	$(222110022)_3$	Contrarian	20% crash model; destabilizer

Table 2: Empirical fits $\ln \tilde{\tau} = A + B/(f - f_0)$ to Kaplan–Meier median onset times, per cash level, with the Model A fit for reference. The Model A row is the mean-lifetime fit of Equation (9); the Model B rows fit Kaplan–Meier medians. Brackets are 95% bootstrap CIs (2,000 resamples of seeds within each f). The cash-500 f_0 lower bound of 0.000 is the physical boundary of the fit domain, not a data quantile. CIs overlap between tight (500) and loose (5000) cash on both B and f_0, but this overlap does not establish parameter equality.

Configuration	n_f	A	B	f_0	R^2
Model A (reference)	12	2.78	0.389	0.101	0.991
Model B, cash 500	5	3.38	0.247 [0.070, 1.135]	0.112 [0.000, 0.153]	0.978
Model B, cash 1000	6	3.80	0.143 [0.071, 0.691]	0.135 [0.048, 0.152]	0.990
Model B, cash 5000	6	4.10	0.098 [0.050, 0.415]	0.143 [0.077, 0.157]	0.980

A Mean-Field and Pair-Approximation Derivations

This appendix records the mean-field and pair-approximation calculations in the form in which they were implemented. The quenched rule mixture enters in its annealed form: Rules 19280 and 19286 share every entry of their totalistic tables except at neighborhood sum $c = 1$, where the majority rule outputs Buy and the contrarian rule outputs Sell, so a site whose rule assignment is drawn afresh at minority fraction f follows the effective stochastic rule

$$R_f(c \to b) \;=\; (1-f)\,\mathbf{1}[\mathcal{R}_{19280}(c) = b] \;+\; f\,\mathbf{1}[\mathcal{R}_{19286}(c) = b]\,, \tag{10}$$

deterministic at every $c \neq 1$ and splitting $c = 1$ between Buy, with probability $1 - f$, and Sell, with probability f.

The mean-field closure is that same effective rule with all spatial structure discarded: each agent's four neighbors are drawn independently from the global state distribution (p_0, p_1, p_2), where p_s is the fraction of agents in state s and $p_0 + p_1 + p_2 = 1$. The neighborhood sum $c = s_N + s_S + s_E + s_W$ is then a multinomial convolution of four i.i.d. draws,

$$\pi_k \;=\; \Pr(c = k) \;=\; \sum_{\substack{n_0+n_1+n_2=4 \\ n_1+2n_2=k}} \binom{4}{n_0,\, n_1,\, n_2}\, p_0^{n_0}\, p_1^{n_1}\, p_2^{n_2}, \qquad k = 0, 1, \ldots, 8, \tag{11}$$

and the verified rule tables (Table 1) send $c \in \{2, 3\}$ to Buy, $c \in \{4, 5\}$ to Hold, and $c \in \{0, 6, 7, 8\}$ to Sell for both rules, leaving $c = 1$ as the only entry on which they differ. The map is therefore

$$p_0' = (1-f)\,\pi_1 + \pi_2 + \pi_3, \tag{12}$$

$$p_1' = \pi_4 + \pi_5, \tag{13}$$

$$p_2' = \pi_0 + f\,\pi_1 + \pi_6 + \pi_7 + \pi_8, \tag{14}$$

with primes denoting next-step fractions and each π_k evaluated from Equation (11) at the current (p_0, p_1, p_2). We iterated Equations (12)–(14) to convergence, $\|\mathbf{p}' - \mathbf{p}\|_1 < 10^{-14}$, from multiple initial conditions and across $f \in [0, 1]$. The two stable fixed points reported in the main text are those of this map, and the Jacobian at the buy-dominated branch has complex eigenvalues of modulus $|\lambda| \approx 0.95$, a stable spiral. The pair approximation restores one level of spatial structure.

The state of the pair approximation is the symmetric matrix of nearest-neighbor pair densities ρ_{ab}, $a, b \in \{0, 1, 2\}$, normalized by $\sum_{ab} \rho_{ab} = 1$, with site marginals $\rho_a = \sum_b \rho_{ab}$ and conditional neighbor distribution $q(s \mid a) = \rho_{as}/\rho_a$. The closure is the local-structure ansatz of Gutowitz, Victor, and Knight [46]: given a focal site in state a, its three neighbors *other* than the paired partner are drawn independently from $q(\cdot \mid a)$. Their sum m therefore follows the trinomial convolution

$$\Pr(m = k \mid a) \;=\; \sum_{\substack{n_0+n_1+n_2=3 \\ n_1+2n_2=k}} \binom{3}{n_0,\, n_1,\, n_2}\, q(0 \mid a)^{n_0}\, q(1 \mid a)^{n_1}\, q(2 \mid a)^{n_2}, \qquad k = 0, \ldots, 6, \tag{15}$$

and the focal site's full neighborhood sum is the partner's state plus that draw, $c = a' + m$, which never exceeds 8. Averaging the effective rule over the trinomial weights gives the transition kernel

$$T_b(a, a') \;=\; \sum_{k=0}^{6} \Pr(m = k \mid a)\; R_f(a' + k \to b), \tag{16}$$

the probability that a site in state a with paired neighbor in state a' lands in state b after one synchronous step. Substituting this kernel for both members of a pair (each site conditioning on its own partner) yields the synchronous pair update $\rho'_{bb'} = \sum_{a,a'} \rho_{aa'}\, T_b(a, a')\, T_{b'}(a', a)$ quoted as Equation (8) of the main text. The implementation symmetrizes and renormalizes ρ' after every step to guard against numerical drift.

We iterated the map to its fixed points, declaring convergence when the largest componentwise change $\max_{ab} |\rho'_{ab} - \rho_{ab}|$ fell below 10^{-12}, and every trajectory reported here converged within a few hundred iterations. Stability of the uncorrelated fixed point $\rho^*_{ab} = \rho^*_a \rho^*_b$, the product state built from the mean-field marginals, was assessed through the Jacobian of the map restricted to the six independent components $(\rho_{00}, \rho_{01}, \rho_{02}, \rho_{11}, \rho_{12}, \rho_{22})$ that survive symmetry and normalization,

computed by finite differences with step $\delta = 10^{-7}$, each perturbed matrix projected back onto the probability simplex before the map was applied. The spectral radius sat strictly below unity at every value of f we sampled (twelve values from 0 to 0.5), declining monotonically from 0.95 at $f = 0$ to 0.71 at $f = 0.5$ (Figure 4B).

The search for correlated fixed points used initial conditions

$$\rho_{ab} \;=\; p_a p_b + \varepsilon\,(\delta_{ab}\,p_a - p_a p_b)\,, \tag{17}$$

which interpolate between the uncorrelated product state at $\varepsilon = 0$ and perfect same-state correlation at $\varepsilon = 1$. We ran $\varepsilon \in \{0.1,\ 0.3,\ 0.5\}$ on buy-biased marginals $(0.50,\ 0.30,\ 0.20)$, $\varepsilon \in \{0.1,\ 0.3\}$ on sell-biased marginals $(0.20,\ 0.10,\ 0.70)$, and $\varepsilon = 0.5$ on $(0.10,\ 0.05,\ 0.85)$, together with an uncorrelated uniform control, at thirteen values of f between 0 and 0.5. Every trajectory converged either to the uncorrelated buy-dominated fixed point or to the all-Sell absorbing state, with the spatial mutual information between adjacent sites decaying to zero in all cases. That mutual information is $I = 2S_1 - S_2$, built from the single-site and pair Shannon entropies

$$S_1 \;=\; -\sum_a \rho_a \ln \rho_a, \qquad S_2 \;=\; -\sum_{a,b} \rho_{ab} \ln \rho_{ab}, \tag{18}$$

in which S_1 is the site entropy plotted for the lattice in the main text (Figure 4A), maximal at $\ln 3$ for a uniform mixture of the three states and zero in a single-state configuration. Because the two site marginals of a symmetric ρ_{ab} are equal, I vanishes exactly when the pair densities factorize as $\rho_{ab} = \rho_a \rho_b$. The pair approximation therefore supports no correlated fixed point, which is the negative result on which the nucleation argument rests.

B Model B Stage 2 Factorial

The Stage 2 factorial sweep crossed six factors in full, giving the design from which every Model B summary statistic in this study is drawn (Table 3).

Every run was assigned to one of five candidate regimes by a classifier evaluated in strict order, each test applying only to runs that fail all earlier ones. A run with zero return standard deviation is *inert*. Failing that, a run is *crash-prone* when its crash frequency, the fraction of steps at which the grid's Sell fraction stands at or above 0.5, reaches at least 0.01. Failing that, a run is *oscillatory* when the net order imbalance φ changes sign at least 20 times per 2,000 steps. That threshold is a rate, one flip per 100 steps, stated here against the sweep's horizon so that it transfers to other run lengths. Failing that, a run whose total price range stays below $0.10\,P_0$ is a *plateau*, and any remaining run is a *drift*. Of the 10,800 runs, 6,589 (61%) were oscillatory and 4,211 (39%) crash-prone; no run was ever inert, plateau, or drift. The sell-dominated (crash-prone) regime appeared abruptly between $f = 0.15$ and $f = 0.20$ at every price impact λ in the design.

Grouping runs by their parameter configuration split each summary metric's variance into a between-configuration share, attributable to the swept parameters jointly, and a within-configuration share, attributable to the disorder seed. The parameters accounted for 94–97% of the variance across the four summary metrics, at between-configuration shares of 0.946 for crash frequency, 0.952 for return kurtosis, 0.959 for return standard deviation, and 0.974 for the φ sign-flip count, leaving the seed at most 5.4%. Across matched design cells, enabling the disposition rule increased return standard deviation by a factor of 2.40 on average.

The crash-prone class divided into the modal freeze and a collapsing tail, the decomposition on which the transition analysis builds: the median crash-prone run had a maximum drawdown of about 5% and a price minimum of 100.00, so half the class never traded below the initial quote at all, while 897 of the 4,211 runs drove the price below 90 and 451 below 80. The deepest collapses, drawdowns of 40–69% arriving after bubbles of 2.5–2.7× the initial price, were concentrated at $f = 0.2$, $\lambda = 0.05$, cash 5,000, disposition off, the loosest budget regime in the design.

C Escape-Time Censoring and Survival Detail

The escape-time campaign ran each seed to a horizon of 2,000 steps, and any run that has not nucleated by then is right-censored: its onset time $\tau = t_{\text{on}}$ is treated as exceeding the horizon, and

the run enters the Kaplan–Meier estimator carrying that censoring mark, since censored runs are never dropped. The resulting counts are recorded for every (f, cash) cell of the campaign (Table 4); rows at 20 seeds belong to the pilot pass and rows at 40 seeds to the refinement pass. The companion velocity–suppression dose-response added two further cells at $f = 0.20$, cash 250 and cash 2,000, each with 20 seeds, one censored run, and an onset rate of 0.95. Those cells feed the dose-response analysis, not the barrier fits.

Because censoring occurs only at the horizon, the Kaplan–Meier survival curve of a cell steps down at each observed onset and its median is finite exactly when the onset rate reaches one half. The fits therefore rest on five points at cash 500, at $f \in \{0.177, 0.18, 0.20, 0.25, 0.30\}$, and on six points at cash 1,000 and 5,000, which add $f = 0.174$ (Table 2). Cash 500 has one fewer because its onset rate at $f = 0.174$ was 0.45, just below the one-half a finite median requires. It was this gap that prompted extending the refinement grid to $f = 0.177$, so that the tightest budget still contributed a point hard against the cliff. The reported confidence intervals (Table 2, Figure S4) come from 2,000 bootstrap resamples drawn over seeds within each f: each resample recomputes the Kaplan–Meier medians and refits the barrier law, and the intervals are the 2.5th and 97.5th percentiles of the resampled parameters.

Finally, the onset-rate observation near the cliff is visible directly in Table 4, and it carries the same label as in the main text: exploratory, reported uncorrected for multiple comparisons, and weighed accordingly. At $f = 0.174$ the onset rate rose monotonically in cash, 0.45, 0.62, 0.70; at $f = 0.177$ the Kaplan–Meier medians fell in the same order, 1,598, 1,427, and 1,208 steps.

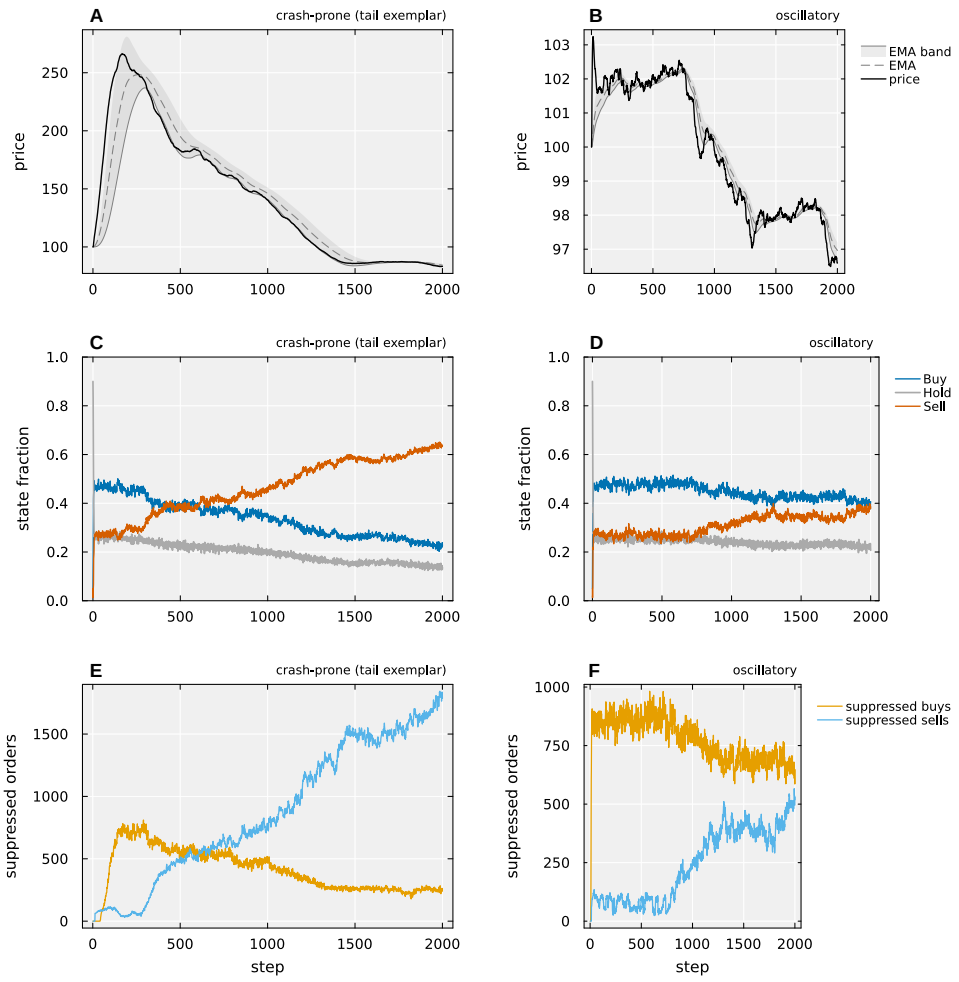


Figure S1: Model B regime exemplars from the Stage 2 factorial: a crash-prone run (run 9652) and an oscillatory run (run 5916). **(A)** Price, its exponential moving average, and the EMA band for the crash-prone exemplar, a bubble-and-collapse trajectory ($100 \to 266 \to 83$). **(B)** The same for the oscillatory exemplar, whose price stayed within a few percent of its starting value. **(C)** Interior Buy, Hold, and Sell fractions for the crash-prone exemplar. **(D)** The same for the oscillatory exemplar. **(E)** Suppressed buy and sell orders per step for the crash-prone exemplar. **(F)** The same for the oscillatory exemplar. The crash-prone exemplar is the dramatic tail of its regime, shown because it makes the mechanism visible; the *typical* crash-prone run is a liquidity freeze in which the grid locks near a Sell fraction of 0.88 while the price barely moves. "Crash-prone" throughout denotes majority-sell lock-in (Sell fraction ≥ 0.5 on at least 1% of steps), not a price discontinuity.

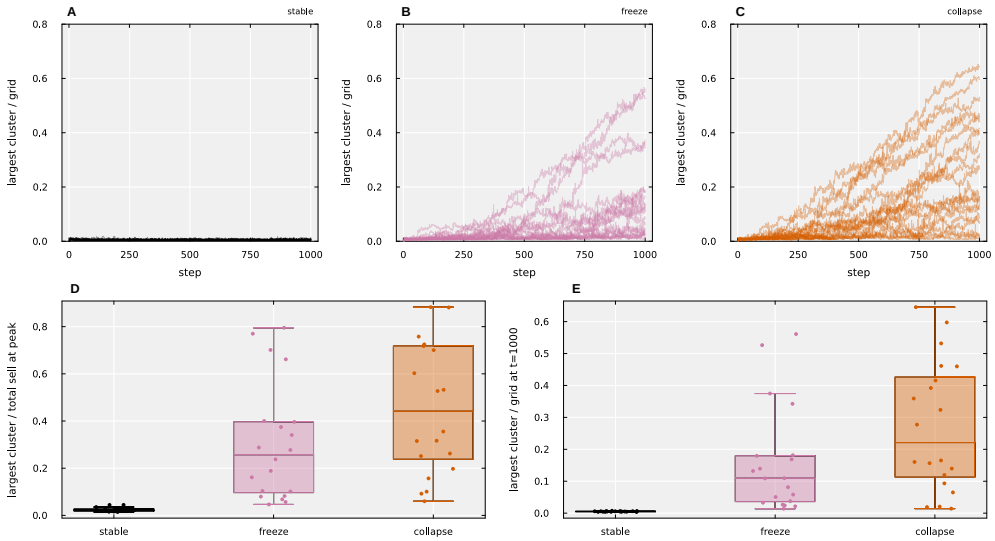


Figure S2: Model B seed sweep: largest-sell-cluster trajectories per seed, 20 disorder seeds per budget configuration. **(A)** Largest sell cluster, expressed as a fraction of the grid, drawn step by step for every seed of the stable configuration ($f = 0.10$, cash 1,000). **(B)** The same for the freeze configuration ($f = 0.20$, cash 500). **(C)** The same for the collapse configuration ($f = 0.20$, cash 5,000). The time-shifted-sigmoid fan within each above-threshold configuration reflects stochastic nucleation waiting times rather than a spread of eventual outcomes. **(D)** Per-seed ratio of the largest cluster to the total sell population at its peak. The stable configuration sits at a median ratio of 0.022, dispersed fluctuations with no dominant domain, while the freeze and collapse configurations concentrate into a single growing domain at median ratios of 0.26 and 0.44, and the stable band shows zero overlap against either above-threshold configuration. **(E)** Largest-cluster fraction of the grid at the end of the run.

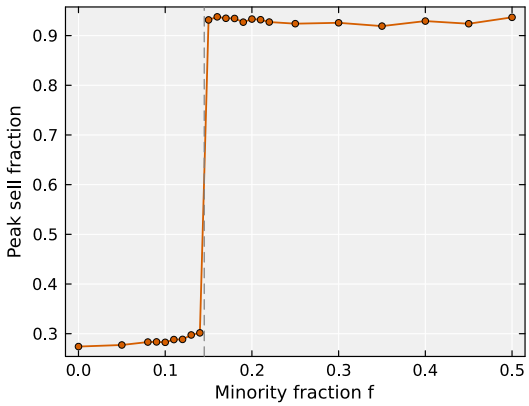


Figure S3: Peak sell fraction against the minority fraction f for Model A (Rules 19280/19286, 128×128 grid, fixed seed), the third summary of the sweep whose mean sell fraction and crash frequency are Figure 2B and 2C. The dashed vertical line at $f = 0.145$ marks the midpoint of the observed $[0.14, 0.15]$ crossover bracket, not a fitted thermodynamic critical point.

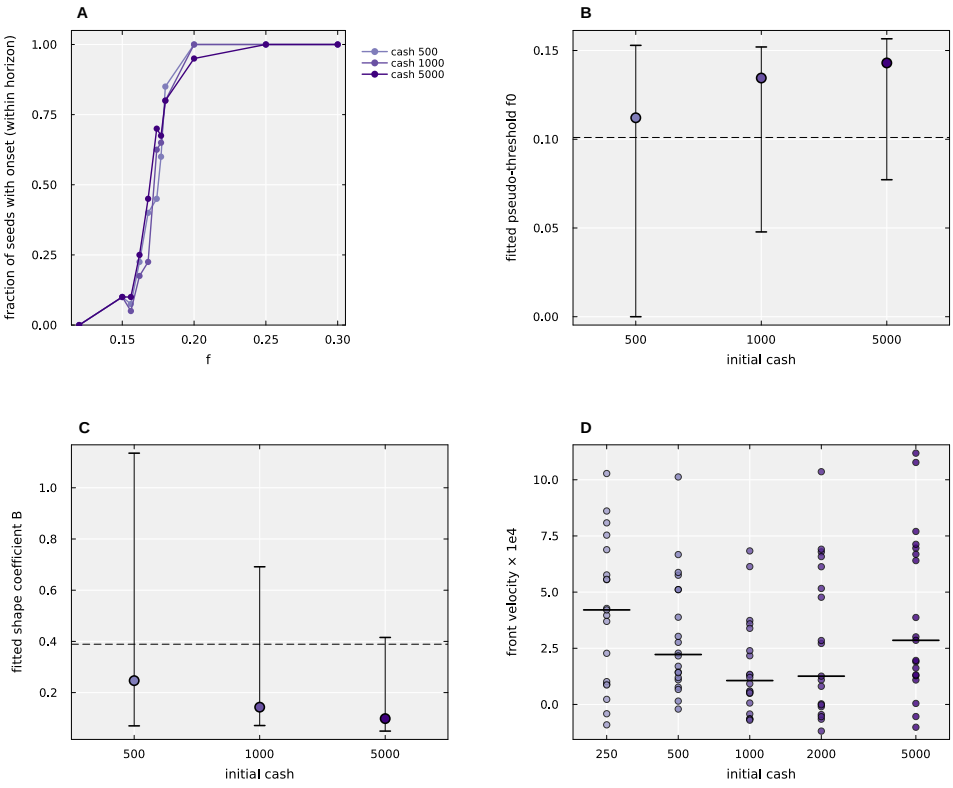


Figure S4: Onset cliff and fitted barrier parameters across the cash ladder. **(A)** Fraction of seeds that nucleated within the 2,000-step horizon, against the minority fraction f at cash 500, 1,000, and 5,000: the onset cliff sits between $f = 0.15$ and $f = 0.18$ at every cash level. **(B)** Fitted pseudo-threshold f_0 per cash level with 95% bootstrap confidence intervals; the dashed line marks the Model A value $f_0 = 0.101$. **(C)** Fitted shape coefficient B on the same footing, against the Model A value $B = 0.389$. The intervals overlap between tight and loose cash on both parameters. That overlap is not itself a test; the paired equivalence analysis reported in the text is, and it bounds the between-budget difference in predicted median onset time to within a factor of two. The resampling behind these intervals is described in Appendix C. **(D)** Front velocity at $f = 0.20$ arranged by initial cash, one point per included run and a bar at each per-dose median: the medians are U-shaped in cash, highest at cash 250, lowest near 1,000–2,000, and rising again at 5,000, while the per-dose median suppression spans only 517–595 orders per step. The supported claim is therefore the pooled and within-dose velocity–suppression relationship of Figure 6C, not a monotone ordering by dose.

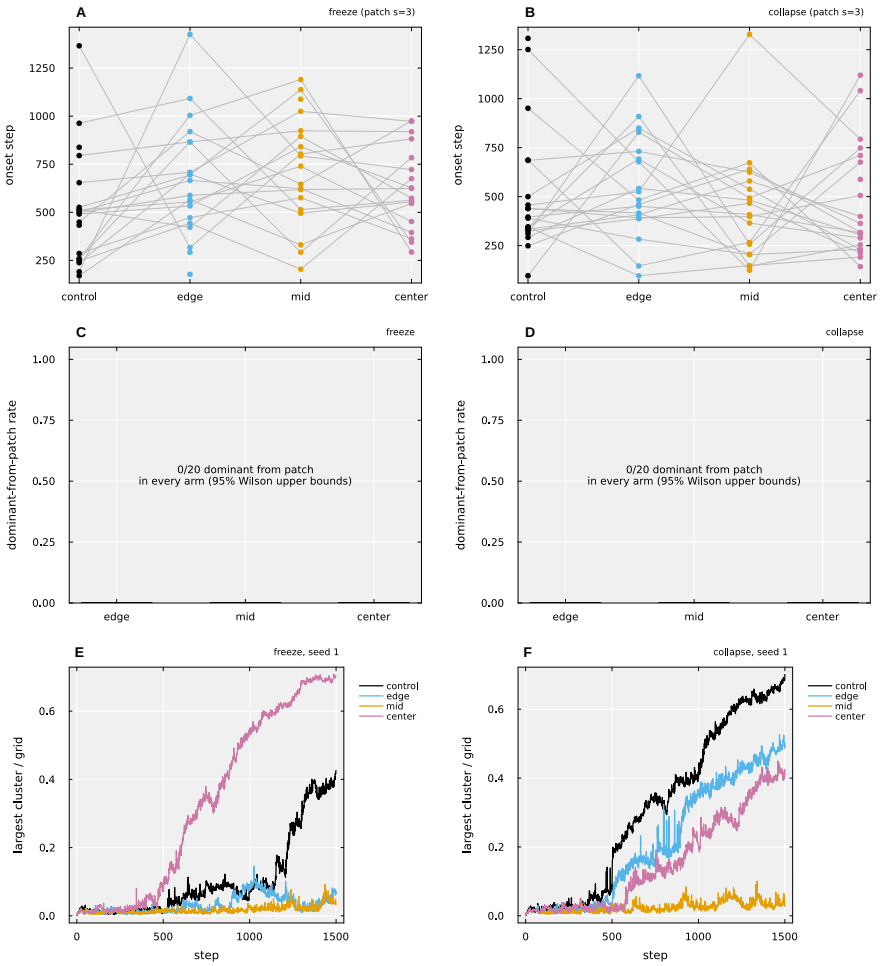


Figure S5: Stamped-patch seeding detail across placements and configurations. **(A)** Onset step per seed under each arm (an unstamped control and a 3×3 sell patch stamped at edge, mid, and center placements) for the freeze configuration, with gray lines pairing the same seed across arms; the pairings cross in both directions and no placement advanced onset. The apparent delay under the freeze mid placement (median paired onset advance -257 steps, sign-flip $p = 0.0238$) is exploratory, uncorrected across the six-test family and failing the Bonferroni threshold of 0.0083, a lead rather than a result. **(B)** The same for the collapse configuration. **(C)** Dominance rate per placement for the freeze configuration: the fraction of runs in which the cluster standing largest at onset descends from the stamped patch is $0/20$ in every arm, so the bars are true but sit at zero height and the panel carries the count as text. **(D)** The same for the collapse configuration, also $0/20$ in every arm, which completes the six configuration-by-placement cells of the design. **(E)** Largest-cluster growth trajectories for a single exemplar seed under each arm in the freeze configuration. **(F)** The same in the collapse configuration. Which arm reaches domain scale first differs between the two configurations rather than tracking the stamped placement, and the domain that grows is in every run one founded away from the patch.

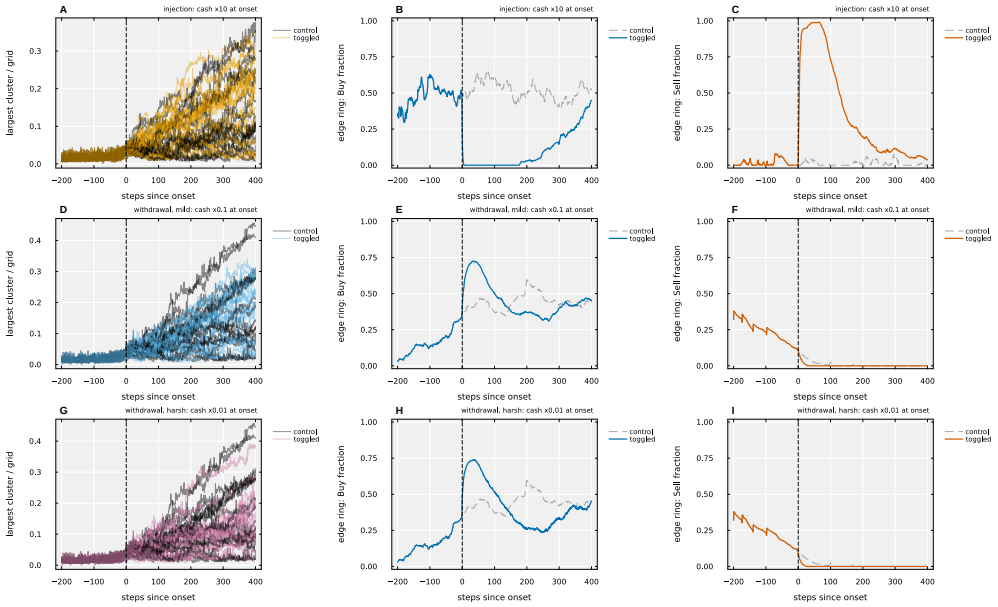


Figure S6: Full budget-toggle detail: paired, onset-aligned counterfactuals at $f = 0.20$ over the 19 included seed pairs per arm, with the dashed line marking the toggle at t_{on} and the pre-toggle traces coinciding by construction. The edge-ring traces are averaged over the seed pairs, toggled against control. **(A)** Largest sell cluster as a fraction of the grid against steps since onset under injection (cash $\times 10$ on the tight cash-500 base), each toggled trajectory drawn over its control. **(B)** Edge-ring Buy fraction under injection. **(C)** Edge-ring Sell fraction under injection; the ring flips from the control's ~ 0.5 Buy and 0 Sell to 0 Buy and ~ 1.0 Sell at the toggle. **(D)** Largest sell cluster under the mild withdrawal (cash $\times 0.1$ on the loose cash-5,000 base). **(E)** Edge-ring Buy fraction under the mild withdrawal. **(F)** Edge-ring Sell fraction under the mild withdrawal. **(G)** Largest sell cluster under the harsh withdrawal (cash $\times 0.01$). **(H)** Edge-ring Buy fraction under the harsh withdrawal. **(I)** Edge-ring Sell fraction under the harsh withdrawal. Either cut lifts the ring's Buy fraction from the control's ~ 0.46 after onset to ~ 0.72–0.74 before the transient excess dissipates by ~ 200 steps, and the Sell fraction, which in this loose configuration declines from ~ 0.37 well before onset to zero just after, remains indistinguishable between control and toggled throughout.

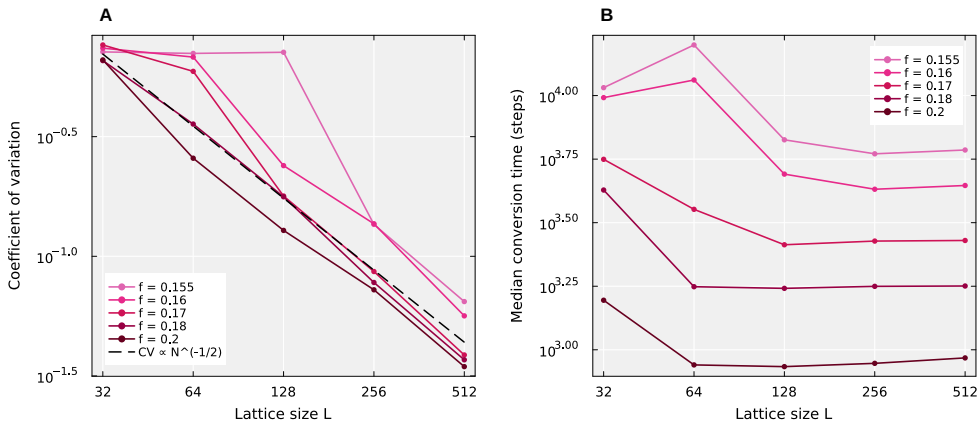


Figure S7: Timing of the Model A conversion across lattice sizes (Rules 19280/19286, periodic boundaries, no price forcing, twenty disorder seeds per cell, restricted to the cells in which every seed converted so that no horizon-pinned value enters a statistic). **(A)** Coefficient of variation of the conversion time against lattice size at five minority fractions, against the $N^{-1/2}$ self-averaging reference (dashed, anchored at $L = 32$). The lowest curve, $f = 0.20$, tracks that reference, while curves nearer the threshold stay flat at small sizes and join it only at larger ones. **(B)** Median conversion time against lattice size at the same five minority fractions, saturating above $f \approx 0.17$. The clock the conversion runs on is therefore size-independent while the outcome inside a fixed window is not.

Table 3: Model B Stage 2 factorial design. The six factors are crossed in full: $5 \times 4 \times 3 \times 2 \times 3 \times 3 = 1{,}080$ configurations, each run at 10 seeds for 2,000 steps (10,800 runs).

Factor	Levels	Count
Minority fraction f	$0,\ 0.1,\ 0.15,\ 0.2,\ 0.3$	5
Price impact λ	$0.005,\ 0.01,\ 0.02,\ 0.05$	4
Initial cash	$500,\ 1{,}000,\ 5{,}000$	3
Disposition rule	off, on	2
Band spread δ	$0,\ 0.5,\ 0.9$	3
Band half-width k	$0.5,\ 1,\ 2$	3

Table 4: Escape-time campaign: run count n, censored count c, and onset rate $(n - c)/n$ for every (f, cash) cell. Rows at $n = 20$ come from the pilot pass, rows at $n = 40$ from the refinement pass. A Kaplan–Meier median is finite only where the onset rate reaches 0.5. The qualifying cells are exactly the $n_f = 5/6/6$ fit points of Table 2.

	Cash 500			Cash 1,000			Cash 5,000		
f	n	c	rate	n	c	rate	n	c	rate
0.12	20	20	0.00	20	20	0.00	20	20	0.00
0.15	20	18	0.10	20	18	0.10	20	18	0.10
0.156	40	37	0.07	40	38	0.05	40	36	0.10
0.162	40	31	0.23	40	33	0.17	40	30	0.25
0.168	40	24	0.40	40	31	0.23	40	22	0.45
0.174	40	22	0.45	40	15	0.62	40	12	0.70
0.177	40	16	0.60	40	14	0.65	40	13	0.68
0.18	20	3	0.85	20	4	0.80	20	4	0.80
0.20	20	0	1.00	20	0	1.00	20	1	0.95
0.25	20	0	1.00	20	0	1.00	20	0	1.00
0.30	20	0	1.00	20	0	1.00	20	0	1.00